

Jeanne Boyarsky and Scott Selikoff

OCA

Oracle Certified Associate Java® SE 8 Programmer I STUDY GUIDE

EXAM 1Z0-808

Covers 100% of exam objectives, including developing Java applications, becoming proficient in Java data types, mastering operators and decision control structures, understanding encapsulation, class inheritance, polymorphism, and much more...

Includes interactive online learning environment and study tools with:

- + 3 custom practice exams
- + More than 200 Electronic Flashcards
- + Searchable key term glossary

 **SYBEX**
A Wiley Brand

www.allitebooks.com

OCA: Oracle® Certified Associate Java® SE 8 Programmer I

Study Guide

Exam 1Z0-808



OCA: Oracle® Certified Associate Java® SE 8 Programmer I

Study Guide

Exam 1Z0-808



Jeanne Boyarsky

Scott Selikoff

Senior Acquisitions Editor: Kenyon Brown
Development Editor: Alexa Murphy
Technical Editors: Ernest Friedman-Hill, Matt Dalen
Production Editor: Rebecca Anderson
Copy Editor: Liz Welch
Editorial Manager: Pete Gaughan
Vice President and Executive Group Publisher: Richard Swadley
Associate Publisher: Jim Minatel
Production Manager: Kathleen Wisor
Media Supervising Producer: Rich Graves
Book Designers: Judy Fung and Bill Gibson
Proofreader: Scott Klemp, Word One New York
Indexer: Ted Laux
Project Coordinator, Cover: Patrick Redmond
Cover Designer: Wiley
Cover Image: ©Getty Images Inc./Jeremy Woodhouse
Copyright © 2015 by John Wiley & Sons, Inc., Indianapolis, Indiana
Published simultaneously in Canada

ISBN: 978-1-118-95740-0

ISBN: 978-1-118-95741-7 (ebk.)

ISBN: 978-1-118-95742-4 (ebk.)

No part of this publication may be reproduced, stored in a retrieval system or transmitted in any form or by any means, electronic, mechanical, photocopying, recording, scanning or otherwise, except as permitted under Sections 107 or 108 of the 1976 United States Copyright Act, without either the prior written permission of the Publisher, or authorization through payment of the appropriate per-copy fee to the Copyright Clearance Center, 222 Rosewood Drive, Danvers, MA 01923, (978) 750-8400, fax (978) 646-8600. Requests to the Publisher for permission should be addressed to the Permissions Department, John Wiley & Sons, Inc., 111 River Street, Hoboken, NJ 07030, (201) 748-6011, fax (201) 748-6008, or online at <http://www.wiley.com/go/permissions>.

Limit of Liability/Disclaimer of Warranty: The publisher and the author make no representations or warranties with respect to the accuracy or completeness of the contents of this work and specifically disclaim all warranties, including without limitation warranties of fitness for a particular purpose. No warranty may be created or extended by sales or promotional materials. The advice and strategies contained herein may not be suitable for every situation. This work is sold with the understanding that the publisher is not engaged in rendering legal, accounting, or other professional services. If professional assistance is required, the services of a competent professional person should be sought. Neither the publisher nor the author shall be liable for damages arising herefrom. The fact that an organization or Web site is referred to in this work as a citation and/or a potential source of further information does not mean that the author or the publisher endorses the information the organization or Web site may provide or recommendations it may make. Further, readers should be aware that Internet Web sites listed in this work may have changed or disappeared between when this work was written and when it is read.

For general information on our other products and services or to obtain technical support, please contact our Customer Care Department within the U.S. at (877) 762-2974, outside the U.S. at (317) 572-3993 or fax (317) 572-4002.

Wiley publishes in a variety of print and electronic formats and by print-on-demand. Some material included with standard print versions of this book may not be included in e-books or in print-on-demand. If this book refers to media such as a CD or DVD that is not included in the version you purchased, you may download this material at <http://booksupport.wiley.com>. For more information about Wiley products, visit www.wiley.com.

Library of Congress Control Number: 2014954685

TRADEMARKS: Wiley, the Wiley logo, and the Sybex logo are trademarks or registered trademarks of John Wiley & Sons, Inc. and/or its affiliates, in the United States and other countries, and may not be used without written permission. Oracle and Java are registered trademarks of Oracle America, Inc. All other trademarks are the property of their respective owners. John Wiley & Sons, Inc. is not associated with any product or vendor mentioned in this book.

10 9 8 7 6 5 4 3 2 1

Dear Reader,

Thank you for choosing *OCA: Oracle Certified Associate Java SE 8 Programmer I Study Guide*. This book is part of a family of premium-quality Sybex books, all of which are written by outstanding authors who combine practical experience with a gift for teaching.

Sybex was founded in 1976. More than 30 years later, we're still committed to producing consistently exceptional books. With each of our titles, we're working hard to set a new standard for the industry. From the paper we print on, to the authors we work with, our goal is to bring you the best books available.

I hope you see all that reflected in these pages. I'd be very interested to hear your comments and get your feedback on how we're doing. Feel free to let me know what you think about this or any other Sybex book by sending me an email at contactus@wiley.com. If you think you've found a technical error in this book, please visit <http://sybex.custhelp.com>. Customer feedback is critical to our efforts at Sybex.

Best regards,

A handwritten signature in black ink, appearing to read "Chris Webb", written in a cursive style.

Chris Webb
Associate Publisher
Sybex, an Imprint of Wiley

To the programmers on FIRST robotics team 694.

—Jeanne

To my wife and the two little bundles of joy she is carrying.

—Scott

Acknowledgments

Jeanne and Scott would like to thank numerous individuals for their contribution to this book. Thank you to Developmental Editor Alexa Murphy for teaching us about Wiley's publishing process and making the book better in so many ways. Thank you to Ernest Friedman-Hill for being our Technical Editor as we wrote our first book. Ernest pointed out many subtle errors in addition to the big ones. And thank you to Matt Dalen for being our Technical Proofer and finding the errors we managed to sneak by Ernest. This book also wouldn't be possible without many people at Wiley, including Jeff Kellum, Kenyon Brown, Pete Gaughan, Rebecca Anderson, and so many others.

Jeanne would personally like to thank Chris Kreussling for knowing almost a decade ago that she would someday write a book. Erik Kariyev motivated her to write her first table of contents ever. Countless CodeRanch.com moderators warned Jeanne about how much work writing a book is to get her to the point where she was ready. Michael Ernest gave her extra advice on the Wiley process. Bert Bates let Jeanne dip her toe in by contributing to his Java 7 book and she learned a ton in the process. Scott was a great co-author and was available to bounce ideas off of or remind her to follow her own advice. Finally, Jeanne would like to thank all of the new programmers at CodeRanch.com and FIRST robotics team 694 for the constant reminders of how new programmers think.

Scott could not have reached this point without the help of a small army of people, led by his perpetually understanding wife Patti, without whose love and support this book would never have been possible. Professor Johannes Gehrke of Cornell University always believed in him and knew he would excel in his career. Jeanne's patience and guidance as co-author was invaluable while Scott adjusted to the learning curve of writing a book. Matt Dalen has been a wonderful friend and sounding board over the last year. Joel McNary introduced him to CodeRanch.com and encouraged him to post regularly, a step that changed his life. Finally, Scott would like to thank his mother and retired teacher Barbara Selikoff for teaching him the value of education and his father Mark Selikoff, for instilling in him the benefits of working hard.

About the Authors

Jeanne Boyarsky has worked as a Java developer for over 12 years at a bank in New York City where she develops, mentors, and conducts training. Besides being a senior moderator at CodeRanch.com in her free time, she works on the forum codebase. Jeanne also mentors the programming division of a FIRST robotics team, where she works with students just getting started with Java.

Jeanne got her Bachelor of Arts in 2002 and her Master's in Computer Information Technology in 2005. She enjoyed getting her Master's degree in an online program while working full time. This was before online education was cool! Jeanne is also a Distinguished Toastmaster and a Scrum Master. You can find out more about Jeanne at www.coderanch.com/how-to/java/BioJeanneBoyarsky.

Scott Selikoff is a professional software consultant, author, and owner of Selikoff Solutions, LLC, which provides software development solutions to businesses in the tri-state New York City area. Skilled in a plethora of software languages and platforms, Scott specializes in database-driven systems, web-based applications, and service-oriented architectures.

A native of Toms River, NJ, Scott achieved his Bachelor of Arts from Cornell University in Mathematics and Computer Science in 2002, after 3 years of study. In 2003, he received his Master's of Engineering in Computer Science, also from Cornell University.

As someone with a deep love of education, Scott has always enjoyed teaching others new concepts. He's given lectures at Cornell University and Rutgers University, as well as conferences including The Server Side Java Symposium. Scott lives in New Jersey with his loving wife and two very playful dogs, a Siberian husky named Webby and standard poodle named Georgette. You can find out more about Scott at www.linkedin.com/in/selikoff.

Jeanne and Scott are both moderators on the CodeRanch.com forums and can be reached there for questions and comments. They also co-author a technical blog called Down Home Country Coding at www.selikoff.net.

Contents at a Glance

<i>Introduction</i>		<i>xxi</i>
<i>Assessment Test</i>		<i>xxxi</i>
Chapter 1	Java Building Blocks	1
Chapter 2	Operators and Statements	51
Chapter 3	Core Java APIs	101
Chapter 4	Methods and Encapsulation	165
Chapter 5	Class Design	233
Chapter 6	Exceptions	299
Appendix A	Answers to Review Questions	333
Appendix B	Study Tips	353
<i>Index</i>		<i>367</i>

Contents

<i>Introduction</i>	<i>xxi</i>
<i>Assessment Test</i>	<i>xxxii</i>
Chapter 1	Java Building Blocks
	1
Understanding the Java Class Structure	2
Fields and Methods	2
Comments	4
Classes vs. Files	5
Writing a <i>main()</i> Method	6
Understanding Package Declarations and Imports	9
Wildcards	10
Redundant Imports	11
Naming Conflicts	12
Creating a New Package	13
Code Formatting on the Exam	16
Creating Objects	16
Constructors	17
Reading and Writing Object Fields	18
Instance_INITIALIZER Blocks	18
Order of Initialization	19
Distinguishing Between Object References and Primitives	20
Primitive Types	20
Reference Types	24
Key Differences	25
Declaring and Initializing Variables	25
Declaring Multiple Variables	26
Identifiers	27
Understanding Default Initialization of Variables	29
Local Variables	29
Instance and Class Variables	30
Understanding Variable Scope	31
Ordering Elements in a Class	34
Destroying Objects	36
Garbage Collection	36
<i>finalize()</i>	38
Benefits of Java	39
Summary	40
Exam Essentials	41
Review Questions	42

Chapter 2	Operators and Statements	51
	Understanding Java Operators	52
	Working with Binary Arithmetic Operators	53
	Arithmetic Operators	53
	Numeric Promotion	55
	Working with Unary Operators	57
	Logical Complement and Negation Operators	57
	Increment and Decrement Operators	58
	Using Additional Binary Operators	60
	Assignment Operators	60
	Compound Assignment Operators	62
	Relational Operators	63
	Logical Operators	64
	Equality Operators	65
	Understanding Java Statements	66
	The <i>if-then</i> Statement	67
	The <i>if-then-else</i> Statement	68
	The <i>switch</i> Statement	72
	The <i>while</i> Statement	76
	The <i>do-while</i> Statement	78
	The <i>for</i> Statement	80
	Understanding Advanced Flow Control	86
	Nested Loops	87
	Adding Optional Labels	87
	The <i>break</i> Statement	88
	The <i>continue</i> Statement	90
	Summary	92
	Exam Essentials	92
	Review Questions	94
Chapter 3	Core Java APIs	101
	Creating and Manipulating Strings	102
	Concatenation	102
	Immutability	104
	The String Pool	105
	Important String Methods	105
	Method Chaining	110
	Using the <i>StringBuilder</i> Class	111
	Mutability and Chaining	112
	Creating a <i>StringBuilder</i>	113
	Important <i>StringBuilder</i> Methods	114
	<i>StringBuilder</i> vs. <i>StringBuffer</i>	117

Understanding Equality	117
Understanding Java Arrays	119
Creating an Array of Primitives	119
Creating an Array with Reference Variables	121
Using an Array	123
Sorting	124
Searching	125
Varargs	126
Multidimensional Arrays	126
Understanding an <i>ArrayList</i>	129
Creating an <i>ArrayList</i>	129
Using an <i>ArrayList</i>	130
Wrapper Classes	134
Autoboxing	136
Converting Between array and <i>List</i>	136
Sorting	138
Working with Dates and Times	138
Creating Dates and Times	138
Manipulating Dates and Times	142
Working with Periods	145
Formatting Dates and Times	148
Parsing Dates and Times	151
Summary	151
Exam Essentials	152
Review Questions	153

Chapter 4 Methods and Encapsulation 165

Designing Methods	166
Optional Specifiers	168
Return Type	169
Method Name	170
Parameter List	171
Optional Exception List	171
Method Body	171
Working with Varargs	172
Applying Access Modifiers	173
Private Access	173
Default (Package Private) Access	175
Protected Access	176
Public Access	180
Designing Static Methods and Fields	181
Calling a Static Variable or Method	182
Static vs. Instance	183
Static Variables	185


Static Initialization	186
Static Imports	187
Passing Data Among Methods	188
Overloading Methods	191
Creating Constructors	196
Default Constructor	197
Overloading Constructors	199
Final Fields	202
Order of Initialization	202
Encapsulating Data	205
Creating Immutable Classes	207
Writing Simple Lambdas	208
Lambda Example	209
Lambda Syntax	211
Predicates	214
Summary	215
Exam Essentials	216
Review Questions	218
 Chapter 5	 233
Class Design	
Introducing Class Inheritance	234
Extending a Class	235
Applying Class Access Modifiers	237
Creating Java Objects	237
Defining Constructors	238
Calling Inherited Class Members	244
Inheriting Methods	246
Inheriting Variables	257
Creating Abstract Classes	259
Defining an Abstract Class	260
Creating a Concrete Class	262
Extending an Abstract Class	263
Implementing Interfaces	266
Defining an Interface	267
Inheriting an Interface	269
Interface Variables	273
Default Interface Methods	274
Static Interface Methods	278
Understanding Polymorphism	279
Object vs. Reference	281
Casting Objects	282
Virtual Methods	284
Polymorphic Parameters	285
Polymorphism and Method Overriding	287

	Summary	288
	Exam Essentials	289
	Review Questions	291
Chapter 6	Exceptions	299
	Understanding Exceptions	300
	The Role of Exceptions	300
	Understanding Exception Types	302
	Throwing an Exception	304
	Using a <i>try</i> Statement	305
	Adding a <i>finally</i> Block	307
	Catching Various Types of Exceptions	309
	Throwing a Second Exception	311
	Recognizing Common Exception Types	313
	Runtime Exceptions	314
	Checked Exceptions	317
	Errors	317
	Calling Methods That Throw Exceptions	318
	Subclasses	319
	Printing an Exception	321
	Summary	323
	Exam Essentials	324
	Review Questions	325
Appendix A	Answers to Review Questions	333
	Chapter 1: Java Building Blocks	334
	Chapter 2: Operators and Statements	336
	Chapter 3: Core Java APIs	339
	Chapter 4: Methods and Encapsulation	342
	Chapter 5: Class Design	346
	Chapter 6: Exceptions	349
Appendix B	Study Tips	353
	Studying for the Test	354
	Creating a Study Plan	354
	Creating and Running Sample Applications	355
	Taking the Test	359
	Understanding the Question	359
	Applying Process of Elimination	362
	Optimizing Your Time	364
	Getting a Good Night's Rest	366
	<i>Index</i>	367

Introduction

Java, “born” in 1995, is now just about 20 years old. As with anything 20 years old, there is a good amount of history and variation between versions of Java. Over the years, the certification exams have changed to cover different topics. The names of the exams have even changed. This book covers the Java 8 Oracle Certified Associate (OCA) exam.

If you read about “the exam” on the Web, you may see information about the older names for the exam. We’ve showed the changes in name. Here’s what happened. Sun Microsystems used to have two exams. The SCJP (Sun Certified Java Programmer) was meant for programmers and the SCJA (Sun Certified Java Associate) was meant for those who wanted broader knowledge. When Oracle bought Sun Microsystems, they changed all the names from Sun to Oracle, giving us the OCJP and OCJA.



Then Oracle made two strategic decisions with Java 7. They decided to stop updating the OCJA exam. They also decided to cover more on in the programmer space and split it into two exams. Now you first take the OCAJP (Oracle Certified Associate Java Programmer), also known as Java Programmer I, or OCA. That’s what this book is about. Then you take the OCPJP (Oracle Certified Professional Java Programmer), also known as Java Programmer II, or OCP. There’s also an upgrade exam in case you took an older version of the SCJP or OCPJP and want to upgrade. Most people refer to the current exams as OCA 8, OCP 8, and the Java 8 upgrade exam. We mention when a topic is split between the OCA and OCP so you know which parts are more advanced.

We try to keep the history to a minimum in this book. There are some places on the exam where you need to know both an “old way” and a “new way” of doing things. When that happens, we will be sure to tell you what version of Java introduced it. We will also let you know about topics that are not on the exam anymore in case you see questions in the older free online mock exams.

The OCA Exam

All you need to do to earn the Oracle Certified Associate Java SE 8 Programmer certification is to pass the exam! That's it.

Oracle has a tendency to fiddle with the length of the exam and the passing score once it comes out. Since it's pretty much a guarantee that whatever we tell you here will become obsolete, we will give you a feel for the range of variation. The OCA exam has varied between 60 and 90 questions since it was introduced. The score to pass has varied between 60 percent and 80 percent. The time allowed to take the exam has varied from two hours to two-and-a-half hours.

Oracle has a tendency to “tweak” the exam objectives over time as well. They do make minor additions and removals from what is covered on the exam. Although this tends to affect the OCP exam more than the OCA exam, there are a few topics that were added to the OCA for Java 8. It wouldn't be a surprise for Oracle to make changes.

Although there will likely be minor changes to the scope of the exam, it certainly isn't a secret. We've created a book page on our blog: www.selikoff.net/oca. If there are any changes to the topics on the exam after this book is published, we will note them there.

That book page also contains a link to the official exam page so that you can check the length and passing score that Oracle has chosen for the moment.

Scheduling the Exam

The exam is administered by Pearson VUE and can be taken at any Pearson VUE testing center. To find a testing center or register for the exam, go to www.pearsonvue.com. Choose IT and then Oracle. If you haven't been to the test center before, we recommend visiting in advance. Some testing centers are nice and professionally run. Others stick you in a closet with lots of people talking around you. You don't want to be taking the test with someone complaining about their broken laptop nearby!

At this time, you can reschedule the exam without penalty until up to 24 hours before. This means that you can register for a convenient time slot well in advance, knowing that you can delay if you aren't ready by that time. Rescheduling is easy and can be done on the Pearson VUE website. This may change, so check the rules before paying.

The Day of the Exam

When you go to take the exam, remember to bring two forms of ID, including one that is government issued. See Pearson's list of what is acceptable ID at <http://www.pearsonvue.com/policies/1S.pdf>. Try not to bring too much extra with you as it will not be allowed

into the exam room. While you will be allowed to check your belongings, it is better to leave extra items at home or in the car.

You will not be allowed to bring paper, your phone, and so forth into the exam room with you. Some centers are stricter than others. At one center, tissues were even taken away from us! Most centers allow keeping your ID and money. They watch you taking the exam, though, so don't even think about writing notes on money.

The exam center will give you writing materials to use during the exam. These are used as scratch paper during the exam to figure out answers and keep track of your thought process. The exam center will dispose of them at the end. Notice how we said "writing materials" rather than "pen and paper." Some centers still give pen and paper. Most give a small erasable board and a dry erase marker. If you have a preference to which you receive, call the testing center in advance to inquire.

Finding Out Your Score

In the past, you would find out right after finishing the exam if you passed. Now you have to wait nervously until you can check your score online.

If you go onto the Pearson VUE website, it will just have a status of "Taken" rather than your result. Oracle uses a separate system for scores. You'll need to go to <http://certview.oracle.com> to find out whether you passed and your score. It doesn't update immediately upon taking the test, but we haven't heard of it taking more than an hour. In addition to your score, you'll also see objectives for which you got a question wrong and instructions on how to get a hardcopy certificate.

At some point, you'll get an electronic certificate and some more time after that you'll receive your printed certificate. Sound vague? It is. The times reported to receive certificates vary widely.

Exam Questions

The OCA exam consists of multiple-choice questions. There are typically five or six possible answers. If a question has more than one answer, the question specifically states exactly how many correct answers there are. This book does not do that. We say "choose all that apply" to make the questions harder. This means the questions in this book are generally harder than those on the exam. The idea is to give you more practice so you can spot the correct answer more easily on the real exam.

Note that exam questions will sometimes have line numbers that begin with numbers higher than 1. This is to indicate that you are looking at a code snippet rather than a complete class. We follow this convention as well to get you used to it.

If you read about older versions of the exam online, you might see references to drag-and-drop questions. These questions had you do a puzzle on how to complete a piece of

code. There was also a bug in the exam software that caused your answers to get lost if you reviewed them again. Luckily, these are no longer on the exam.

Getting Started

We recommend reading Appendix B, “Study Tips,” before diving into the technical material in this book. Knowing how to approach studying will help you make better use of your study time.

Next, make sure you have downloaded version 8 of the JDK. If you learned Java some time ago, you might have version 7 or even earlier. There have been both big and small changes to the language. You could get a question wrong if you study with the wrong version.

Also, please check our book page to make sure Oracle hasn’t changed the objectives. For example, if Oracle decided that lambdas weren’t on the exam, you’d want to know that before studying. We will post any updates that you should know about at www.selikoff.net/oca.

Getting Help

Both of the authors are moderators at CodeRanch.com. CodeRanch.com is a very large and active programming forum that is very friendly toward Java beginners. It has a forum just for this exam called OCAJP. It also has a forum called Beginning Java for non-exam-specific questions. As you read the book, feel free to ask your questions in either of those forums. It could be you are having trouble compiling a class or that you are just plain confused about something. You’ll get an answer from a knowledgeable Java programmer. It might even be one of us.

Who Should Buy This Book

If you want to become certified as a Java programmer, this book is definitely for you. If you want to acquire a solid foundation in Java and your goal is to prepare for the exam, this book is also for you. You’ll find clear explanations of the concepts you need to grasp and plenty of help to achieve the high level of professional competency you need in order to succeed in your chosen field.

This book is intended to be understandable to anyone who has a tiny bit of Java knowledge. If you've never read a Java book before, we recommend starting with a book that teaches programming from the beginning and then returning to this study guide.

This book is for anyone from high school students to those beginning their programming journey to experienced professionals who need a review for the certification.

How This Book Is Organized

This book consists of six chapters, plus supplementary information: a glossary, this introduction, three appendices, and the assessment test after the introduction. You might have noticed that there are more than six exam objectives. We split up what you need to know to make it easy to learn and remember. Each chapter begins with a list of the objectives that are covered in that chapter.

The chapters are organized as follows:

- Chapter 1, “Java Building Blocks,” covers the basics of Java such as scoping variables and how to run a program. It also includes calling methods and types of variables.
- Chapter 2, “Operators and Statements,” focuses on the core logical constructs such as conditionals and loops. It also talks about the meaning and precedence of operators.
- Chapter 3, “Core Java APIs,” introduces you to array, ArrayList, String, StringBuilder, and various date classes.
- Chapter 4, “Methods and Encapsulation,” explains how to write methods, including access modifiers. It also shows how to call lambdas.
- Chapter 5, “Class Design,” adds interfaces and superclasses. It also includes casting and polymorphism.
- Chapter 6, “Exceptions,” shows the different types of exception classes and how to use them.

At the end of each chapter, you'll find a few elements you can use to prepare for the exam:

Summary This section reviews the most important topics that were covered in the chapter and serves as a good review.

Exam Essentials This section summarizes highlights that were covered in the chapter. You should be readily familiar with the key points of each chapter and be able to explain them in detail.

Review Questions Each chapter concludes with at least 20 review questions. You should answer these questions and check your answers against the ones provided in Appendix A.

If you can't answer at least 80 percent of these questions correctly, go back and review the chapter, or at least those sections that seem to be giving you difficulty.



The review questions, assessment test, and other testing elements included in this book are *not* derived from the real exam questions, so don't memorize the answers to these questions and assume that doing so will enable you to pass the exam. You should learn the underlying topic, as described in the text of the book. This will let you answer the questions provided with this book *and* pass the exam. Learning the underlying topic is also the approach that will serve you best in the workplace—the ultimate goal of a certification.

To get the most out of this book, you should read each chapter from start to finish before going to the chapter-end elements. They are most useful for checking and reinforcing your understanding. Even if you're already familiar with a topic, you should skim the chapter. There are a number of subtleties to Java that you could easily not encounter even when working with Java, even for years.

Free Online Learning Environment

This book provides a free online interactive learning environment and test bank with several additional elements. The online test bank includes:

Sample Tests All of the questions in this book, including the 20-question assessment test at the end of this introduction and over 130 questions that make up the Review Question sections for each chapter. In addition, there are three 60-question Practice Exams to test your knowledge of the material. The online test bank runs on multiple devices.

Electronic Flashcards Over 200 questions in flashcard format (a question followed by a single correct answer). You can use these to reinforce your learning and provide last-minute test prep before the exam.

Glossary The key terms from this book and their definitions are available as a fully searchable PDF.



Go to www.sybex.com/go/ocejavase8 to register and gain access to this comprehensive study tool package.

Conventions Used in This Book

This book uses certain typographic styles in order to help you quickly identify important information and to avoid confusion about the meaning of words, such as onscreen prompts. In particular, look for the following styles:

- *Italicized text* indicates key terms that are described at length for the first time in a chapter. (Italics are also used for emphasis.)
- A monospaced font indicates code or command-line text.
- *Italicized monospaced text* indicates a variable.

In addition to these text conventions, which can apply to individual words or entire paragraphs, a few conventions highlight segments of text:



A note indicates information that's useful or interesting. It is often something to pay special attention to for the exam.

Sidebars

A sidebar is like a note but longer. The information in a sidebar is useful, but it doesn't fit into the main flow of the text.



Real World Scenario

Real World Scenario

A real world scenario describes a task or an example that's particularly grounded in the real world. Although interesting, the scenario will not show up on the exam.

OCA Exam Objectives

OCA: Oracle Certified Associate Java SE 8 Programmer I Study Guide: Exam 1Z0-808 has been written to cover every OCA exam objective. The following table provides a breakdown of this book's exam coverage, showing you the chapter where each objective or sub-objective is covered:

Exam Objective	Chapter
■ Java Basics	
Define the scope of variables	1
Define the structure of a Java class	1
Create executable Java applications with a main method; run a Java program from the command line, including console output	1
Import other Java packages to make them accessible in your code	1
Compare and contrast the features and components of Java such as platform independence, object orientation, encapsulation, etc.	1
■ Working with Java Data Types	
Declare and initialize variables (including casting of primitive data types)	1
Differentiate between object reference variables and primitive variables	1
Know how to read or write to object fields	1
Explain an Object's Lifecycle (creation, "dereference by reassignment," and garbage collection)	1
Develop code that uses wrapper classes such as Boolean, Double, and Integer	1
■ Using Operators and Decision Constructs	
Use Java operators, including parentheses to override operator precedence	2
Test equality between Strings and other objects using == and equals ()	3
Create if and if/else and ternary constructs	2
Use a switch statement	2
■ Creating and Using Arrays	
Declare, instantiate, initialize, and use a one-dimensional array	3
Declare, instantiate, initialize, and use multi-dimensional array	3

Exam Objective	Chapter
■ Using Loop Constructs	
Create and use while loops	2
Create and use for loops including the enhanced for loop	2
Create and use do/while loops	2
Compare loop constructs	2
Use break and continue	2
■ Working with Methods and Encapsulation	
Create methods with arguments and return values, including overloaded methods	4
Apply the static keyword to methods and fields	4
Create and overload constructors, including impact on default constructors	4
Apply access modifiers	4
Apply encapsulation principles to a class	4
Determine the effect upon object references and primitive values when they are passed into methods that change the values	4
■ Working with Inheritance	
Describe inheritance and its benefits	5
Develop code that demonstrates the use of polymorphism, including overriding and object type versus reference type	5
Determine when casting is necessary	5
Use super and this to access objects and constructors	5
Use abstract classes and interfaces	5
■ Handling Exceptions	
Differentiate among checked exceptions, unchecked exceptions, and Errors	6
Create a try-catch block and determine how exceptions alter normal program flow	6

(continued)

Exam Objective	Chapter
Describe the advantages of Exception handling	6
Create and invoke a method that throws an exception	6
Recognize common exception classes (such as NullPointerException, ArithmeticException, ArrayIndexOutOfBoundsException, ClassCastException)	6
■ Working with Selected Classes from the Java API	
Manipulate data using the StringBuilder class and its methods	3
Creating and manipulating Strings	3
Create and manipulate calendar data using classes from java.time.LocalDate-Time, java.time.LocalDate, java.time.LocalTime, java.time.format.DateTime-Formatter, java.time.Period	3
Declare and use an ArrayList of a given type	3
Write a simple Lambda expression that consumes a Lambda Predicate expression	4

Assessment Test

1. What is the result of the following class? (Choose all that apply)

```
1: public class _C {  
2:     private static int $;  
3:     public static void main(String[] main) {  
4:         String a_b;  
5:         System.out.print($);  
6:         System.out.print(a_b);  
7:     } }
```

- A. Compiler error on line 1.
- B. Compiler error on line 2.
- C. Compiler error on line 4.
- D. Compiler error on line 5.
- E. Compiler error on line 6.
- F. `0null`
- G. `nullnull`

2. What is the result of the following code?

```
String s1 = "Java";  
String s2 = "Java";  
StringBuilder sb1 = new StringBuilder();  
sb1.append("Ja").append("va");  
System.out.println(s1 == s2);  
System.out.println(s1.equals(s2));  
System.out.println(sb1.toString() == s1);  
System.out.println(sb1.toString().equals(s1));
```

- A. `true` is printed out exactly once.
- B. `true` is printed out exactly twice.
- C. `true` is printed out exactly three times.
- D. `true` is printed out exactly four times.
- E. The code does not compile.

3. What is the output of the following code? (Choose all that apply)

```
1: interface HasTail { int getTailLength(); }  
2: abstract class Puma implements HasTail {  
3:     protected int getTailLength() {return 4;}  
4: }  
5: public class Cougar extends Puma {
```

```
6:   public static void main(String[] args) {
7:       Puma puma = new Puma();
8:       System.out.println(puma.getTailLength());
9:   }
10:
11:   public int getTailLength(int length) {return 2;}
12: }
```

- A.** 2
- B.** 4
- C.** The code will not compile because of line 3.
- D.** The code will not compile because of line 5.
- E.** The code will not compile because of line 7.
- F.** The code will not compile because of line 11.
- G.** The output cannot be determined from the code provided.

4. What is the output of the following program?

```
1: public class FeedingSchedule {
2:   public static void main(String[] args) {
3:       boolean keepGoing = true;
4:       int count = 0;
5:       int x = 3;
6:       while(count++ < 3) {
7:           int y = (1 + 2 * count) % 3;
8:           switch(y) {
9:               default:
10:                  case 0: x -= 1; break;
11:                  case 1: x += 5;
12:           }
13:       }
14:       System.out.println(x);
15: } }
```

- A.** 4
- B.** 5
- C.** 6
- D.** 7
- E.** 13
- F.** The code will not compile because of line 7.

5. What is the output of the following code snippet?

```
13: System.out.print("a");
14: try {
15:     System.out.print("b");
16:     throw new IllegalArgumentException();
17: } catch (RuntimeException e) {
18:     System.out.print("c");
19: } finally {
20:     System.out.print("d");
21: }
22: System.out.print("e");
```

- A. abe
- B. abce
- C. abde
- D. abcde
- E. The code does not compile.
- F. An uncaught exception is thrown.

6. What is the result of the following program?

```
1: public class MathFunctions {
2:     public static void addToInt(int x, int amountToAdd) {
3:         x = x + amountToAdd;
4:     }
5:     public static void main(String[] args) {
6:         int a = 15;
7:         int b = 10;
8:         MathFunctions.addToInt(a, b);
9:         System.out.println(a);    } }
```

- A. 10
- B. 15
- C. 25
- D. Compiler error on line 3.
- E. Compiler error on line 8.
- F. None of the above.

7. What is the result of the following code?

```
int[] array = {6,9,8};
List<Integer> list = new ArrayList<>();
```

```
list.add(array[0]);  
list.add(array[2]);  
list.set(1, array[1]);  
list.remove(0);  
System.out.println(list);
```

- A. [8]
 - B. [9]
 - C. Something like [Ljava.lang.String;@160bc7c0
 - D. An exception is thrown.
 - E. The code does not compile.
8. What is the output of the following code?
- ```
1: public class Deer {
2: public Deer() { System.out.print("Deer"); }
3: public Deer(int age) { System.out.print("DeerAge"); }
4: private boolean hasHorns() { return false; }
5: public static void main(String[] args) {
6: Deer deer = new Reindeer(5);
7: System.out.println(", "+deer.hasHorns());
8: }
9: }
10: class Reindeer extends Deer {
11: public Reindeer(int age) { System.out.print("Reindeer"); }
12: public boolean hasHorns() { return true; }
13: }
```
- A. DeerReindeer,false
  - B. DeerReindeer,true
  - C. ReindeerDeer,false
  - D. ReindeerDeer,true
  - E. DeerAgeReindeer,false
  - F. DeerAgeReindeer,true
  - G. The code will not compile because of line 7.
  - H. The code will not compile because of line 12.
9. Which of the following statements are true? (Choose all that apply)
- A. Checked exceptions are intended to be thrown by the JVM (and not the programmer).
  - B. Checked exceptions are required to be caught or declared.
  - C. Errors are intended to be thrown by the JVM (and not the programmer).
  - D. Errors are required to be caught or declared.
  - E. Runtime exceptions are intended to be thrown by the JVM (and not the programmer).
  - F. Runtime exceptions are required to be caught or declared.

**10.** Which are true of the following code? (Choose all that apply)

```
1: import java.util.*;
2: public class Grasshopper {
3: public Grasshopper(String n) {
4: name = n;
5: }
6: public static void main(String[] args) {
7: Grasshopper one = new Grasshopper("g1");
8: Grasshopper two = new Grasshopper("g2");
9: one = two;
10: two = null;
11: one = null;
12: }
13: private String name; }
```

- A.** Immediately after line 9, no grasshopper objects are eligible for garbage collection.
- B.** Immediately after line 10, no grasshopper objects are eligible for garbage collection.
- C.** Immediately after line 9, only one grasshopper object is eligible for garbage collection.
- D.** Immediately after line 10, only one grasshopper object is eligible for garbage collection.
- E.** Immediately after line 11, only one grasshopper object is eligible for garbage collection.
- F.** The code compiles.
- G.** The code does not compile.

**11.** What is the output of the following program?

```
1: public class FeedingSchedule {
2: public static void main(String[] args) {
3: int x = 5, j = 0;
4: OUTER: for(int i=0; i<3;)
5: INNER: do {
6: i++; x++;
7: if(x > 10) break INNER;
8: x += 4;
9: j++;
10: } while(j <= 2);
11: System.out.println(x);
12: } }
```

- A.** 10
- B.** 12
- C.** 13
- D.** 17
- E.** The code will not compile because of line 4.
- F.** The code will not compile because of line 6.

**12.** What is the result of the following program?

```
1: public class Egret {
2: private String color;
3: public Egret() {
4: this("white");
5: }
6: public Egret(String color) {
7: color = color;
8: }
9: public static void main(String[] args) {
10: Egret e = new Egret();
11: System.out.println("Color:" + e.color);
12: }
13: }
```

- A. Color:
- B. Color:null
- C. Color:White
- D. Compiler error on line 4.
- E. Compiler error on line 10.
- F. Compiler error on line 11.

**13.** What is the output of the following program?

```
1: public class BearOrShark {
2: public static void main(String[] args) {
3: int luck = 10;
4: if((luck>10 ? luck++: --luck)<10) {
5: System.out.print("Bear");
6: } if(luck<10) System.out.print("Shark");
7: } }
```

- A. Bear
- B. Shark
- C. BearShark
- D. The code will not compile because of line 4.
- E. The code will not compile because of line 6.
- F. The code compiles without issue but does not produce any output.

**14.** Assuming we have a valid, non-null HenHouse object whose value is initialized by the blank line shown here, which of the following are possible outputs of this application? (Choose all that apply)

```
1: class Chicken {}
2: interface HenHouse { public java.util.List<Chicken> getChickens(); }
3: public class ChickenSong {
```

```

4: public static void main(String[] args) {
5: HenHouse house = -----
6: Chicken chicken = house.getChickens().get(0);
7: for(int i=0; i<house.getChickens().size();
8: chicken = house.getChickens().get(i++)) {
9: System.out.println("Cluck");
10: } } }

```

- A. The code will not compile because of line 6.
  - B. The code will not compile because of lines 7–8.
  - C. The application will compile but not produce any output.
  - D. The application will output Cluck exactly once.
  - E. The application will output Cluck more than once.
  - F. The application will compile but produce an exception at runtime.
15. Which of the following statements can be inserted in the blank line so that the code will compile successfully? (Choose all that apply)

```

public interface CanSwim {}
public class Amphibian implements CanSwim {}
class Tadpole extends Amphibian {}
public class FindAllTadPole {
 public static void main(String[] args) {
 List<Tadpole> tadpoles = new ArrayList<Tadpole>();
 for(Amphibian amphibian : tadpoles) {
 ----- tadpole = amphibian;
 } } }

```

- A. CanSwim
  - B. Long
  - C. Amphibian
  - D. Tadpole
  - E. Object
16. What individual changes, if any, would allow the following code to compile? (Choose all that apply)
- ```

1: public interface Animal { public default String getName() { return null; } }
2: interface Mammal { public default String getName() { return null; } }
3: abstract class Otter implements Mammal, Animal {}

```
- A. The code compiles without issue.
 - B. Remove the default method modifier and method implementation on line 1.
 - C. Remove the default method modifier and method implementation on line 2.
 - D. Remove the default method modifier and method implementation on lines 1 and 2.
 - E. Change the return value on line 1 from null to "Animal".

- F. Override the `getName()` method with an abstract method in the `Otter` class.
 - G. Override the `getName()` method with a concrete method in the `Otter` class.
17. Which of the following lines can be inserted at line 11 to print true? (Choose all that apply)
- ```
10: public static void main(String[] args) {
11: // INSERT CODE HERE
12: }
13: private static boolean test(Predicate<Integer> p) {
14: return p.test(5);
15: }
```
- A. `System.out.println(test(i -> i == 5));`
  - B. `System.out.println(test(i -> {i == 5;}));`
  - C. `System.out.println(test((i) -> i == 5));`
  - D. `System.out.println(test((int i) -> i == 5);`
  - E. `System.out.println(test((int i) -> {return i == 5;}));`
  - F. `System.out.println(test((i) -> {return i == 5;}));`
18. Which of the following print out a date representing April 1, 2015? (Choose all that apply)
- A. `System.out.println(LocalDate.of(2015, Calendar.APRIL, 1));`
  - B. `System.out.println(LocalDate.of(2015, Month.APRIL, 1));`
  - C. `System.out.println(LocalDate.of(2015, 3, 1));`
  - D. `System.out.println(LocalDate.of(2015, 4, 1));`
  - E. `System.out.println(new LocalDate(2015, 3, 1));`
  - F. `System.out.println(new LocalDate(2015, 4, 1));`
19. Bytecode is in a file with which extension?
- A. `.bytecode`
  - B. `.bytes`
  - C. `.class`
  - D. `.exe`
  - E. `.javac`
  - F. `.java`
20. Which of the following are checked exceptions? (Choose all that apply)
- A. `Exception`
  - B. `IllegalArgumentException`
  - C. `IOException`
  - D. `NullPointerException`
  - E. `NumberFormatException`
  - F. `StackOverflowError`



# Answers to Assessment Test

1. E. Option E is correct because local variables require assignment before referencing them. Option D is incorrect because class and instance variables have default values and allow referencing. `a_b` defaults to a null value. Options A, B, and C are incorrect because identifiers may begin with a letter, underscore, or dollar sign. Options F and G are incorrect because the code does not compile. If `a_b` was an instance variable, the code would compile and output `0null`. For more information, see Chapter 1.
2. C. String literals are used from the string pool. This means that `s1` and `s2` refer to the same object and are equal. Therefore, the first two print statements print true. The third print statement prints false because `toString()` uses a method to compute the value and it is not from the string pool. The final print statement again prints true because `equals()` looks at the values of String objects. For more information, see Chapter 3.
3. C, D, E. First, the method `getTailLength()` in the interface `HasTail` is assumed to be public, since it is part of an interface. The implementation of the method on line 3 is therefore an invalid override, as `protected` is a more restrictive access modifier than public, so option C is correct. Next, the class `Cougar` implements an overloaded version of `getTailLength()`, but since the declaration in the parent class `Puma` is invalid, it needs to implement a public version of the method. Since it does not, the declaration of `Puma` is invalid, so option D is correct. Option E is incorrect, since `Puma` is marked abstract and cannot be instantiated. The overloaded method on line 11 is declared correctly, so option F is not correct. Finally, as the code has multiple compiler errors, options A, B, and G can be eliminated. For more information, see Chapter 5.
4. C. The code compiles and runs without issue; therefore, option F is incorrect. This type of problem is best examined one loop iteration at a time:
  - The loop continues as count loop expression evaluates to  $0 < 3$ , which is true, with `y` taking a new value of 1. The value of `y` is set to:
 

```
y = (1 + 2 * 1) % 3
 = (1 + 2) % 3
 = 3 % 3
 = 0
```
  - The first case block is called and the value of `x` is then set to:
 

```
x = 3 - 1 = 2
```
  - The loop continues as count loop expression evaluates to  $1 < 3$ , which is true, with `y` taking a new value of 2. The value of `y` is set to:
 

```
y = (1 + 2 * 2) % 3
 = (1 + 4) % 3
 = 4 % 3
 = 2
```

- The default block is called and the value of  $x$  is set to:  

$$x = 2 - 1 = 1$$
  - The loop continues as the count loop expression evaluates to  $2 < 3$ , which is true, with  $y$  taking a new value of 3. The value of  $y$  is set to:  

$$\begin{aligned} y &= (1 + 2 * 3) \% 3 \\ &= (1 + 6) \% 3 \\ &= 7 \% 3 \\ &= 1 \end{aligned}$$
  - The second case block is called and the value of  $x$  is then set to:  

$$x = 1 + 5 = 6$$
  - The loop ends as the count loop expression evaluates to  $3 < 3$ , with  $y$  also taking a new value of 4. The most recent value of  $x$ , 6, is output, so the answer is option C. For more information, see Chapter 2.
5. D. The code starts running and prints  $a$  and  $b$  on lines 13 and 15. Line 16 throws an exception, which is caught on line 17. After line 18 prints  $c$ , the `finally` block is run and  $d$  is printed. Then the `try` statement ends and  $e$  is printed on line 22. For more information, see Chapter 6.
  6. B. The code compiles successfully, so options D and E are incorrect. The value of  $a$  cannot be changed by the `addToInt` method, no matter what the method does, because only a copy of the variable is passed into the parameter  $x$ . Therefore,  $a$  does not change and the output on line 9 is 15. For more information, see Chapter 4.
  7. B. The array is allowed to use an anonymous initializer because it is in the same line as the declaration. The `ArrayList` uses the diamond operator allowed since Java 7. This specifies the type matches the one on the left without having to re-type it. After adding the two elements, `list` contains `[6, 8]`. We then replace the element at index 1 with 9, resulting in `[6, 9]`. Finally, we remove the element at index 0, leaving `[9]`. Option C is incorrect because arrays output something like that rather than an `ArrayList`. For more information, see Chapter 3.
  8. A. The code compiles and runs without issue, so options G and H are incorrect. First, the `Reindeer` object is instantiated using the constructor that takes an `int` value. Since there is no explicit call to the parent constructor, the default no-argument `super()` is inserted as the first line of the constructor. The output is then `Deer`, followed by `Reindeer` in the child constructor, so only options A and B can be correct. Next, the method `hasHorns()` looks like an overridden method, but it is actually a hidden method since it is declared `private` in the parent class. Because the hidden method is referenced in the parent class, the parent version is used, so the code outputs `false`, and option A is the correct answer.
  9. B, C. Only checked exceptions are required to be handled (caught) or declared. Runtime exceptions are commonly thrown by both the JVM and programmer code. Checked exceptions are usually thrown by programmer code. Errors are intended to be

thrown by the JVM. While a programmer could throw one, this would be a horrible practice. For more information, see Chapter 6.

10. C, D, F. Immediately after line 9, only Grasshopper `g1` is eligible for garbage collection since both one and two point to Grasshopper `g2`. Immediately after line 10, we still only have Grasshopper `g1` eligible for garbage collection. Reference one points to `g1` and reference two is `null`. Immediately after line 11, both Grasshopper objects are eligible for garbage collection since both one and two point to `null`. The code does compile. Although it is traditional to declare instance variables early in the class, you don't have to. For more information, see Chapter 1.
11. B. The code compiles and runs without issue; therefore, options E and F are incorrect. This type of problem is best examined one loop iteration at a time:
  - On the first iteration of the outer loop `i` is 0, so the loop continues.
  - On the first iteration of the inner loop, `i` is updated to 1 and `x` to 6. The `if-then` statement branch is not executed, and `x` is increased to 10 and `j` to 1.
  - On the second iteration of the inner loop (since `j` = 1 and `1` <= 2), `i` is updated to 2 and `x` to 11. At this point, the `if-then` branch will evaluate to true for the remainder of the program run, which causes the flow to break out of the inner loop each time it is reached.
  - On the second iteration of the outer loop (since `i` = 2), `i` is updated to 3 and `x` to 12. As before, the inner loop is broken since `x` is still greater than 10.
  - On the third iteration of the outer loop, the outer loop is broken, as `i` is already not less than 3. The most recent value of `x`, 12, is output, so the answer is option B. For more information, see Chapter 2.
12. B. Line 10 calls the constructor on lines 3–5. That constructor calls the other constructor. However, the constructor on lines 6–8 assigns the method parameter to itself, which leaves the `color` instance variable on line 2 set to its default value of `null`. For more information, see Chapter 4.
13. C. The code compiles and runs without issue, so options D and E are correct. Remember that only one of the right-hand ternary expressions will be evaluated at runtime. Since `luck` is not less than 10, the second expression, `--luck`, will be evaluated, and since the pre-increment operator was used, the value returned will be 9, which is less than 10. So the first `if-then` statement will be visited and Bear will be output. Notice there is no `else` statement on line 6. Since `luck` is still less than 10, the second `if-then` statement will also be reached and Shark will be output; therefore, the correct answer is option C. For more information, see Chapter 2.
14. D, E, F. The code compiles without issue, so options A and B are incorrect. If `house.getChickens()` returns an array of one element, the code will output Cluck once, so option D is correct. If `house.getChickens()` returns an array of multiple elements, the code will output Cluck once for each element in the array, so option E is correct. Alternatively, if `house.getChickens()` returns an array of zero elements, then the code will

throw an `IndexOutOfBoundsException` on the call to `house.getChickens().get(0)`; therefore, option C is not possible and option F is correct. The code will also throw an exception if the array returned by `house.getChickens()` is null, so option F is possible under multiple circumstances. For more information, see Chapter 2.

15. A, C, E. The for-each loop automatically casts each `Tadpole` object to an `Amphibian` reference, which does not require an explicit cast because `Tadpole` is a subclass of `Amphibian`. From there, any parent class or interface that `Amphibian` inherits from is permitted without an explicit cast. This includes `CanSwim`, the interface `Amphibian` implements, and `Object`, which all classes extend from, so options A and E are correct. Option C is also correct since the reference is being cast to the same type, so no explicit cast is required. Option B is incorrect, since `Long` is not a parent of `Amphibian`. Option D is incorrect as well, although an explicit cast to `Tadpole` on the right-hand side of the expression would be required to allow the code to compile. For more information, see Chapter 5.
16. D, F, G. The code does not compile, **since a class cannot inherit two interfaces that both define default methods with the same signature, unless the class implementing the interfaces overrides it with an abstract or concrete method.** Therefore, option A is incorrect and options F and G are correct. The alternate approach is to make the `getName()` method abstract in the interfaces, because an interface may inherit two abstract methods with the same signature. The change must be made to both interfaces, though, so options B and C are incorrect if taken individually, and option D is correct since the changes are taken together. For more information, see Chapter 5.
17. A, C, F. **The only functional programming interface you need to memorize for the exam is `Predicate`. It takes a single parameter and returns a boolean. Lambda expressions with one parameter are allowed to omit the parentheses around the parameter list, making options A and C correct. The return statement is optional when a single statement is in the body, making option F correct.** Option B is incorrect because a return statement must be used if braces are included around the body. Options D and E are incorrect because the type is `Integer` in the predicate and `int` in the lambda. **Autoboxing works for collections not inferring predicates.** If these two were changed to `Integer`, they would be correct. For more information, see Chapter 4.
18. B, D. The new date APIs added in Java 8 use static methods rather than a constructor to create a new date, making options E and F incorrect. The months are indexed starting with 1 in these APIs, making options A and C incorrect. Option A uses the old `Calendar` constants which are indexed from 0. Therefore, options B and D are correct. For more information, see Chapter 3.
19. C. Files with the `.java` extension contain the Java source code and are compiled to files with the `.class` extension that contain the bytecode. For more information, see Chapter 1.
20. A, C. Option A is the exception base class, which is a checked exception. Options B, D, and E extend `RuntimeException` directly or indirectly and therefore are unchecked exceptions. Option F is a throwable and not an exception, and so should not be caught or declared. For more information, see Chapter 6.

# **OCA: Oracle® Certified Associate Java® SE 8 Programmer I**

**Study Guide**

**Exam 1Z0-808**





# Chapter 1

## Java Building Blocks

---

### OCA EXAM OBJECTIVES COVERED IN THIS CHAPTER:

#### ✓ Java Basics

- Define the scope of variables
- Define the structure of a Java class
- Create executable Java applications with a main method; run a Java program from the command line; including console output
- Import other Java packages to make them accessible in your code
- Compare and contrast the features and components of Java such as platform independence, object orientation, encapsulation, etc.

#### ✓ Working with Java Data Types

- Declare and initialize variables (including casting or primitive types)
- Differentiate between object reference variables and primitive variables
- Know how to read or write to object fields
- Explain an Object's Lifecycle (creation, "dereference by reassignment" and garbage collection)



Welcome to the beginning of your journey to become certified on Java. We assume this isn't the first Java programming book you've read. Although we do talk about the basics, we do so only because we want to make sure you have all the terminology and detail you'll need for the OCA exam. If you've never written a Java program before, we recommend you pick up an introductory book on any version of Java—something like *Head First Java, 2nd Edition* (O'Reilly Media, 2005); *Java for Dummies* (For Dummies, 2014), or *Thinking in Java, 4th Edition* (Prentice Hall, 2006). (It's okay if the book covers an older version of Java—even Java 1.3 is fine.) Then come back to this certification study guide.

This chapter covers the fundamentals of Java. You'll see how to define and run a Java class, and learn about packages, variables, and the object life cycle.

## Understanding the Java Class Structure

In Java programs, classes are the basic building blocks. When defining a *class*, you describe all the parts and characteristics of one of those building blocks. To use most classes, you have to create objects. An *object* is a runtime instance of a class in memory. All the various objects of all the different classes represent the state of your program.

In the following sections, we'll look at fields, methods, and comments. We'll also explore the relationship between classes and files.

### Fields and Methods

Java classes have two primary elements: *methods*, often called functions or procedures in other languages, and *fields*, more generally known as variables. Together these are called the *members* of the class. Variables hold the state of the program, and methods operate on that state. If the change is important to remember, a variable stores that change. That's all classes really do. It's the programmer who creates and arranges these elements in such a way that the resulting code is useful and, ideally, easy for other programmers to understand.

Other building blocks include interfaces, which you'll learn about in Chapter 5, "Class Design," and enums, which you'll learn about when you start studying for the OCP exam.



The simplest Java class you can write looks like this:

```
1: public class Animal {
2: }
```

Java calls a word with special meaning a *keyword*. The `public` keyword on line 1 means the class can be used by other classes. The `class` keyword indicates you're defining a class. `Animal` gives the name of the class. Granted, this isn't a very interesting class, so add your first field:

```
1: public class Animal {
2: String name;
3: }
```



The line numbers aren't part of the program; they're just there to make the code easier to talk about.

On line 2, we define a variable named `name`. We also define the type of that variable to be a `String`. A `String` is a value that we can put text into, such as **"this is a string"**. `String` is also a class supplied with Java. Next you can add methods:

```
1: public class Animal {
2: String name;
3: public String getName() {
4: return name;
5: }
6: public void setName(String newName) {
7: name = newName;
8: }
9: }
```

On lines 3–5, you've defined your first method. A method is an operation that can be called. Again, `public` is used to signify that this method may be called from other classes. Next comes the return type—in this case, the method returns a `String`. On lines 6–8 is another method. This one has a special return type called *void*. *void* means that no value at all is returned. This method requires information be supplied to it from the calling method; this information is called a *parameter*. `setName` has one parameter named `newName`, and it is of type `String`. This means the caller should pass in one `String` parameter and expect nothing to be returned.

The full declaration of a method is called a *method signature*. In this example, can you identify the return type and parameters?

```
public int numberVisitors(int month)
```

The return type is `int`, which is a numeric type. There's one parameter named *month*, which is of type `int` as well.

## Comments

Another common part of the code is called a *comment*. Because comments aren't executable code, you can place them anywhere. Comments make your code easier to read. You won't see many comments on the exam—the exam creators are trying to make the code difficult to read—but you'll see them in this book as we explain the code. And we hope you use them in your own code. There are three types of comments in Java. The first is called a single-line comment:

```
// comment until end of line
```

A single-line comment begins with two slashes. Anything you type after that on the same line is ignored by the compiler. Next comes the multiple-line comment:

```
/* Multiple
 * line comment
 */
```

A multiple-line comment (also known as a multiline comment) includes anything starting from the symbol `/*` until the symbol `*/`. People often type an asterisk (\*) at the beginning of each line of a multiline comment to make it easier to read, but you don't have to. Finally, we have a Javadoc comment:

```
/**
 * Javadoc multiple-line comment
 * @author Jeanne and Scott
 */
```

This comment is similar to a multiline comment except it starts with `/**`. This special syntax tells the Javadoc tool to pay attention to the comment. Javadoc comments have a specific structure that the Javadoc tool knows how to read. You won't see a Javadoc comment on the exam—just remember it exists so you can read up on it online when you start writing programs for others to use.

As a bit of practice, can you identify which type of comment each of these five words is in? Is it a single-line or a multiline comment?

```
/*
 * // anteater
 */
// bear
// // cat
```

```
// /* dog */
/* elephant */
/*
 * /* ferret */
 */
```

Did you look closely? Some of these are tricky. Even though comments technically aren't on the exam, it is good to practice to look at code carefully.

Okay, on to the answers. `anteater` is in a multiline comment. Everything between `/*` and `*/` is part of a multiline comment—even if it includes a single-line comment within it! `bear` is your basic single-line comment. `cat` and `dog` are also single-line comments. Everything from `//` to the end of the line is part of the comment, even if it is another type of comment. `elephant` is your basic multiline comment.

The line with `ferret` is interesting in that it doesn't compile. Everything from the first `/*` to the first `*/` is part of the comment, which means the compiler sees something like this:

```
/* */ */
```

We have a problem. There is an extra `*/`. That's not valid syntax—a fact the compiler is happy to inform you about.

## Classes vs. Files

Most of the time, each Java class is defined in its own `*.java` file. It is usually `public`, which means any code can call it. Interestingly, **Java does not require that the class be `public`**. For example, this class is just fine:

```
1: class Animal {
2: String name;
3: }
```

**You can even put two classes in the same file. When you do so, at most one of the classes in the file is allowed to be `public`**. That means a file containing the following is also fine:

```
1: public class Animal {
2: private String name;
3: }
4: class Animal2 {
5: }
```

If you do have a public class, it needs to match the filename. `public class Animal2` would not compile in a file named `Animal.java`. In Chapter 5, we will discuss what non-public access means.

## Writing a *main()* Method

A Java program begins execution with its *main()* *method*. A *main()* method is the gateway between the startup of a Java process, which is managed by the *Java Virtual Machine* (JVM), and the beginning of the programmer's code. The JVM calls on the underlying system to allocate memory and CPU time, access files, and so on.

The *main()* method lets us hook our code into this process, keeping it alive long enough to do the work we've coded. The simplest possible class with a *main()* method looks like this:

```
1: public class Zoo {
2: public static void main(String[] args) {
3:
4: }
5:}
```

This code doesn't do anything useful (or harmful). It has no instructions other than to declare the entry point. It does illustrate, in a sense, that what you can put in a *main()* method is arbitrary. Any legal Java code will do. In fact, the only reason we even need a class structure to start a Java program is because the language requires it. To compile and execute this code, type it into a file called *Zoo.java* and execute the following:

```
$ javac Zoo.java
$ java Zoo
```

If you don't get any error messages, you were successful. If you do get error messages, check that you've installed a Java Development Kit (JDK) and not a Java Runtime Environment (JRE), that you have added it to the *PATH*, and that you didn't make any typos in the example. If you have any of these problems and don't know what to do, post a question with the error message you received in the Beginning Java forum at CodeRanch ([www.coderanch.com/forums/f-33/java](http://www.coderanch.com/forums/f-33/java)).

To compile Java code, the file must have the extension *.java*. The name of the file must match the name of the class. The result is a file of *bytecode* by the same name, but with a *.class* filename extension. Bytecode consists of instructions that the JVM knows how to execute. Notice that we must omit the *.class* extension to run *Zoo.java* because the period has a reserved meaning in the JVM.

The rules for what a Java code file contains, and in what order, are more detailed than what we have explained so far (there is more on this topic later in the chapter). To keep things simple for now, we'll follow a subset of the rules:

- Each file can contain only one class.
- The filename must match the class name, including case, and have a *.java* extension.

Suppose we replace line 3 in `Zoo.java` with `System.out.println("Welcome!");`. When we compile and run the code again, we'll get the line of output that matches what's between the quotes. In other words, the program will output `Welcome!`.

Let's first review the words in the `main()` method's signature, one at a time. The keyword `public` is what's called an *access modifier*. It declares this method's level of exposure to potential callers in the program. Naturally, `public` means anywhere in the program. You'll learn about access modifiers in Chapter 4, "Methods and Encapsulation."

The keyword `static` binds a method to its class so it can be called by just the class name, as in, for example, `Zoo.main()`. Java doesn't need to create an object to call the `main()` method—which is good since you haven't learned about creating objects yet! In fact, the JVM does this, more or less, when loading the class name given to it. If a `main()` method isn't present in the class we name with the `.java` executable, the process will throw an error and terminate. Even if a `main()` method is present, Java will throw an exception if it isn't static. A nonstatic `main()` method might as well be invisible from the point of view of the JVM. We'll see `static` again in Chapter 4.

The keyword `void` represents the *return type*. A method that returns no data returns control to the caller silently. In general, it's good practice to use `void` for methods that change an object's state. In that sense, the `main()` method changes the program state from started to finished. We will explore return types in Chapter 4 as well. Excited for Chapter 4 yet?

Finally we arrive at the `main()` method's parameter list, represented as an array of `java.lang.String` objects. In practice, you can write `String[] args`, `String args[]` or `String... args`; the compiler accepts any of these. The variable name `args` hints that this list contains values that were read in (arguments) when the JVM started. You can use any name you like, though. The characters `[]` are brackets and represent an array. An array is a fixed-size list of items that are all of the same type. The characters `...` are called *varargs* (variable argument lists). You will learn about `String` in Chapter 2, "Operators and Statements." Arrays and *varargs* will follow in Chapter 3, "Core Java APIs."

Let's see how to use the `args` parameter. First we modify the `Zoo` program to print out the first two arguments passed in:

```
public class Zoo {
 public static void main(String[] args) {
 System.out.println(args[0]);
 System.out.println(args[1]);
 } }
```

`args[0]` accesses the first element of the array. That's right: array indexes begin with 0 in Java. To run it, type this:

```
$ javac Zoo.java
$ java Zoo Bronx Zoo
```

The output is what you might expect:

```
Bronx
Zoo
```

The program correctly identifies the first two “words” as the arguments. Spaces are used to separate the arguments. If you want spaces inside an argument, you need to use quotes as in this example:

```
$ javac Zoo.java
$ java Zoo "San Diego" Zoo
```

Now we have a space in the output:

```
San Diego
Zoo
```

All command-line arguments are treated as `String` objects, even if they represent another data type:

```
$ javac Zoo.java
$ java Zoo Zoo 2
```

No matter. You still get the values output as `Strings`. In Chapter 2, you’ll learn how to convert `Strings` to numbers.

```
Zoo
2
```

Finally, what happens if you don’t pass in enough arguments?

```
$ javac Zoo.java
$ java Zoo Zoo
```

Reading `args[0]` goes fine and `Zoo` is printed out. Then Java panics. There’s no second argument! What to do? Java prints out an exception telling you it has no idea what to do with this argument at position 1. (You’ll learn about exceptions in Chapter 6, “Exceptions.”)

```
 ZooException in thread "main"
java.lang.ArrayIndexOutOfBoundsException: 1
 at mainmethod.Zoo.main(Zoo.java:7)
```

To review, you need to have a JDK to compile because it includes a compiler. You do not need to have a JDK to run the code—a JRE is enough. Java class files run on the JVM and therefore run on any machine with Java rather than just the machine or operating system they happened to have been compiled on.

# Understanding Package Declarations and Imports

Java comes with thousands of built-in classes, and there are countless more from developers like you. With all those classes, Java needs a way to organize them. It handles this in a way similar to a file cabinet. You put all your pieces of paper in folders. Java puts classes in *packages*. These are logical groupings for classes.

We wouldn't put you in front of a file cabinet and tell you to find a specific paper. Instead, we'd tell you which folder to look in. Java works the same way. It needs you to tell it which packages to look in to find code.

Suppose you try to compile this code:

```
public class ImportExample {
 public static void main(String[] args) {
 Random r = new Random(); // DOES NOT COMPILE
 System.out.println(r.nextInt(10));
 }
}
```

The Java compiler helpfully gives you an error that looks like this:

Random cannot be resolved to a type

This error could mean you made a typo in the name of the class. You double-check and discover that you didn't. The other cause of this error is omitting a needed *import* statement. Import statements tell Java which packages to look in for classes. Since you didn't tell Java where to look for `Random`, it has no clue.

Trying this again with the import allows you to compile:

```
import java.util.Random; // import tells us where to find Random
public class ImportExample {
 public static void main(String[] args) {
 Random r = new Random();
 System.out.println(r.nextInt(10)); // print a number between 0 and 9
 }
}
```

Now the code runs; it prints out a random number between 0 and 9. Just like arrays, Java likes to begin counting with 0.

Java classes are grouped into packages. The import statement tells the compiler which package to look in to find a class. This is similar to how mailing a letter works.

Imagine you are mailing a letter to 123 Main St., Apartment 9. The mail carrier first brings the letter to 123 Main St. Then she looks for the mailbox for apartment number 9. The address is like the package name in Java. The apartment number is like the class name in Java. Just as the mail carrier only looks at apartment numbers in the building, Java only looks for class names in the package.

Package names are hierarchical like the mail as well. The postal service starts with the top level, looking at your country first. You start reading a package name at the beginning too. If it begins with `java` or `javax`, this means it came with the JDK. If it starts with something else, it likely shows where it came from using the website name in reverse. From example, `com.amazon.java8book` tells us the code came from `amazon.com`. After the website name, you can add whatever you want. For example, `com.amazon.java8.my.name` also came from `amazon.com`. Java calls more detailed packages *child packages*. `com.amazon.java8book` is a child package of `com.amazon`. You can tell because it's longer and thus more specific.

You'll see package names on the exam that don't follow this convention. Don't be surprised to see package names like `a.b.c`. The rule for package names is that they are mostly letters or numbers separated by dots. Technically, you're allowed a couple of other characters between the dots. The rules are the same as for variable names, which you'll see later in the chapter. The exam may try to trick you with invalid variable names. Luckily, it doesn't try to trick you by giving invalid package names.

In the following sections, we'll look at imports with wildcards, naming conflicts with imports, how to create a package of your own, and how the exam formats code.

## Wildcards

Classes in the same package are often imported together. You can use a shortcut to import all the classes in a package:

```
import java.util.*; // imports java.util.Random among other things
public class ImportExample {
 public static void main(String[] args) {
 Random r = new Random();
 System.out.println(r.nextInt(10));
 }
}
```

In this example, we imported `java.util.Random` and a pile of other classes. The `*` is a wildcard that matches all classes in the package. Every class in the `java.util` package is available to this program when Java compiles it. It doesn't import child packages, fields, or methods; it imports only classes. (Okay, it's only classes for now, but there's a special type of import called the "static import" that imports other types. You'll learn more about that in Chapter 4.)

You might think that including so many classes slows down your program, but it doesn't. The compiler figures out what's actually needed. Which approach you choose is personal preference.



Listing the classes used makes the code easier to read, especially for new programmers. Using the wildcard can shorten the import list. You'll see both approaches on the exam.

## Redundant Imports

Wait a minute! We've been referring to `System` without an import and Java found it just fine. There's one special package in the Java world called `java.lang`. This package is special in that it is automatically imported. You can still type this package in an import statement, but you don't have to. In the following code, how many of the imports do you think are redundant?

```
1: import java.lang.System;
2: import java.lang.*;
3: import java.util.Random;
4: import java.util.*;
5: public class ImportExample {
6: public static void main(String[] args) {
7: Random r = new Random();
8: System.out.println(r.nextInt(10));
9: }
10: }
```

The answer is that three of the imports are redundant. Lines 1 and 2 are redundant because everything in `java.lang` is automatically considered to be imported. Line 4 is also redundant in this example because `Random` is already imported from `java.util.Random`. If line 3 wasn't present, `java.util.*` wouldn't be redundant, though, since it would cover importing `Random`.

Another case of redundancy involves importing a class that is in the same package as the class importing it. Java automatically looks in the current package for other classes.

Let's take a look at one more example to make sure you understand the edge cases for imports. For this example, `Files` and `Paths` are both in the package `java.nio.file`. You don't need to memorize this package for the OCA exam (but you should know it for the OCP exam). When testing your understanding of packages and imports, the OCA exam will use packages you may never have seen before. The question will let you know which package the class is in if you need to know that in order to answer the question.

What imports do you think would work to get this code to compile?

```
public class InputImports {
 public void read(Files files) {
 Paths.get("name");
 }
}
```

There are two possible answers. The shorter one is to use a wildcard to import both at the same time:

```
import java.nio.file.*;
```

The other answer is to import both classes explicitly:

```
import java.nio.file.Files;
import java.nio.file.Paths;
```

Now let's consider some imports that don't work:

```
import java.nio.*; // NO GOOD - a wildcard only matches
 //class names, not "file.*Files"
import java.nio.*.*; // NO GOOD - you can only have one wildcard
 //and it must be at the end
import java.nio.files.Paths.*; // NO GOOD - you cannot import methods
 //only class names
```

## Naming Conflicts

One of the reasons for using packages is so that class names don't have to be unique across all of Java. This means you'll sometimes want to import a class that can be found in multiple places. A common example of this is the `Date` class. Java provides implementations of `java.util.Date` and `java.sql.Date`. This is another example where you don't need to know the package names for the OCA exam—they will be provided to you. What import could we use if we want the `java.util.Date` version?

```
public class Conflicts {
 Date date;
 // some more code
}
```

The answer should be easy by now. You can write either **`import java.util.*;`** or **`import java.util.Date;`**. The tricky cases come about when other imports are present:

```
import java.util.*;
import java.sql.*; // DOES NOT COMPILE
```

When the class is found in multiple packages, Java gives you the compiler error:

The type `Date` is ambiguous

In our example, the solution is easy—remove the `java.sql.Date` import that we don't need. But what do we do if we need a whole pile of other classes in the `java.sql` package?

```
import java.util.Date;
import java.sql.*;
```

Ah, now it works. If you explicitly import a class name, it takes precedence over any wildcards present. Java thinks, “Okay! The programmer really wants me to assume use of the `java.util.Date` class.”

One more example. What does Java do with “ties” for precedence?

```
import java.util.Date;
import java.sql.Date;
```

Java is smart enough to detect that this code is no good. As a programmer, you’ve claimed to explicitly want the default to be both the `java.util.Date` and `java.sql.Date` implementations. Because there can’t be two defaults, the compiler tells you:

The import `java.sql.Date` collides with another import statement

### If You Really Need to Use Two Classes with the Same Name...

Sometimes you really do want to use `Date` from two different packages. When this happens, you can pick one to use in the import and use the other’s fully qualified class name (the package name, a dot, and the class name) to specify that it’s special. For example:

```
import java.util.Date;

public class Conflicts {
 Date date;
 java.sql.Date sqlDate;

}
```

Or you could have neither with an import and always use the fully qualified class name:

```
public class Conflicts {
 java.util.Date date;
 java.sql.Date sqlDate;

}
```

## Creating a New Package

Up to now, all the code we’ve written in this chapter has been in the *default package*. This is a special unnamed package that you should use only for throwaway code. You can tell the code is in the default package, because there’s no package name. On the exam, you’ll see the default package used a lot to save space in code listings. In real life, always name your packages to avoid naming conflicts and to allow others to reuse your code.

Now it's time to create a new package. The directory structure on your computer is related to the package name. Suppose we have these two classes:

```
C:\temp\packagea\ClassA.java
```

```
package packagea;
public class ClassA {
}
```

```
C:\temp\packageb\ClassB.java
```

```
package packageb;
import packagea.ClassA;
public class ClassB {
 public static void main(String[] args) {
 ClassA a;
 System.out.println("Got it");
 }
}
```

When you run a Java program, Java knows where to look for those package names. In this case, running from C:\temp works because both packagea and packageb are underneath it.

### Compiling Code with Packages

You'll learn Java much more easily by using the command line to compile and test your examples. Once you know the Java syntax well, you can switch to an integrated development environment (IDE) like Eclipse. An IDE will save you time in coding. But for the exam, your goal is to know details about the language and not have the IDE hide them for you.

Follow this example to make sure you know how to use the command line. If you have any problems following this procedure, post a question in the Beginning Java forum at CodeRanch ([www.coderanch.com/forums/f-33/java](http://www.coderanch.com/forums/f-33/java)). Describe what you tried and what the error said.

#### Windows Setup

Create the two files:

- C:\temp\packagea\ClassA.java
- C:\temp\packageb\ClassB.java

Then type this command:

```
cd C:\temp
```

**Mac/Linux Setup**

Create the two files:

- /tmp/packagea/ClassA.java
- /tmp/packageb/ClassB.java

Then type this command:

```
cd /tmp
```

**To Compile**

Type this command:

```
javac packagea/ClassA.java packageb/ClassB.java
```

If this command doesn't work, you'll get an error message. Check your files carefully for typos against the provided files. If the command does work, two new files will be created:

packagea/ClassA.class and packageb/ClassB.class.

**To Run**

Type this command:

```
java packageb.ClassB
```

If it works, you'll see `Got it` printed. You might have noticed we typed `ClassB` rather than `ClassB.class`. In Java you don't pass the extension when running a program.

**Class Paths and JARs**

You can also specify the location of the other files explicitly using a class path. This technique is useful when the class files are located elsewhere or in special JAR files. A JAR file is like a zip file of mainly Java class files. This goes beyond what you'll need to do on version 8 of the exam, although it appears on older versions.

On Windows, you type the following:

```
java -cp ".;C:\temp\someOtherLocation;c:\temp\myJar.jar" myPackage.MyClass
```

And on Mac OS/Linux, you type this:

```
java -cp ".: /tmp/someOtherLocation:/tmp/myJar.jar" myPackage.MyClass
```

The dot indicates you want to include the current directory in the class path. The rest of the command says to look for loose class files (or packages) in `someOtherLocation` and within `myJar.jar`. Windows uses semicolons to separate parts of the class path; other operating systems use colons.

Finally, you can use a wildcard (\*) to match all the JARs in a directory. Here's an example:

```
java -cp "C:\temp\directoryWithJars*" myPackage.MyClass
```

This command will add all the JARs to the class path that are in `directoryWithJars`. It won't include any JARs in the class path that are in a subdirectory of `directoryWithJars`.

## Code Formatting on the Exam

Not all questions will include the imports. If the exam isn't asking about imports in the question, it will often omit the imports to save space. You'll see examples with line numbers that don't begin with 1 in this case. The question is telling you, "Don't worry—imagine the code we omitted is correct; just focus on what I'm giving you." This means when you do see the line number 1 or no line numbers at all, you have to make sure imports aren't missing. Another thing the exam does to save space is to merge code on the same line. You should expect to see code like the following and to be asked whether it compiles. (You'll learn about `ArrayList` in Chapter 3—assume that part is good for now.)

```
6: public void method(ArrayList list) {
7: if (list.isEmpty()) { System.out.println("e");
8: } else { System.out.println("n");
9: } }
```

The answer here is that it does compile because the code starts below the imports. Now, what about this one? Does it compile?

```
1: public class LineNumbers {
2: public void method(ArrayList list) {
3: if (list.isEmpty()) { System.out.println("e");
4: } else { System.out.println("n");
5: } }
```

For this one, you would answer "Does not compile." Since the code begins with line 1, you don't get to assume that valid imports were provided earlier. The exam will let you know what package classes are in unless they're covered in the objectives. You'll be expected to know that `ArrayList` is in `java.util`—at least you will once you get to Chapter 3 of this book!

You'll also see code that doesn't have a `main()` method. When this happens, assume the `main()` method, class definition, and all necessary imports are present. You're just being asked if the part of the code you're shown compiles when dropped into valid surrounding code.

## Creating Objects

Our programs wouldn't be able to do anything useful if we didn't have the ability to create new objects. Remember that an **object is an instance of a class**. In the following sections, we'll look at constructors, object fields, instance initializers, and the order in which values are initialized.

## Constructors

To create an instance of a class, all you have to do is write `new` before it. For example:

```
Random r = new Random();
```

First you declare the type that you'll be creating (`Random`) and give the variable a name (`r`). This gives Java a place to store a reference to the object. Then you write `new Random()` to actually create the object.

`Random()` looks like a method since it is followed by parentheses. It's called a *constructor*, which is a special type of method that creates a new object. Now it's time to define a constructor of your own:

```
public class Chick {
 public Chick() {
 System.out.println("in constructor");
 }
}
```

There are two key points to note about the constructor: the name of the constructor matches the name of the class, and there's no return type. You'll likely see a method like this on the exam:

```
public void Chick() { } // NOT A CONSTRUCTOR
```

When you see a method name beginning with a capital letter and having a return type, pay special attention to it. It is *not* a constructor since there's a return type. It's a regular method that won't be called when you write `new Chick()`.

The purpose of a constructor is to initialize fields, although you can put any code in there. Another way to initialize fields is to do so directly on the line on which they're declared. This example shows both approaches:

```
public class Chicken {
 int numEggs = 0; // initialize on line
 String name;
 public Chicken() {
 name = "Duke"; // initialize in constructor
 } }
}
```

For most classes, you don't have to code a constructor—the compiler will supply a “do nothing” default constructor for you. There's one scenario that requires you to declare a constructor that you'll learn about in Chapter 5.

## Reading and Writing Object Fields

It's possible to read and write instance variables directly from the caller. In this example, a mother swan lays eggs:

```
public class Swan {
 int numberEggs; // instance variable
 public static void main(String[] args) {
 Swan mother = new Swan();
 mother.numberEggs = 1; // set variable
 System.out.println(mother.numberEggs); // read variable
 }
}
```

Reading a variable is known as *getting* it. The class gets *numberEggs* directly to print it out. Writing to a variable is known as *setting* it. This class sets *numberEggs* to 1.

In Chapter 4, you'll learn how to protect the *Swan* class from having someone set a negative number of eggs.

You can even read and write fields directly on the line declaring them:

```
1: public class Name {
2: String first = "Theodore";
3: String last = "Moose";
4: String full = first + last;
5: }
```

Lines 2 and 3 both write to fields. Line 4 does both. It reads the fields *first* and *last*. It then writes the field *full*.

## Instance Initializer Blocks

When you learned about methods, you saw braces (`{}`). The code between the braces is called a *code block*. Sometimes this code is called being inside the braces. Anywhere you see braces is a code block.

Sometimes code blocks are inside a method. These are run when the method is called. Other times, code blocks appear outside a method. These are called *instance initializers*. In Chapter 5, you'll learn how to use a static initializer.

How many blocks do you see in this example? How many instance initializers do you see?

```
3: public static void main(String[] args) {
4: { System.out.println("Feathers"); }
5: }
6: { System.out.println("Snowy"); }
```



There are three code blocks and one instance initializer. Counting code blocks is easy: you just count the number of pairs of braces. If there aren't the same number of open ({) and close (}) braces, the code doesn't compile. It doesn't matter that one set of braces is inside the `main()` method—it still counts.

When counting instance initializers, keep in mind that it does matter whether the braces are inside a method. There's only one pair of braces outside a method. Line 6 is an instance initializer.

## Order of Initialization

When writing code that initializes fields in multiple places, you have to keep track of the order of initialization. We'll add some more rules to the order of initialization in Chapters 4 and 5. In the meantime, you need to remember:

- Fields and instance initializer blocks are run in the order in which they appear in the file.
- The constructor runs after all fields and instance initializer blocks have run.

Let's look at an example:

```
1: public class Chick {
2: private String name = "Fluffy";
3: { System.out.println("setting field"); }
4: public Chick() {
5: name = "Tiny";
6: System.out.println("setting constructor");
7: }
8: public static void main(String[] args) {
9: Chick chick = new Chick();
10: System.out.println(chick.name); } }
```

Running this example prints this:

```
setting field
setting constructor
Tiny
```

Let's look at what's happening here. We start with the `main()` method because that's where Java starts execution. On line 9, we call the constructor of `Chick`. Java creates a new object. First it initializes `name` to "Fluffy" on line 2. Next it executes the print statement in the instance initializer on line 3. Once all the fields and instance initializers have run, Java returns to the constructor. Line 5 changes the value of `name` to "Tiny" and line 6 prints another statement. At this point, the constructor is done executing and goes back to the print statement on line 10.

Order matters for the fields and blocks of code. You can't refer to a variable before it has been initialized:

```
{ System.out.println(name); } // DOES NOT COMPILE
private String name = "Fluffy";
```

You should expect to see a question about initialization on the exam. Let's try one more. What do you think this code prints out?

```
public class Egg {
 public Egg() {
 number = 5;
 }
 public static void main(String[] args) {
 Egg egg = new Egg();
 System.out.println(egg.number);
 }
 private int number = 3;
 { number = 4; } }
```

If you answered 5, you got it right. Fields and blocks are run first in order, setting *number* to 3 and then 4. Then the constructor runs, setting *number* to 5.

## Distinguishing Between Object References and Primitives

Java applications contain two types of data: primitive types and reference types. In this section, we'll discuss the differences between a primitive type and a reference type.

### Primitive Types

Java has eight built-in data types, referred to as the Java *primitive types*. These eight data types represent the building blocks for Java objects, because all Java objects are just a complex collection of these primitive data types. The exam assumes you are well versed in the eight primitive data types, their relative sizes, and what can be stored in them.

Table 1.1 shows the Java primitive types together with their size in bytes and the range of values that each holds.

**TABLE 1.1** Java primitive types

| Keyword | Type                        | Example |
|---------|-----------------------------|---------|
| boolean | true or false               | true    |
| byte    | 8-bit integral value        | 123     |
| short   | 16-bit integral value       | 123     |
| int     | 32-bit integral value       | 123     |
| long    | 64-bit integral value       | 123     |
| float   | 32-bit floating-point value | 123.45f |
| double  | 64-bit floating-point value | 123.456 |
| char    | 16-bit Unicode value        | 'a'     |

There's a lot of information in Table 1.1. Let's look at some key points:

- float and double are used for floating-point (decimal) values.
- A float requires the letter f following the number so Java knows it is a float.
- byte, short, int, and long are used for numbers without decimal points.
- Each numeric type uses twice as many bits as the smaller similar type. For example, short uses twice as many bits as byte does.

You won't be asked about the exact sizes of most of these types. You should know that a byte can hold a value from -128 to 127. So you aren't stuck memorizing this, let's look at how Java gets that. A byte is 8 bits. A bit has two possible values. (These are basic computer science definitions that you should memorize.)  $2^8$  is  $2 \times 2 = 4 \times 2 = 8 \times 2 = 16 \times 2 = 32 \times 2 = 64 \times 2 = 128 \times 2 = 256$ . Since 0 needs to be included in the range, Java takes it away from the positive side. Or if you don't like math, you can just memorize it.

The number of bits is used by Java when it figures out how much memory to reserve for your variable. For example, Java allocates 32 bits if you write this:

```
int num;
```



## Real World Scenario

### What Is the Largest int?

You do not have to know this for the exam, but the maximum number an `int` can hold is 2,147,483,647. How do we know this? One way is to have Java tell us:

```
System.out.println(Integer.MAX_VALUE);
```

The other way is with math. An `int` is 32 bits.  $2^{32}$  is 4,294,967,296. Divide that by 2 and you get 2,147,483,648. Then subtract 1 as we did with bytes and you get 2,147,483,647. It's easier to just ask Java to print the value, isn't it?

There are a few more things you should know about numeric primitives. When a number is present in the code, it is called a *literal*. By default, Java assumes you are defining an `int` value with a literal. In this example, the number listed is bigger than what fits in an `int`. Remember, you aren't expected to memorize the maximum value for an `int`. The exam will include it in the question if it comes up.

```
long max = 3123456789; // DOES NOT COMPILE
```

Java complains the number is out of range. And it is—for an `int`. However, we don't have an `int`. The solution is to add the character `L` to the number:

```
long max = 3123456789L; // now Java knows it is a long
```

Alternatively, you could add a lowercase `l` to the number. But please use the uppercase `L`. The lowercase `l` looks like the number 1.

Another way to specify numbers is to change the “base.” When you learned how to count, you studied the digits 0–9. This numbering system is called base 10 since there are 10 numbers. It is also known as the decimal number system. Java allows you to specify digits in several other formats:

- octal (digits 0–7), which uses the number 0 as a prefix—for example, 017
- hexadecimal (digits 0–9 and letters A–F), which uses the number 0 followed by `x` or `X` as a prefix—for example, 0xFF
- binary (digits 0–1), which uses the number 0 followed by `b` or `B` as a prefix—for example, 0b10

You won't need to convert between number systems on the exam. You'll have to recognize valid literal values that can be assigned to numbers.

### Converting Back to Binary

Although you don't need to convert between number systems on the exam, we'll look at one example in case you're curious:

```
System.out.println(56); // 56
System.out.println(0b11); // 3
System.out.println(017); // 15
System.out.println(0x1F); // 31
```

First we have our normal base 10 value. We know you already know how to read that, but bear with us. The rightmost digit is 6, so it's "worth" 6. The second-to-rightmost digit is 5, so it's "worth" 50 (5 times 10.) Adding these together, we get 56.

Next we have binary, or base 2. The rightmost digit is 1 and is "worth" 1. The second-to-rightmost digit is also 1. In this case, it's "worth" 2 (1 times 2) because the base is 2. Adding these gets us 3.

Then comes octal, or base 8. The rightmost digit is 7 and is "worth" 7. The second-to-rightmost digit is 1. In this case, it's "worth" 8 (1 times 8) because the base is 8. Adding these gets us 15.

Finally, we have hexadecimal, or base 16, which is also known as hex. The rightmost "digit" is F and it's "worth" 15 (9 is "worth" 9, A is "worth" 10, B is "worth" 11, and so forth). The second-to-rightmost digit is 1. In this case, it's "worth" 16 (1 times 16) because the base is 16. Adding these gets us 31.

The last thing you need to know about numeric literals is a feature added in Java 7. You can have underscores in numbers to make them easier to read:

```
int million1 = 1000000;
int million2 = 1_000_000;
```

We'd rather be reading the latter one because the zeroes don't run together. You can add underscores anywhere except at the beginning of a literal, the end of a literal, right before a decimal point, or right after a decimal point. Let's look at a few examples:

```
double notAtStart = _1000.00; // DOES NOT COMPILE
double notAtEnd = 1000.00_; // DOES NOT COMPILE
double notByDecimal = 1000_.00; // DOES NOT COMPILE
double annoyingButLegal = 1_00_0.0_0; // this one compiles
```

## Reference Types

A *reference type* refers to an object (an instance of a class). Unlike primitive types that hold their values in the memory where the variable is allocated, references do not hold the value of the object they refer to. Instead, a reference “points” to an object by storing the memory address where the object is located, a concept referred to as a *pointer*. Unlike other languages, Java does not allow you to learn what the physical memory address is. You can only use the reference to refer to the object.

Let’s take a look at some examples that declare and initialize reference types. Suppose we declare a reference of type `java.util.Date` and a reference of type `String`:

```
java.util.Date today;
String greeting;
```

The `today` variable is a reference of type `Date` and can only point to a `Date` object. The `greeting` variable is a reference that can only point to a `String` object. A value is assigned to a reference in one of two ways:


- A reference can be assigned to another object of the same type.
- A reference can be assigned to a new object using the `new` keyword.

For example, the following statements assign these references to new objects:

```
today = new java.util.Date();
greeting = "How are you?";
```

The `today` reference now points to a new `Date` object in memory, and `today` can be used to access the various fields and methods of this `Date` object. Similarly, the `greeting` reference points to a new `String` object, “How are you?”. The `String` and `Date` objects do not have names and can be accessed only via their corresponding reference. Figure 1.1 shows how the reference types appear in memory.

**FIGURE 1.1** An object in memory can be accessed only via a reference.



## Key Differences

There are a few important differences you should know between primitives and reference types. First, reference types can be assigned `null`, which means they do not currently refer to an object. Primitive types will give you a compiler error if you attempt to assign them `null`. In this example, `value` cannot point to `null` because it is of type `int`:

```
int value = null; // DOES NOT COMPILE
String s = null;
```

Next, reference types can be used to call methods when they do not point to `null`. Primitives do not have methods declared on them. In this example, we can call a method on `reference` since it is of a reference type. You can tell `length` is a method because it has `()` after it. The following line is gibberish. No methods exist on `len` because it is an `int` primitive. Primitives do not have methods.

```
String reference = "hello";
int len = reference.length();
int bad = len.length(); // DOES NOT COMPILE
```

Finally, notice that all the primitive types have lowercase type names. All classes that come with Java begin with uppercase. You should follow this convention for classes you create as well.

## Declaring and Initializing Variables

We've seen some variables already. A *variable* is a name for a piece of memory that stores data. When you declare a variable, you need to state the variable type along with giving it a name. For example, the following code declares two variables. One is named `zooName` and is of type `String`. The other is named `numberAnimals` and is of type `int`.

```
String zooName;
int numberAnimals;
```

Now that we've declared a variable, we can give it a value. This is called initializing a variable. To initialize a variable, you just type the variable name followed by an equal sign, followed by the desired value:

```
zooName = "The Best Zoo";
numberAnimals = 100;
```

Since you often want to initialize a variable right away, you can do so in the same statement as the declaration. For example, here we merge the previous declarations and initializations into more concise code:

```
String zooName = "The Best Zoo";
int numberAnimals = 100;
```

In the following sections, we'll look at how to declare multiple variables in one-line and legal identifiers.

## Declaring Multiple Variables

You can also declare and initialize multiple variables in the same statement. How many variables do you think are declared and initialized in the following two lines?

```
String s1, s2;
String s3 = "yes", s4 = "no";
```

Four String variables were declared: *s1*, *s2*, *s3*, and *s4*. You can declare many variables in the same declaration as long as they are all of the same type. You can also initialize any or all of those values inline. In the previous example, we have two initialized variables: *s3* and *s4*. The other two variables remain declared but not yet initialized.

This is where it gets tricky. Pay attention to tricky things! The exam will attempt to trick you. Again, how many variables do you think are declared and initialized in this code?

```
int i1, i2, i3 = 0;
```

As you should expect, three variables were declared: *i1*, *i2*, and *i3*. However, only one of those values was initialized: *i3*. The other two remain declared but not yet initialized. That's the trick. Each snippet separated by a comma is a little declaration of its own. The initialization of *i3* only applies to *i3*. It doesn't have anything to do with *i1* or *i2* despite being in the same statement.

Another way the exam could try to trick you is to show you code like this line:

```
int num, String value; // DOES NOT COMPILE
```

This code doesn't compile because it tries to declare multiple variables of *different* types in the same statement. The shortcut to declare multiple variables in the same statement only works when they share a type.

To make sure you understand this, see if you can figure out which of the following are legal declarations. “Legal,” “valid,” and “compiles” are all synonyms in the Java exam world. We try to use all the terminology you could encounter on the exam.

```
boolean b1, b2;
String s1 = "1", s2;
double d1, double d2;
int i1; int i2;
int i3; i4;
```

The first statement is legal. It declares two variables without initializing them. The second statement is also legal. It declares two variables and initializes only one of them.

The third statement is *not* legal. Java does not allow you to declare two different types in the same statement. Wait a minute! Variables *d1* and *d2* are the same type. They are both



of type `double`. Although that's true, it still isn't allowed. If you want to declare multiple variables in the same statement, they must share the same type declaration and not repeat it. `double d1, d2;` would have been legal.

The fourth statement is legal. Although `int` does appear twice, each one is in a separate statement. A semicolon (;) separates statements in Java. It just so happens there are two completely different statements on the same line. The fifth statement is *not* legal. Again, we have two completely different statements on the same line. The second one is not a valid declaration because it omits the type. When you see an oddly placed semicolon on the exam, pretend the code is on separate lines and think about whether the code compiles that way. In this case, we have the following:

```
int i1;
int i2;
int i3;
i4;// DOES NOT COMPILE
```

Looking at the last line on its own, you can easily see that the declaration is invalid. And yes, the exam really does cram multiple statements onto the same line—partly to try to trick you and partly to fit more code on the screen. In the real world, please limit yourself to one declaration per statement and line. Your teammates will thank you for the readable code.

## Identifiers

It probably comes as no surprise that Java has precise rules about *identifier* names. Luckily, the same rules for identifiers apply to anything you are free to name, including variables, methods, classes, and fields.

**There are only three rules to remember for legal identifiers:**

- The name must begin with a letter or the symbol `$` or `_`.
- Subsequent characters may also be numbers.
- You cannot use the same name as a Java *reserved word*. As you might imagine, a reserved word is a keyword that Java has reserved so that you are not allowed to use it. Remember that Java is case sensitive, so you can use versions of the keywords that only differ in case. Please don't, though.

Don't worry—you won't need to memorize the full list of reserved words. The exam will only ask you about ones you've already learned, such as `class`. The following is a list of all the reserved words in Java. `const` and `goto` aren't actually used in Java. They are reserved so that people coming from other languages don't use them by accident—and in theory, in case Java wants to use them one day.

|           |            |           |              |            |
|-----------|------------|-----------|--------------|------------|
| abstract  | assert     | boolean   | break        | byte       |
| case      | catch      | char      | class        | const*     |
| continue  | default    | do        | double       | else       |
| enum      | extends    | false     | final        | finally    |
| float     | for        | goto*     | if           | implements |
| import    | instanceof | int       | interface    | long       |
| native    | new        | null      | package      | private    |
| protected | public     | return    | short        | static     |
| strictfp  | super      | switch    | synchronized | this       |
| throw     | throws     | transient | true         | try        |
| void      | volatile   | while     |              |            |

Prepare to be tested on these rules. The following examples are legal:

```
okidentifier
$OK2Identifier
_alsoOK1d3ntifi3r
__SStillOkbutKnotsonice$
```

These examples are not legal:

```
3DPointClass // identifiers cannot begin with a number
hollywood@vine // @ is not a letter, digit, $ or _
*$coffee // * is not a letter, digit, $ or _
public // public is a reserved word
```

Although you can do crazy things with identifier names, you shouldn't. Java has conventions so that code is readable and consistent. This consistency includes CamelCase. In CamelCase, each word begins with an uppercase letter. This makes multiple-word variable names easier to read. Which would you rather read: *Thisismyclass name* or *ThisIsMyClass name*? The exam will mostly use common conventions for identifiers, but not always. When you see a nonstandard identifier, be sure to check if it is legal. If not, you get to mark the answer “does not compile” and skip analyzing everything else in the question.



## Real World Scenario

### Identifiers in the Real World

Most Java developers follow these conventions for identifier names:

- Method and variables names begin with a lowercase letter followed by CamelCase.
- Class names begin with an uppercase letter followed by CamelCase. Don't start any identifiers with \$. The compiler uses this symbol for some files.

Also, know that valid letters in Java are not just characters in the English alphabet. Java supports the Unicode character set, so there are more than 45,000 characters that can start a legal Java identifier. A few hundred more are non-Arabic numerals that may appear after the first character in a legal identifier. Luckily, you don't have to worry about memorizing those for the exam. If you are in a country that doesn't use the English alphabet, this is useful to know for a job.

# Understanding Default Initialization of Variables

Before you can use a variable, it needs a value. Some types of variables get this value set automatically, and others require the programmer to specify it. In the following sections, we'll look at the differences between the defaults for local, instance, and class variables.

## Local Variables

**A *local variable* is a variable defined within a method. Local variables must be initialized before use. They do not have a default value and contain garbage data until initialized. The compiler will not let you read an uninitialized value.** For example, the following code generates a compiler error:

```
4: public int notValid() {
5: int y = 10;
6: int x;
7: int reply = x + y; // DOES NOT COMPILE
8: return reply;
9: }
```

`y` is initialized to 10. However, because `x` is not initialized before it is used in the expression on line 7, the compiler generates the following error:

```
Test.java:5: variable x might not have been initialized
```

```
 int reply = x + y;
```

```
 ^
```

Until `x` is assigned a value, it cannot appear within an expression, and the compiler will gladly remind you of this rule. The compiler knows your code has control of what happens inside the method and can be expected to initialize values.

The compiler is smart enough to recognize variables that have been initialized after their declaration but before they are used. Here's an example:

```
public int valid() {
 int y = 10;
 int x; // x is declared here
 x = 3; // and initialized here
 int reply = x + y;
 return reply;
}
```

The compiler is also smart enough to recognize initializations that are more complex. In this example, there are two branches of code. `answer` is initialized in both of them so the compiler is perfectly happy. `onlyOneBranch` is only initialized if `check` happens to be true. The compiler knows there is the possibility for `check` to be false, resulting in uninitialized code, and gives a compiler error. You'll learn more about the `if` statement in the next chapter.

```
public void findAnswer(boolean check) {
 int answer;
 int onlyOneBranch;
 if (check) {
 onlyOneBranch = 1;
 answer = 1;
 } else {
 answer = 2;
 }
 System.out.println(answer);
 System.out.println(onlyOneBranch); // DOES NOT COMPILE
}
```

## Instance and Class Variables

Variables that are not local variables are known as *instance variables* or *class variables*. Instance variables are also called fields. Class variables are shared across multiple objects.

You can tell a variable is a class variable because it has the keyword `static` before it. You'll learn about this in Chapter 4. For now, just know that a variable is a class variable if it has the `static` keyword in its declaration.

Instance and class variables do not require you to initialize them. As soon as you declare these variables, they are given a default value. You'll need to memorize everything in table 1.2 except the default value of `char`. To make this easier, remember that the compiler doesn't know what value to use and so wants the simplest type it can give the value: `null` for an object and `0/false` for a primitive.

**TABLE 1.2** Default initialization values by type

| Variable type                                                                 | Default initialization value                |
|-------------------------------------------------------------------------------|---------------------------------------------|
| <code>boolean</code>                                                          | <code>false</code>                          |
| <code>byte</code> , <code>short</code> , <code>int</code> , <code>long</code> | <code>0</code> (in the type's bit-length)   |
| <code>float</code> , <code>double</code>                                      | <code>0.0</code> (in the type's bit-length) |
| <code>char</code>                                                             | <code>'\u0000'</code> (NUL)                 |
| All object references (everything else)                                       | <code>null</code>                           |

## Understanding Variable Scope

You've learned that local variables are declared within a method. How many local variables do you see in this example?

```
public void eat(int piecesOfCheese) {
 int bitesOfCheese = 1;
}
```

There are two local variables in this method. *bitesOfCheese* is declared inside the method. *piecesOfCheese* is called a method parameter. It is also local to the method. Both of these variables are said to have a *scope* local to the method. This means they cannot be used outside the method.

Local variables can never have a scope larger than the method they are defined in. However, they can have a smaller scope. Consider this example:

```
3: public void eatIfHungry(boolean hungry) {
4: if (hungry) {
5: int bitesOfCheese = 1;
```

```

6: } // bitesOfCheese goes out of scope here
7: System.out.println(bitesOfCheese); // DOES NOT COMPILE
8: }

```

*hungry* has a scope of the entire method. *bitesOfCheese* has a smaller scope. It is only available for use in the *if* statement because it is declared inside of it. When you see a set of braces (*{ }*) in the code, it means you have entered a new block of code. Each block of code has its own scope. When there are multiple blocks, you match them from the inside out. In our case, the *if* statement block begins at line 4 and ends at line 6. The method's block begins at line 3 and ends at line 8.

Since *bitesOfCheese* is declared in such a block, the scope is limited to that block. When the compiler gets to line 7, it complains that it doesn't know anything about this *bitesOfCheese* thing and gives an error:

bitesOfCheese cannot be resolved to a variable

Remember that blocks can contain other blocks. These smaller contained blocks can reference variables defined in the larger scoped blocks, but not vice versa. For example:

```

16: public void eatIfHungry(boolean hungry) {
17: if (hungry) {
18: int bitesOfCheese = 1;
19: {
20: boolean teenyBit = true;
21: System.out.println(bitesOfCheese);
22: }
23: }
24: System.out.println(teenyBit); // DOES NOT COMPILE
25: }

```

The variable defined on line 18 is in scope until the block ends on line 23. Using it in the smaller block from lines 19 to 22 is fine. The variable defined on line 20 goes out of scope on line 22. Using it on line 24 is not allowed.

The exam may attempt to trick you with questions on scope. You'll probably see a question that appears to be about something complex and fails to compile because one of the variables is out of scope. Let's try one. Don't worry if you aren't familiar with *if* statements or *while* loops yet. It doesn't matter what the code does since we are talking about scope. See if you can figure out on which line each of the five local variables goes into and out of scope:

```

11: public void eatMore(boolean hungry, int amountOfFood) {
12: int roomInBelly = 5;
13: if (hungry) {
14: boolean timeToEat = true;
15: while (amountOfFood > 0) {
16: int amountEaten = 2;

```

```
17: roomInBelly = roomInBelly - amountEaten;
18: amountOfFood = amountOfFood - amountEaten;
19: }
20: }
21: System.out.println(amountOfFood);
22: }
```

The first step in figuring out the scope is to identify the blocks of code. In this case, there are three blocks. You can tell this because there are three sets of braces. Starting from the innermost set, we can see where the `while` loop’s block starts and ends. Repeat this as we go out for the `if` statement block and method block. Table 1.3 shows the line numbers that each block starts and ends on.

**TABLE 1.3** Blocks for scope

| Line   | First line in block | Last line in block |
|--------|---------------------|--------------------|
| while  | 15                  | 19                 |
| if     | 13                  | 20                 |
| Method | 11                  | 22                 |

You’ll want to practice this skill a lot. Identifying blocks needs to be second nature for the exam. The good news is that there are lots of code examples to practice on. You can look at any code example in this book on any topic and match up braces.

Now that we know where the blocks are, we can look at the scope of each variable. *hungry* and *amountOfFood* are method parameters, so they are available for the entire method. This means their scope is lines 11 to 22. *roomInBelly* goes into scope on line 12 because that is where it is declared. It stays in scope for the rest of the method and so goes out of scope on line 22. *timeToEat* goes into scope on line 14 where it is declared. It goes out of scope on line 20 where the `if` block ends. *amountEaten* goes into scope on line 16 where it is declared. It goes out of scope on line 19 where the `while` block ends.

All that was for local variables. Luckily the rule for instance variables is easier: they are available as soon as they are defined and last for the entire lifetime of the object itself. The rule for class (static) variables is even easier: they go into scope when declared like the other variables types. However, they stay in scope for the entire life of the program.

Let’s do one more example to make sure you have a handle on this. Again, try to figure out the type of the four variables and when they go into and out of scope.

```
1: public class Mouse {
2: static int MAX_LENGTH = 5;
3: int length;
```

```
4: public void grow(int inches) {
5: if (length < MAX_LENGTH) {
6: int newSize = length + inches;
7: length = newSize;
8: }
9: }
10: }
```

In this class, we have one class variable (*MAX\_LENGTH*), one instance variable (*length*), and two local variables (*inches* and *newSize*.) *MAX\_LENGTH* is a class variable because it has the static keyword in its declaration. *MAX\_LENGTH* goes into scope on line 2 where it is declared. It stays in scope until the program ends. *length* goes into scope on line 3 where it is declared. It stays in scope as long as this *Mouse* object exists. *inches* goes into scope where it is declared on line 4. It goes out of scope at the end of the method on line 9. *newSize* goes into scope where it is declared on line 6. Since it is defined inside the *if* statement block, it goes out of scope when that block ends on line 8.

Got all that? Let's review the rules on scope:

- Local variables—in scope from declaration to end of block
- Instance variables—in scope from declaration until object garbage collected
- Class variables—in scope from declaration until program ends

# Ordering Elements in a Class

Now that you've seen the most common parts of a class, let's take a look at the correct order to type them into a file. Comments can go anywhere in the code. Beyond that, you need to memorize the rules in Table 1.4.

**TABLE 1.4** Elements of a class

| Element             | Example             | Required? | Where does it go?             |
|---------------------|---------------------|-----------|-------------------------------|
| Package declaration | package abc;        | No        | First line in the file        |
| Import statements   | import java.util.*; | No        | Immediately after the package |
| Class declaration   | public class C      | Yes       | Immediately after the import  |
| Field declarations  | int value;          | No        | Anywhere inside a class       |
| Method declarations | void method()       | No        | Anywhere inside a class       |



Let's look at a few examples to help you remember this. The first example contains one of each element:

```
package structure; // package must be first non-comment
import java.util.*; // import must come after package
public class Meerkat { // then comes the class
 double weight; // fields and methods can go in either order
 public double getWeight() {
 return weight; }
 double height; // another field - they don't need to be together
}
```

So far so good. This is a common pattern that you should be familiar with. How about this one?

```
/* header */
package structure;
// class Meerkat
public class Meerkat { }
```

Still good. We can put comments anywhere, and imports are optional. In the next example, we have a problem:

```
import java.util.*;
package structure; // DOES NOT COMPILE
String name; // DOES NOT COMPILE
public class Meerkat { }
```

There are two problems here. One is that the package and import statements are reversed. Though both are optional, package must come before import if present. The other issue is that a field attempts declaration outside a class. This is not allowed. Fields and methods must be within a class.

Got all that? Think of the acronym PIC (picture): package, import, and class. Fields and methods are easier to remember because they merely have to be inside of a class.

You need to know one more thing about class structure for the OCA exam: multiple classes can be defined in the same file, but only one of them is allowed to be public. The public class matches the name of the file. For example, these two classes must be in a file named `Meerkat.java`:

```
1: public class Meerkat { }
2: class Paw { }
```

A file is also allowed to have neither class be public. As long as there isn't more than one public class in a file, it is okay. On the OCP exam, you'll also need to understand inner classes, which are classes within a class.

# Destroying Objects

Now that we've played with our objects, it is time to put them away. Luckily, Java automatically takes care of that for you. Java provides a garbage collector to automatically look for objects that aren't needed anymore.

All Java objects are stored in your program memory's *heap*. The heap, which is also referred to as the free store, represents a large pool of unused memory allocated to your Java application. The heap may be quite large, depending on your environment, but there is always a limit to its size. If your program keeps instantiating objects and leaving them on the heap, eventually it will run out of memory.

In the following sections, we'll look at garbage collection and the `finalize()` method.

## Garbage Collection

Garbage collection refers to the process of automatically freeing memory on the heap by deleting objects that are no longer reachable in your program. There are many different algorithms for garbage collection, but you don't need to know any of them for the exam. You *do* need to know that `System.gc()` is not guaranteed to run, and you should be able to recognize when objects become eligible for garbage collection.

Let's start with the first one. Java provides a method called `System.gc()`. Now you might think from the name that this tells Java to run garbage collection. Nope! It meekly *suggests* that now might be a good time for Java to kick off a garbage collection run. Java is free to ignore the request.

The more interesting part of garbage collection is when the memory belonging to an object can be reclaimed. Java waits patiently until the code no longer needs that memory. An object will remain on the heap until it is no longer reachable. An object is no longer reachable when one of two situations occurs:


- The object no longer has any references pointing to it.
- All references to the object have gone out of scope.

### Objects vs. References

Do not confuse a reference with the object that it refers to; they are two different entities. The reference is a variable that has a name and can be used to access the contents of an object. A reference can be assigned to another reference, passed to a method, or returned from a method. All references are the same size, no matter what their type is.

An object sits on the heap and does not have a name. Therefore, you have no way to access an object except through a reference. Objects come in all different shapes and sizes and consume varying amounts of memory. An object cannot be assigned to another

object, nor can an object be passed to a method or returned from a method. It is the object that gets garbage collected, not its reference.



Realizing the difference between a reference and an object goes a long way toward understanding garbage collection, the new operator, and many other facets of the Java language. Look at this code and see if you can figure out when each object first becomes eligible for garbage collection:


```

1: public class Scope {
2: public static void main(String[] args) {
3: String one, two;
4: one = new String("a");
5: two = new String("b");
6: one = two;
7: String three = one;
8: one = null;
9: } }


```

When you get asked a question about garbage collection on the exam, we recommend you draw what's going on. There's a lot to keep track of in your head and it's easy to make a silly mistake trying to keep it all in your memory. Let's try it together now. Really. Get a pencil and paper. We'll wait.

Got that paper? Okay, let's get started. On line 3, we write *one* and *two*. Just the words. No need for boxes or arrows yet since no objects have gone on the heap yet. On line 4, we have our first object. Draw a box with the string "a" in it and draw an arrow from the word *one* to that box. Line 5 is similar. Draw another box with the string "b" in it this time and an arrow from the word *two*. At this point, your work should look like Figure 1.2.

**FIGURE 1.2** Your drawing after line 5

On line 6, the variable *one* changes to point to "b". Either erase or cross out the arrow from *one* and draw a new arrow from *one* to "b". On line 7, we have a new variable, so write the word *three* and draw an arrow from *three* to "b". Notice that *three* points to what *one* is pointing to right now and not what it was pointing to at the beginning. This is why we are drawing pictures. It's easy to forget something like that. At this point, your work should look like Figure 1.3.

**FIGURE 1.3** Your drawing after line 7

Finally, cross out the line between *one* and "b" since line 8 sets this variable to null. Now, we were trying to find out when the objects were first eligible for garbage collection. On line 6, we got rid of the only arrow pointing to "a", making that object eligible for garbage collection. "b" has arrows pointing to it until it goes out of scope. This means "b" doesn't go out of scope until the end of the method on line 9.

## ***finalize()***

Java allows objects to implement a method called `finalize()` that might get called. This method gets called if the garbage collector tries to collect the object. If the garbage collector doesn't run, the method doesn't get called. If the garbage collector fails to collect the object and tries to run it again later, the method doesn't get called a second time.

In practice, this means you are highly unlikely to use it in real projects. Luckily, there isn't much to remember about `finalize()` for the exam. Just keep in mind that it might not get called and that it definitely won't be called twice.

With that said, this call produces no output when we run it:

```
public class Finalizer {
 protected void finalize() {
```

```
 System.out.println("Calling finalize");
 }
 public static void main(String[] args) {
 Finalizer f = new Finalizer();
 } }
```

The reason is that the program exits before there is any need to run the garbage collector. While *f* is eligible for garbage collection, Java has better things to do than take out the trash constantly. For the exam, you need to know that this `finalize()` call could run zero or one time. Now for a more interesting example:

```
public class Finalizer {
 private static List objects = new ArrayList();
 protected void finalize() {
 objects.add(this); // Don't do this
 } }
}
```

Remember, `finalize()` is only run when the object is eligible for garbage collection. The problem here is that by the end of the method, the object is no longer eligible for garbage collection because a static variable is referring to it and static variables stay in scope until the program ends. Java is smart enough to realize this and aborts the attempt to throw out the object. Now suppose later in the program *objects* is set to null. Oh, good, we can finally remove the object from memory. Java remembers already running `finalize()` on this object and will not do so again. The lesson is that the `finalize()` call could run zero or one time. This is the exact same lesson as the simple example—that's why it's so easy to remember.

## Benefits of Java

Java has some key benefits that you'll need to know for the exam:

**Object Oriented** Java is an object-oriented language, which means all code is defined in classes and most of those classes can be instantiated into objects. We'll discuss this more throughout the book. Many languages before Java were procedural, which meant there were routines or methods but no classes. Another common approach is functional programming. Java allows for functional programming within a class, but object oriented is still the main organization of code.

**Encapsulation** Java supports access modifiers to protect data from unintended access and modification. Most people consider encapsulation to be an aspect of object-oriented languages. Since the exam objectives call attention to it specifically, so do we.

**Platform Independent** Java is an interpreted language because it gets compiled to bytecode. A key benefit is that Java code gets compiled once rather than needing to be

recompiled for different operating systems. This is known as “write once, run everywhere.” On the OCP exam, you’ll learn that it is possible to write code that does not run everywhere. For example, you might refer to a file in a specific directory. If you get asked on the OCA exam, the answer is that the same class files run everywhere.

**Robust** One of the major advantages of Java over C++ is that it prevents memory leaks. Java manages memory on its own and does garbage collection automatically. Bad memory management in C++ is a big source of errors in programs.

**Simple** Java was intended to be simpler than C++. In addition to eliminating pointers, it got rid of operator overloading. In C++, you could write `a + b` and have it mean almost anything.

**Secure** Java code runs inside the JVM. This creates a sandbox that makes it hard for Java code to do evil things to the computer it is running on.

## Summary

In this chapter, you saw that Java classes consist of members called fields and methods. An object is an instance of a Java class. There are three styles of comment: a single-line comment (`//`), a multiline comment (`/* */`), and a Javadoc comment (`/** */`).

Java begins program execution with a `main()` method. The most common signature for this method run from the command line is `public static void main(String[] args)`. Arguments are passed in after the class name, as in `java NameOfClass firstArgument`. Arguments are indexed starting with 0.

Java code is organized into folders called packages. To reference classes in other packages, you use an import statement. A wildcard ending an import statement means you want to import all classes in that package. It does not include packages that are inside that one. `java.lang` is a special package that does not need to be imported.

**Constructors create Java objects. A constructor is a method matching the class name and omitting the return type. When an object is instantiated, fields and blocks of code are initialized first. Then the constructor is run.**

Primitive types are the basic building blocks of Java types. They are assembled into reference types. Reference types can have methods and be assigned to `null`. In addition to “normal” numbers, numeric literals are allowed to begin with 0 (octal), 0x (hex), 0X (hex), 0b (binary), or 0B (binary). Numeric literals are also allowed to contain underscores as long as they are directly between two other numbers.

Declaring a variable involves stating the data type and giving the variable a name. Variables that represent fields in a class are automatically initialized to their corresponding “zero” or null value during object instantiation. Local variables must be specifically initialized. Identifiers may contain letters, numbers, \$, or `_`. Identifiers may not begin with numbers.

Scope refers to that portion of code where a variable can be accessed. There are three kinds of variables in Java, depending on their scope: instance variables, class variables, and

local variables. Instance variables are the nonstatic fields of your class. Class variables are the static fields within a class. Local variables are declared within a method.

For some class elements, order matters within the file. The package statement comes first if present. Then comes the import statements if present. Then comes the class declaration. Fields and methods are allowed to be in any order within the class.

Garbage collection is responsible for removing objects from memory when they can never be used again. An object becomes eligible for garbage collection when there are no more references to it or its references have all gone out of scope. The `finalize()` method will run once for each object if/when it is first garbage collected.

Java code is object oriented, meaning all code is defined in classes. Access modifiers allow classes to encapsulate data. Java is platform independent, compiling to bytecode. It is robust and simple by not providing pointers or operator overloading. Finally, Java is secure because it runs inside a virtual machine.

## Exam Essentials

**Be able to write code using a `main()` method.** A `main()` method is usually written as `public static void main(String[] args)`. Arguments are referenced starting with `args[0]`. Accessing an argument that wasn't passed in will cause the code to throw an exception.

**Understand the effect of using packages and imports.** Packages contain Java classes. Classes can be imported by class name or wildcard. Wildcards do not look at subdirectories. In the event of a conflict, class name imports take precedence.

**Be able to recognize a constructor.** A constructor has the same name as the class. It looks like a method without a return type.

**Be able to identify legal and illegal declarations and initialization.** Multiple variables can be declared and initialized in the same statement when they share a type. Local variables require an explicit initialization; others use the default value for that type. Identifiers may contain letters, numbers, \$, or \_. Identifiers may not begin with numbers. Numeric literals may contain underscores between two digits and begin with 1–9, 0, 0x, 0X, 0b, and 0B.

**Be able to determine where variables go into and out of scope.** All variables go into scope when they are declared. Local variables go out of scope when the block they are declared in ends. Instance variables go out of scope when the object is garbage collected. Class variables remain in scope as long as the program is running.

**Be able to recognize misplaced statements in a class.** Package and import statements are optional. If present, both go before the class declaration in that order. Fields and methods are also optional and are allowed in any order within the class declaration.

**Know how to identify when an object is eligible for garbage collection.** Draw a diagram to keep track of references and objects as you trace the code. When no arrows point to a box (object), it is eligible for garbage collection.

# Review Questions

1. Which of the following are valid Java identifiers? (Choose all that apply)

- A. A\$B
- B. \_helloWorld
- C. true
- D. java.lang
- E. Public
- F. 1980\_s

2. What is the output of the following program?

```
1: public class WaterBottle {
2: private String brand;
3: private boolean empty;
4: public static void main(String[] args) {
5: WaterBottle wb = new WaterBottle();
6: System.out.print("Empty = " + wb.empty);
7: System.out.print(", Brand = " + wb.brand);
8: } }
```

- A. Line 6 generates a compiler error.
- B. Line 7 generates a compiler error.
- C. There is no output.
- D. Empty = false, Brand = null
- E. Empty = false, Brand =
- F. Empty = null, Brand = null

3. Which of the following are true? (Choose all that apply)

```
4: short numPets = 5;
5: int numGrains = 5.6;
6: String name = "Scruffy";
7: numPets.length();
8: numGrains.length();
9: name.length();
```

- A. Line 4 generates a compiler error.
- B. Line 5 generates a compiler error.
- C. Line 6 generates a compiler error.
- D. Line 7 generates a compiler error.
- E. Line 8 generates a compiler error.



**F.** Line 9 generates a compiler error.

**G.** The code compiles as is.

- 4.** Given the following class, which of the following is true? (Choose all that apply)

```
1: public class Snake {
2:
3: public void shed(boolean time) {
4:
5: if (time) {
6:
7: }
8: System.out.println(result);
9:
10: }
11: }
```

**A.** If `String result = "done";` is inserted on line 2, the code will compile.

**B.** If `String result = "done";` is inserted on line 4, the code will compile.

**C.** If `String result = "done";` is inserted on line 6, the code will compile.

**D.** If `String result = "done";` is inserted on line 9, the code will compile.

**E.** None of the above changes will make the code compile.

- 5.** Given the following classes, which of the following can independently replace `INSERT IMPORTS HERE` to make the code compile? (Choose all that apply)

```
package aquarium;
public class Tank { }
```

```
package aquarium.jellies;
public class Jelly { }
```

```
package visitor;
INSERT IMPORTS HERE
public class AquariumVisitor {
 public void admire(Jelly jelly) { } }
```

**A.** `import aquarium.*;`

**B.** `import aquarium.*.Jelly;`

**C.** `import aquarium.jellies.Jelly;`

**D.** `import aquarium.jellies.*;`

**E.** `import aquarium.jellies.Jelly.*;`

**F.** None of these can make the code compile.

6. Given the following classes, what is the maximum number of imports that can be removed and have the code still compile?

```
package aquarium; public class Water { }
```

```
package aquarium;
import java.lang.*;
import java.lang.System;
import aquarium.Water;
import aquarium.*;
public class Tank {
 public void print(Water water) {
 System.out.println(water); } }
```

- A. 0
  - B. 1
  - C. 2
  - D. 3
  - E. 4
  - F. Does not compile.
7. Given the following classes, which of the following snippets can be inserted in place of INSERT IMPORTS HERE and have the code compile? (Choose all that apply)

```
package aquarium;
public class Water {
 boolean salty = false;
}
package aquarium.jellies;
public class Water {
 boolean salty = true;
}
package employee;
 INSERT IMPORTS HERE
public class WaterFiller {
 Water water;
}
```

- A. `import aquarium.*;`
- B. `import aquarium.Water;`  
`import aquarium.jellies.*;`
- C. `import aquarium.*;`  
`import aquarium.jellies.Water;`

- D. `import aquarium.*;`  
`import aquarium.jellies.*;`
  - E. `import aquarium.Water;`  
`import aquarium.jellies.Water;`
  - F. None of these imports can make the code compile.
8. Given the following class, which of the following calls print out Blue Jay? (Choose all that apply)
- ```
public class BirdDisplay {  
    public static void main(String[] name) {  
        System.out.println(name[1]);  
    }  
}
```
- A. `java BirdDisplay Sparrow Blue Jay`
 - B. `java BirdDisplay Sparrow "Blue Jay"`
 - C. `java BirdDisplay Blue Jay Sparrow`
 - D. `java BirdDisplay "Blue Jay" Sparrow`
 - E. `java BirdDisplay.class Sparrow "Blue Jay"`
 - F. `java BirdDisplay.class "Blue Jay" Sparrow`
 - G. Does not compile.
9. Which of the following legally fill in the blank so you can run the `main()` method from the command line? (Choose all that apply)
- ```
public static void main(_____)
```
- A. `String[] _names`
  - B. `String[] 123`
  - C. `String abc[]`
  - D. `String _Names[]`
  - E. `String... $n`
  - F. `String names`
  - G. None of the above.
10. Which of the following are legal entry point methods that can be run from the command line? (Choose all that apply)
- A. `private static void main(String[] args)`
  - B. `public static final main(String[] args)`
  - C. `public void main(String[] args)`
  - D. `public static void test(String[] args)`
  - E. `public static void main(String[] args)`
  - F. `public static main(String[] args)`
  - G. None of the above.

11. Which of the following are true? (Choose all that apply)
- A. An instance variable of type `double` defaults to `null`.
  - B. An instance variable of type `int` defaults to `null`.
  - C. An instance variable of type `String` defaults to `null`.
  - D. An instance variable of type `double` defaults to `0.0`.
  - E. An instance variable of type `int` defaults to `0.0`.
  - F. An instance variable of type `String` defaults to `0.0`.
  - G. None of the above.
12. Which of the following are true? (Choose all that apply)
- A. A local variable of type `boolean` defaults to `null`.
  - B. A local variable of type `float` defaults to `0`.
  - C. A local variable of type `Object` defaults to `null`.
  - D. A local variable of type `boolean` defaults to `false`.
  - E. A local variable of type `boolean` defaults to `true`.
  - F. A local variable of type `float` defaults to `0.0`.
  - G. None of the above.
13. Which of the following are true? (Choose all that apply)
- A. An instance variable of type `boolean` defaults to `false`.
  - B. An instance variable of type `boolean` defaults to `true`.
  - C. An instance variable of type `boolean` defaults to `null`.
  - D. An instance variable of type `int` defaults to `0`.
  - E. An instance variable of type `int` defaults to `0.0`.
  - F. An instance variable of type `int` defaults to `null`.
  - G. None of the above.
14. Given the following class in the file `/my/directory/named/A/Bird.java`:
- ```
INSERT CODE HERE
public class Bird { }
```
- Which of the following replaces `INSERT CODE HERE` if we compile from `/my/directory/`? (Choose all that apply)
- A. `package my.directory.named.a;`
 - B. `package my.directory.named.A;`
 - C. `package named.a;`
 - D. `package named.A;`
 - E. `package a;`
 - F. `package A;`
 - G. Does not compile.

15. Which of the following lines of code compile? (Choose all that apply)
- A. `int i1 = 1_234;`
 - B. `double d1 = 1_234_.0;`
 - C. `double d2 = 1_234._0;`
 - D. `double d3 = 1_234.0_;`
 - E. `double d4 = 1_234.0;`
 - F. None of the above.
16. Given the following class, which of the following lines of code can replace INSERT CODE HERE to make the code compile? (Choose all that apply)
- ```
public class Price {
 public void admission() {
 INSERT CODE HERE
 System.out.println(amount);
 }
}
```
- A. `int amount = 9L;`
  - B. `int amount = 0b101;`
  - C. `int amount = 0xE;`
  - D. `double amount = 0xE;`
  - E. `double amount = 1_2_.0_0;`
  - F. `int amount = 1_2_;`
  - G. None of the above.
17. Which of the following are true? (Choose all that apply)
- ```
public class Bunny {  
    public static void main(String[] args) {  
        Bunny bun = new Bunny();  
    }  
}
```
- A. Bunny is a class.
 - B. bun is a class.
 - C. main is a class.
 - D. Bunny is a reference to an object.
 - E. bun is a reference to an object.
 - F. main is a reference to an object.
 - G. None of the above.
18. Which represent the order in which the following statements can be assembled into a program that will compile successfully? (Choose all that apply)
- A: `class Rabbit {}`
 - B: `import java.util.*;`
 - C: `package animals;`

- A. A, B, C
- B. B, C, A
- C. C, B, A
- D. B, A
- E. C, A
- F. A, C
- G. A, B

19. Suppose we have a class named `Rabbit`. Which of the following statements are true? (Choose all that apply)

```

1: public class Rabbit {
2:     public static void main(String[] args) {
3:         Rabbit one = new Rabbit();
4:         Rabbit two = new Rabbit();
5:         Rabbit three = one;
6:         one = null;
7:         Rabbit four = one;
8:         three = null;
9:         two = null;
10:        two = new Rabbit();
11:        System.gc();
12:    } }

```

- A. The `Rabbit` object from line 3 is first eligible for garbage collection immediately following line 6.
- B. The `Rabbit` object from line 3 is first eligible for garbage collection immediately following line 8.
- C. The `Rabbit` object from line 3 is first eligible for garbage collection immediately following line 12.
- D. The `Rabbit` object from line 4 is first eligible for garbage collection immediately following line 9.
- E. The `Rabbit` object from line 4 is first eligible for garbage collection immediately following line 11.
- F. The `Rabbit` object from line 4 is first eligible for garbage collection immediately following line 12.

20. What is true about the following code? (Choose all that apply)

```

public class Bear {
    protected void finalize() {
        System.out.println("Roar!");
    }
}

```

```
public static void main(String[] args) {  
    Bear bear = new Bear();  
    bear = null;  
    System.gc();  
} }
```

- A. `finalize()` is guaranteed to be called.
- B. `finalize()` might or might not be called.
- C. `finalize()` is guaranteed not to be called.
- D. Garbage collection is guaranteed to run.
- E. Garbage collection might or might not run.
- F. Garbage collection is guaranteed not to run.
- G. The code does not compile.

21. What does the following code output?

```
1: public class Salmon {  
2:     int count;  
3:     public void Salmon() {  
4:         count = 4;  
5:     }  
6:     public static void main(String[] args) {  
7:         Salmon s = new Salmon();  
8:         System.out.println(s.count);  
9:     } }
```

- A. 0
- B. 4
- C. Compilation fails on line 3.
- D. Compilation fails on line 4.
- E. Compilation fails on line 7.
- F. Compilation fails on line 8.

22. Which of the following are true statements? (Choose all that apply)

- A. Java allows operator overloading.
- B. Java code compiled on Windows can run on Linux.
- C. Java has pointers to specific locations in memory.
- D. Java is a procedural language.
- E. Java is an object-oriented language.
- F. Java is a functional programming language.

- 23.** Which of the following are true? (Choose all that apply)
- A.** javac compiles a .class file into a .java file.
 - B.** javac compiles a .java file into a .bytecode file.
 - C.** javac compiles a .java file into a .class file.
 - D.** Java takes the name of the class as a parameter.
 - E.** Java takes the name of the .bytecode file as a parameter.
 - F.** Java takes the name of the .class file as a parameter.

Chapter 2

Operators and Statements

OCA EXAM OBJECTIVES COVERED IN THIS CHAPTER:

✓ Using Operators and Decision Constructs

- Use Java operators; including parentheses to override operator precedence
- Create if and if/else and ternary constructs
- Use a switch statement

✓ Using Loop Constructs

- Create and use while loops
- Create and use for loops including the enhanced for loop
- Create and use do/while loops
- Compare loop constructs
- Use break and continue





Like many programming languages, Java is composed primarily of variables, operators, and statements put together in some logical order. In the previous chapter, we discussed variables and gave some examples; in this chapter we'll discuss the various operators and statements available to you within the language. This knowledge will allow you to build complex functions and class structures that you'll see in later chapters.

Understanding Java Operators

A Java *operator* is a special symbol that can be applied to a set of variables, values, or literals—referred to as operands—and that returns a result. Three flavors of operators are available in Java: unary, binary, and ternary. These types of operators can be applied to one, two, or three operands, respectively. For the OCA exam, you'll need know a specific subset of Java operators, how to apply them, and the order in which they should be applied.

Java operators are not necessarily evaluated from left-to-right order. For example, the following Java expression is actually evaluated from right-to-left given the specific operators involved:

```
int y = 4;
double x = 3 + 2 * --y;
```

In this example, you would first decrement *y* to 3, and then multiply the resulting value by 2, and finally add 3. The value would then be automatically upcast from 9 to 9.0 and assigned to *x*. The final values of *x* and *y* would be 9.0 and 3, respectively. If you didn't follow that evaluation, don't worry. By the end of this chapter, solving problems like this should be second nature.

Unless overridden with parentheses, Java operators follow *order of operation*, listed in Table 2.1, by decreasing order of *operator precedence*. If two operators have the same level of precedence, then Java guarantees left-to-right evaluation. You need to know only those operators in bold for the OCA exam.

TABLE 2.1 Order of operator precedence

Operator	Symbols and examples
Post-unary operators	<i>expression</i> ++ , <i>expression</i> --
Pre-unary operators	++ <i>expression</i> , -- <i>expression</i>

Operator	Symbols and examples
Other unary operators	<code>+, -, !</code>
Multiplication/Division/Modulus	<code>*, /, %</code>
Addition/Subtraction	<code>+, -</code>
Shift operators	<code><<, >>, >>></code>
Relational operators	<code><, >, <=, >=, instanceof</code>
Equal to/not equal to	<code>==, !=</code>
Logical operators	<code>&, ^, </code>
Short-circuit logical operators	<code>&&, </code>
Ternary operators	<code>boolean expression ? expression1 : expression2</code>
Assignment operators	<code>=, +=, -=, *=, /=, %=, &=, ^=, !=, <<=, >>=, >>>=</code>

We'll spend the first half of this chapter discussing many of the operators in this list as well as how operator precedence determines which operators should be applied first. Note that you won't be tested on some operators, although we recommend that you be aware of their existence.

Working with Binary Arithmetic Operators

We'll begin our discussion with *binary operators*, by far the most common operators in the Java language. They can be used to perform mathematical operations on variables, create logical expressions, as well as perform basic variable assignments. Binary operators are commonly combined in complex expressions with more than two variables; therefore, operator precedence is very important in evaluating expressions.

Arithmetic Operators

Arithmetic operators are often encountered in early mathematics and include addition (+), subtraction (-), multiplication (*), division (/), and modulus (%). They also include the unary operators, ++ and --, although we cover them later in this chapter. As you may have

noticed in Table 2.1, the *multiplicative* operators (*, /, %) have a higher order of precedence than the *additive* operators (+, -). That means when you see an expression such as this:

```
int x = 2 * 5 + 3 * 4 - 8;
```

you first evaluate the $2 * 5$ and $3 * 4$, which reduces the expression to the following:

```
int x = 10 + 12 - 8;
```

Then, you evaluate the remaining terms in left-to-right order, resulting in a value of x of 14. Make sure you understand why the result is 24 as you'll likely see this kind of operator precedence question on the exam.

Notice that we said “Unless overridden with parentheses...” prior to Table 2.1. That's because you can change the order of operation explicitly by wrapping parentheses around the sections you want evaluated first. Compare the previous example with the following one containing the same values and operators, in the same order, but with two sets of parentheses:

```
int x = 2 * ((5 + 3) * 4 - 8);
```

This time you would evaluate the addition operator $10 + 3$, which reduces the expression to the following:

```
int x = 2 * (8 * 4 - 8);
```

You can further reduce this expression by multiplying the first two values within the parentheses:

```
int x = 2 * (32 - 8);
```

Next, you subtract the values within the parentheses before applying terms outside the parentheses:

```
int x = 2 * 24;
```

Finally, you would multiply the result by 2, resulting in a value of 48 for x .

All of the arithmetic operators may be applied to any Java primitives, except boolean and String. Furthermore, only the addition operators + and += may be applied to String values, which results in String concatenation.

Although we are sure you have seen most of the arithmetic operators before, the modulus operator, %, may be new to you. The modulus, or remainder operator, is simply the remainder when two numbers are divided. For example, 9 divided by 3 divides evenly and has no remainder; therefore, the remainder, or $9 \% 3$, is 0. On the other hand, 11 divided by 3 does not divide evenly; therefore, the remainder, or $11 \% 3$, is 2.

Be sure to understand the difference between arithmetic division and modulus. For integer values, division results in the floor value of the nearest integer that fulfills the operation, whereas modulus is the remainder value. The following examples illustrate this distinction:

```
System.out.print(9 / 3); // Outputs 3
System.out.print(9 % 3); // Outputs 0
```

```
System.out.print(10 / 3); // Outputs 3
System.out.print(10 % 3); // Outputs 1

System.out.print(11 / 3); // Outputs 3
System.out.print(11 % 3); // Outputs 2

System.out.print(12 / 3); // Outputs 4
System.out.print(12 % 3); // Outputs 0
```

Note that the division results only increase when the value on the left-hand side goes from 9 to 12, whereas the modulus remainder value increases by 1 each time the left-hand side is increased until it wraps around to zero. For a given divisor y , which is 3 in these examples, the modulus operation results in a value between 0 and $(y - 1)$ for positive dividends. This means that the result of a modulus operation is always 0, 1, or 2.



The modulus operation is not limited to positive integer values in Java and may also be applied to negative integers and floating-point integers. For a given divisor y and negative dividend, the resulting modulus value is between $(-y + 1)$ and 0. For the OCA exam, though, you are not required to be able to take the modulus of a negative integer or a floating-point number.

Numeric Promotion

Now that you understand the basics of arithmetic operators, it is vital we talk about primitive *numeric promotion*, as Java may do things that seem unusual to you at first. If you recall in Chapter 1, “Java Building Blocks,” where we listed the primitive numeric types, each primitive has a bit-length. **You don’t need to know the exact size of these types for the exam, but you should know which are bigger than others.** For example, you should know that a long takes up more space than an int, which in turn takes up more space than a short, and so on.

You should memorize certain rules Java will follow when applying operators to data types:

Numeric Promotion Rules

1. If two values have different data types, Java will automatically promote one of the values to the larger of the two data types.
2. If one of the values is integral and the other is floating-point, Java will automatically promote the integral value to the floating-point value’s data type.

3. Smaller data types, namely byte, short, and char, are first promoted to int any time they're used with a Java binary arithmetic operator, even if neither of the operands is int.
4. After all promotion has occurred and the operands have the same data type, the resulting value will have the same data type as its promoted operands.

The last two rules are the ones most people have trouble with, and the ones likely to trip you up on the exam. For the third rule, note that unary operators are excluded from this rule. For example, applying ++ to a short value results in a short value. We'll discuss unary operators in the next section.

Let's tackle some examples for illustrative purposes:

- What is the data type of $x * y$?

```
int x = 1;
long y = 33;
```

If we follow the first rule, since one of the values is long and the other is int, and long is larger than int, then the int value is promoted to a long, and the resulting value is long.

- What is the data type of $x + y$?

```
double x = 39.21;
float y = 2.1;
```

This is actually a trick question, as this code will not compile! As you may remember from Chapter 1, floating-point literals are assumed to be double, unless postfix with an f, as in 2.1f. If the value was set properly to 2.1f, then the promotion would be similar to the last example, with both operands being promoted to a double, and the result would be a double value.

- What is the data type of x / y ?

```
short x = 10;
short y = 3;
```

In this case, we must apply the third rule, namely that x and y will both be promoted to int before the operation, resulting in an output of type int. Pay close attention to the fact that the resulting output is not a short, as we'll come back to this example in the upcoming section on assignment operators.

- What is the data type of $x * y / z$?

```
short x = 14;
float y = 13;
double z = 30;
```

In this case, we must apply all of the rules. First, x will automatically be promoted to int solely because it is a short and it is being used in an arithmetic binary operation.

The promoted `x` value will then be automatically promoted to a `float` so that it can be multiplied with `y`. The result of `x * y` will then be automatically promoted to a `double`, so that it can be multiplied with `z`, resulting in a `double` value.

Working with Unary Operators

By definition, a *unary* operator is one that requires exactly one operand, or variable, to function. As shown in Table 2.2, they often perform simple tasks, such as increasing a numeric variable by one, or negating a boolean value.

TABLE 2.2 Java unary operators

Unary operator	Description
+	Indicates a number is positive, although numbers are assumed to be positive in Java unless accompanied by a negative unary operator
-	Indicates a literal number is negative or negates an expression
++	Increments a value by 1
--	Decrements a value by 1
!	Inverts a Boolean's logical value

Logical Complement and Negation Operators

The *logical complement operator*, `!`, flips the value of a boolean expression. For example, if the value is `true`, it will be converted to `false`, and vice versa. To illustrate this, compare the outputs of the following statements:

```
boolean x = false;
System.out.println(x); // false
x = !x;
System.out.println(x); // true
```

Likewise, the *negation operator*, `-`, reverses the sign of a numeric expression, as shown in these statements:

```
double x = 1.21;
```

```
System.out.println(x); // 1.21
x = -x;
System.out.println(x); // -1.21
x = -x;
System.out.println(x); // 1.21
```

Based on the description, it might be obvious that some operators require the variable or expression they're acting upon to be of a specific type. For example, you cannot apply a negation operator, `-`, to a boolean expression, nor can you apply a logical complement operator, `!`, to a numeric expression. Be wary of questions on the exam that try to do this, as they'll cause the code to fail to compile. For example, none of the following lines of code will compile:

```
int x = !5; // DOES NOT COMPILE
boolean y = -true; // DOES NOT COMPILE
boolean z = !0; // DOES NOT COMPILE
```

The first statement will not compile due the fact that in Java you cannot perform a logical inversion of a numeric value. The second statement does not compile because you cannot numerically negate a boolean value; you need to use the logical inverse operator. Finally, the last statement does not compile because you cannot take the logical complement of a numeric value, nor can you assign an integer to a boolean variable.



Keep an eye out for questions on the exam that use the logical complement operator or numeric values with boolean expressions or variables. Unlike some other programming languages, in Java `1` and `true` are not related in any way, just as `0` and `false` are not related.

Increment and Decrement Operators

Increment and decrement operators, `++` and `--`, respectively, can be applied to numeric operands and have the higher order or precedence, as compared to binary operators. In other words, they often get applied first to an expression.

Increment and decrement operators require special care because the order they are applied to their associated operand can make a difference in how an expression is processed. If the operator is placed before the operand, referred to as the *pre-increment operator* and the *pre-decrement operator*, then the operator is applied first and the value returned is the *new value* of the expression. Alternatively, if the operator is placed after the operand, referred to as the *post-increment operator* and the *post-decrement operator*, then the *original value* of the expression is returned, with operator applied after the value is returned.

The following code snippet illustrates this distinction:

```
int counter = 0;
System.out.println(counter); // Outputs 0
System.out.println(++counter); // Outputs 1
System.out.println(counter); // Outputs 1
System.out.println(counter--); // Outputs 1
System.out.println(counter); // Outputs 0
```

The first pre-increment operator updates the value for counter and outputs the new value of 1. The next post-decrement operator also updates the value of counter but outputs the value before the decrement occurs.

One common practice in a certification exam, albeit less common in the real world, is to apply multiple increment or decrement operators to a single variable on the same line:

```
int x = 3;
int y = ++x * 5 / x-- + --x;
System.out.println("x is " + x);
System.out.println("y is " + y);
```

This one is more complicated than the previous example because x is modified three times on the same line. Each time it is modified, as the expression moves from left to right, the value of x changes, with different values being assigned to the variable. As you'll recall from our discussion on operator precedence, order of operation plays an important part in evaluating this example.

So how do you read this code? First, the x is incremented and returned to the expression, which is multiplied by 5. We can simplify this:

```
int y = 4 * 5 / x-- + --x; // x assigned value of 4
```

Next, x is decremented, but the original value of 4 is used in the expression, leading to this:

```
int y = 4 * 5 / 4 + --x; // x assigned value of 3
```

The final assignment of x reduces the value to 2, and since this is a pre-increment operator, that value is returned to the expression:

```
int y = 4 * 5 / 4 + 2; // x assigned value of 2
```

Finally, we evaluate the multiple and division from left-to-right, and finish with the addition. The result is then printed:

```
x is 2
y is 7
```

Using Additional Binary Operators

We'll now expand our discussion of binary operators to include all other binary operators that you'll need to know for the exam. This includes operators that perform assignments, those that compare arithmetic values and return boolean results, and those that compare boolean and object values and return boolean results.

Assignment Operators

An assignment operator is a binary operator that modifies, or assigns, the variable on the left-hand side of the operator, with the result of the value on the right-hand side of the equation. The simplest assignment operator is the = assignment, which you have seen already:

```
int x = 1;
```

This statement assigns `x` the value of 1.

Java will automatically promote from smaller to larger data types, as we saw in the previous section on arithmetic operators, but it will throw a compiler exception if it detects you are trying to convert from larger to smaller data types.

Let's return to some examples similar to what you saw in Chapter 1 in order to show how casting can resolve these issues:

```
int x = 1.0; // DOES NOT COMPILE
short y = 1921222; // DOES NOT COMPILE
int z = 9f; // DOES NOT COMPILE
long t = 192301398193810323; // DOES NOT COMPILE
```

The first statement does not compile because you are trying to assign a double `1.0` to an integer value. Even though the value is a mathematic integer, by adding `.0`, you're instructing the compiler to treat it as a double. The second statement does not compile because the literal value `1921222` is outside the range of `short` and the compiler detects this. The third statement does not compile because of the `f` added to the end of the number that instructs the compiler to treat the number as floating-point value. Finally, the last statement does not compile because Java interprets the literal as an `int` and notices that the value is larger than `int` allows. The literal would need a postfix `L` to be considered a `long`.

Casting Primitive Values

We can fix the examples in the previous section by casting the results to a smaller data type. Casting primitives is required any time you are going from a larger numerical data type to a smaller numerical data type, or converting from a floating-point number to an integral value.

```
int x = (int)1.0;
short y = (short)1921222; // Stored as 20678
```

```
int z = (int)9L;  
long t = 192301398193810323L;
```

Overflow and Underflow

The expressions in the previous example now compile, although there's a cost. The second value, 1,921,222, is too large to be stored as a short, so numeric overflow occurs and it becomes 20,678. *Overflow* is when a number is so large that it will no longer fit within the data type, so the system "wraps around" to the next lowest value and counts up from there. There's also an analogous *underflow*, when the number is too low to fit in the data type.

This is beyond the scope of the exam, but something to be careful of in your own code. For example, the following statement outputs a negative number:

```
System.out.print(2147483647+1); // -2147483648
```

Since 2147483647 is the maximum `int` value, adding any strictly positive value to it will cause it to wrap to the next negative number.

Let's return to one of our earlier examples for a moment:

```
short x = 10;  
short y = 3;  
short z = x * y; // DOES NOT COMPILE
```

Based on everything you have learned up until now, can you understand why the last line of this statement will not compile? If you remember, short values are automatically promoted to `int` when applying any arithmetic operator, with the resulting value being of type `int`. Trying to set a short variable to an `int` results in a compiler error, as Java thinks you are trying to implicitly convert from a larger data type to a smaller one.

There are times that you may want to override the default behavior of the compiler. For example, in the preceding example, we know the result of `10 * 3` is 30, which can easily fit into a short variable. If you need the result to be a short, though, you can override this behavior by casting the result of the multiplication:

```
short x = 10;  
short y = 3;  
short z = (short)(x * y);
```

By performing this explicit cast of a larger value into a smaller data type, you are instructing the compiler to ignore its default behavior. In other words, you are telling the compiler that you have taken additional steps to prevent overflow or underflow. It is also possible that in your particular application and scenario, overflow or underflow would result in acceptable values.

Compound Assignment Operators

Besides the simple assignment operator, `=`, there are also numerous *compound assignment operators*. Only two of the compound operators listed in Table 2.1 are required for the exam, `+=` and `-=`. Complex operators are really just glorified forms of the simple assignment operator, with a built-in arithmetic or logical operation that applies the left- and right-hand sides of the statement and stores the resulting value in a variable in the left-hand side of the statement. For example, the following two statements after the declaration of `x` and `z` are equivalent:

```
int x = 2, z = 3;
x = x * z; // Simple assignment operator
x *= z; // Compound assignment operator
```

The left-hand side of the compound operator can only be applied to a variable that is already defined and cannot be used to declare a new variable. In the previous example, if `x` was not already defined, then the expression `x *= z` would not compile.

Compound operators are useful for more than just shorthand—they can also save us from having to explicitly cast a value. For example, consider the following example, in which the last line will not compile due to the result being promoted to a `long` and assigned to an `int` variable:

```
long x = 10;
int y = 5;
y = y * x; // DOES NOT COMPILE
```

Based on the last two sections, you should be able to spot the problem in the last line. This last line could be fixed with an explicit cast to `(int)`, but there's a better way using the compound assignment operator:

```
long x = 10;
int y = 5;
y *= x;
```

The compound operator will first cast `x` to a `long`, apply the multiplication of two `long` values, and then cast the result to an `int`. Unlike the previous example, in which the compiler threw an exception, in this example we see that the compiler will automatically cast the resulting value to the data type of the value on the left-hand side of the compound operator.

One final thing to know about the assignment operator is that the result of the assignment is an expression in and of itself, equal to the value of the assignment. For example, the following snippet of code is perfectly valid, if not a little odd looking:

```
long x = 5;
long y = (x=3);
System.out.println(x); // Outputs 3
System.out.println(y); // Also, outputs 3
```

The key here is that `(x=3)` does two things. First, it sets the value of the variable `x` to be 3. Second, it returns a value of the assignment, which is also 3. The exam creators are fond of inserting the assignment operator `=` in the middle of an expression and using the value of the assignment as part of a more complex expression.

Relational Operators

We now move on to *relational operators*, which compare two expressions and return a boolean value. The first four relational operators (see Table 2.3) are applied to numeric primitive data types only. If the two numeric operands are not of the same data type, the smaller one is promoted in the manner as previously discussed.

TABLE 2.3 Relational operators

<	Strictly less than
<=	Less than or equal to
>	Strictly greater than
>=	Greater than or equal to

Let’s look at examples of these operators in action:

```
int x = 10, y = 20, z = 10;
System.out.println(x < y); // Outputs true
System.out.println(x <= y); // Outputs true
System.out.println(x >= z); // Outputs true
System.out.println(x > z); // Outputs false
```

Notice that the last example outputs `false`, because although `x` and `z` are the same value, `x` is not strictly greater than `z`.

The fifth relational operator (Table 2.4) is applied to object references and classes or interfaces.

TABLE 2.4 Relational instanceof operator

<code>a instanceof b</code>	True if the reference that <code>a</code> points to is an instance of a class, subclass, or class that implements a particular interface, as named in <code>b</code>
-----------------------------	--

The `instanceof` operator, while useful for determining whether an arbitrary object is a member of a particular class or interface, is out of scope for the OCA exam.

Logical Operators

If you have studied computer science, you may have already come across logical operators before. If not, no need to panic—we'll be covering them in detail in this section.

The *logical operators*, (&), (|), and (^), may be applied to both numeric and boolean data types. When they're applied to boolean data types, they're referred to as *logical operators*. Alternatively, when they're applied to numeric data types, they're referred to as *bitwise operators*, as they perform bitwise comparisons of the bits that compose the number. For the exam, though, you don't need to know anything about numeric bitwise comparisons, so we'll leave that educational aspect to other books.

You should familiarize with the **truth tables** in Figure 2.1, where x and y are assumed to be boolean data types.

FIGURE 2.1 The logical true tables for &, |, and ^

x & y (AND)			x y (INCLUSIVE OR)			x ^ y (EXCLUSIVE OR)		
	y = true	y = false		y = true	y = false		y = true	y = false
x = true	true	false	x = true	true	true	x = true	false	true
x = false	false	false	x = false	true	false	x = false	true	false

Here are some tips to help remember this table:

- AND is only true if both operands are true.
- Inclusive OR is only false if both operands are false.
- Exclusive OR is only true if the operands are different.

Finally, we present the conditional operators, && and ||, which are often referred to as short-circuit operators. The *short-circuit operators* are nearly identical to the logical operators, & and |, respectively, except that the right-hand side of the expression may never be evaluated if the final result can be determined by the left-hand side of the expression. For example, consider the following statement:

```
boolean x = true || (y < 4);
```

Referring to the truth tables, the value x can only be false if both sides of the expression are false. Since we know the left-hand side is true, there's no need to evaluate the right-hand side, since no value of y will ever make the value of x anything other than true. It may help you to illustrate this concept by executing the previous line of code for various values of y.

A more common example of where short-circuit operators are used is checking for `null` objects before performing an operation, such as this:

```
if(x != null && x.getValue() < 5) {  
    // Do something  
}
```

In this example, if `x` was `null`, then the short-circuit prevents a `NullPointerException` from ever being thrown, since the evaluation of `x.getValue() < 5` is never reached. Alternatively, if we used a logical `&`, then both sides would always be evaluated and when `x` was `null` this would throw an exception:

```
if(x != null & x.getValue() < 5) { // Throws an exception if x is null  
    // Do something  
}
```

Be wary of short-circuit behavior on the exam, as questions are known to alter a variable on the right-hand side of the expression that may never be reached. For example, what is the output of the following code?

```
int x = 6;  
boolean y = (x >= 6) || (++x <= 7);  
System.out.println(x);
```

Because `x >= 6` is true, the increment operator on the right-hand side of the expression is never evaluated, so the output is 6.

Equality Operators

Determining equality in Java can be a nontrivial endeavor as there's a semantic difference between “two objects are the same” and “two objects are equivalent.” It is further complicated by the fact that for numeric and boolean primitives, there is no such distinction.

Let's start with the basics, the *equals* operator `==` and *not equals* operator `!=`. Like the relational operators, they compare two operands and return a boolean value about whether the expressions or values are equal, or not equal, respectively.

The equality operators are used in one of three scenarios:

1. Comparing two numeric primitive types. If the numeric values are of different data types, the values are automatically promoted as previously described. For example, `5 == 5.00` returns true since the left side is promoted to a double.
2. Comparing two boolean values.
3. Comparing two objects, including `null` and `String` values.

The comparisons for equality are limited to these three cases, so you cannot mix and match types. For example, each of the following would result in a compiler error:

```
boolean x = true == 3; // DOES NOT COMPILE
boolean y = false != "Giraffe"; // DOES NOT COMPILE
boolean z = 3 == "Kangaroo"; // DOES NOT COMPILE
```

Pay close attention to the data types when you see an equality operator on the exam. The exam creators also have a habit of mixing assignment operators and equality operators, as in the following snippet:

```
boolean y = false;
boolean x = (y = true);
System.out.println(x); // Outputs true
```

At first glance, you might think the output should be false, and if the expression was (`y == true`), then you would be correct. In this example, though, the expression is assigning the value of `true` to `y`, and as you saw in the section on assignment operators, the assignment itself has the value of the assignment. Therefore, the output would be `true`.

For object comparison, the equality operator is applied to the references to the objects, not the objects they point to. Two references are equal if and only if they point to the same object, or both point to `null`. Let's take a look at some examples:

```
File x = new File("myFile.txt");
File y = new File("myFile.txt");
File z = x;
System.out.println(x == y); // Outputs false
System.out.println(x == z); // Outputs true
```

Even though all of the variables point to the same file information, only two, `x` and `z`, are equal in terms of `==`. In this example, as well as during the OCA exam, you may be presented with classnames that are unfamiliar, such as `File`. Many times you can answer questions about these classes without knowing the specific details of these classes. In particular, you should be able to answer questions that indicate `x` and `y` are two separate and distinct objects, even if you do not know the data types of these objects.

In Chapter 3, “Core Java APIs,” we’ll continue the discussion of object equality by introducing what it means for two different objects to be equivalent. We’ll also cover `String` equality and show how this can be a nontrivial topic.

Understanding Java Statements

Java operators allow you to create a lot of complex expressions, but they’re limited in the manner in which they can control program flow. For example, imagine you want a section of code to only be executed under certain conditions that cannot be evaluated until

runtime. Or suppose you want a particular segment of code to repeat once for every item in some list.


As you may recall from Chapter 1, a Java *statement* is a complete unit of execution in Java, terminated with a semicolon (;). For the remainder of the chapter, we'll be introducing you to various Java control flow statements. *Control flow statements* break up the flow of execution by using decision making, looping, and branching, allowing the application to selectively execute particular segments of code.

These statements can be applied to single expressions as well as a block of Java code. As described in the previous chapter, a *block* of code in Java is a group of zero or more statements between balanced braces, ({}), and can be used anywhere a single statement is allowed.

The *if-then* Statement

Often, we only want to execute a block of code under certain circumstances. The *if-then* statement, as shown in Figure 2.2, accomplishes this by allowing our application to execute a particular block of code if and only if a boolean expression evaluates to true at runtime.

FIGURE 2.2 The structure of an *if-then* statement



For example, imagine we had a function that used the hour of day, an integer value from 0 to 23, to display a message to the user:

```
if(hourOfDay < 11)
    System.out.println("Good Morning");
```

If the hour of the day is less than 11, then the message will be displayed. Now let's say we also wanted to increment some value, `morningGreetingCount`, every time the greeting is printed. We could write the *if-then* statement twice, but luckily Java offers us a more natural approach using a block:

```
if(hourOfDay < 11) {
    System.out.println("Good Morning");
    morningGreetingCount++;
}
```

The block allows multiple statements to be executed based on the *if*-then evaluation. Notice that the first statement didn't contain a block around the print section, but it easily could have. For readability, it is considered good coding practice to put blocks around the execution component of *if*-then statements, as well as many other control flow statements, although it is not required.

Watch Indentation and Braces

One area that the exam writers will try to trip you up is on *if*-then statements without braces (`{}`). For example, take a look at this slightly modified form of our example:

```
if(hourOfDay < 11)
    System.out.println("Good Morning");
    morningGreetingCount++;
```

Based on the indentation, you might be inclined to think the variable `morningGreetingCount` is only going to be incremented if the `hourOfDay` is less than 11, but that's not what this code does. It will execute the print statement only if the condition is met, but it will always execute the increment operation.


Remember that in Java, unlike some other programming languages, tabs are just whitespace and are not evaluated as part of the execution. When you see a control flow statement in a question, be sure to trace the open and close braces of the block and ignore any indentation you may come across.

The *if-then-else* Statement

Let's expand our example a little. What if we want to display a different message if it is 11 a.m. or later? Could we do it using only the tools we have? Of course we can!

```
if(hourOfDay < 11) {
    System.out.println("Good Morning");
}
if(hourOfDay >= 11) {
    System.out.println("Good Afternoon");
}
```

This seems a bit redundant, though, since we're performing an evaluation on `hourOfDay` twice. It's also wasteful because in some circumstances the cost of the boolean expression we're evaluating could be computationally expensive. Luckily, Java offers us a more useful approach in the form of an *if-then-else* statement, as shown in Figure 2.3.

FIGURE 2.3 The structure of an if-then-else statement

Let's return to this example:

```

if(hourOfDay < 11) {
    System.out.println("Good Morning");
} else {
    System.out.println("Good Afternoon");
}

```

Now our code is truly branching between one of the two possible options, with the boolean evaluation happening only once. The `else` operator takes a statement or block of statement, in the same manner as the `if` statement does. In this manner, we can append additional if-then statements to an `else` block to arrive at a more refined example:

```

if(hourOfDay < 11) {
    System.out.println("Good Morning");
} else if(hourOfDay < 15) {
    System.out.println("Good Afternoon");
} else {
    System.out.println("Good Evening");
}

```

In this example, the Java process will continue execution until it encounters an if-then statement that evaluates to true. If neither of the first two expressions are true, it will execute the final code of the `else` block. One thing to keep in mind in creating complex

if-then-else statements is that order is important. For example, see what happens if we reorder the previous snippet of code as follows:

```
if(hourOfDay < 15) {
    System.out.println("Good Afternoon");
} else if(hourOfDay < 11) {
    System.out.println("Good Morning"); // UNREACHABLE CODE
} else {
    System.out.println("Good Evening");
}
```

For hours of the day less than 11, this code behaves very differently than the previous set of code. See if you can determine why the second block can never be executed regardless of the value of `hourOfDay`.

If a value is less than 11, then it must be also less than 15 by definition. Therefore, if the second branch in the example can be reached, the first branch can also be reached. Since execution of each branch is mutually exclusive in this example—that is, only one branch can be executed—if the first branch is executed, then the second cannot be executed. Therefore, there is no way the second branch will ever be executed, and the code is deemed unreachable.

Verifying the *if* Statement Evaluates to a Boolean Expression

Another common place the exam may try to lead you astray is by providing code where the boolean expression inside the if-then statement is not actually a boolean expression. For example, take a look at the following lines of code:

```
int x = 1;
if(x) { // DOES NOT COMPILE
    ...
}
```

This statement may be valid in some other programming and scripting languages, but not in Java, where 0 and 1 are not considered boolean values. Also, be wary of assignment operators being used as if they were equals `==` operators in if-then statements:

```
int x = 1;
if(x = 5) { // DOES NOT COMPILE
    ...
}
```

Ternary Operator

Now that we have discussed if-then-else statements, we can briefly return to our discussion of operators and present the final operator that you need to learn for the exam. The conditional operator, `?` `:`, otherwise known as the *ternary operator*, is the only operator that takes three operands and is of the form:

```
booleanExpression ? expression1 : expression2
```

The first operand must be a boolean expression, and the second and third can be any expression that returns a value. The ternary operation is really a condensed form of an if-then-else statement that returns a value. For example, the following two snippets of code are equivalent:

```
int y = 10;
final int x;
if(y > 5) {
    x = 2 * y;
} else {
    x = 3 * y;
}
```

Compare the previous code snippet with the following equivalent ternary operator code snippet:

```
int y = 10;
int x = (y > 5) ? (2 * y) : (3 * y);
```

Note that it is often helpful for readability to add parentheses around the expressions in ternary operations, although it is certainly not required.

There is no requirement that second and third expressions in ternary operations have the same data types, although it may come into play when combined with the assignment operator. Compare the following two statements:

```
System.out.println((y > 5) ? 21 : "Zebra");
int animal = (y < 91) ? 9 : "Horse"; // DOES NOT COMPILE
```

Both expressions evaluate similar boolean values and return an `int` and a `String`, although only the first line will compile. The `System.out.println()` does not care that the statements are completely different types, because it can convert both to `String`. On the other hand, the compiler does know that "Horse" is of the wrong data type and cannot be assigned to an `int`; therefore, it will not allow the code to be compiled.

Ternary Expression Evaluation

As of Java 7, only one of the right-hand expressions of the ternary operator will be evaluated at runtime. In a manner similar to the short-circuit operators, if one of the two right-hand expressions in a ternary operator performs a side effect, then it may not be applied at runtime. Let's illustrate this principle with the following example:

```
int y = 1;
int z = 1;
final int x = y < 10 ? y++ : z++;
System.out.println(y + ", " + z); // Outputs 2,1
```

Notice that since the left-hand boolean expression was true, only *y* was incremented. Contrast the preceding example with the following modification:

```
int y = 1;
int z = 1;
final int x = y >= 10 ? y++ : z++;
System.out.println(y + ", " + z); // Outputs 1,2
```

Now that the left-hand boolean expression evaluates to false, only *z* was incremented. In this manner, we see how the expressions in a ternary operator may not be applied if the particular expression is not used.

For the exam, be wary of any question that includes a ternary expression in which a variable is modified in one of the right-hand side expressions.

The *switch* Statement


We now expand on our discussion of if-then-else statements by discussing a switch statement. A *switch* statement, as shown in Figure 2.4, is a complex decision-making structure in which a single value is evaluated and flow is redirected to the first matching branch, known as a *case* statement. If no such case statement is found that matches the value, an optional *default* statement will be called. If no such default option is available, the entire switch statement will be skipped.

Supported Data Types

As shown in Figure 2.4, a switch statement has a target variable that is not evaluated until runtime. Prior to Java 5.0, this variable could only be int values or those values that could be promoted to int, specifically byte, short, char, or int. When enum was added in Java 5.0, support was added to switch statements to support enum values. In Java 7, switch

statements were further updated to allow matching on String values. Finally, the switch statement also supports any of the primitive numeric wrapper classes, such as Byte, Short, Character, or Integer.

FIGURE 2.4 The structure of a switch statement



Data types supported by switch statements include the following:

- int and Integer
- byte and Byte
- short and Short
- char and Character
- int and Integer
- String
- enum values

For the exam, we recommend you memorize this list. Note that boolean and long, and their associated wrapper classes, are not supported by switch statements.

Compile-time Constant Values

The values in each case statement must be compile-time constant values of the same data type as the switch value. This means you can use only literals, enum constants, or final

constant variables of the same data type. By final constant, we mean that the variable must be marked with the final modifier and initialized with a literal value in the same expression in which it is declared.

Let's look at a simple example using the day of the week, with 0 for Sunday, 1 for Monday, and so on:

```
int dayOfWeek = 5;
switch(dayOfWeek) {
    default:
        System.out.println("Weekday");
        break;
    case 0:
        System.out.println("Sunday");
        break;
    case 6:
        System.out.println("Saturday");
        break;
}
```

With a value of dayOfWeek of 5, this code will output:

```
Weekday
```

The first thing you may notice is that there is a break statement at the end of each case and default section. We'll discuss break statements in detail when we discuss loops, but for now all you need to know is that they terminate the switch statement and return flow control to the enclosing statement. As we'll soon see, if you leave out the break statement, flow will continue to the next proceeding case or default block automatically.

Another thing you might notice is that the default block is not at the end of the switch statement. There is no requirement that the case or default statements be in a particular order, unless you are going to have pathways that reach multiple sections of the switch block in a single execution.

To illustrate both of the preceding points, consider the following variation:

```
int dayOfWeek = 5;
switch(dayOfWeek) {
    case 0:
        System.out.println("Sunday");
    default:
        System.out.println("Weekday");
    case 6:
        System.out.println("Saturday");
        break;
}
```


This code looks a lot like the previous example except two of the break statements have been removed and the order has been changed. This means that for the given value of `dayOfWeek`, 5, the code will jump to the default block and then execute all of the proceeding case statements in order until it finds a break statement or finishes the structure:

```
Weekday
Saturday
```

The order of the case and default statements is now important since placing the default statement at the end of the switch statement would cause only one word to be output.

What if the value of `dayOfWeek` was 6 in this example? Would the default block still be executed? The output of this example with `dayOfWeek` set to 6 would be:

```
Saturday
```

Even though the default block was before the case block, only the case block was executed. If you recall the definition of the default block, it is only branched to if there is no matching case value for the switch statement, regardless of its position within the switch statement.

Finally, if the value of `dayOfWeek` was 0, all three statements would be output:

```
Sunday
Weekday
Saturday
```

Notice that in this last example, the default is executed since there was no break statement at the end of the preceding case block. While the code will not branch to the default statement if there is a matching case value within the switch statement, it will execute the default statement if it encounters it after a case statement for which there is no terminating break statement.

The exam creators are fond of switch examples that are missing break statements! When evaluating switch statements on the exam, always consider that multiple branches may be visited in a single execution.

We conclude our discussion on switch statements by acknowledging that the data type for case statements must all match the data type of the switch variable. As already discussed, the case statement value must also be a literal, enum constant, or final constant variable. For example, given the following switch statement, notice which case statements will compile and which will not:

```
private int getSortOrder(String firstName, final String lastName) {
    String middleName = "Patricia";
    final String suffix = "JR";
    int id = 0;
    switch(firstName) {
        case "Test":
            return 52;
```

```
case middleName: // DOES NOT COMPILE
    id = 5;
    break;
case suffix:
    id = 0;
    break;
case lastName: // DOES NOT COMPILE
    id = 8;
    break;
case 5: // DOES NOT COMPILE
    id = 7;
    break;
case 'J': // DOES NOT COMPILE
    id = 10;
    break;
case java.time.DayOfWeek.SUNDAY: // DOES NOT COMPILE
    id=15;
    break;
}
return id;
}
```

The first case statement compiles without issue using a `String` literal and is a good example of how a `return` statement, like a `break` statement, can be used to exit the `switch` statement early. The second case statement does not compile because `middleName` is not a `final` variable, despite having a known value at this particular line of execution. The third case statement compiles without issue because `suffix` is a `final` constant variable.


In the fourth case statement, despite `lastName` being `final`, it is not constant as it is passed to the function; therefore, this line does not compile as well. Finally, the last three case statements don't compile because none of them have a matching type of `String`; the last one is an `enum` value.

The *while* Statement

A repetition control structure, which we refer to as a *loop*, executes a statement of code multiple times in succession. By using nonconstant variables, each repetition of the statement may be different. For example, a statement that iterates over a list of unique names and outputs them would encounter a new name on every execution of the loop.

The simplest such repetition control structure in Java is the *while* statement, described in Figure 2.5. Like all repetition control structures, it has a termination condition, implemented as a boolean expression, that will continue as long as the expression evaluates to true.

FIGURE 2.5 The structure of a while statement



As shown in Figure 2.5, a *while* loop is similar to an *if-then* statement in that it is composed of a boolean expression and a statement, or block of statements. During execution, the boolean expression is evaluated before each iteration of the loop and exits if the evaluation returns false. It is important to note that a *while* loop may terminate after its first evaluation of the boolean expression. In this manner, the statement block may never be executed.

Let's return to our mouse example from Chapter 1 and show a loop can be used to model a mouse eating a meal:

```

int roomInBelly = 5;

public void eatCheese(int bitesOfCheese) {
    while (bitesOfCheese > 0 && roomInBelly > 0) {
        bitesOfCheese--;
        roomInBelly--;
    }
    System.out.println(bitesOfCheese+" pieces of cheese left");
}
  
```

This method takes an amount of food, in this case cheese, and continues until the mouse has no room in its belly or there is no food left to eat. With each iteration of the loop, the mouse “eats” one bite of food and loses one spot in its belly. By using a compound boolean statement, you ensure that the *while* loop can end for either of the conditions.



Real World Scenario

Infinite Loops

Consider the following segment of code:

```
int x = 2;
int y = 5;
while(x < 10)
    y++;
```


You may notice one glaring problem with this statement: it will never end! The boolean expression that is evaluated prior to each loop iteration is never modified, so the expression ($x < 10$) will always evaluate to true. The result is that the loop will never end, creating what is commonly referred to as an *infinite loop*.

Infinite loops are something you should be aware of any time you create a loop in your application. You should be absolutely certain that the loop will eventually terminate under some condition. First, make sure the loop variable is modified. Then, ensure that the termination condition will be eventually reached in all circumstances. As you'll see in the upcoming section "Understanding Advanced Flow Control," a loop may also exit under other conditions such as a `break` statement.

The *do-while* Statement

Java also allows for the creation of a *do-while* loop, which like a `while` loop, is a repetition control structure with a termination condition and statement, or block of statements, as shown in Figure 2.6. Unlike a `while` loop, though, a *do-while* loop guarantees that the statement or block will be executed at least once.

FIGURE 2.6 The structure of a *do-while* statement



The primary difference between the syntactic structure of a do-while loop and a while loop is that a do-while loop purposely orders the statement or block of statements before the conditional expression, in order to reinforce that the statement will be executed before the expression is ever evaluated. For example, take a look at the output of the following statements:

```
int x = 0;
do {
    x++;
} while(false);
System.out.println(x); // Outputs 1
```

Java will execute the statement block first, and then check the loop condition. Even though the loop exits right away, the statement block was still executed once and the program outputs a 1.

When to Use *while* vs. *do-while* Loops

In practice, it might be difficult to determine when you should use a while loop and when you should use a do-while loop. The short answer is that it does not actually matter. Any while loop can be converted to a do-while loop, and vice versa. For example, compare this while loop:

```
while(x > 10) {
    x--;
}
```

and this do-while loop:

```
if(x > 10) {
    do {
        x--;
    } while(x > 10);
}
```

Though one of the loops is certainly easier to read, they are functionally equivalent. Java recommends you use a while loop when a loop might not be executed at all and a do-while loop when the loop is executed at least once. But determining whether you should use a while loop or a do-while loop in practice is sometimes about personal preference and code readability.

For example, although the first statement is shorter, the second has the advantage that you could leverage the existing if-then statement and perform some other operation in a new else branch, as shown in the following example:

continues

continued

```

if(x > 10) {
    do {
        x--;
    } while(x > 10);
} else {
    x++;
}

```

The *for* Statement


Now that you can build applications with simple while and do-while statements, we expand our discussion of loops to a more complex repetition control structure called a for loop.

Starting in Java 5.0, there are now two types of for statements. The first is referred to as the basic for loop, and the second is often called the enhanced for loop. For clarity, we'll refer to the enhanced for loop as the for-each statement throughout the book.

The Basic *for* Statement

A basic *for loop* has the same conditional boolean expression and statement, or block of statements, as the other loops you have seen, as well as two new sections: an *initialization block* and an *update* statement. Figure 2.7 shows how these components are laid out.

FIGURE 2.7 The structure of a basic for statement



- ① Initialization statement executes
- ② If `booleanExpression` is true continue, else exit loop
- ③ Body executes
- ④ Execute `updateStatements`
- ⑤ Return to Step 2

Although Figure 2.7 might seem a little confusing and almost arbitrary at first, the organization of the components and flow allow us to create extremely powerful statements in a very small amount of space that otherwise would take multiple lines with a standard while loop. Note that each section is separated by a semicolon. The initialization and update sections may contain multiple statements, separated by commas.

Variables declared in the initialization block of a for loop have limited scope and are only accessible within the for loop. Be wary of any exam questions in which a variable declared within the initialization block of a for loop is available outside the loop. Alternatively, variables declared before the for loop and assigned a value in the initialization block may be used outside the for loop because their scope precedes the for loop creation.

Let's take a look at an example that prints the numbers 0 to 9:

```
for(int i = 0; i < 10; i++) {  
    System.out.print(i + " ");  
}
```

The local variable `i` is initialized first to 0. The variable `i` is only in scope for the duration of the loop and is not available outside the loop once the loop has completed. Like a while loop, the boolean condition is evaluated on every iteration of the loop *before* the loop executes. Since it returns true, the loop executes and outputs the 0 followed by a space. Next, the loop executes the update section, which in this case increases the value of `i` to 1. The loop then evaluates the boolean expression a second time, and the process repeats multiple times, printing:

0 1 2 3 4 5 6 7 8 9

On the 10th iteration of the loop, the value of `i` reaches 9 and is incremented by 1 to reach 10. On the 11th iteration of the loop, the boolean expression is evaluated and since $(10 < 10)$ returns false, the loop terminates without executing the statement loop body.

Although most for loops you are likely to encounter in practice will be well defined and similar to the previous example, there are a number of variations and edge cases you could see on the exam. You should familiarize yourself with the following five examples: variations of these are likely to be seen on the exam.

Let's tackle some examples for illustrative purposes:

1. Creating an Infinite Loop

```
for( ; ; ) {  
    System.out.println("Hello World");  
}
```

Although this for loop may look like it will throw a compiler error, it will in fact compile and run without issue. It is actually an infinite loop that will print the same statement repeatedly. This example reinforces the fact that the components of the for loop are each optional. Note that the semicolons separating the three sections are required, as `for( ; )` and `for( )` will not compile.

2. Adding Multiple Terms to the `for` Statement

```
int x = 0;
for(long y = 0, z = 4; x < 5 && y < 10; x++, y++) {
    System.out.print(y + " ");
}
System.out.print(x);
```

This code demonstrates three variations of the `for` loop you may not have seen. First, you can declare a variable, such as `x` in this example, before the loop begins and use it after it completes. Second, your initialization block, boolean expression, and update statements can include extra variables that may not reference each other. For example, `z` is defined in the initialization block and is never used. Finally, the update statement can modify multiple variables. This code will print the following when executed:

```
0 1 2 3 4
```

Keep this example in mind when we look at the next three examples, none of which compile.

3. Redeclaring a Variable in the Initialization Block

```
int x = 0;
for(long y = 0, x = 4; x < 5 && y < 10; x++, y++) {    // DOES NOT COMPILE
    System.out.print(x + " ");
}
```

This example looks similar to the previous one, but it does not compile because of the initialization block. The difference is that `x` is repeated in the initialization block after already being declared before the loop, resulting in the compiler stopping because of a duplicate variable declaration. We can fix this loop by changing the declaration of `x` and `y` as follows:

```
int x = 0;
long y = 10;
for(y = 0, x = 4; x < 5 && y < 10; x++, y++) {
    System.out.print(x + " ");
}
```

Note that this variation will now compile because the initialization block simply assigns a value to `x` and does not declare it.

4. Using Incompatible Data Types in the Initialization Block

```
for(long y = 0, int x = 4; x < 5 && y < 10; x++, y++) {    // DOES NOT COMPILE
    System.out.print(x + " ");
}
```

This example also looks a lot like our second example, but like the third example will not compile, although this time for a different reason. The variables in the initialization block must all be of the same type. In the first example, `y` and `z` were both `long`, so the code compiled without issue, but in this example they have differing types, so the code will not compile.

5. Using Loop Variables Outside the Loop


```
for(long y = 0, x = 4; x < 5 && y < 10; x++, y++) {
    System.out.print(y + " ");
}
System.out.print(x); // DOES NOT COMPILE
```

The final variation on the second example will not compile for a different reason than the previous examples. If you notice, `x` is defined in the initialization block of the loop, and then used after the loop terminates. Since `x` was only scoped for the loop, using it outside the loop will throw a compiler error.

The *for-each* Statement

Starting with Java 5.0, Java developers have had a new type of enhanced for loop at their disposal, one specifically designed for iterating over arrays and Collection objects. This enhanced for loop, which for clarity we'll refer to as a *for-each* loop, is shown in Figure 2.8.

FIGURE 2.8 The structure of an enhancement for statement



The for-each loop declaration is composed of an initialization section and an object to be iterated over. The right-hand side of the for-each loop statement must be a built-in Java array or an object whose class implements `java.lang.Iterable`, which includes most of the Java Collections framework. The left-hand side of the for-each loop must include a declaration for an instance of a variable, whose type matches the type of a member of the array or collection in the right-hand side of the statement. On each iteration of the loop, the named variable on the left-hand side of the statement is assigned a new value from the array or collection on the right-hand side of the statement.



For the OCA exam, the only members of the Collections framework that you need to be aware of are `List` and `ArrayList`. In this chapter, we'll show how to iterate over `List` objects, and in Chapter 3 we'll go into detail about how to create `List` objects and how they differ from traditional Java arrays.

Let's review some examples:

- What will this code output?

```
final String[] names = new String[3];
names[0] = "Lisa";
names[1] = "Kevin";
names[2] = "Roger";
for(String name : names) {
    System.out.print(name + ", ");
}
```

This code will compile and print:

Lisa, Kevin, Roger,

- What will this code output?

```
java.util.List<String> values = new java.util.ArrayList<String>();
values.add("Lisa");
values.add("Kevin");
values.add("Roger");
for(String value : values) {
    System.out.print(value + ", ");
}
```

This code will compile and print the same values:

Lisa, Kevin, Roger,

When you see a for-each loop on the exam, make sure the right-hand side is an array or Iterable object and the left-hand side has a matching type. For example, the two examples that follow will not compile.

- Why will the following fail to compile?

```
String names = "Lisa";
for(String name : names) {    // DOES NOT COMPILE
    System.out.print(name + " ");
}
```

In this example, the String names is not an array, nor does it implement java.lang.Iterable, so the compiler will throw an exception since it does not know how to iterate over the String.

- Why will the following fail to compile?

```
String[] names = new String[3];
for(int name : names) {    // DOES NOT COMPILE
    System.out.print(name + " ");
}
```

This code will fail to compile because the left-hand side of the for-each statement does not define an instance of `String`. Notice that in this last example, the array is initialized with three null pointer values. In and of itself, that will not cause the code to not compile, as a corrected loop would just output null three times.



Real World Scenario

Comparing *for* and *for-each* Loops

Since `for` and `for-each` both use the same keyword, you might be wondering how they are related. While this discussion is out of scope for the exam, let's take a moment to explore how `for-each` loops are converted to `for` loops by the compiler.

When `for-each` was introduced in Java 5, it was added as a compile-time enhancement. This means that Java actually converts the `for-each` loop into a standard `for` loop during compilation. For example, assuming `names` is an array of `String[]` as we saw in the first example, the following two loops are equivalent:

```
for(String name : names) {  
    System.out.print(name + ", ");  
}  
for(int i=0; i < names.length; i++) {  
    String name = names[i];  
    System.out.print(name + ", ");  
}
```

For objects that inherit `java.lang.Iterable`, there is a different, but similar, conversion. For example, assuming `values` is an instance of `List<Integer>`, as we saw in the second example, the following two loops are equivalent:

```
for(int value : values) {  
    System.out.print(value + ", ");  
}  
for(java.util.Iterator<Integer> i = values.iterator(); i.hasNext(); ) {  
    int value = i.next();  
    System.out.print(value + ", ");  
}
```

Notice that in the second version, there is no update statement as it is not required when using the `java.util.Iterator` class.

You may have noticed that in the previous `for-each` examples, there was an extra comma printed at the end of the list:

Lisa, Kevin, Roger,

While the `for-each` statement is convenient for working with lists in many cases, it does hide access to the loop iterator variable. If we wanted to print only the comma between names, we could convert the example into a standard `for` loop, as in the following example:

```
java.util.List<String> names = new java.util.ArrayList<String>();
names.add("Lisa");
names.add("Kevin");
names.add("Roger");
for(int i=0; i<names.size(); i++) {
    String name = names.get(i);
    if(i>0) {
        System.out.print(", ");
    }
    System.out.print(name);
}
```

This sample code would output the following:

Lisa, Kevin, Roger

It is also common to use a standard `for` loop over a `for-each` loop if comparing multiple elements in a loop within a single iteration, as in the following example. Notice that we skip the first loop's execution, since `value[-1]` is not defined and would throw an `IndexOutOfBoundsException` error.

```
int[] values = new int[3];
values[0] = 10;
values[1] = new Integer(5);
values[2] = 15;
for(int i=1; i<values.length; i++) {
    System.out.print(values[i]-values[i-1]);
}
```

This sample code would output the following:

-5, 10,

Despite these examples, enhanced `for-each` loops are quite useful in Java in a variety of circumstances. As a developer, though, you can always revert to a standard `for` loop if you need fine-grain control.

Understanding Advanced Flow Control

Up to now, we have been dealing with single loops that only ended when their boolean expression evaluated to false. We'll now show you other ways loops could end, or branch, and you'll see that the path taken during runtime may not be as straightforward as in previous examples.

Nested Loops

First off, loops can contain other loops. For example, consider the following code that iterates over a two-dimensional array, an array that contains other arrays as its members. We'll cover multidimensional arrays in detail in Chapter 3, but for now assume the following is how you would declare a two-dimensional array.

```
int[][] myComplexArray = {{5,2,1,3},{3,9,8,9},{5,7,12,7}};
for(int[] mySimpleArray : myComplexArray) {
    for(int i=0; i<mySimpleArray.length; i++) {
        System.out.print(mySimpleArray[i]+"\\t");
    }
    System.out.println();
}
```

Notice that we intentionally mix a `for` and `for-each` loop in this example. The outer loops will execute a total of three times. Each time the outer loop executes, the inner loop is executed four times. When we execute this code, we see the following output:

```
5      2      1      3
3      9      8      9
5      7     12      7
```

Nested loops can include `while` and `do-while`, as shown in this example. See if you can determine what this code will output.

```
int x = 20;
while(x>0) {
    do {
        x -= 2
    } while (x>5);
    x--;
    System.out.print(x+"\\t");
}
```

The first time this loop executes, the inner loop repeats until the value of `x` is 4. The value will then be decremented to 3 and that will be the output at the end of the first iteration of the outer loop. On the second iteration of the outer loop, the inner `do-while` will be executed once, even though `x` is already not greater than 5. As you may recall, `do-while` statements always execute the body at least once. This will reduce the value to 1, which will be further lowered by the decrement operator in the outer loop to 0. Once the value reaches 0, the outer loop will terminate. The result is that the code will output the following:

```
3      0
```

Adding Optional Labels

One thing we skipped when we presented `if-then` statements, `switch` statements, and loops is that they can all have optional labels. A *label* is an optional pointer to the head of a

currently in the process of executing. The optional label parameter allows us to break out of a higher level outer loop. In the following example, we search for the first (x,y) array index position of a number within an unsorted two-dimensional array:

```
public class SearchSample {
    public static void main(String[] args) {
        int[][] list = {{1,13,5},{1,2,5},{2,7,2}};
        int searchValue = 2;
        int positionX = -1;
        int positionY = -1;
        PARENT_LOOP: for(int i=0; i<list.length; i++) {
            for(int j=0; j<list[i].length; j++) {
                if(list[i][j]==searchValue) {
                    positionX = i;
                    positionY = j;
                    break PARENT_LOOP;
                }
            }
        }
        if(positionX== -1 || positionY== -1) {
            System.out.println("Value "+searchValue+" not found");
        } else {
            System.out.println("Value "+searchValue+" found at: " +
                "("+positionX+","+positionY+")");
        }
    }
}
```

When executed, this code will output:

```
Value 2 found at: (1,1)
```

In particular, take a look at the statement `break PARENT_LOOP`. This statement will break out of the entire loop structure as soon as the first matching value is found. Now, imagine what would happen if we replaced the body of the inner loop with the following:

```
if(list[i][j]==searchValue) {
    positionX = i;
    positionY = j;
    break;
}
```

How would this change our flow and would the output change? Instead of exiting when the first matching value is found, the program will now only exit the inner loop when the

condition is met. In other words, the structure will now find the first matching value of the last inner loop to contain the value, resulting in the following output:

Value 2 found at: (2,0)

Finally, what if we removed the break altogether?

```
if(list[i][j]==searchValue) {
    positionX = i;
    positionY = j;
}
```

In this case, the code will search for the last value in the entire structure that has the matching value. The output will look like this:


Value 2 found at: (2,2)

You can see from this example that using a label on a break statement in a nested loop, or not using the break statement at all, can cause the loop structure to behave quite differently.

The *continue* Statement

Let's now complete our discussion of advanced loop control with the *continue* statement, a statement that causes flow to finish the execution of the current loop, as shown in Figure 2.10.

FIGURE 2.10 The structure of a continue statement



You may notice the syntax of the continue statement mirrors that of the break statement. In fact, the statements are similar in how they are used, but with different results. While the break statement transfers control to the enclosing statement, the continue statement transfers control to the boolean expression that determines if the loop should continue. In other words, it ends the current iteration of the loop. Also like the break statement, the continue statement is applied to the nearest inner loop under execution using optional label statements to override this behavior. Let's take a look at the following example:


```

public class SwitchSample {
    public static void main(String[] args) {
        FIRST_CHAR_LOOP: for (int a = 1; a <= 4; a++) {
            for (char x = 'a'; x <= 'c'; x++) {
                if (a == 2 || x == 'b')
                    continue FIRST_CHAR_LOOP;
                System.out.print(" " + a + x);
            }
        }
    }
}

```

With the structure as defined, the loop will return control to the parent loop any time the first value is 2 or the second value is b. This results in one execution of the inner loop for each of three outer loop calls. The output looks like this:

1a 3a 4a

Now, imagine we removed the `FIRST_CHAR_LOOP` label in the `continue` statement so that control is returned to the inner loop instead of the outer. See if you can understand how the output will be changed to:

1a 1c 3a 3c 4a 4c

Finally, if we remove the `continue` statement and associated `if-then` statement altogether, we arrive at a structure that outputs all the values, such as:

1a 1b 1c 2a 2b 2c 3a 3b 3c 4a 4b 4c

Table 2.5 will help remind you when labels, `break`, and `continue` statements are permitted in Java. Although for illustrative purposes our examples have included using these statements in nested loops, they can be used inside single loops as well.

TABLE 2.5 Advanced flow control usage

	Allows optional labels	Allows <i>break</i> statement	Allows <i>continue</i> statement
if	Yes *	No	No
while	Yes	Yes	Yes
do while	Yes	Yes	Yes
for	Yes	Yes	Yes
switch	Yes	Yes	No

* Labels are allowed for any block statement, including those that are preceded with an `if-then` statement.

Summary

This chapter covered a wide variety of topics, including dozens of Java operators, along with numerous control flow statements. Many of these operators and statements may have been new to you.

It is important that you understand how to use all of the required Java operators covered in this chapter and know how operator precedence influences the way a particular expression is interpreted. There will likely be numerous questions on the exam that appear to test one thing, such as `StringBuilder` or exception handling, when in fact the answer is related to the misuse of a particular operator that causes the application to fail to compile. When you see an operator on the exam, always check that the appropriate data types are used and that they match each other where applicable.

For statements, this chapter covered two types of control structures: decision-making controls structures, including `if-then`, `if-then-else`, and `switch` statements, as well as repetition control structures including `for`, `for-each`, `while`, and `do-while`. Remember that most of these structures require the evaluation of a particular boolean expression either for branching decisions or once per repetition. The `switch` statement is the only one that supports a variety of data types, including `String` variables as of Java 7.

With a `for-each` statement you don't need to explicitly write a boolean expression, since the compiler builds them implicitly. For clarity, we referred to an enhanced `for` loop as a `for-each` loop, but syntactically they are written as a `for` statement.

We concluded this chapter by discussing advanced control options and how flow can be enhanced through nested loops, `break` statements, and `continue` statements. Be wary of questions on the exam that use nested statements, especially ones with labels, and verify they are being used correctly.

This chapter is especially important because at least one component of this chapter will likely appear in every exam question with sample code. Many of the questions on the exam focus on proper syntactic use of the structures, as they will be a large source of questions that end in "Does not compile." You should be able to answer all of the review questions correctly or fully understand those that you answered incorrectly before moving on to later chapters.

Exam Essentials

Be able to write code that uses Java operators. This chapter covered a wide variety of operator symbols. Go back and review them several times so that you are familiar with them throughout the rest of the book.

Be able to recognize which operators are associated with which data types. Some operators may be applied only to numeric primitives, some only to boolean values, and some only to objects. It is important that you notice when an operator and operand(s) are mismatched, as this issue is likely to come up in a couple of exam questions.

Understand Java operator precedence. Most Java operators you'll work with are binary, but the number of expressions is often greater than two. Therefore, you must understand the order in which Java will evaluate each operator symbol.

Be able to write code that uses parentheses to override operator precedence. You can use parentheses in your code to manually change the order of precedence.

Understand if and switch decision control statements. The if-then and switch statements come up frequently throughout the exam in questions unrelated to decision control, so make sure you fully understand these basic building blocks of Java.

Understand loop statements. Know the syntactical structure of all loops, including while, do-while, and for. Each loop has its own special properties and structures. Also, be familiar with the enhanced for-each loops that iterate over lists.

Understand how break and continue can change flow control. Know how to change the flow control within a statement by applying a break or continue command. Also know which control statements can accept break statements and which can accept continue statements. Finally, understand how these statements work inside embedded loops or switch statements.

Review Questions

1. Which of the following Java operators can be used with boolean variables? (Choose all that apply)
 - A. `==`
 - B. `+`
 - C. `--`
 - D. `!`
 - E. `%`
 - F. `<=`
2. What data type (or types) will allow the following code snippet to compile? (Choose all that apply)

```
byte x = 5;
byte y = 10;
----- z = x + y;
```

 - A. `int`
 - B. `long`
 - C. `boolean`
 - D. `double`
 - E. `short`
 - F. `byte`
3. What is the output of the following application?

```
1: public class CompareValues {
2:     public static void main(String[] args) {
3:         int x = 0;
4:         while(x++ < 10) {}
5:         String message = x > 10 ? "Greater than" : false;
6:         System.out.println(message+" "+x);
7:     }
8: }
```

 - A. Greater than,10
 - B. false,10
 - C. Greater than,11
 - D. false,11
 - E. The code will not compile because of line 4.
 - F. The code will not compile because of line 5.

4. What change would allow the following code snippet to compile? (Choose all that apply)

```
3: long x = 10;  
4: int y = 2 * x;
```

- A. No change; it compiles as is.
- B. Cast x on line 4 to int.
- C. Change the data type of x on line 3 to short.
- D. Cast 2 * x on line 4 to int.
- E. Change the data type of y on line 4 to short.
- F. Change the data type of y on line 4 to long.

5. What is the output of the following code snippet?

```
3: java.util.List<Integer> list = new java.util.ArrayList<Integer>();  
4: list.add(10);  
5: list.add(14);  
6: for(int x : list) {  
7:     System.out.print(x + ", ");  
8:     break;  
9: }
```

- A. 10, 14,
- B. 10, 14
- C. 10,
- D. The code will not compile because of line 7.
- E. The code will not compile because of line 8.
- F. The code contains an infinite loop and does not terminate.

6. What is the output of the following code snippet?

```
3: int x = 4;  
4: long y = x * 4 - x++;  
5: if(y<10) System.out.println("Too Low");  
6: else System.out.println("Just right");  
7: else System.out.println("Too High");
```

- A. Too Low
- B. Just Right
- C. Too High
- D. Compiles but throws a NullPointerException.
- E. The code will not compile because of line 6.
- F. The code will not compile because of line 7.

7. What is the output of the following code?
- ```
1: public class TernaryTester {
2: public static void main(String[] args) {
3: int x = 5;
4: System.out.println(x > 2 ? x < 4 ? 10 : 8 : 7);
5: }}
```
- A. 5  
B. 4  
C. 10  
D. 8  
E. 7  
F. The code will not compile because of line 4.
8. What is the output of the following code snippet?
- ```
3: boolean x = true, z = true;  
4: int y = 20;  
5: x = (y != 10) ^ (z=false);  
6: System.out.println(x+", "+y+", "+z);
```
- A. true, 10, true
B. true, 20, false
C. false, 20, true
D. false, 20, false
E. false, 20, true
F. The code will not compile because of line 5.
9. How many times will the following code print "Hello World"?
- ```
3: for(int i=0; i<10 ;) {
4: i = i++;
5: System.out.println("Hello World");
6: }
```
- A. 9  
B. 10  
C. 11  
D. The code will not compile because of line 3.  
E. The code will not compile because of line 5.  
F. The code contains an infinite loop and does not terminate.
10. What is the output of the following code?
- ```
3: byte a = 40, b = 50;  
4: byte sum = (byte) a + b;  
5: System.out.println(sum);
```

- A. 40
 - B. 50
 - C. 90
 - D. The code will not compile because of line 4.
 - E. An undefined value.
11. What is the output of the following code?
- ```
1: public class ArithmeticSample {
2: public static void main(String[] args) {
3: int x = 5 * 4 % 3;
4: System.out.println(x);
5: }}
```
- A. 2
  - B. 3
  - C. 5
  - D. 6
  - E. The code will not compile because of line 3.
12. What is the output of the following code snippet?
- ```
3: int x = 0;  
4: String s = null;  
5: if(x == s) System.out.println("Success");  
6: else System.out.println("Failure");
```
- A. Success
 - B. Failure
 - C. The code will not compile because of line 4.
 - D. The code will not compile because of line 5.
13. What is the output of the following code snippet?
- ```
3: int x1 = 50, x2 = 75;
4: boolean b = x1 >= x2;
5: if(b = true) System.out.println("Success");
6: else System.out.println("Failure");
```
- A. Success
  - B. Failure
  - C. The code will not compile because of line 4.
  - D. The code will not compile because of line 5.
14. What is the output of the following code snippet?
- ```
3: int c = 7;  
4: int result = 4;
```

```
5: result += ++c;
6: System.out.println(result);
```

- A. 8
- B. 11
- C. 12
- D. 15
- E. 16
- F. The code will not compile because of line 5.

15. What is the output of the following code snippet?

```
3: int x = 1, y = 15;
4: while x < 10
5:     y--;
6:     x++;
7: System.out.println(x+", "+y);
```

- A. 10, 5
- B. 10, 6
- C. 11, 5
- D. The code will not compile because of line 3.
- E. The code will not compile because of line 4.
- F. The code contains an infinite loop and does not terminate.

16. What is the output of the following code snippet?

```
3: do {
4:     int y = 1;
5:     System.out.print(y++ + " ");
6: } while(y <= 10);
```

- A. 1 2 3 4 5 6 7 8 9
- B. 1 2 3 4 5 6 7 8 9 10
- C. 1 2 3 4 5 6 7 8 9 10 11
- D. The code will not compile because of line 6.
- E. The code contains an infinite loop and does not terminate.

17. What is the output of the following code snippet?

```
3: boolean keepGoing = true;
4: int result = 15, i = 10;
5: do {
6:     i--;
7:     if(i==8) keepGoing = false;
8:     result -= 2;
9: } while(keepGoing);
```



```
10: System.out.println(result);
```

- A. 7
- B. 9
- C. 10
- D. 11
- E. 15
- F. The code will not compile because of line 8.

18. What is the output of the following code snippet?

```
3: int count = 0;
4: ROW_LOOP: for(int row = 1; row <=3; row++)
5:   for(int col = 1; col <=2 ; col++) {
6:     if(row * col % 2 == 0) continue ROW_LOOP;
7:     count++;
8:   }
9: System.out.println(count);
```

- A. 1
- B. 2
- C. 3
- D. 4
- E. 6
- F. The code will not compile because of line 6.

19. What is the result of the following code snippet?

```
3: int m = 9, n = 1, x = 0;
4: while(m > n) {
5:   m--;
6:   n += 2;
7:   x += m + n;
8: }
9: System.out.println(x);
```

- A. 11
- B. 13
- C. 23
- D. 36
- E. 50
- F. The code will not compile because of line 7.

20. What is the result of the following code snippet?

```
3: final char a = 'A', d = 'D';
4: char grade = 'B';
```

```
5: switch(grade) {  
6:   case a:  
7:     case 'B': System.out.print("great");  
8:     case 'C': System.out.print("good"); break;  
9:     case d:  
10:    case 'F': System.out.print("not good");  
11: }
```

- A.** great
- B.** greatgood
- C.** The code will not compile because of line 3.
- D.** The code will not compile because of line 6.
- E.** The code will not compile because of lines 6 and 9.

Chapter 3

Core Java APIs

OCA EXAM OBJECTIVES COVERED IN THIS CHAPTER:

✓ Using Operators and Decision Constructs

- Test equality between Strings and other objects using `==` and `equals()`

✓ Creating and Using Arrays

- Declare, instantiate, initialize and use a one-dimensional array
- Declare, instantiate, initialize and use a multi-dimensional array

✓ Working with Selected classes from the Java API

- Creating and manipulating Strings
- Manipulate data using the `StringBuilder` class and its methods
- Declare and use an `ArrayList` of a given type
- Create and manipulate calendar data using classes from `java.time.LocalDateTime`, `java.time.LocalDate`, `java.time.LocalTime`, `java.time.format.DateTimeFormatter`, `java.time.Period`

✓ Working with Java Data Types

- Develop code that uses wrapper classes such as `Boolean`, `Double`, and `Integer`.





The OCA exam expects you to know the core data structures and classes used in Java, and in this chapter readers will learn about the most common methods available. For example, `String` and `StringBuilder` are used for text data. An array and an `ArrayList` are used when you have multiple values. A variety of classes are used for working with dates. In this chapter, you'll also learn how to determine whether two objects are equal.

API stands for application programming interface. In Java, an interface is something special. In the context of an API, it can be a group of class or interface definitions that gives you access to a service or functionality. You will learn about the most common APIs for each of the classes covered in this chapter.

Creating and Manipulating Strings

The `String` class is such a fundamental class that you'd be hard-pressed to write code without it. After all, you can't even write a `main()` method without using the `String` class. A *string* is basically a sequence of characters; here's an example:

```
String name = "Fluffy";
```

As you learned in Chapter 1, “Java Building Blocks,” this is an example of a reference type. You also learned that reference types are created using the `new` keyword. Wait a minute. Something is missing from the previous example: it doesn't have `new` in it! In Java, these two snippets both create a `String`:

```
String name = "Fluffy";  
String name = new String("Fluffy");
```

Both give you a reference variable of type *name* pointing to the `String` object "Fluffy". They are subtly different, as you'll see in the section “String Pool,” later in this chapter. For now, just remember that the **String class is special and doesn't need to be instantiated with `new`.**

In this section, we'll look at concatenation, immutability, the string pool, common methods, and method chaining.

Concatenation

In Chapter 2, “Operators and Statements,” you learned how to add numbers. $1 + 2$ is clearly 3. But what is `"1" + "2"`? It's actually `"12"` because Java combines the two `String` objects.

Placing one String before the other String and combining them together is called string *concatenation*. The OCA exam creators like string concatenation because the + operator can be used in two ways within the same line of code. **There aren't a lot of rules to know for this, but you have to know them well:**

- 1. If both operands are numeric, + means numeric addition.**
- 2. If either operand is a String, + means concatenation.**
- 3. The expression is evaluated left to right.**

Now let's look at some examples:

```
System.out.println(1 + 2);           // 3
System.out.println("a" + "b");       // ab
System.out.println("a" + "b" + 3);    // ab3
System.out.println(1 + 2 + "c");      // 3c
```

The first example uses the first rule. Both operands are numbers, so we use normal addition. The second example is simple string concatenation, described in the second rule. The quotes for the String are only used in code—they don't get output.

The third example combines both the second and third rules. Since we start on the left, Java figures out what "a" + "b" evaluates to. You already know that one: it's "ab". Then Java looks at the remaining expression of "ab" + 3. The second rule tells us to concatenate since one of the operands is a String.

In the fourth example, we start with the third rule, which tells us to consider 1 + 2. Both operands are numeric, so the first rule tells us the answer is 3. Then we have 3 + "c", which uses the second rule to give us "3c". Notice all three rules get used in one line? The exam takes this a step further and will try to trick you with something like this:

```
int three = 3;
String four = "4";
System.out.println(1 + 2 + three + four);
```

When you see this, just take it slow and remember the three rules—and be sure to check the variable types. In this example, we start with the third rule, which tells us to consider 1 + 2. The first rule gives us 3. Next we have 3 + three. Since three is of type int, we still use the first rule, giving us 6. Next we have 6 + four. Since four is of type String, we switch to the second rule and get a final answer of "64". When you see questions like this, just take your time and check the types. Being methodical pays off.

There is only one more thing to know about concatenation, but it is an easy one. In this example, you just have to remember what += does. s += "2" means the same thing as s = s + "2".

```
4: String s = "1";           // s currently holds "1"
5: s += "2";                 // s currently holds "12"
6: s += 3;                   // s currently holds "123"
7: System.out.println(s);    // 123
```

On line 5, we are “adding” two strings, which means we concatenate them. Line 6 tries to trick you by adding a number, but it’s just like we wrote `s = s + 3`. We know that a string “plus” anything else means to use concatenation.

To review the rules one more time: use numeric addition if two numbers are involved, use concatenation otherwise, and evaluate from left to right. Have you memorized these three rules yet? Be sure to do so before the exam!

Immutability

Once a `String` object is created, it is not allowed to change. It cannot be made larger or smaller, and you cannot change one of the characters inside it.

You can think of a string as a storage box you have perfectly full and whose sides can’t bulge. There’s no way to add objects, nor can you replace objects without disturbing the entire arrangement. The trade-off for the optimal packing is zero flexibility.

Mutable is another word for changeable. *Immutable* is the opposite—an object that can’t be changed once it’s created. On the OCA exam, you need to know that `String` is immutable.

More on Immutability

You won’t be asked to identify whether custom classes are immutable on the exam, but it’s helpful to see an example. Consider the following code:

```
class Mutable {
    private String s;
    public void setS(String newS){ s = newS; } // Setter makes it mutable
    public String getS() { return s; }
}
final class Immutable {
    private String s = "name";
    public String getS() { return s; }
}
```

`Immutable` only has a getter. There’s no way to change the value of `s` once it’s set. `Mutable` has a setter as well. This allows the reference `s` to change to point to a different `String` later. Note that even though the `String` class is immutable, it can still be used in a mutable class. You can even make the instance variable `final` so the compiler reminds you if you accidentally change `s`.

Also, immutable classes in Java are `final`, and subclasses can’t add mutable behavior.

You learned that `+` is used to do `String` concatenation in Java. There's another way, which isn't used much on real projects but is great for tricking people on the exam. What does this print out?

```
String s1 = "1";
String s2 = s1.concat("2");
s2.concat("3");
System.out.println(s2);
```

Did you say "12"? Good. The trick is to see if you forget that the `String` class is immutable by throwing a method at you.

The String Pool

Since strings are everywhere in Java, they use up a lot of memory. In some production applications, they can use up 25–40 percent of the memory in the entire program. Java realizes that many strings repeat in the program and solves this issue by reusing common ones. **The *string pool*, also known as the intern pool, is a location in the Java virtual machine (JVM) that collects all these strings.**

The string pool contains literal values that appear in your program. For example, "name" is a literal and therefore goes into the string pool. `myObject.toString()` is a string but not a literal, so it does not go into the string pool. **Strings not in the string pool are garbage collected just like any other object.**

Remember back when we said these two lines are subtly different?

```
String name = "Fluffy";
String name = new String("Fluffy");
```

The former says to use the string pool normally. The second says "No, JVM. I really don't want you to use the string pool. Please create a new object for me even though it is less efficient." When you write programs, you wouldn't want to do this. For the exam, you need to know that it is allowed.

Important String Methods

The `String` class has dozens of methods. Luckily, you need to know only a handful for the exam. The exam creators pick most of the methods developers use in the real world.

For all these methods, you need to remember that a string is a sequence of characters and Java counts from 0 when indexed. Figure 3.1 shows how each character in the string "animals" is indexed.

FIGURE 3.1 Indexing for a string

a	n	i	m	a	l	s
0	1	2	3	4	5	6

Let's look at thirteen methods from the `String` class. Many of them are straightforward so we won't discuss them at length. You need to know how to use these methods.

length()

The method `length()` returns the number of characters in the `String`. The method signature is as follows:

```
int length()
```

The following code shows how to use `length()`:

```
String string = "animals";  
System.out.println(string.length()); // 7
```

Wait. It outputs 7? Didn't we just tell you that Java counts from 0? The difference is that zero counting happens only when you're using indexes or positions within a list. When determining the total size or length, Java uses normal counting again.

charAt()

The method `charAt()` lets you query the string to find out what character is at a specific index. The method signature is as follows:

```
char charAt(int index)
```

The following code shows how to use `charAt()`:

```
String string = "animals";  
System.out.println(string.charAt(0)); // a  
System.out.println(string.charAt(6)); // s  
System.out.println(string.charAt(7)); // throws exception
```

Since indexes start counting with 0, `charAt(0)` returns the “first” character in the sequence. Similarly, `charAt(6)` returns the “seventh” character in the sequence. `charAt(7)` is a problem. It asks for the “eighth” character in the sequence, but there are only seven characters present. When something goes wrong that Java doesn't know how to deal with, it throws an exception, as shown here. You'll learn more about exceptions in Chapter 6, “Exceptions.”

```
java.lang.StringIndexOutOfBoundsException: String index out of range: 7
```

indexOf()

The method `indexOf()` looks at the characters in the string and finds the first index that matches the desired value. `indexOf` can work with an individual character or a whole `String` as input. It can also start from a requested position. The method signatures are as follows:


```
int indexOf(char ch)
int indexOf(char ch, index fromIndex)
int indexOf(String str)
int indexOf(String str, index fromIndex)
```

The following code shows how to use `indexOf()`:

```
String string = "animals";
System.out.println(string.indexOf('a'));           // 0
System.out.println(string.indexOf("al"));          // 4
System.out.println(string.indexOf('a', 4));        // 4
System.out.println(string.indexOf("al", 5));       // -1
```

Since indexes begin with 0, the first 'a' matches at that position. The second statement looks for a more specific string and so matches later on. The third statement says Java shouldn't even look at the characters until it gets to index 4. The final statement doesn't find anything because it starts looking after the match occurred. Unlike `charAt()`, the `indexOf()` method doesn't throw an exception if it can't find a match. `indexOf()` returns -1 when no match is found. Because indexes start with 0, the caller knows that -1 couldn't be a valid index. This makes it a common value for a method to signify to the caller that no match is found.

substring()

The method `substring()` also looks for characters in a string. It returns parts of the string. The first parameter is the index to start with for the returned string. As usual, this is a zero-based index. There is an optional second parameter, which is the end index you want to stop at.

Notice we said “stop at” rather than “include.” **This means the `endIndex` parameter is allowed to be 1 past the end of the sequence if you want to stop at the end of the sequence.** That would be redundant, though, since you could omit the second parameter entirely in that case. In your own code, you want to avoid this redundancy. Don't be surprised if the exam uses it though. The method signatures are as follows:

```
int substring(int beginIndex)
int substring(int beginIndex, int endIndex)
```

The following code shows how to use `substring()`:

```
String string = "animals";
System.out.println(string.substring(3)); // mals
System.out.println(string.substring(string.indexOf('m'))); // mals
System.out.println(string.substring(3, 4)); // m
System.out.println(string.substring(3, 7)); // mals
```

The `substring()` method is the trickiest `String` method on the exam. The first example says to take the characters starting with index 3 through the end, which gives us "mals". The second example does the same thing: it calls `indexOf()` to get the index rather than hard-coding it. This is a common practice when coding because you may not know the index in advance.

The third example says to take the characters starting with index 3 until, but not including, the character at index 4—which is a complicated way of saying we want a `String` with one character: the one at index 3. This results in "m". The final example says to take the characters starting with index 3 until we get to index 7. Since index 7 is the same as the end of the string, it is equivalent to the first example.

We hope that wasn't too confusing. The next examples are less obvious:

```
System.out.println(string.substring(3, 3)); // empty string
System.out.println(string.substring(3, 2)); // throws exception
System.out.println(string.substring(3, 8)); // throws exception
```

The first example in this set prints an empty string. The request is for the characters starting with index 3 until you get to index 3. Since we start and end with the same index, there are *no* characters in between. The second example in this set throws an exception because the indexes can be backward. Java knows perfectly well that it will never get to index 2 if it starts with index 3. The third example says to continue until the eighth character. There is no eighth position, so Java throws an exception. Granted, there is no seventh character either, but at least there is the "end of string" invisible position.

Let's review this one more time since `substring()` is so tricky. The method returns the string starting from the requested index. If an end index is requested, it stops right before that index. Otherwise, it goes to the end of the string.

toLowerCase()* and *toUpperCase()

Whew. After that mental exercise, it is nice to have methods that do exactly what they sound like! These methods make it easy to convert your data. The method signatures are as follows:

```
String toLowerCase(String str)
String toUpperCase(String str)
```

The following code shows how to use these methods:

```
String string = "animals";
System.out.println(string.toUpperCase()); // ANIMALS
System.out.println("Abc123".toLowerCase()); // abc123
```

These methods do what they say. `toUpperCase()` converts any lowercase characters to uppercase in the returned string. `toLowerCase()` converts any uppercase characters to lowercase in the returned string. These methods leave alone any characters other than letters. Also, remember that strings are immutable, so the original string stays the same.

equals()* and *equalsIgnoreCase()

The `equals()` method checks whether two `String` objects contain exactly the same characters in the same order. The `equalsIgnoreCase()` method checks whether two `String` objects contain the same characters with the exception that it will convert the characters' case if needed. The method signatures are as follows:

```
boolean equals(String str)
boolean equalsIgnoreCase(String str)
```

The following code shows how to use these methods:

```
System.out.println("abc".equals("ABC")); // false
System.out.println("ABC".equals("ABC")); // true
System.out.println("abc".equalsIgnoreCase("ABC")); // true
```

This example should be fairly intuitive. In the first example, the values aren't exactly the same. In the second, they are exactly the same. In the third, they differ only by case, but it is okay because we called the method that ignores differences in case.

startsWith()* and *endsWith()

The `startsWith()` and `endsWith()` methods look at whether the provided value matches part of the `String`. The method signatures are as follows:

```
boolean startsWith(String prefix)
boolean endsWith(String suffix)
```

The following code shows how to use these methods:

```
System.out.println("abc".startsWith("a")); // true
System.out.println("abc".startsWith("A")); // false
System.out.println("abc".endsWith("c")); // true
System.out.println("abc".endsWith("a")); // false
```

Again, nothing surprising here. Java is doing a case-sensitive check on the values provided.

contains()

The `contains()` method also looks for matches in the `String`. It isn't as particular as `startsWith()` and `endsWith()`—the match can be anywhere in the `String`. The method signature is as follows:

```
boolean contains(String str)
```

The following code shows how to use these methods:

```
System.out.println("abc".contains("b")); // true
System.out.println("abc".contains("B")); // false
```

Again, we have a case-sensitive search in the `String`. The `contains()` method is a convenience method so you don't have to write `str.indexOf(otherString) != -1`.

replace()

The `replace()` method does a simple search and replace on the string. There's a version that takes `char` parameters as well as a version that takes `CharSequence` parameters. A `CharSequence` is a general way of representing several classes, including `String` and `StringBuilder`. It's called an interface, which we'll cover in Chapter 5, "Class Design." The method signatures are as follows:

```
String replace(char oldChar, char newChar)
String replace(CharSequence oldChar, CharSequence newChar)
```

The following code shows how to use these methods:

```
System.out.println("abcabc".replace('a', 'A')); // AbcAbc
System.out.println("abcabc".replace("a", "A")); // AbcAbc
```

The first example uses the first method signature, passing in `char` parameters. The second example uses the second method signature, passing in `String` parameters.

trim()

You've made it through all the `String` methods you need to know except one. We left the easy one for last. The `trim()` method removes whitespace from the beginning and end of a `String`. In terms of the exam, whitespace consists of spaces along with the `\t` (tab) and `\n` (newline) characters. Other characters, such as `\r` (carriage return), are also included in what gets trimmed. The method signature is as follows:

```
public String trim()
```

The following code shows how to use this method:

```
System.out.println("abc".trim());           // abc
System.out.println("\t a b c\n".trim()); // a b c
```

The first example prints the original string because there are no whitespace characters at the beginning or end. The second example gets rid of the leading tab, subsequent spaces, and the trailing newline. It leaves the spaces that are in the middle of the string.

Method Chaining

It is common to call multiple methods on the same `String`, as shown here:

```
String start = "AniMaL ";
String trimmed = start.trim();           // "AniMaL"
```

```
String lowercase = trimmed.toLowerCase();    // "animal"
String result = lowercase.replace('a', 'A');  // "Animal"
System.out.println(result);
```

This is just a series of `String` methods. Each time one is called, the returned value is put in a new variable. There are four `String` values along the way, and `Animal` is output.

However, on the exam there is a tendency to cram as much code as possible into a small space. You'll see code using a technique called method chaining. Here's an example:

```
String result = "AniMaL  ".trim().toLowerCase().replace('a', 'A');
System.out.println(result);
```

This code is equivalent to the previous example. It also creates four `String` objects and outputs `Animal`. To read code that uses method chaining, start at the left and evaluate the first method. Then call the next method on the returned value of the first method. Keep going until you get to the semicolon.

Remember that `String` is immutable. What do you think the result of this code is?

```
5: String a = "abc";
6: String b = a.toUpperCase();
7: b = b.replace("B", "2").replace('C', '3');
8: System.out.println("a=" + a);
9: System.out.println("b=" + b);
```

On line 5, we set `a` to point to `"abc"` and never pointed `a` to anything else. Since we are dealing with an immutable object, none of the code on lines 6 or 7 changes `a`.

`b` is a little trickier. Line 6 has `b` pointing to `"ABC"`, which is straightforward. On line 7, we have method chaining. First, `"ABC".replace("B", "2")` is called. This returns `"A2C"`. Next, `"A2C".replace('C', '3')` is called. This returns `"A23"`. Finally, `b` changes to point to this returned `String`. When line 9 executes, `b` is `"A23"`.

Using the *StringBuilder* Class

A small program can create a lot of `String` objects very quickly. For example, how many do you think this piece of code creates?

```
10: String alpha = "";
11: for(char current = 'a'; current <= 'z'; current++)
12:   alpha += current;
13: System.out.println(alpha);
```

The empty `String` on line 10 is instantiated, and then line 12 appends an `"a"`. However, because the `String` object is immutable, a new `String` object is assigned to `alpha` and the

“” object becomes eligible for garbage collection. The next time through the loop, *alpha* is assigned a new String object, "ab", and the "a" object becomes eligible for garbage collection. The next iteration assigns *alpha* to "abc" and the "ab" object becomes eligible for garbage collection, and so on.

This sequence of events continues, and after 26 iterations through the loop, a total of 27 objects are instantiated, most of which are immediately eligible for garbage collection.

This is very inefficient. Luckily, Java has a solution. The `StringBuilder` class creates a String without storing all those interim String values. Unlike the String class, `StringBuilder` is not immutable.

```
15: StringBuilder alpha = new StringBuilder();
16: for(char current = 'a'; current <= 'z'; current++)
17:   alpha.append(current);
18: System.out.println(alpha);
```

On line 15, a new `StringBuilder` object is instantiated. The call to `append()` on line 17 adds a character to the `StringBuilder` object each time through the for loop and appends the value of `current` to the end of *alpha*. This code reuses the same `StringBuilder` without creating an interim String each time.

In this section, we'll look at creating a `StringBuilder`, common methods, and a comparison to `StringBuffer`.

Mutability and Chaining

We're sure you noticed this from the previous example, but `StringBuilder` is not immutable. In fact, we gave it 27 different values in the example (blank plus adding each letter in the alphabet). The exam will likely try to trick you with respect to String and `StringBuilder` being mutable.

Chaining makes this even more interesting. When we chained String method calls, the result was a new String with the answer. Chaining `StringBuilder` objects doesn't work this way. Instead, the `StringBuilder` changes its own state and returns a reference to itself! Let's look at an example to make this clearer:

```
4: StringBuilder sb = new StringBuilder("start");
5: sb.append("+middle");           // sb = "start+middle"
6: StringBuilder same = sb.append("+end"); // "start+middle+end"
```

Line 5 adds text to the end of *sb*. It also returns a reference to *sb*, which is ignored. Line 6 also adds text to the end of *sb* and returns a reference to *sb*. This time the reference is stored in *same*—which means *sb* and *same* point to the exact same object and would print out the same value.

The exam won't always make the code easy to read by only having one method per line. What do you think this example prints?

```
4: StringBuilder a = new StringBuilder("abc");
5: StringBuilder b = a.append("de");
```

```
6: b = b.append("f").append("g");
7: System.out.println("a=" + a);
8: System.out.println("b=" + b);
```

Did you say both print "abcdefg"? Good. There's only one *StringBuilder* object here. We know that because `new StringBuilder()` was called only once. On line 5, there are two variables referring to that object, which has a value of "abcde". On line 6, those two variables are still referring to that same object, which now has a value of "abcdefg". Incidentally, the assignment back to *b* does absolutely nothing. *b* is already pointing to that *StringBuilder*.

Creating a *StringBuilder*

There are three ways to construct a *StringBuilder*:

```
StringBuilder sb1 = new StringBuilder();
StringBuilder sb2 = new StringBuilder("animal");
StringBuilder sb3 = new StringBuilder(10);
```

The first says to create a *StringBuilder* containing an empty sequence of characters and assign *sb1* to point to it. The second says to create a *StringBuilder* containing a specific value and assign *sb2* to point to it. For the first two, it tells Java to manage the implementation details. The final example tells Java that we have some idea of how big the eventual value will be and would like the *StringBuilder* to reserve a certain number of slots for characters.

Size vs. Capacity

The behind-the-scenes process of how objects are stored isn't on the exam, but some knowledge of this process may help you better understand and remember *StringBuilder*.

Size is the number of characters currently in the sequence, and capacity is the number of characters the sequence can currently hold. Since a *String* is immutable, the size and capacity are the same. The number of characters appearing in the *String* is both the size and capacity.

For *StringBuilder*, Java knows the size is likely to change as the object is used. When *StringBuilder* is constructed, it may start at the default capacity (which happens to be 16) or one of the programmer's choosing. In the example, we request a capacity of 5. At this point, the size is 0 since no characters have been added yet, but we have space for 5.

continues

(continued)

Next we add four characters. At this point, the size is 4 since four slots are taken. The capacity is still 5. Then we add three more characters. The size is now 7 since we have used up seven slots. Because the capacity wasn't large enough to store seven characters, Java automatically increased it for us.

```
StringBuilder sb = new StringBuilder(5);
```

0	1	2	3	4

```
sb.append("anim");
```

a	n	i	m	
0	1	2	3	4

```
sb.append("als");
```

a	n	i	m	a	l	s		
0	1	2	3	4	5	6	7	...

Important *StringBuilder* Methods

As with *String*, we aren't going to cover every single method in the *StringBuilder* class. These are the ones you might see on the exam.

charAt()*, *indexOf()*, *length()*, and *substring()

These four methods work exactly the same as in the *String* class. Be sure you can identify the output of this example:

```
StringBuilder sb = new StringBuilder("animals");
String sub = sb.substring(sb.indexOf("a"), sb.indexOf("al"));
int len = sb.length();
char ch = sb.charAt(6);
System.out.println(sub + " " + len + " " + ch);
```

The correct answer is `anim 7 s`. The `indexOf()` method calls return 0 and 4, respectively. `substring()` returns the *String* starting with index 0 and ending right before index 4.

`length()` returns 7 because it is the number of characters in the *StringBuilder* rather than an index. Finally, `charAt()` returns the character at index 6. Here we do start with 0 because we are referring to indexes. If any of this doesn't sound familiar, go back and read the section on *String* again.

Notice that `substring()` returns a `String` rather than a `StringBuilder`. That is why `sb` is not changed. `substring()` is really just a method that inquires about where the substring happens to be.

append()

The `append()` method is by far the most frequently used method in `StringBuilder`. In fact, it is so frequently used that we just started using it without comment. Luckily, this method does just what it sounds like: it adds the parameter to the `StringBuilder` and returns a reference to the current `StringBuilder`. One of the method signatures is as follows:

```
StringBuilder append(String str)
```

Notice that we said *one* of the method signatures. **There are more than 10 method signatures that look similar but that take different data types as parameters.** All those methods are provided so you can write code like this:

```
StringBuilder sb = new StringBuilder().append(1).append('c');
sb.append("-").append(true);
System.out.println(sb);      // 1c-true
```

Nice method chaining, isn't it? `append()` is called directly after the constructor. By having all these method signatures, you can just call `append()` without having to convert your parameter to a `String` first.

insert()

The `insert()` method adds characters to the `StringBuilder` at the requested index and returns a reference to the current `StringBuilder`. Just like `append()`, there are lots of method signatures for different types. Here's one:

```
StringBuilder insert(int offset, String str)
```

Pay attention to the offset in these examples. It is the index where we want to insert the requested parameter.

```
3: StringBuilder sb = new StringBuilder("animals");
4: sb.insert(7, "-");           // sb = animals-
5: sb.insert(0, "-");          // sb = -animals-
6: sb.insert(4, "-");          // sb = -ani-mals
7: System.out.println(sb);
```

Line 4 says to insert a dash at index 7, which happens to be the end of sequence of characters. Line 5 says to insert a dash at index 0, which happens to be the very beginning. Finally, line 6 says to insert a dash right before index 4. The exam creators will try to trip you up on this. As we add and remove characters, their indexes change. When you see a question dealing with such operations, draw what is going on so you won't be confused.

delete()* and *deleteCharAt()

The `delete()` method is the opposite of the `insert()` method. It removes characters from the sequence and returns a reference to the current `StringBuilder`. The `deleteCharAt()` method is convenient when you want to delete only one character. The method signatures are as follows:

```
StringBuilder delete(int start, int end)
StringBuilder deleteCharAt(int index)
```

The following code shows how to use these methods:

```
StringBuilder sb = new StringBuilder("abcdef");
sb.delete(1, 3);                // sb = adef
sb.deleteCharAt(5);             // throws an exception
```

First, we delete the characters starting with index 1 and ending right before index 3. This gives us `adef`. Next, we ask Java to delete the character at position 5. However, the remaining value is only four characters long, so it throws a `StringIndexOutOfBoundsException`.

reverse()

After all that, it's time for a nice, easy method. The `reverse()` method does just what it sounds like: it reverses the characters in the sequences and returns a reference to the current `StringBuilder`. The method signature is as follows:

```
StringBuilder reverse()
```

The following code shows how to use this method:

```
StringBuilder sb = new StringBuilder("ABC");
sb.reverse();
System.out.println(sb);
```

As expected, this prints `CBA`. This method isn't that interesting. Maybe the exam creators like to include it to encourage you to write down the value rather than relying on memory for indexes.

toString()

The last method converts a `StringBuilder` into a `String`. The method signature is as follows:

```
String toString()
```

The following code shows how to use this method:

```
String s = sb.toString();
```

Often `StringBuilder` is used internally for performance purposes but the end result needs to be a `String`. For example, maybe it needs to be passed to another method that is expecting a `String`.

StringBuilder vs. StringBuffer

When writing new code that concatenates a lot of `String` objects together, you should use `StringBuilder`. `StringBuilder` was added to Java in Java 5. If you come across older code, you will see `StringBuffer` used for this purpose. `StringBuffer` does the same thing but more slowly because it is thread safe. You'll learn about threads for the OCP exam. In theory, you don't need to know about `StringBuffer` on the exam at all. However, we bring this up anyway, since an older question might still be left on the exam.

Understanding Equality

In Chapter 2, you learned how to use `==` to compare numbers and that object references refer to the same object.

```
StringBuilder one = new StringBuilder();
StringBuilder two = new StringBuilder();
StringBuilder three = one.append("a");
System.out.println(one == two); // false
System.out.println(one == three); // true
```

Since this example isn't dealing with primitives, we know to look for whether the references are referring to the same object. *one* and *two* are both completely separate `StringBuilders`, giving us two objects. Therefore, the first print statement gives us `false`. *three* is more interesting. Remember how `StringBuilder` methods like to return the current reference for chaining? This means *one* and *three* both point to the same object and the second print statement gives us `true`.

Let's now visit the more complex and confusing scenario, `String` equality, made so in part because of the way the JVM reuses `String` literals:

```
String x = "Hello World";
String y = "Hello World";
System.out.println(x == y);    // true
```

Remember that `Strings` are immutable and literals are pooled. The JVM created only one literal in memory. *x* and *y* both point to the same location in memory; therefore, the statement outputs `true`. It gets even trickier. Consider this code:

```
String x = "Hello World";
String z = " Hello World".trim();
System.out.println(x == z);    // false
```

In this example, we don't have two of the same `String` literal. Although `x` and `z` happen to evaluate to the same string, one is computed at runtime. Since it isn't the same at compile-time, a new `String` object is created.

You can even force the issue by creating a new `String`:

```
String x = new String("Hello World");
String y = "Hello World";
System.out.println(x == y); // false
```

Since you have specifically requested a different `String` object, the pooled value isn't shared.



The lesson is to never use `==` to compare `String` objects. The only time you should have to deal with `==` for `Strings` is on the exam.

You saw earlier that you can say you want logical equality rather than object equality for `String` objects:

```
String x = "Hello World";
String z = " Hello World".trim();
System.out.println(x.equals(z)); // true
```

This works because the authors of the `String` class implemented a standard method called `equals` to check the values inside the `String` rather than the `String` itself. If a class doesn't have an `equals` method, Java determines whether the references point to the same object—which is exactly what `==` does. In case you are wondering, the authors of `StringBuilder` did not implement `equals()`. If you call `equals()` on two `StringBuilder` instances, it will check reference equality.

The exam will test you on your understanding of equality with objects they define too. For example:

```
1: public class Tiger {
2:     String name;
3:     public static void main(String[] args) {
4:         Tiger t1 = new Tiger();
5:         Tiger t2 = new Tiger();
6:         Tiger t3 = t1;
7:         System.out.println(t1 == t1); // true
8:         System.out.println(t1 == t2); // false
9:         System.out.println(t1.equals(t2)); // false
10:    } }
```

The first two statements check object reference equality. Line 7 prints `true` because we are comparing references to the same object. Line 8 prints `false` because the two object

references are different. Line 9 prints `false` since `Tiger` does not implement `equals()`. Don't worry—you aren't expected to know how to implement `equals()` for the OCA exam.

Understanding Java Arrays

Up to now, we've been referring to the `String` and `StringBuilder` classes as a “sequence of characters.” This is true. They are implemented using an *array* of characters. **An array is an area of memory on the heap with space for a designated number of elements.** A `String` is implemented as an array with some methods that you might want to use when dealing with characters specifically. A `StringBuilder` is implemented as an array where the array object is replaced with a new bigger array object when it runs out of space to store all the characters. A big difference is that an array can be of any other Java type. If we didn't want to use a `String` for some reason, we could use an array of `char` primitives directly:

```
char[] letters;
```

This wouldn't be very convenient because we'd lose all the special properties `String` gives us, such as writing “Java”. Keep in mind that `letters` is a reference variable and not a primitive. `char` is a primitive. But `char` is what goes into the array and not the type of the array itself. The array itself is of type `char[]`. You can mentally read the brackets `[]` as “array.”

In other words, an array is an ordered list. It can contain duplicates. You will learn about data structures that cannot contain duplicates for the OCP exam. In this section, we'll look at creating an array of primitives and objects, sorting, searching, `varargs`, and multidimensional arrays.


Creating an Array of Primitives

The most common way to create an array looks like this:

```
int[] numbers1 = new int[3];
```


The basic parts are shown in Figure 3.2. It specifies the type of the array (`int`) and the size (3). The brackets tell you this is an array.

FIGURE 3.2 **The basic structure of an array**



When using this form to instantiate an array, set all the elements to the default value for that type. As you learned in Chapter 1, the default value of an `int` is 0. Since `numbers1` is a reference variable, it points to the array object, as shown in Figure 3.3. As you can see, the default value for all the elements is 0. Also, the indexes start with 0 and count up, just as they did for a `String`.

FIGURE 3.3 An empty array




Another way to create an array is to specify all the elements it should start out with:

```
int[] numbers2 = new int[] {42, 55, 99};
```

In this example, we also create an `int` array of size 3. This time, we specify the initial values of those three elements instead of using the defaults. Figure 3.4 shows what this array looks like.

FIGURE 3.4 An initialized array



Java recognizes that this expression is redundant. Since you are specifying the type of the array on the left side of the equal sign, Java already knows the type. And since you are specifying the initial values, it already knows the size. As a shortcut, Java lets you write this:

```
int[] numbers2 = {42, 55, 99};
```

This approach is called an anonymous array. It is anonymous because you don't specify the type and size.

Finally, you can type the `[]` before or after the name, and adding a space is optional. This means that all four of these statements do the exact same thing:

```
int[] numAnimals;  
int [] numAnimals2;  
int numAnimals3[];  
int numAnimals4 [];
```

Most people use the first one. You could see any of these on the exam, though, so get used to seeing the brackets in odd places.

Multiple “Arrays” in Declarations

What types of reference variables do you think the following code creates?

```
int[] ids, types;
```

The correct answer is two variables of type `int[]`. This seems logical enough. After all, `int a, b;` created two `int` variables. What about this example?

```
int ids[], types;
```

All we did was move the brackets, but it changed the behavior. This time we get one variable of type `int[]` and one variable of type `int`. Java sees this line of code and thinks something like this: “They want two variables of type `int`. The first one is called `ids[]`. This one is a `int[]` called `ids`. The second one is just called `types`. No brackets, so it is a regular integer.”

Needless to say, you shouldn’t write code that looks like this. But you do still need to understand it for the exam.

Creating an Array with Reference Variables

You can choose any Java type to be the type of the array. This includes classes you create yourself. Let’s take a look at a built-in type with `String`:

```
public class ArrayType {  
    public static void main(String args[]) {  
        String [] bugs = { "cricket", "beetle", "ladybug" };  
        String [] alias = bugs;  
        System.out.println(bugs.equals(alias));    // true  
        System.out.println(bugs.toString()); // [Ljava.lang.String;@160bc7c0  
    } }  
}
```

We can call `equals()` because an array is an object. It returns `true` because of reference equality. The `equals()` method on arrays does not look at the elements of the array. Remember, this would work even on an `int[]` too. `int` is a primitive; `int[]` is an object.


The second print statement is even more interesting. What on earth is `[Ljava.lang.String;@160bc7c0`? You don’t have to know this for the exam, but `[L` means it is an array, `java.lang.String` is the reference type, and `160bc7c0` is the hash code.



Since Java 5, Java has provided a method that prints an array nicely: `java.util.Arrays.toString(bugs)` would print `[cricket, beetle, ladybug]`. The exam tends not to use it because most of the questions on arrays were written a long time ago. Regardless, this is a useful method when testing your own code.

Make sure you understand Figure 3.5. The array does not allocate space for the `String` objects. Instead, it allocates space for a reference to where the objects are really stored.

FIGURE 3.5 An array pointing to strings



As a quick review, what do you think this array points to?

```
class Names {
    String names[];
}
```

You got us. It was a review of Chapter 1 and not our discussion on arrays. The answer is `null`. The code never instantiated the array so it is just a reference variable to `null`. Let's try that again—what do you think this array points to?

```
class Names {
    String names[] = new String[2];
}
```

It is an array because it has brackets. It is an array of type `String` since that is the type mentioned in the declaration. It has two elements because the length is 2. Each of those two slots currently is `null`, but has the potential to point to a `String` object.

Remember casting from the previous chapter when you wanted to force a bigger type into a smaller type? You can do that with arrays too:

```
3: String[] strings = { "stringValue" };
4: Object[] objects = strings;
5: String[] againStrings = (String[]) objects;
6: againStrings[0] = new StringBuilder(); // DOES NOT COMPILE
7: objects[0] = new StringBuilder();      // careful!
```


Line 3 creates an array of type `String`. Line 4 doesn't require a cast because `Object` is a broader type than `String`. On line 5, a cast is needed because we are moving to a more specific type. Line 6 doesn't compile because a `String[]` only allows `String` objects and `StringBuilder` is not a `String`.

Line 7 is where this gets interesting. From the point of view of the compiler, this is just fine. A `StringBuilder` object can clearly go in an `Object[]`. The problem is that we don't actually have an `Object[]`. We have a `String[]` referred to from an `Object[]` variable. At runtime, the code throws an `ArrayStoreException`. You don't need to memorize the name of this exception, but you do need to know that the code will throw an exception.

Using an Array

Now that we know how to create an array, let's try accessing one:

```
4: String[] mammals = {"monkey", "chimp", "donkey"};
5: System.out.println(mammals.length);           // 3
6: System.out.println(mammals[0]);               // monkey
7: System.out.println(mammals[1]);               // chimp
8: System.out.println(mammals[2]);               // donkey
```

Line 4 declares and initializes the array. Line 5 tells us how many elements the array can hold. The rest of the code prints the array. Notice elements are indexed starting with 0. This should be familiar from `String` and `StringBuilder`, which also start counting with 0. Those classes also counted length as the number of elements.

To make sure you understand how length works, what do you think this prints?

```
String[] birds = new String[6];
System.out.println(birds.length);
```

The answer is 6. Even though all 6 elements of the array are `null`, there are still 6 of them. `length` does not consider what is in the array; it only considers how many slots have been allocated.

It is very common to use a loop when reading from or writing to an array. This loop sets each element of number to 5 higher than the current index:

```
5: int[] numbers = new int[10];
6: for (int i = 0; i < numbers.length; i++)
7:   numbers[i] = i + 5;
```

Line 5 simply instantiates an array with 10 slots. Line 6 is a `for` loop using an extremely common pattern. It starts at index 0, which is where an array begins as well. It keeps going, one at a time, until it hits the end of the array. Line 7 sets the current element of *numbers*.

The exam will test whether you are being observant by trying to access elements that are not in the array. Can you tell why each of these throws an `ArrayIndexOutOfBoundsException` for our array of size 10?

```
numbers[10] = 3;
numbers[numbers.length] = 5;
for (int i = 0; i <= numbers.length; i++) numbers[i] = i + 5;
```

The first one is trying to see if you know that indexes start with 0. Since we have 10 elements in our array, this means only `numbers[0]` through `numbers[9]` are valid. The second example assumes you are clever enough to know 10 is invalid and disguises it by using the `length` variable. However, the `length` is always one more than the maximum valid index. Finally, the `for` loop incorrectly uses `<=` instead of `<`, which is also a way of referring to that 10th element.

Sorting

Java makes it easy to sort an array by providing a sort method—or rather, a bunch of sort methods. Just like `StringBuilder` allowed you to pass almost anything to `append()`, you can pass almost any array to `Arrays.sort()`.

`Arrays` is the first class provided by Java we have used that requires an import. To use it, you must have either of the following two statements in your class:

```
import java.util.*           // import whole package including Arrays
import java.util.Arrays;     // import just Arrays
```

There is one exception, although it doesn't come up often on the exam. You can write `java.util.Arrays` every time it is used in the class instead of specifying it as an import.

Remember that if you are shown a code snippet with a line number that doesn't begin with 1, you can assume the necessary imports are there. Similarly, you can assume the imports are present if you are shown a snippet of a method.

This simple example sorts three numbers:

```
int[] numbers = { 6, 9, 1 };
Arrays.sort(numbers);
for (int i = 0; i < numbers.length; i++)
    System.out.print (numbers[i] + " ");
```

The result is 1 6 9, as you should expect it to be. Notice that we had to loop through the output to print the values in the array. Just printing the array variable directly would give the annoying hash of `[I@2bd9c3e7`.

Try this again with `String` types:

```
String[] strings = { "10", "9", "100" };
Arrays.sort(strings);
for (String string : strings)
    System.out.print(string + " ");
```

This time the result might not be what you expect. This code outputs 10 100 9. The problem is that String sorts in alphabetic order, and 1 sorts before 9. (Numbers sort before letters and uppercase sorts before lowercase, in case you were wondering.) For the OCP exam, you’ll learn how to create custom sort orders using something called a comparator.

Did you notice we snuck in the enhanced for loop in this example? Since we aren’t using the index, we don’t need the traditional for loop. That won’t stop the exam creators from using it, though, so we’ll be sure to use both to keep you sharp!

Searching

Java also provides a convenient way to search—but only if the array is already sorted. Table 3.1 covers the rules for binary search.

TABLE 3.1 Binary search rules

Scenario	Result
Target element found in sorted array	Index of match
Target element not found in sorted array	Negative value showing one smaller than the negative of index, where a match needs to be inserted to preserve sorted order
Unsorted array	A surprise—this result isn’t predictable

Let’s try out these rules with an example:

```
3: int[] numbers = {2,4,6,8};
4: System.out.println(Arrays.binarySearch(numbers, 2)); // 0
5: System.out.println(Arrays.binarySearch(numbers, 4)); // 1
6: System.out.println(Arrays.binarySearch(numbers, 1)); // -1
7: System.out.println(Arrays.binarySearch(numbers, 3)); // -2
8: System.out.println(Arrays.binarySearch(numbers, 9)); // -5
```

Take note of the fact that line 3 is a sorted array. If it weren’t, we couldn’t apply either of the other rules. Line 4 searches for the index of 2. The answer is index 0. Line 5 searches for the index of 4, which is 1.

Line 5 searches for the index of 1. Although 1 isn’t in the list, the search can determine that it should be inserted at element 0 to preserve the sorted order. Since 0 already means something for array indexes, Java needs to subtract 1 to give us the answer of -1. Line 7 is similar. Although 3 isn’t in the list, it would need to be inserted at element 1 to preserve the sorted order. We negate and subtract 1 for consistency, getting -1 -1, also known as -2. Finally, line 8 wants to tell us that 9 should be inserted at index 4. We again negate and subtract 1, getting -4 -1, also known as -5.

What do you think happens in this example?

```
5: int numbers = new int[] {3,2,1};
6: System.out.println(Arrays.binarySearch(numbers, 2));
7: System.out.println(Arrays.binarySearch(numbers, 3));
```

Note that on line 5, the array isn't sorted. This means the output will not be predictable. When testing this example, line 6 correctly gave 1 as the output. However, line 3 gave the wrong answer. The exam creators will not expect you to know what incorrect values come out. As soon as you see the array isn't sorted, look for an answer choice about unpredictable output.

On the exam, you need to know what a binary search returns in various scenarios. Oddly, you don't need to know why "binary" is in the name. In case you are curious, a binary search splits the array into two equal pieces (remember 2 is binary) and determines which half the target it is. It repeats this process until only one element is left.

Varargs

When creating an array yourself, it looks like what we've seen thus far. When one is passed to your method, there is another way it can look. Here are three examples with a `main()` method:

```
public static void main(String[] args)
public static void main(String args[])
public static void main(String... args) // varargs
```

The third example uses a syntax called `varargs` (variable arguments), which you saw in Chapter 1. You'll learn how to call a method using `varargs` in Chapter 4, "Methods and Encapsulation." For now, all you need to know is that you can use a variable defined using `varargs` as if it were a normal array. For example `args.length` and `args[0]` are legal.

Multidimensional Arrays

Arrays are objects, and of course array components can be objects. It doesn't take much time, rubbing those two facts together, to wonder if arrays can hold other arrays, and of course they can.

Creating a Multidimensional Array

Multiple array separators are all it takes to declare arrays with multiple dimensions. You can locate them with the type or variable name in the declaration, just as before:

```
int[][] vars1;           // 2D array
int vars2 [][];          // 2D array
int[] vars3[];           // 2D array
int[] vars4 [], space [][]; // a 2D AND a 3D array
```

The first two examples are nothing surprising and declare a two-dimensional (2D) array. The third example also declares a 2D array. There's no good reason to use this style other than to confuse readers of your code. The final example declares two arrays on the same line. Adding up the brackets, we see that the `vars4` is a 2D array and `space` is a 3D array. Again, there's no reason to use this style other than to confuse readers of your code. The exam creators like to try to confuse you, though. Luckily you are on to them and won't let this happen to you!

You can specify the size of your multidimensional array in the declaration if you like:

```
String [][] rectangle = new String[3][2];
```


The result of this statement is an array `rectangle` with three elements, each of which refers to an array of two elements. You can think of the addressable range as `[0][0]` through `[2][1]`, but don't think of it as a structure of addresses like `[0,0]` or `[2,1]`.

Now suppose we set one of these values:

```
rectangle[0][1] = "set";
```

You can visualize the result as shown in Figure 3.6. This array is sparsely populated because it has a lot of null values. You can see that `rectangle` still points to an array of three elements and that we have three arrays of two elements. You can also follow the trail from reference to the one value pointing to a `String`. First you start at index 0 in the top array. Then you go to index 1 in the next array.


FIGURE 3.6 A sparsely populated multidimensional array



While that array happens to be rectangular in shape, an array doesn't need to be. Consider this one:

```
int[][] differentSize = {{1, 4}, {3}, {9,8,7}};
```

We still start with an array of three elements. However, this time the elements in the next level are all different sizes. One is of length 2, the next length 1, and the last length 3 (see Figure 3.7). This time the array is of primitives, so they are shown as if they are in the array themselves.

FIGURE 3.7 An asymmetric multidimensional array

Another way to create an asymmetric array is to initialize just an array's first dimension, and define the size of each array component in a separate statement:

```
int [][] args = new int[4][];
args[0] = new int[5];
args[1] = new int[3];
```

This technique reveals what you really get with Java: arrays of arrays that, properly managed, offer a multidimensional effect.

Using a Multidimensional Array

The most common operation on a multidimensional array is to loop through it. This example prints out a 2D array:

```
int[][] twoD = new int[3][2];
for (int i = 0; i < twoD.length; i++) {
    for (int j = 0; j < twoD[i].length; j++)
        System.out.print(twoD[i][j] + " "); // print element
    System.out.println();                    // time for a new row
}
```

We have two loops here. The first uses index *i* and goes through the first subarray for *twoD*. The second uses a different loop variable *j*. It is very important these be different variable names so the loops don't get mixed up. The inner loop looks at how many elements are in the second-level array. The inner loop prints the element and leaves a space for readability. When the inner loop completes, the outer loop goes to a new line and repeats the process for the next element.

This entire exercise would be easier to read with the enhanced for loop.

```
for (int[] inner : twoD) {
    for (int num : inner)
```

```
        System.out.print(num + " ");  
        System.out.println();  
    }
```

We'll grant you that it isn't fewer lines, but each line is less complex and there aren't any loop variables or terminating conditions to mix up.

Understanding an *ArrayList*

An array has one glaring shortcoming: you have to know how many elements will be in the array when you create it and then you are stuck with that choice. Just like a *StringBuilder*, *ArrayList* can change size at runtime as needed. Like an array, an *ArrayList* is an ordered sequence that allows duplicates.

As when we used *Arrays.sort*, *ArrayList* requires an import. To use it, you must have either of the following two statements in your class:

```
import java.util.*           // import whole package including ArrayList  
import java.util.ArrayList;  // import just ArrayList
```

Remember that if you are shown a code snippet with a line number that doesn't begin with 1, you can assume the necessary imports are there. Similarly, you can assume the imports are present if you are shown a snippet of a method.

In this section, we'll look at creating an *ArrayList*, common methods, autoboxing, conversion, and sorting.

Experienced programmers, take note: This section is simplified and doesn't cover a number of topics that are out of scope for the OCA exam.

Creating an *ArrayList*

As with *StringBuilder*, there are three ways to create an *ArrayList*:

```
ArrayList list1 = new ArrayList();  
ArrayList list2 = new ArrayList(10);  
ArrayList list3 = new ArrayList(list2);
```

The first says to create an *ArrayList* containing space for the default number of elements but not to fill any slots yet. The second says to create an *ArrayList* containing a specific number of slots, but again not to assign any. The final example tells Java that we want to make a copy of another *ArrayList*. We copy both the size and contents of that *ArrayList*. Granted, *list2* is empty in this example so it isn't particularly interesting.

Although these are the only three constructors you need to know, you do need to learn some variants of it. The previous examples were the old pre-Java 5 way of creating an *ArrayList*. They still work and you still need to know they work. You also need to know

the new and improved way. Java 5 introduced *generics*, which allow you to specify the type of class that the `ArrayList` will contain.

```
ArrayList<String> list4 = new ArrayList<String>();  
ArrayList<String> list5 = new ArrayList<>();
```

Java 5 allows you to tell the compiler what the type would be by specifying it between `<` and `>`. Starting in Java 7, you can even omit that type from the right side. The `<` and `>` are still required, though. This is called the diamond operator because `<>` looks like a diamond.

Just when you thought you knew everything about creating an `ArrayList`, there is one more thing you need to know. `ArrayList` implements an interface called `List`. In other words, an `ArrayList` is a `List`. You will learn about interfaces in Chapter 5. In the meantime, just know that you can store an `ArrayList` in a `List` reference variable but not vice versa. The reason is that `List` is an interface and interfaces can't be instantiated.

```
List<String> list6 = new ArrayList<>();  
ArrayList<String> list7 = new List<>(); // DOES NOT COMPILE
```

Using an *ArrayList*

`ArrayList` has many methods, but you only need to know a handful of them—even fewer than you did for `String` and `StringBuilder`.

Before reading any further, you are going to see something new in the method signatures: a “class” named `E`. Don't worry—it isn't really a class. `E` is used by convention in generics to mean “any class that this array can hold.” If you didn't specify a type when creating the `ArrayList`, `E` means `Object`. Otherwise, it means the class you put between `<` and `>`.

You should also know that `ArrayList` implements `toString()` so you can easily see the contents just by printing it. Arrays do not do produce such pretty output.

add()

The `add()` methods insert a new value in the `ArrayList`. The method signatures are as follows:

```
boolean add(E element)  
void add(int index, E element)
```

Don't worry about the boolean return value. It always returns `true`. It is there because other classes in the collections family need a return value in the signature when adding an element.

Since `add()` is the most critical `ArrayList` method you need to know for the exam, we are going to show a few sets of examples for it. Let's start with the most straightforward case:

```
ArrayList list = new ArrayList();  
list.add("hawk");           // [hawk]
```



```
list.add(Boolean.TRUE);    // [hawk, true]
System.out.println(list);  // [hawk, true]
```

`add()` does exactly what we expect: it stores the `String` in the no longer empty `ArrayList`. It then does the same thing for the `boolean`. This is okay because we didn't specify a type for `ArrayList`; therefore, the type is `Object`, which includes everything except primitives. It may not have been what we intended, but the compiler doesn't know that. Now, let's use generics to tell the compiler we only want to allow `String` objects in our `ArrayList`:

```
ArrayList<String> safer = new ArrayList<>();
safer.add("sparrow");
safer.add(Boolean.TRUE);    // DOES NOT COMPILE
```

This time the compiler knows that only `String` objects are allowed in and prevents the attempt to add a `boolean`. Now let's try adding multiple values to different positions.

```
4: List<String> birds = new ArrayList<>();
5: birds.add("hawk");           // [hawk]
6: birds.add(1, "robin");       // [hawk, robin]
7: birds.add(0, "blue jay");    // [blue jay, hawk, robin]
8: birds.add(1, "cardinal");    // [blue jay, cardinal, hawk, robin]
9: System.out.println(birds);   // [blue jay, cardinal, hawk, robin]
```

When a question has code that adds objects at indexed positions, draw it so that you won't lose track of which value is at which index. In this example, line 5 adds "hawk" to the end of `birds`. Then line 6 adds "robin" to index 1 of `birds`, which happens to be the end. Line 7 adds "blue jay" to index 0, which happens to be the beginning of `birds`. Finally, line 8 adds "cardinal" to index 1, which is now near the middle of `birds`.

remove()

The `remove()` methods remove the first matching value in the `ArrayList` or remove the element at a specified index. The method signatures are as follows:

```
boolean remove(Object object)
E remove(int index)
```

This time the `boolean` return value tells us whether a match was removed. The `E` return type is the element that actually got removed. The following shows how to use these methods:

```
3: List<String> birds = new ArrayList<>();
4: birds.add("hawk");           // [hawk]
5: birds.add("hawk");           // [hawk, hawk]
6: System.out.println(birds.remove("cardinal")); // prints false
7: System.out.println(birds.remove("hawk"));    // prints true
```

```
8: System.out.println(birds.remove(0)); // prints hawk
9: System.out.println(birds);          // []
```

Line 6 tries to remove an element that is not in `birds`. It returns `false` because no such element is found. Line 7 tries to remove an element that is in `birds` and so returns `true`. Notice that it removes only one match. Line 8 removes the element at index 0, which is the last remaining element in the `ArrayList`.

Since calling `remove()` with an `int` uses the index, an index that doesn't exist will throw an exception. For example, `birds.remove(100)` throws an `IndexOutOfBoundsException`.

There is also a `removeIf()` method. We'll cover it in the next chapter because it uses lambda expressions (a topic in that chapter).

set()

The `set()` method changes one of the elements of the `ArrayList` without changing the size. The method signature is as follows:

```
E set(int index, E newElement)
```

The `E` return type is the element that got replaced. The following shows how to use this method:

```
15: List<String> birds = new ArrayList<>();
16: birds.add("hawk");           // [hawk]
17: System.out.println(birds.size()); // 1
18: birds.set(0, "robin");       // [robin]
19: System.out.println(birds.size()); // 1
20: birds.set(1, "robin");       // IndexOutOfBoundsException
```

Line 16 adds one element to the array, making the size 1. Line 18 replaces that one element and the size stays at 1. Line 20 tries to replace an element that isn't in the `ArrayList`. Since the size is 1, the only valid index is 0. Java throws an exception because this isn't allowed.

isEmpty() and size()

The `isEmpty()` and `size()` methods look at how many of the slots are in use. The method signatures are as follows:

```
boolean isEmpty()
int size()
```

The following shows how to use these methods:

```
System.out.println(birds.isEmpty()); // true
System.out.println(birds.size());    // 0
birds.add("hawk");                   // [hawk]
birds.add("hawk");                   // [hawk, hawk]
```

```
System.out.println(birds.isEmpty());    // false
System.out.println(birds.size());      // 2
```

At the beginning, `birds` has a size of 0 and is empty. It has a capacity that is greater than 0. However, as with `StringBuilder`, we don't use the capacity in determining size or length. After adding elements, the size becomes positive and it is no longer empty.

clear()

The `clear()` method provides an easy way to discard all elements of the `ArrayList`. The method signature is as follows:

```
void clear()
```

The following shows how to use this method:

```
List<String> birds = new ArrayList<>();
birds.add("hawk");           // [hawk]
birds.add("hawk");           // [hawk, hawk]
System.out.println(birds.isEmpty());    // false
System.out.println(birds.size());      // 2
birds.clear();               // []
System.out.println(birds.isEmpty());    // true
System.out.println(birds.size());      // 0
```

After we call `clear()`, `birds` is back to being an empty `ArrayList` of size 0.

contains()

The `contains()` method checks whether a certain value is in the `ArrayList`. The method signature is as follows:

```
boolean contains(Object object)
```

The following shows how to use this method:

```
List<String> birds = new ArrayList<>();
birds.add("hawk");           // [hawk]
System.out.println(birds.contains("hawk")); // true
System.out.println(birds.contains("robin")); // false
```

This method calls `equals()` on each element of the `ArrayList` to see whether there are any matches. Since `String` implements `equals()`, this works out well.

equals()

Finally, `ArrayList` has a custom implementation of `equals()` so you can compare two lists to see if they contain the same elements in the same order.

```
boolean equals(Object object)
```

The following shows an example:

```
31: List<String> one = new ArrayList<>();
32: List<String> two = new ArrayList<>();
33: System.out.println(one.equals(two));    // true
34: one.add("a");                          // [a]
35: System.out.println(one.equals(two));    // false
36: two.add("a");                          // [a]
37: System.out.println(one.equals(two));    // true
38: one.add("b");                          // [a,b]
39: two.add(0, "b");                       // [b,a]
40: System.out.println(one.equals(two));    // false
```

On line 33, the two `ArrayList` objects are equal. An empty list is certainly the same elements in the same order. On line 35, the `ArrayList` objects are not equal because the size is different. On line 37, they are equal again because the same one element is in each. On line 40, they are not equal. The size is the same and the values are the same, but they are not in the same order.

Wrapper Classes

Up to now, we've only put `String` objects in the `ArrayList`. What happens if we want to put primitives in? Each primitive type has a wrapper class, which is an object type that corresponds to the primitive. Table 3.2 lists all the wrapper classes along with the constructor for each.

TABLE 3.2 Wrapper classes

Primitive type	Wrapper class	Example of constructing
boolean	Boolean	<code>new Boolean(true)</code>
byte	Byte	<code>new Byte((byte) 1)</code>
short	Short	<code>new Short((short) 1)</code>
int	Integer	<code>new Integer(1)</code>
long	Long	<code>new Long(1)</code>
float	Float	<code>new Float(1.0)</code>
double	Double	<code>new Double(1.0)</code>
char	Character	<code>new Character('c')</code>

The wrapper classes also have a method that converts back to a primitive. You don't need to know much about the constructors or `intValue()` type methods for the exam because autoboxing has removed the need for them (see the next section). You might encounter this syntax on questions that have been on the exam for many years. However, you just need to be able to read the code and not look for tricks in it.

There are also methods for converting a `String` to a primitive or wrapper class. You do need to know these methods. The parse methods, such as `parseInt()`, return a primitive, and the `valueOf()` method returns a wrapper class. This is easy to remember because the name of the returned primitive is in the method name. For example:

```
int primitive = Integer.parseInt("123");
Integer wrapper = Integer.valueOf("123");
```

The first line converts a `String` to an `int` primitive. The second converts a `String` to an `Integer` wrapper class. If the `String` passed in is not valid for the given type, Java throws an exception. In these examples, letters and dots are not valid for an integer value:

```
int bad1 = Integer.parseInt("a");           // throws NumberFormatException
Integer bad2 = Integer.valueOf("123.45");    // throws NumberFormatException
```

Before you worry, the exam won't make you recognize that the method `parseInt()` is used rather than `parseInteger()`. You simply need to be able to recognize the methods when put in front of you. Also, the `Character` class doesn't participate in the `parse/valueOf` methods. Since a `String` is made up of characters, you can just call `charAt()` normally.

Table 3.3 lists the methods you need to recognize for creating a primitive or wrapper class object from a `String`. In real coding, you won't be so concerned which is returned from each method due to autoboxing.

TABLE 3.3 Converting from a String

Wrapper class	Converting String to primitive	Converting String to wrapper class
Boolean	<code>Boolean.parseBoolean("true");</code>	<code>Boolean.valueOf("TRUE");</code>
Byte	<code>Byte.parseByte("1");</code>	<code>Byte.valueOf("2");</code>
Short	<code>Short.parseShort("1");</code>	<code>Short.valueOf("2");</code>
Integer	<code>Integer.parseInt("1");</code>	<code>Integer.valueOf("2");</code>
Long	<code>Long.parseLong("1");</code>	<code>Long.valueOf("2");</code>
Float	<code>Float.parseFloat("1");</code>	<code>Float.valueOf("2.2");</code>
Double	<code>Double.parseDouble("1");</code>	<code>Double.valueOf("2.2");</code>
Character	None	None

Autoboxing

Why won't you need to be concerned with whether a primitive or wrapper class is returned, you ask? Since Java 5, you can just type the primitive value and Java will convert it to the relevant wrapper class for you. This is called *autoboxing*. Let's look at an example:

```
4: List<Double> weights = new ArrayList<>();
5: weights.add(50.5);           // [50.5]
6: weights.add(new Double(60)); // [50.5, 60.0]
7: weights.remove(50.5);        // [60.0]
8: double first = weights.get(0); // 60.0
```

Line 5 autoboxes the double primitive into a Double object and adds that to the List. Line 6 shows that you can still write code the long way and pass in a wrapper object. Line 7 again autoboxes into the wrapper object and passes it to remove(). Line 8 retrieves the Double and unboxes it into a double primitive.

What do you think happens if you try to unbox a null?

```
3: List<Integer> heights = new ArrayList<>();
4: heights.add(null);
5: int h = heights.get(0);           // NullPointerException
```

On line 4, we add a null to the list. This is legal because a null reference can be assigned to any reference variable. On line 5, we try to unbox that null to an int primitive. This is a problem. Java tries to get the int value of null. Since calling any method on null gives a NullPointerException, that is just what we get. Be careful when you see null in relation to autoboxing.

Be careful when autoboxing into Integer. What do you think this code outputs?

```
List<Integer> numbers = new ArrayList<>();
numbers.add(1);
numbers.add(2);
numbers.remove(1);
System.out.println(numbers);
```

It actually outputs 1. After adding the two values, the List contains [1, 2]. We then request the element with index 1 be removed. That's right: index 1. Because there's already a remove() method that takes an int parameter, Java calls that method rather than autoboxing. If you want to remove the 2, you can write numbers.remove(new Integer(2)) to force wrapper class use.

Converting Between *array* and *List*

You should know how to convert between an array and an ArrayList. Let's start with turning an ArrayList into an array:

```
3: List<String> list = new ArrayList<>();
4: list.add("hawk");
```

```
5: list.add("robin");
6: Object[] objectArray = list.toArray();
7: System.out.println(objectArray.length);    // 2
8: String[] stringArray = list.toArray(new String[0]);
9: System.out.println(stringArray.length);    // 2
```

Line 6 shows that an ArrayList knows how to convert itself to an array. The only problem is that it defaults to an array of class Object. This isn't usually what you want. Line 8 specifies the type of the array and does what we actually want. The advantage of specifying a size of 0 for the parameter is that Java will create a new array of the proper size for the return value. If you like, you can suggest a larger array to be used instead. **If the ArrayList fits in that array, it will be returned. Otherwise, a new one will be created.**

Converting from an array to a List is more interesting. The original array and created array backed List are linked. When a change is made to one, it is available in the other. It is a fixed-size list and is also known a *backed* List because the array changes with it. Pay careful attention to the values here:

```
20: String[] array = { "hawk", "robin" };    // [hawk, robin]
21: List<String> list = Arrays.asList(array); // returns fixed size list
22: System.out.println(list.size());        // 2
23: list.set(1, "test");                    // [hawk, test]
24: array[0] = "new";                       // [new, test]
25: for (String b : array) System.out.print(b + " "); // new test
26: list.remove(1);                         // throws UnsupportedOperationException
```

Line 21 converts the array to a List. Note that it isn't the `java.util.ArrayList` we've grown used to. It is a fixed-size, backed version of a List. Line 23 is okay because `set()` merely replaces an existing value. It updates both `array` and `list` because they point to the same data store. Line 24 also changes both `array` and `list`. Line 25 shows the array has changed to `new test`. Line 26 throws an exception because we are not allowed to change the size of the list.

A Cool Trick with Varargs

This topic isn't on the exam, but merging varargs with ArrayList conversion allows you to create an ArrayList in a cool way:

```
List<String> list = Arrays.asList("one", "two");
```

`asList()` takes varargs, which let you pass in an array or just type out the String values. This is handy when testing because you can easily create and populate a List on one line.

Sorting

Sorting an `ArrayList` is very similar to sorting an array. You just use a different helper class:

```
List<Integer> numbers = new ArrayList<>();
numbers.add(99);
numbers.add(5);
numbers.add(81);
Collections.sort(numbers);
System.out.println(numbers); [5, 81, 99]
```

As you can see, the numbers got sorted, just like you'd expect. Isn't it nice to have something that works just like you think it will?

Working with Dates and Times

In Java 8, Oracle completely revamped how we work with dates and times. You can still write code the “old way,” but those classes aren't on the exam. We'll mention the “old way” in real-world scenarios so that you can learn the “new way” more easily if you first learned Java before version 8. Even if you are learning Java starting with version 8, this will help you when you need to read older code. Just know that the “old way” is not on the exam.

As with an `ArrayList`, you need an import statement to work with the date and time classes. Most of them are in the `java.time` package. To use it, add this import to your program:

```
import java.time.*;           // import time classes
```

In the following sections, we'll look at creating, manipulating, and formatting dates and times.

Creating Dates and Times

In the real world, we usually talk about dates and time zones as if the other person is located near us. For example, if you say “I'll call you at 11:00 on Tuesday morning,” we assume that 11:00 means the same thing to both of us. But if you live in New York and we live in California, you need to be more specific. California is three hours earlier than New York because the states are in different time zones. You would instead say, “I'll call you at 11:00 EST on Tuesday morning.” Luckily, the exam doesn't cover time zones, so discussing dates and times is easier.

When working with dates and times, the first thing to do is decide how much information you need. The exam gives you three choices:

LocalDate Contains just a date—no time and no time zone. A good example of `LocalDate` is your birthday this year. It is your birthday for a full day regardless of what time it is.

LocalTime Contains just a time—no date and no time zone. A good example of `LocalTime` is midnight. It is midnight at the same time every day.

LocalDateTime Contains both a date and time but no time zone. A good example of `LocalDateTime` is “the stroke of midnight on New Year’s.” Midnight on January 2 isn’t nearly as special, and clearly an hour after midnight isn’t as special either.

Oracle recommends avoiding time zones unless you really need them. Try to act as if everyone is in the same time zone when you can. If you do need to communicate across time zones, `ZonedDateTime` handles them.

Ready to create your first date and time objects?

```
System.out.println(LocalDate.now());  
System.out.println(LocalTime.now());  
System.out.println(LocalDateTime.now());
```

Each of the three classes has a static method called `now()` that gives the current date and time. Your output is going to depend on what date/time you run it and where you live. The authors live in the United States, making the output look like the following when run on January 20 at 12:45 p.m.:

```
2015-01-20  
12:45:18.401  
2015-01-20T12:45:18.401
```

The key is to notice the type of information in the output. The first one contains only a date and no time. The second contains only a time and no date. This time displays hours, minutes, seconds, and nanoseconds. The third contains both date and time. Java uses `T` to separate the date and time when converting `LocalDateTime` to a `String`.

Wait—I Don’t Live in the United States

The exam recognizes that exam takers live all over the world and will not ask you about the details of United States date and time formats.

In the United States, the month is written before the date. The exam won’t ask you about the difference between 02/03/2015 and 03/02/2015. That would be mean and not internationally friendly, and it would be testing your knowledge of United States dates rather than your knowledge of Java. That said, our examples do use United States date and time formats as will the questions on the exam. Just remember that the month comes before the date. Also, Java tends to use a 24-hour clock even though the United States uses a 12-hour clock with a.m./p.m.

Now that you know how to create the current date and time, let's look at other specific dates and times. To begin, let's create just a date with no time. Both of these examples create the same date:

```
LocalDate date1 = LocalDate.of(2015, Month.JANUARY, 20);
LocalDate date2 = LocalDate.of(2015, 1, 20);
```

Both pass in the year, month, and date. Although it is good to use the Month constants (to make the code easier to read), you can pass the int number of the month directly. Just use the number of the month the same way you would if you were writing the date in real life.

The method signatures are as follows:

```
public static LocalDate of(int year, int month, int dayOfMonth)
public static LocalDate of(int year, Month month, int dayOfMonth)
```

Month is a special type of class called an enum. You don't need to know about enums on the OCA exam and can just treat them as constants.



Up to now, we've been like a broken record telling you that Java counts starting with 0. Well, months are an exception. For months in the new date and time methods, Java counts starting from 1 like we human beings do.

When creating a time, you can choose how detailed you want to be. You can specify just the hour and minute, or you can add the number of seconds. You can even add nanoseconds if you want to be very precise. (A nanosecond is a billionth of a second—you probably won't need to be that specific.)

```
LocalTime time1 = LocalTime.of(6, 15);           // hour and minute
LocalTime time2 = LocalTime.of(6, 15, 30);        // + seconds
LocalTime time3 = LocalTime.of(6, 15, 30, 200);    // + nanoseconds
```

These three times are all different but within a minute of each other. The method signatures are as follows:

```
public static LocalTime of(int hour, int minute)
public static LocalTime of(int hour, int minute, int second)
public static LocalTime of(int hour, int minute, int second, int nanos)
```

Finally, we can combine dates and times:

```
LocalDateTime dateTime1 = LocalDateTime.of(2015, Month.JANUARY, 20, 6, 15, 30);
LocalDateTime dateTime2 = LocalDateTime.of(dateTime1, time1);
```

The first line of code shows how you can specify all the information about the LocalDateTime right in the same line. There are many method signatures allowing you

to specify different things. Having that many numbers in a row gets to be hard to read, though. The second line of code shows how you can create `LocalDate` and `LocalTime` objects separately first and then combine them to create a `LocalDateTime` object.

This time there are a lot of method signatures since there are more combinations. The method signatures are as follows:

```
public static LocalDateTime of(int year, int month,
    int dayOfMonth, int hour, int minute)
public static LocalDateTime of(int year, int month,
    int dayOfMonth, int hour, int minute, int second)
public static LocalDateTime of(int year, int month,
    int dayOfMonth, int hour, int minute, int second, int nanos)
public static LocalDateTime of(int year, Month month,
    int dayOfMonth, int hour, int minute)
public static LocalDateTime of(int year, Month month,
    int dayOfMonth, int hour, int minute, int second)
public static LocalDateTime of(int year, Month month,
    int dayOfMonth, int hour, int minute, int second, int nanos)
public static LocalDateTime of(LocalDate date, LocalTime)
```

Did you notice that we did not use a constructor in any of the examples? The date and time classes have private constructors to force you to use the static methods. The exam creators may try to throw something like this at you:

```
LocalDate d = new LocalDate(); // DOES NOT COMPILE
```

Don't fall for this. You are not allowed to construct a date or time object directly. Another trick is to see what happens when you pass invalid numbers to `of()`. For example:

```
LocalDate.of(2015, Month.JANUARY, 32) // throws DateTimeException
```

You don't need to know the exact exception that's thrown, but it's a clear one:

```
java.time.DateTimeException: Invalid value for DayOfMonth
(valid values 1 - 28/31): 32
```



Real World Scenario

Creating Dates in Java 7 and Earlier

You can see some of the problems with the "old way" in the following table. There wasn't a way to specify just a date without the time. The `Date` class represented both the date and time whether you wanted it to or not. Trying to create a specific date required more code than it should have. Month indexes were 0 based instead of 1 based, which was confusing.

continues

continued

There’s an old way to create a date. In Java 1.1, you created a specific `Date` with this: `Date jan = new Date(2015, Calendar.JANUARY, 1);`. You could use the `Calendar` class beginning with Java 1.2. `Date` exists mainly for backward compatibility and so that `Calendar` can work with code—making the “new way” the third way. The “new way” is much better so it looks like this is a case of the third time is the charm!

	Old way	New way (Java 8 and later)
Importing	<code>import java.util.*;</code>	<code>import java.time.*;</code>
Creating an object with the current date	<code>Date d = new Date();</code>	<code>LocalDate d = LocalDate.now();</code>
Creating an object with the current date and time	<code>Date d = new Date();</code>	<code>LocalDateTime dt = LocalDateTime.now();</code>
Creating an object representing January 1, 2015	<code>Calendar c = Calendar.getInstance(); c.set(2015, Calendar.JANUARY, 1); Date jan = c.getTime();</code> or <code>Calendar c = new GregorianCalendar(2015, Calendar.JANUARY, 1); Date jan = c.getTime();</code>	<code>LocalDate jan = LocalDate.of(2015, Month.JANUARY, 1);</code>
Creating January 1, 2015 without the constant	<code>Calendar c = Calendar.getInstance(); c.set(2015, 0, 1); Date jan = c.getTime();</code>	<code>LocalDate jan = LocalDate.of(2015, 1, 1)</code>

Manipulating Dates and Times

Adding to a date is easy. The date and time classes are immutable, just like `String` was. This means that we need to remember to assign the results of these methods to a reference variable so they are not lost.

```
12: LocalDate date = LocalDate.of(2014, Month.JANUARY, 20);
13: System.out.println(date);           // 2014-01-20
14: date = date.plusDays(2);
15: System.out.println(date);           // 2014-01-22
16: date = date.plusWeeks(1);
```

```
17: System.out.println(date);           // 2014-01-29
18: date = date.plusMonths(1);
19: System.out.println(date);           // 2014-02-28
20: date = date.plusYears(5);
21: System.out.println(date);           // 2019-02-28
```

This code is nice because it does just what it sounds like. We start out with January 20, 2014. On line 14, we add two days to it and reassign it to our reference variable. On line 16, we add a week. This method allows us to write clearer code than `plusDays(7)`. Now date is January 29, 2014. On line 18, we add a month. This would bring us to February 29, 2014. However, 2014 is not a leap year. (2012 and 2016 are leap years.) Java is smart enough to realize February 29, 2014 does not exist and gives us February 28, 2014 instead. Finally, line 20 adds five years.

There are also nice easy methods to go backward in time. This time, let's work with `LocalDateTime`.

```
22: LocalDate date = LocalDate.of(2020, Month.JANUARY, 20);
23: LocalTime time = LocalTime.of(5, 15);
24: LocalDateTime dateTime = LocalDateTime.of(date, time);
25: System.out.println(dateTime);        // 2020-01-20T05:15
26: dateTime = dateTime.minusDays(1);
27: System.out.println(dateTime);        // 2020-01-19T05:15
28: dateTime = dateTime.minusHours(10);
29: System.out.println(dateTime);        // 2020-01-18T19:15
30: dateTime = dateTime.minusSeconds(30);
31: System.out.println(dateTime);        // 2020-01-18T19:14:30
```

Line 25 prints the original date of January 20, 2020 at 5:15 a.m. Line 26 subtracts a full day, bringing us to January 19, 2020 at 5:15 a.m. Line 28 subtracts 10 hours, showing that the date will change if the hours cause it to and brings us to January 18, 2020 at 19:15 (7:15 p.m.). Finally, line 30 subtracts 30 seconds. We see that all of a sudden the display value starts displaying seconds. Java is smart enough to hide the seconds and nanoseconds when we aren't using them.

It is common for date and time methods to be chained. For example, without the print statements, the previous example could be rewritten as follows:

```
LocalDate date = LocalDate.of(2020, Month.JANUARY, 20);
LocalTime time = LocalTime.of(5, 15);
LocalDateTime dateTime = LocalDateTime.of(date, time)
    .minusDays(1).minusHours(10).minusSeconds(30);
```

When you have a lot of manipulations to make, this chaining comes in handy. There are two ways the exam creators can try to trick you. What do you think this prints?

```
LocalDate date = LocalDate.of(2020, Month.JANUARY, 20);
date.plusDays(10);
System.out.println(date);
```

It prints January 20, 2020. Adding 10 days was useless because we ignored the result. Whenever you see immutable types, pay attention to make sure the return value of a method call isn't ignored.

The exam also may test to see if you remember what each of the date and time objects includes. Do you see what is wrong here?

```
LocalDate date = LocalDate.of(2020, Month.JANUARY, 20);
date = date.plusMinutes(1);    // DOES NOT COMPILE
```

`LocalDate` does not contain time. This means you cannot add minutes to it. This can be tricky in a chained sequence of additions/subtraction operations, so make sure you know which methods in Table 3.4 can be called on which of the three objects.

TABLE 3.4 Methods in `LocalDate`, `LocalTime`, and `LocalDateTime`

	Can call on <code>LocalDate</code> ?	Can call on <code>LocalTime</code> ?	Can call on <code>LocalDateTime</code> ?
<code>plusYears/minusYears</code>	Yes	No	Yes
<code>plusMonths/minusMonths</code>	Yes	No	Yes
<code>plusWeeks/minusWeeks</code>	Yes	No	Yes
<code>plusDays/minusDays</code>	Yes	No	Yes
<code>plusHours/minusHours</code>	No	Yes	Yes
<code>plusMinutes/minusMinutes</code>	No	Yes	Yes
<code>plusSeconds/minusSeconds</code>	No	Yes	Yes
<code>plusNanos/minusNanos</code>	No	Yes	Yes

Manipulating Dates in Java 7 and Earlier

As you look at all the code in the following table to do time calculations in the “old way,” you can see why Java needed to revamp the date and time APIs! The “old way” took a lot of code to do something simple.

	Old way	New way (Java 8 and later)
Adding a day	<pre>public Date addDay(Date date) { Calendar cal = Calendar .getInstance(); cal.setTime(date); cal.add(Calendar.DATE, 1); return cal.getTime(); }</pre>	<pre>public LocalDate addDay(LocalDate date) { return date. plusDays(1); }</pre>
Subtracting a day	<pre>public Date subtractDay(Date date) { Calendar cal = Calendar .getInstance(); cal.setTime(date); cal.add(Calendar.DATE, -1); return cal.getTime(); }</pre>	<pre>public LocalDate subtractDay(LocalDate date) { return date. minusDays(1); }</pre>

Working with Periods

Now we know enough to do something fun with dates! Our zoo performs animal enrichment activities to give the animals something fun to do. The head zookeeper has decided to switch the toys every month. This system will continue for three months to see how it works out.

```
public static void main(String[] args) {
    LocalDate start = LocalDate.of(2015, Month.JANUARY, 1);
    LocalDate end = LocalDate.of(2015, Month.MARCH, 30);
    performAnimalEnrichment(start, end);
}

private static void performAnimalEnrichment(LocalDate start, LocalDate end) {
    LocalDate upTo = start;
    while (upTo.isBefore(end)) {           // check if still before end
        System.out.println("give new toy: " + upTo);
        upTo = upTo.plusMonths(1);         // add a month
    }
}
```

This code works fine. It adds a month to the date until it hits the end date. The problem is that this method can't be reused. Our zookeeper wants to try different schedules to see which works best.

Converting to a long

`LocalDate` and `LocalDateTime` have a method to convert them into long equivalents in relation to 1970. What's special about 1970? That's what UNIX started using for date standards, so Java reused it. And don't worry—you don't have to memorize the names for the exam.

- `LocalDate` has `toEpochDay()`, which is the number of days since January 1, 1970.
- `LocalDateTime` has `toEpochTime()`, which is the number of seconds since January 1, 1970.
- `LocalTime` does not have an epoch method. Since it represents a time that occurs on any date, it doesn't make sense to compare it in 1970. Although the exam pretends time zones don't exist, you may be wondering if this special January 1, 1970 is in a specific time zone. The answer is yes. This special time refers to when it was January 1, 1970 in GMT (Greenwich Mean Time). Greenwich is in England and GMT does not participate in daylight savings time. This makes it a good reference point. (Again, you don't have to know about GMT for the exam.)

Luckily, Java has a `Period` class that we can pass in. This code does the same thing as the previous example:

```
public static void main(String[] args) {
    LocalDate start = LocalDate.of(2015, Month.JANUARY, 1);
    LocalDate end = LocalDate.of(2015, Month.MARCH, 30);
    Period period = Period.ofMonths(1);           // create a period
    performAnimalEnrichment(start, end, period);
}
private static void performAnimalEnrichment(LocalDate start, LocalDate end,
    Period period) {                             // uses the generic period
    LocalDate upTo = start;
    while (upTo.isBefore(end)) {
        System.out.println("give new toy: " + upTo);
        upTo = upTo.plus(period);               // adds the period
    }
}
```

The method can add an arbitrary period of time that gets passed in. This allows us to reuse the same method for different periods of time as our zookeeper changes her mind.

There are five ways to create a `Period` class:

```
Period annually = Period.ofYears(1);           // every 1 year
Period quarterly = Period.ofMonths(3);         // every 3 months
```



```

Period everyThreeWeeks = Period.ofWeeks(3);           // every 3 weeks
Period everyOtherDay = Period.ofDays(2);             // every 2 days
Period everyYearAndAWeek = Period.of(1, 0, 7);       // every year and 7 days

```

There's one catch. You cannot chain methods when creating a `Period`. The following code looks like it is equivalent to the `everyYearAndAWeek` example, but it's not. Only the last method is used because the `Period.ofXXX` methods are static methods.

```

Period wrong = Period.ofYears(1).ofWeeks(1);         // every week

```

This tricky code is really like writing the following:

```

Period wrong = Period.ofYears(1);
wrong = Period.ofWeeks(7);

```

This is clearly not what you intended! That's why the `of()` method allows us to pass in the number of years, months, and days. They are all included in the same period. You will get a compiler warning about this. Compiler warnings tell you something is wrong or suspicious without failing compilation.

You've probably noticed by now that a `Period` is a day or more of time. There is also `Duration`, which is intended for smaller units of time. For `Duration`, you can specify the number of days, hours, minutes, seconds, or nanoseconds. And yes, you could pass 365 days to make a year, but you really shouldn't—that's what `Period` is for. `Duration` isn't on the exam since it roughly works the same way as `Period`. It's good to know it exists, though.

The last thing to know about `Period` is what objects it can be used with. Let's look at some code:

```

3: LocalDate date = LocalDate.of(2015, 1, 20);
4: LocalTime time = LocalTime.of(6, 15);
5: LocalDateTime dateTime = LocalDateTime.of(date, time);
6: Period period = Period.ofMonths(1);
7: System.out.println(date.plus(period));           // 2015-02-20
8: System.out.println(dateTime.plus(period));       // 2015-02-20T06:15
9: System.out.println(time.plus(period));           // UnsupportedOperationException

```

Lines 7 and 8 work as expected. They add a month to January 20, 2015, giving us February 20, 2015. The first has only the date, and the second has both the date and time.

Line 9 attempts to add a month to an object that only has a time. This won't work. Java throws an exception and complains that we attempt to use an `Unsupported unit: Months`.

As you can see, you'll have to pay attention to the type of date and time objects every place you see them.

Formatting Dates and Times

The date and time classes support many methods to get data out of them:

```
LocalDate date = LocalDate.of(2020, Month.JANUARY, 20);
System.out.println(date.getDayOfWeek());    // MONDAY
System.out.println(date.getMonth());        // JANUARY
System.out.println(date.getYear());         // 2020
System.out.println(date.getDayOfYear());    // 20
```

We could use this information to display information about the date. However, it would be more work than necessary. Java provides a class called `DateTimeFormatter` to help us out. Unlike the `LocalDateTime` class, `DateTimeFormatter` can be used to format any type of date and/or time object. What changes is the format. `DateTimeFormatter` is in the package `java.time.format`.

```
LocalDate date = LocalDate.of(2020, Month.JANUARY, 20);
LocalTime time = LocalTime.of(11, 12, 34);
LocalDateTime dateTime = LocalDateTime.of(date, time);
System.out.println(date.format(DateTimeFormatter.ISO_LOCAL_DATE));
System.out.println(time.format(DateTimeFormatter.ISO_LOCAL_TIME));
System.out.println(dateTime.format(DateTimeFormatter.ISO_LOCAL_DATE_TIME));
```

ISO is a standard for dates. The output of the previous code looks like this:

```
2020-01-20
11:12:34
2020-01-20T11:12:34
```

This is a reasonable way for computers to communicate, but probably not how you want to output the date and time in your program. Luckily there are some predefined formats that are more useful:

```
DateTimeFormatter shortDateTime =
    DateTimeFormatter.ofLocalizedDate(FormatStyle.SHORT);
System.out.println(shortDateTime.format(dateTime));    // 1/20/20
System.out.println(shortDateTime.format(date));        // 1/20/20
System.out.println(
    shortDateTime.format(time)); // UnsupportedOperationException
```

Here we say we want a localized formatter in the predefined short format. The last line throws an exception because a time cannot be formatted as a date. The `format()` method is declared on both the formatter objects and the date/time objects, allowing you to reference the objects in either order. The following statements print exactly the same thing as the previous code:

```

DateTimeFormatter shortDateTime =
    DateTimeFormatter.ofLocalizedDate(FormatStyle.SHORT);
System.out.println(dateTime.format(shortDateTime));
System.out.println(date.format(shortDateTime));
System.out.println(time.format(shortDateTime));

```

In this book, we'll change around the orders to get you used to seeing it both ways. Table 3.5 shows the legal and illegal localized formatting methods.

TABLE 3.5 ofLocalized methods

DateTimeFormatter f = DateTime Formatter._____ (FormatStyle.SHORT);	Calling f.format (localDate)	Calling f.format (localDateTime)	Calling f.format (localTime)
ofLocalizedDate	Legal – shows whole object	Legal – shows just date part	Throws runtime exception
ofLocalizedDateTime	Throws runtime exception	Legal – shows whole object	Throws runtime exception
ofLocalizedTime	Throws runtime exception	Legal – shows just time part	Legal – shows whole object

There are two predefined formats that can show up on the exam: SHORT and MEDIUM. The other predefined formats involve time zones, which are not on the exam.

```

LocalDate date = LocalDate.of(2020, Month.JANUARY, 20);
LocalTime time = LocalTime.of(11, 12, 34);
LocalDateTime dateTime = LocalDateTime.of(date, time);

```

```

DateTimeFormatter shortF = DateTimeFormatter
    .ofLocalizedDateTime(FormatStyle.SHORT);
DateTimeFormatter mediumF = DateTimeFormatter
    .ofLocalizedDateTime(FormatStyle.MEDIUM);
System.out.println(shortF.format(dateTime));    // 1/20/20 11:12 AM
System.out.println(mediumF.format(dateTime));    // Jan 20, 2020 11:12:34 AM

```

If you don't want to use one of the predefined formats, you can create your own. For example, this code spells out the month:

```

DateTimeFormatter f = DateTimeFormatter.ofPattern("MMMM dd, yyyy, hh:mm");
System.out.println(dateTime.format(f));    // January 20, 2020, 11:12

```

Before we look at the syntax, know you are not expected to memorize what different numbers of each symbol mean. The most you will need to do is recognize the date and time pieces.

MMMM **M** represents the month. The more **M**s you have, the more verbose the Java output. For example, **M** outputs 1, **MM** outputs 01, **MMM** outputs Jan, and **MMMM** outputs January.

dd **d** represents the date in the month. As with month, the more **d**s you have, the more verbose the Java output. **dd** means to include the leading zero for a single-digit month.

• Use **,** if you want to output a comma (this also appears after the year).

yyyy **y** represents the year. **yy** outputs a two-digit year and **yyyy** outputs a four-digit year.

hh **h** represents the hour. Use **hh** to include the leading zero if you're outputting a single-digit hour.

• Use **:** if you want to output a colon.

mm **m** represents the minute.



Real World Scenario

Formatting Dates in Java 7 and Earlier

Formatting is roughly equivalent to the “old way”; it just uses a different class.

	Old way	New way (Java 8 and later)
Formatting the times	<pre>SimpleDateFormat sf = new SimpleDateFormat("hh:mm"); sf.format(jan3);</pre>	<pre>DateTimeFormatter f = DateTimeFormatter. ofPattern("hh:mm"); dt.format(f);</pre>

Let's do a quick review. Can you figure out which of these lines will throw an exception?

```
4: DateTimeFormatter f = DateTimeFormatter.ofPattern("hh:mm");  
5: f.format(dateTime);  
6: f.format(date);  
7: f.format(time);
```

If you get this question on the exam, think about what the symbols represent. We have **h** for hour and **m** for minute. Remember **M** (uppercase) is month and **m** (lowercase) is minute. We can only use this formatter with objects containing times. Therefore, line 6 will throw an exception.

Parsing Dates and Times

Now that you know how to convert a date or time to a formatted `String`, you'll find it easy to convert a `String` to a date or time. Just like the `format()` method, the `parse()` method takes a formatter as well. If you don't specify one, it uses the default for that type.

```
DateTimeFormatter f = DateTimeFormatter.ofPattern("MM dd yyyy");
LocalDate date = LocalDate.parse("01 02 2015", f);
LocalTime time = LocalTime.parse("11:22");
System.out.println(date);           // 2015-01-02
System.out.println(time);           // 11:22
```

Here we show using both a custom formatter and a default value. This isn't common, but you might have to read code that looks like this on the exam. Parsing is consistent in that if anything goes wrong, Java throws a runtime exception. That could be a format that doesn't match the `String` to be parsed or an invalid date.

Summary

In this chapter, you learned that `Strings` are immutable sequences of characters. The new operator is optional. The concatenation operator (+) creates a new `String` with the content of the first `String` followed by the content of the second `String`. If either operand involved in the + expression is a `String`, concatenation is used; otherwise, addition is used. `String` literals are stored in the string pool. The `String` class has many methods. You need to know `charAt()`, `concat()`, `endsWith()`, `equals()`, `equalsIgnoreCase()`, `indexOf()`, `length()`, `replace()`, `startsWith()`, `substring()`, `toLowerCase()`, `toUpperCase()`, and `trim()`.

`StringBuilders` are mutable sequences of characters. Most of the methods return a reference to the current object to allow method chaining. The `StringBuilder` class has many methods. You need to know `append()`, `charAt()`, `delete()`, `deleteCharAt()`, `indexOf()`, `insert()`, `length()`, `reverse()`, `substring()`, and `toString()`. `StringBuffer` is the same as `StringBuilder` except that it is thread safe.

Calling `==` on `String` objects will check whether they point to the same object in the pool. Calling `==` on `StringBuilder` references will check whether they are pointing to the same `StringBuilder` object. Calling `equals()` on `String` objects will check whether the sequence of characters is the same. Calling `equals()` on `StringBuilder` objects will check whether they are pointing to the same object rather than looking at the values inside.

An array is a fixed-size area of memory on the heap that has space for primitives or pointers to objects. You specify the size when creating it—for example, `int[] a = new int[6];`. Indexes begin with 0 and elements are referred to using `a[0]`. The `Arrays.sort()` method sorts an array. `Arrays.binarySearch()` searches a sorted array and returns the index of a match. If no match is found, it negates the position where the element would

need to be inserted and subtracts 1. Methods that are passed varargs (...) can be used as if a normal array was passed in. In a multidimensional array, the second-level arrays and beyond can be different sizes.

An `ArrayList` can change size over its life. It can be stored in an `ArrayList` or `List` reference. Generics can specify the type that goes in the `ArrayList`. You need to know the methods `add()`, `clear()`, `contains()`, `equals()`, `isEmpty()`, `remove()`, `set()`, and `size()`. Although an `ArrayList` is not allowed to contain primitives, Java will autobox parameters passed in to the proper wrapper type. `Collections.sort()` sorts an `ArrayList`.

A `LocalDate` contains just a date, a `LocalTime` contains just a time, and a `LocalDateTime` contains both a date and time. All three have private constructors and are created using `LocalDate.now()` or `LocalDate.of()` (or the equivalents for that class). Dates and times can be manipulated using `plusXXX` or `minusXXX` methods. The `Period` class represents a number of days, months, or years to add or subtract from a `LocalDate` or `LocalDateTime`. `DateTimeFormatter` is used to output dates and times in the desired format. The date and time classes are all immutable, which means the return value must be used.

Exam Essentials

Be able to determine the output of code using `String`. Know the rules for concatenating Strings and how to use common `String` methods. Know that Strings are immutable. Pay special attention to the fact that indexes are zero based and that `substring()` gets the string up until right before the index of the second parameter.

Be able to determine the output of code using `StringBuilder`. Know that `StringBuilder` is mutable and how to use common `StringBuilder` methods. Know that `substring()` does not change the value of a `StringBuilder` whereas `append()`, `delete()`, and `insert()` do change it. Also note that most `StringBuilder` methods return a reference to the current instance of `StringBuilder`.

Understand the difference between `==` and `equals`. `==` checks object equality. `equals()` depends on the implementation of the object it is being called on. For Strings, `equals()` checks the characters inside of it.

Be able to determine the output of code using arrays. Know how to declare and instantiate one-dimensional and multidimensional arrays. Be able to access each element and know when an index is out of bounds. Recognize correct and incorrect output when searching and sorting.

Be able to determine the output of code using `ArrayList`. Know that `ArrayList` can increase in size. Be able to identify the different ways of declaring and instantiating an `ArrayList`. Identify correct output from `ArrayList` methods, including the impact of autoboxing.

Recognize invalid uses of dates and times. `LocalDate` does not contain time fields and `LocalTime` does not contain date fields. Watch for operations being performed on the wrong time. Also watch for adding or subtracting time and ignoring the result.

Review Questions

1. What is output by the following code? (Choose all that apply)

```
1: public class Fish {  
2:     public static void main(String[] args) {  
3:         int numFish = 4;  
4:         String fishType = "tuna";  
5:         String anotherFish = numFish + 1;  
6:         System.out.println(anotherFish + " " + fishType);  
7:         System.out.println(numFish + " " + 1);  
8:     } }
```

- A. 4 1
- B. 41
- C. 5
- D. 5 tuna
- E. 5tuna
- F. 51tuna
- G. The code does not compile.

2. Which of the following are output by this code? (Choose all that apply)

```
3: String s = "Hello";  
4: String t = new String(s);  
5: if ("Hello".equals(s)) System.out.println("one");  
6: if (t == s) System.out.println("two");  
7: if (t.equals(s)) System.out.println("three");  
8: if ("Hello" == s) System.out.println("four");  
9: if ("Hello" == t) System.out.println("five");
```

- A. one
- B. two
- C. three
- D. four
- E. five
- F. The code does not compile.

3. Which are true statements? (Choose all that apply)

- A. An immutable object can be modified.
- B. An immutable object cannot be modified.
- C. An immutable object can be garbage collected.

- D. An immutable object cannot be garbage collected.
 - E. String is immutable.
 - F. StringBuffer is immutable.
 - G. StringBuilder is immutable.
4. What is the result of the following code?
- ```
7: StringBuilder sb = new StringBuilder();
8: sb.append("aaa").insert(1, "bb").insert(4, "ccc");
9: System.out.println(sb);
```
- A. abbaaccc
  - B. abbaccca
  - C. bbaaaccc
  - D. bbaaccca
  - E. An exception is thrown.
  - F. The code does not compile.
5. What is the result of the following code?
- ```
2: String s1 = "java";
3: StringBuilder s2 = new StringBuilder("java");
4: if (s1 == s2)
5:   System.out.print("1");
6: if (s1.equals(s2))
7:   System.out.print("2");
```
- A. 1
 - B. 2
 - C. 12
 - D. No output is printed.
 - E. An exception is thrown.
 - F. The code does not compile.
6. What is the result of the following code?
- ```
public class Lion {
 public void roar(String roar1, StringBuilder roar2) {
 roar1.concat("!!!");
 roar2.append("!!!");
 }
 public static void main(String[] args) {
 String roar1 = "roar";
 StringBuilder roar2 = new StringBuilder("roar");
 new Lion().roar(roar1, roar2);
 }
}
```



```
 System.out.println(roar1 + " " + roar2);
} }
```

- A. roar roar
  - B. roar roar!!!
  - C. roar!!! roar
  - D. roar!!! roar!!!
  - E. An exception is thrown.
  - F. The code does not compile.
7. Which are the results of the following code? (Choose all that apply)
- ```
String letters = "abcdef";  
System.out.println(letters.length());  
System.out.println(letters.charAt(3));  
System.out.println(letters.charAt(6));
```
- A. 5
 - B. 6
 - C. c
 - D. d
 - E. An exception is thrown.
 - F. The code does not compile.
8. Which are the results of the following code? (Choose all that apply)
- ```
String numbers = "012345678";
System.out.println(numbers.substring(1, 3));
System.out.println(numbers.substring(7, 7));
System.out.println(numbers.substring(7));
```
- A. 12
  - B. 123
  - C. 7
  - D. 78
  - E. A blank line.
  - F. An exception is thrown.
  - G. The code does not compile.
9. What is the result of the following code?
- ```
3: String s = "purr";  
4: s.toUpperCase();  
5: s.trim();  
6: s.substring(1, 3);
```

```
7: s += " two";  
8: System.out.println(s.length());
```

- A. 2
- B. 4
- C. 8
- D. 10
- E. An exception is thrown.
- F. The code does not compile.

10. What is the result of the following code? (Choose all that apply)

```
13: String a = "";  
14: a += 2;  
15: a += 'c';  
16: a += false;  
17: if ( a == "2cfalse") System.out.println("==");  
18: if ( a.equals("2cfalse")) System.out.println("equals");
```

- A. Compile error on line 14.
- B. Compile error on line 15.
- C. Compile error on line 16.
- D. Compile error on another line.
- E. ==
- F. equals
- G. An exception is thrown.

11. What is the result of the following code?

```
4: int total = 0;  
5: StringBuilder letters = new StringBuilder("abcdefg");  
6: total += letters.substring(1, 2).length();  
7: total += letters.substring(6, 6).length();  
8: total += letters.substring(6, 5).length();  
9: System.out.println(total);
```

- A. 1
- B. 2
- C. 3
- D. 7
- E. An exception is thrown.
- F. The code does not compile.

12. What is the result of the following code? (Choose all that apply)
- ```
StringBuilder numbers = new StringBuilder("0123456789");
numbers.delete(2, 8);
numbers.append("-").insert(2, "+");
System.out.println(numbers);
```
- A. 01+89-
  - B. 012+9-
  - C. 012+-9
  - D. 0123456789
  - E. An exception is thrown.
  - F. The code does not compile.
13. What is the result of the following code?
- ```
StringBuilder b = "rumble";
b.append(4).deleteCharAt(3).delete(3, b.length() - 1);
System.out.println(b);
```
- A. rum
 - B. rum4
 - C. rumb4
 - D. rumble4
 - E. An exception is thrown.
 - F. The code does not compile.
14. Which of the following can replace line 4 to print "avaJ"? (Choose all that apply)
- ```
3: StringBuilder puzzle = new StringBuilder("Java");
4: // INSERT CODE HERE
5: System.out.println(puzzle);
```
- A. puzzle.reverse();
  - B. puzzle.append("vaJ\$").substring(0, 4);
  - C. puzzle.append("vaJ\$").delete(0, 3).deleteCharAt(puzzle.length() - 1);
  - D. puzzle.append("vaJ\$").delete(0, 3).deleteCharAt(puzzle.length());
  - E. None of the above.
15. Which of these array declarations is not legal? (Choose all that apply)
- A. `int[][] scores = new int[5][];`
  - B. `Object[][][] cubbies = new Object[3][0][5];`
  - C. `String beans[] = new beans[6];`
  - D. `java.util.Date[] dates[] = new java.util.Date[2][];`
  - E. `int[][] types = new int[];`
  - F. `int[][] java = new int[][];`

16. Which of these compile when replacing line 8? (Choose all that apply)

```
7: char[] c = new char[2];
```

```
8: // INSERT CODE HERE
```

- A. `int length = c.capacity;`
- B. `int length = c.capacity();`
- C. `int length = c.length;`
- D. `int length = c.length();`
- E. `int length = c.size;`
- F. `int length = c.size();`
- G. None of the above.

17. Which of these compile when replacing line 8? (Choose all that apply)

```
7: ArrayList l = new ArrayList();
```

```
8: // INSERT CODE HERE
```

- A. `int length = l.capacity;`
- B. `int length = l.capacity();`
- C. `int length = l.length;`
- D. `int length = l.length();`
- E. `int length = l.size;`
- F. `int length = l.size();`
- G. None of the above.

18. Which of the following are true? (Choose all that apply)

- A. An array has a fixed size.
- B. An `ArrayList` has a fixed size.
- C. An array allows multiple dimensions.
- D. An array is ordered.
- E. An `ArrayList` is ordered.
- F. An array is immutable.
- G. An `ArrayList` is immutable.

19. Which of the following are true? (Choose all that apply)

- A. Two arrays with the same content are equal.
- B. Two `ArrayLists` with the same content are equal.
- C. If you call `remove(0)` using an empty `ArrayList` object, it will compile successfully.
- D. If you call `remove(0)` using an empty `ArrayList` object, it will run successfully.
- E. None of the above.

**20.** What is the result of the following statements?

```
6: List<String> list = new ArrayList<String>();
7: list.add("one");
8: list.add("two");
9: list.add(7);
10: for(String s : list) System.out.print(s);
```

- A. onetwo
- B. onetwo7
- C. onetwo followed by an exception
- D. Compiler error on line 9.
- E. Compiler error on line 10.

**21.** What is the result of the following statements?

```
3: ArrayList<Integer> values = new ArrayList<>();
4: values.add(4);
5: values.add(5);
6: values.set(1, 6);
7: values.remove(0);
8: for (Integer v : values) System.out.print(v);
```

- A. 4
- B. 5
- C. 6
- D. 46
- E. 45
- F. An exception is thrown.
- G. The code does not compile.

**22.** What is the result of the following?

```
int[] random = { 6, -4, 12, 0, -10 };
int x = 12;
int y = Arrays.binarySearch(random, x);
System.out.println(y);
```

- A. 2
- B. 4
- C. 6
- D. The result is undefined.
- E. An exception is thrown.
- F. The code does not compile.

**23.** What is the result of the following?

```
4: List<Integer> list = Arrays.asList(10, 4, -1, 5);
5: Collections.sort(list);
6: Integer array[] = list.toArray(new Integer[4]);
7: System.out.println(array[0]);
```

- A. -1
- B. 10
- C. Compiler error on line 4.
- D. Compiler error on line 5.
- E. Compiler error on line 6.
- F. An exception is thrown.

**24.** What is the result of the following?

```
6: String [] names = {"Tom", "Dick", "Harry"};
7: List<String> list = names.asList();
8: list.set(0, "Sue");
9: System.out.println(names[0]);
```

- A. Sue
- B. Tom
- C. Compiler error on line 7.
- D. Compiler error on line 8.
- E. An exception is thrown.

**25.** What is the result of the following?

```
List<String> hex = Arrays.asList("30", "8", "3A", "FF");
Collections.sort(hex);
int x = Collections.binarySearch(hex, "8");
int y = Collections.binarySearch(hex, "3A");
int z = Collections.binarySearch(hex, "4F");
System.out.println(x + " " + y + " " + z);
```

- A. 0 1 -2
- B. 0 1 -3
- C. 2 1 -2
- D. 2 1 -3
- E. None of the above.
- F. The code doesn't compile.

**26.** Which of the following are true statements about the following code? (Choose all that apply)

```
4: List<Integer> ages = new ArrayList<>();
5: ages.add(Integer.parseInt("5"));
```

```
6: ages.add(Integer.valueOf("6"));
7: ages.add(7);
8: ages.add(null);
9: for (int age : ages) System.out.print(age);
```

- A. The code compiles.
- B. The code throws a runtime exception.
- C. Exactly one of the add statements uses autoboxing.
- D. Exactly two of the add statements use autoboxing.
- E. Exactly three of the add statements use autoboxing.

27. What is the result of the following?

```
List<String> one = new ArrayList<String>();
one.add("abc");
List<String> two = new ArrayList<>();
two.add("abc");
if (one == two)
 System.out.println("A");
else if (one.equals(two))
 System.out.println("B");
else
 System.out.println("C");
```

- A. A
- B. B
- C. C
- D. An exception is thrown.
- E. The code does not compile.

28. Which of the following can be inserted into the blank to create a date of June 21, 2014?  
(Choose all that apply)

```
import java.time.*;
```

```
public class StartOfSummer {
```

```
 public static void main(String[] args) {
 LocalDate date = _____
 }
```

```
}
```

- A. new LocalDate(2014, 5, 21);
- B. new LocalDate(2014, 6, 21);
- C. LocalDate.of(2014, 5, 21);

- D. `LocalDate.of(2014, 6, 21);`
- E. `LocalDate.of(2014, Calendar.JUNE, 21);`
- F. `LocalDate.of(2014, Month.JUNE, 21);`

**29.** What is the output of the following code?

```
LocalDate date = LocalDate.parse("2018-04-30", DateTimeFormatter.ISO_LOCAL_
DATE);
date.plusDays(2);
date.plusHours(3);
System.out.println(date.getYear() + " " + date.getMonth() + " "
+ date.getDayOfMonth());
```

- A. 2018 APRIL 2
- B. 2018 APRIL 30
- C. 2018 MAY 2
- D. The code does not compile.
- E. A runtime exception is thrown.

**30.** What is the output of the following code?

```
LocalDate date = LocalDate.of(2018, Month.APRIL, 40);
System.out.println(date.getYear() + " " + date.getMonth() + " "
+ date.getDayOfMonth());
```

- A. 2018 APRIL 4
- B. 2018 APRIL 30
- C. 2018 MAY 10
- D. Another date.
- E. The code does not compile.
- F. A runtime exception is thrown.

**31.** What is the output of the following code?

```
LocalDate date = LocalDate.of(2018, Month.APRIL, 30);
date.plusDays(2);
date.plusYears(3);
System.out.println(date.getYear() + " " + date.getMonth() + " "
+ date.getDayOfMonth());
```

- A. 2018 APRIL 2
- B. 2018 APRIL 30
- C. 2018 MAY 2
- D. 2021 APRIL 2



- E. 2021 APRIL 30
- F. 2021 MAY 2
- G. A runtime exception is thrown.

**32.** What is the output of the following code?

```
LocalDateTime d = LocalDateTime.of(2015, 5, 10, 11, 22, 33);
Period p = Period.of(1, 2, 3);
d = d.minus(p);
DateTimeFormatter f = DateTimeFormatter.ofLocalizedTime(FormatStyle.SHORT);
System.out.println(d.format(f));
```

- A. 3/7/14 11:22 AM
- B. 5/10/15 11:22 AM
- C. 3/7/14
- D. 5/10/15
- E. 11:22 AM
- F. The code does not compile.
- G. A runtime exception is thrown.

**33.** What is the output of the following code?

```
LocalDateTime d = LocalDateTime.of(2015, 5, 10, 11, 22, 33);
Period p = Period.ofDays(1).ofYears(2);
d = d.minus(p);
DateTimeFormatter f = DateTimeFormatter.ofLocalizedDateTime(FormatStyle.SHORT);
System.out.println(f.format(d));
```

- A. 5/9/13 11:22 AM
- B. 5/10/13 11:22 AM
- C. 5/9/14
- D. 5/10/14
- E. The code does not compile.
- F. A runtime exception is thrown.



# Chapter 4

## Methods and Encapsulation

---

### OCA EXAM OBJECTIVES COVERED IN THIS CHAPTER:

#### ✓ Working with Methods and Encapsulation

- Create methods with arguments and return values; including overloaded methods
- Apply the static keyword to methods and fields
- Create and overload constructors; include impact on default constructors
- Apply access modifiers
- Apply encapsulation principles to a class
- Determine the effect upon object references and primitive values when they are passed into methods that change the values

#### ✓ Working with Selected classes from the Java API

- Write a simple Lambda expression that consumes a Lambda Predicate expression






In previous chapters, we’ve used methods and constructors without examining them in detail. In this chapter, we’ll explore methods and constructors in depth and cover everything you need to know about them for the OCA exam. (Well, almost—Chapter 5, “Class Design,” will explain the effects of inheritance on both methods and constructors.) This chapter discusses instance variables, the `final` keyword, access modifiers, and initialization. You’ll also learn how to write a simple lambda expression.

# Designing Methods

Every interesting Java program we’ve seen has had a `main()` method. We can write other methods, too. For example, we can write a basic method to take a nap, as shown in Figure 4.1.

**FIGURE 4.1** Method signature



This is called a *method declaration*, which specifies all the information needed to call the method. There are a lot of parts, and we’ll cover each one in more detail. Table 4.1 is a brief reference to the elements of a method declaration. Don’t worry if it seems like a lot of information—by the time you finish this chapter, it will all fit together.

**TABLE 4.1** Parts of a method declaration

| Element            | Value in <code>nap()</code> example | Required? |
|--------------------|-------------------------------------|-----------|
| Access modifier    | <code>public</code>                 | No        |
| Optional specifier | <code>final</code>                  | No        |

| Element                 | Value in <code>nap()</code> example      | Required?                         |
|-------------------------|------------------------------------------|-----------------------------------|
| Return type             | <code>void</code>                        | Yes                               |
| Method name             | <code>nap</code>                         | Yes                               |
| Parameter list          | <code>(int minutes)</code>               | Yes, but can be empty parentheses |
| Optional exception list | <code>throws InterruptedException</code> | No                                |
| Method body             | <pre>{     // take a nap }</pre>         | Yes, but can be empty braces      |

To call this method, just type its name, followed by a single `int` value in parentheses:  
`nap(10);`

Let's start by taking a look at each of these parts of a basic method.

## Access Modifiers

Java offers four choices of access modifier:

***public*** The method can be called from any class.

***private*** The method can only be called from within the same class.

***protected*** The method can only be called from classes in the same package or subclasses. You'll learn about subclasses in Chapter 5.

***Default (Package Private) Access*** The method can only be called from classes in the same package. This one is tricky because there is no keyword for default access. You simply omit the access modifier.



There's a default keyword in Java. You saw it in the `switch` statement in Chapter 2, "Operators and Statements," and you'll see it again in the next chapter when we discuss interfaces. It's not used for access control.

We'll explore the impact of the various access modifiers later in this chapter. For now, just master identifying valid syntax of methods. The exam creators like to trick you by putting method elements in the wrong order or using incorrect values.

We'll have practice examples as we go through each of the method elements in this section. Make sure you understand why each of these is a valid or invalid method declaration. Pay attention to the access modifiers as you figure out what is wrong with the ones that don't compile:

```
public void walk1() {}
default void walk2() {} // DOES NOT COMPILE
```

```
void public walk3() {} // DOES NOT COMPILE
void walk4() {}
```

`walk1()` is a valid method declaration with public access. `walk4()` is a valid method declaration with default access. `walk2()` doesn't compile because default is not a valid access modifier. `walk3()` doesn't compile because the access modifier is specified after the return type.

## Optional Specifiers

There are a number of optional specifiers, but most of them aren't on the exam. Optional specifiers come from the following list. Unlike with access modifiers, you can have multiple specifiers in the same method (although not all combinations are legal). When this happens, you can specify them in any order. And since it is optional, you can't have any of them at all. This means you can have zero or more specifiers in a method declaration.

***static*** Covered later in this chapter. Used for class methods.

***abstract*** Covered in Chapter 5. Used when not providing a method body.

***final*** Covered in Chapter 5. Used when a method is not allowed to be overridden by a subclass.

***synchronized*** On the OCP but not the OCA exam.

***native*** Not on the OCA or OCP exam. Used when interacting with code written in another language such as C++.

***strictfp*** Not on the OCA or OCP exam. Used for making floating-point calculations portable.

Again, just focus on syntax for now. Do you see why these compile or don't compile?

```
public void walk1() {}
public final void walk2() {}
public static final void walk3() {}
public final static void walk4() {}
public modifier void walk5() {} // DOES NOT COMPILE
public void final walk6() {} // DOES NOT COMPILE
final public void walk7() {}
```

`walk1()` is a valid method declaration with no optional specifier. This is okay; it is optional, after all. `walk2()` is a valid method declaration, with `final` as the optional specifier. `walk3()` and `walk4()` are valid method declarations with both `final` and `static` as optional specifiers. The order of these two keywords doesn't matter. `walk5()` doesn't

compile because *modifier* is not a valid optional specifier. `walk6()` doesn't compile because the optional specifier is after the return type.

`walk7()` does compile. Java allows the optional specifiers to appear before the access modifier. This is a weird case and not one you need to know for the exam. We are mentioning it so you don't get confused when practicing.

## Return Type

The next item in a method declaration is the return type. The return type might be an actual Java type such as `String` or `int`. If there is no return type, the `void` keyword is used. This special return type comes from the English language: *void* means without contents. In Java, we have no type there.



Remember that a method must have a return type. If no value is returned, the return type is `void`. You cannot omit the return type.

When checking return types, you also have to look inside the method body. Methods with a return type other than `void` are required to have a return statement inside the method body. This return statement must include the primitive or object to be returned. Methods that have a return type of `void` are permitted to have a return statement with no value returned or omit the return statement entirely.

Ready for some examples? Can you explain why these methods compile or don't?

```
public void walk1() { }
public void walk2() { return; }
public String walk3() { return ""; }
public String walk4() { } // DOES NOT COMPILE
public walk5() { } // DOES NOT COMPILE
String walk6(int a) { if (a == 4) return ""; } // DOES NOT COMPILE
```

Since the return type of `walk1()` is `void`, the return statement is optional. `walk2()` shows the optional return statement that correctly doesn't return anything. `walk3()` is a valid method with a `String` return type and a return statement that returns a `String`. `walk4()` doesn't compile because the return statement is missing. `walk5()` doesn't compile because the return type is missing.

`walk6()` is a little tricky. There is a return statement, but it doesn't always get run. If `a` is 6, the return statement doesn't get executed. Since the `String` always needs to be returned, the compiler complains.

When returning a value, it needs to be assignable to the return type. Imagine there is a local variable of that type to which it is assigned before being returned. Can you think of how to add a line of code with a local variable in these two methods?

```
int integer() {
 return 9;
}
int long() {}
 return 9L; // DOES NOT COMPILE
}
```

It is a fairly mechanical exercise. You just add a line with a local variable. The type of the local variable matches the return type of the method. Then you return that local variable instead of the value directly:

```
int integerExpanded() {
 int temp = 9;
 return temp;
}
int longExpanded() {
 int temp = 9L; // DOES NOT COMPILE
 return temp;
}
```

This shows more clearly why you can't return a `long` primitive in a method that returns an `int`. You can't stuff that `long` into an `int` variable, so you can't return it directly either.

## Method Name

Method names follow the same rules as we practiced with variable names in Chapter 1, “Java Building Blocks.” To review, an identifier may only contain letters, numbers, `$`, or `_`. Also, the first character is not allowed to be a number, and reserved words are not allowed. By convention, methods begin with a lowercase letter but are not required to. Since this is a review of Chapter 1, we can jump right into practicing with some examples:

```
public void walk1() {}
public void 2walk() {} // DOES NOT COMPILE
public walk3 void() {} // DOES NOT COMPILE
public void Walk_$() {}
public void() {} // DOES NOT COMPILE
```

`walk1()` is a valid method declaration with a traditional name. `2walk()` doesn't compile because identifiers are not allowed to begin with numbers. `walk3()` doesn't compile because the method name is before the return type. `Walk_$( )` is a valid method declaration. While it certainly isn't good practice to start a method name with a capital letter and end with punctuation, it is legal. The final line of code doesn't compile because the method name is missing.



## Parameter List

Although the parameter list is required, it doesn't have to contain any parameters. This means you can just have an empty pair of parentheses after the method name, such as `void nap() {}`. If you do have multiple parameters, you separate them with a comma. There are a couple more rules for the parameter list that you'll see when we cover varargs shortly. For now, let's practice looking at method signatures with "regular" parameters:

```
public void walk1() { }
public void walk2 { } // DOES NOT COMPILE
public void walk3(int a) { }
public void walk4(int a; int b) { } // DOES NOT COMPILE
public void walk5(int a, int b) { }
```

`walk1()` is a valid method declaration without any parameters. `walk2()` doesn't compile because it is missing the parentheses around the parameter list. `walk3()` is a valid method declaration with one parameter. `walk4()` doesn't compile because the parameters are separated by a semicolon rather than a comma. Semicolons are for separating statements, not parameter lists. `walk5()` is a valid method declaration with two parameters.

## Optional Exception List

In Java, code can indicate that something went wrong by throwing an exception. We'll cover this in Chapter 6, "Exceptions." For now, you just need to know that it is optional and where in the method signature it goes if present. In the example, `InterruptedException` is a type of `Exception`. You can list as many types of exceptions as you want in this clause separated by commas. For example:

```
public void zeroExceptions() { }
public void oneException() throws IllegalArgumentException { }
public void twoExceptions() throws
 IllegalArgumentException, InterruptedException { }
```

You might be wondering what methods do with these exceptions. The calling method can throw the same exceptions or handle them. You'll learn more about this in Chapter 6.

## Method Body

The final part of a method declaration is the method body (except for abstract methods and interfaces, but you don't need to know about either of those until next chapter). A method body is simply a code block. It has braces that contain zero or more Java statements. We've spent several chapters looking at Java statements by now, so you should find it easy to figure out why these compile or don't:

```
public void walk1() { }
public void walk2; // DOES NOT COMPILE
public void walk3(int a) { int name = 5; }
```

`walk1()` is a valid method declaration without an empty method body. `walk2()` doesn't compile because it is missing the braces around the empty method body. `walk3()` is a valid method declaration with one statement in the method body.

You've made it through the basics of identifying correct and incorrect method declarations. Now we can delve into more detail.

## Working with Varargs

As you saw in the previous chapter, a method may use a vararg parameter (variable argument) as if it is an array. It is a little different than an array, though. A vararg parameter must be the last element in a method's parameter list. This implies you are only allowed to have one vararg parameter per method.

Can you identify why each of these does or doesn't compile? (Yes, there is a lot of practice in this chapter. You have to be really good at identifying valid and invalid methods for the exam.)

```
public void walk1(int... nums) { }
public void walk2(int start, int... nums) { }
public void walk3(int... nums, int start) { } // DOES NOT COMPILE
public void walk4(int... start, int... nums) { } // DOES NOT COMPILE
```

`walk1()` is a valid method declaration with one vararg parameter. `walk2()` is a valid method declaration with one `int` parameter and one vararg parameter. `walk3()` and `walk4()` do not compile because they have a vararg parameter in a position that is not the last one.

When calling a method with a vararg parameter, you have a choice. You can pass in an array, or you can list the elements of the array and let Java create it for you. You can even omit the vararg values in the method call and Java will create an array of length zero for you.

Finally! We get to do something other than identify whether method declarations are valid. Instead we get to look at method calls. Can you figure out why each method call outputs what it does?

```
15: public static void walk(int start, int... nums) {
16: System.out.println(nums.length);
17: }
18: public static void main(String[] args) {
19: walk(1); // 0
20: walk(1, 2); // 1
```

```

21: walk(1, 2, 3); // 2
22: walk(1, new int[] {4, 5}); // 2
23: }

```

Line 19 passes 1 as start but nothing else. This means Java creates an array of length 0 for `nums`. Line 20 passes 1 as start and one more value. Java converts this one value to an array of length 1. Line 21 passes 1 as start and two more values. Java converts these two values to an array of length 2. Line 22 passes 1 as start and an array of length 2 directly as `nums`.

You've seen that Java will create an empty array if no parameters are passed for a vararg. However, it is still possible to pass `null` explicitly:

```
walk(1, null); // throws a NullPointerException
```

Since `null` isn't an `int`, Java treats it as an array reference that happens to be `null`. It just passes on the `null` array object to `walk`. Then the `walk()` method throws an exception because it tries to determine the length of `null`.

Accessing a vararg parameter is also just like accessing an array. It uses array indexing. For example:

```

16: public static void run(int... nums) {
17: System.out.println(nums[1]);
18: }
19: public static void main(String[] args) {
20: run(11, 22); // 22
21: }

```

Line 20 calls a vararg parameter two parameters. When the method gets called, it sees an array of size 2. Since indexes are 0 based, 22 is printed.

## Applying Access Modifiers

You already saw that there are four access modifiers: `public`, `private`, `protected`, and default access. We are going to discuss them in order from most restrictive to least restrictive:


- `private`: Only accessible within the same class
- default (package private) access: `private` and other classes in the same package
- `protected`: default access and child classes
- `public`: `protected` and classes in the other packages

### Private Access

Private access is easy. Only code in the same class can call private methods or access private fields.

Before we start, take a look at Figure 4.2. It shows the classes we'll use to explore private and default access. The big boxes are the names of the packages. The smaller boxes inside them are the classes in each package. You can refer back to this figure if you want to quickly see how the classes relate.

**FIGURE 4.2** Classes used to show private and default access



This is perfectly legal code because everything is one class:

```

1: package pond.duck;
2: public class FatherDuck {
3: private String noise = "quack";
4: private void quack() {
5: System.out.println(noise); // private access is ok
6: }
7: private void makeNoise() {
8: quack(); // private access is ok
9: } }

```

So far, so good. `FatherDuck` makes a call to private method `quack()` on line 8 and uses private instance variable `noise` on line 5.

Now we add another class:

```

1: package pond.duck;
2: public class BadDuckling {
3: public void makeNoise() {
4: FatherDuck duck = new FatherDuck();
5: duck.quack(); // DOES NOT COMPILE
6: System.out.println(duck.noise); // DOES NOT COMPILE
7: } }

```

BadDuckling is trying to access members it has no business touching. On line 5, it tries to access a private method in another class. On line 6, it tries to access a private instance variable in another class. Both generate compiler errors. Bad duckling!

Our bad duckling is only a few days old and doesn't know better yet. Luckily, you know that accessing private members of other classes is not allowed and you need to use a different type of access.

## Default (Package Private) Access

Luckily, MotherDuck is more accommodating about what her ducklings can do. She allows classes in the same package to access her members. When there is no access modifier, Java uses the default, which is package private access. This means that the member is “private” to classes in the same package. In other words, only classes in the package may access it.

```
package pond.duck;

public class MotherDuck {
 String noise = "quack";
 void quack() {
 System.out.println(noise); // default access is ok
 }
 private void makeNoise() {
 quack(); // default access is ok
 } }
```

MotherDuck can call `quack()` and refer to *noise*. After all, members in the same class are certainly in the same package. The big difference is MotherDuck lets other classes in the same package access members (due to being package private) whereas FatherDuck doesn't (due to being private). GoodDuckling has a much better experience than BadDuckling:

```
package pond.duck;

public class GoodDuckling {
 public void makeNoise() {
 MotherDuck duck = new MotherDuck();
 duck.quack(); // default access
 System.out.println(duck.noise); // default access
 } }
```

GoodDuckling succeeds in learning to `quack()` and make *noise* by copying its mother. Notice that all the classes we've covered so far are in the same package `pond.duck`. This allows default (package private) access to work.

In this same pond, a swan just gave birth to a baby swan. A baby swan is called a cygnet. The cygnet sees the ducklings learning to quack and decides to learn from MotherDuck as well.

```
package pond.swan;

import pond.duck.MotherDuck; // import another package
```

```
public class BadCygnet {
 public void makeNoise() {
 MotherDuck duck = new MotherDuck();
 duck.quack(); // DOES NOT COMPILE
 System.out.println(duck.noise); // DOES NOT COMPILE
 }
}
```

Oh no! `MotherDuck` only allows lessons to other ducks by restricting access to the `pond.duck` package. Poor little `BadCygnet` is in the `pond.swan` package and the code doesn't compile.


Remember that when there is no access modifier, only classes in the same package can access it.

### Protected Access

Protected access allows everything that default (package private) access allows and more. The protected access modifier adds the ability to access members of a parent class. We'll cover creating subclasses in depth in Chapter 5. For now, we'll cover the simplest possible use of a child class.

Figure 4.3 shows the many classes we'll create in this section. There are a number of classes and packages, so don't worry about keeping them all in your head. Just check back with this figure as you go.

**FIGURE 4.3** Classes used to show protected access



First, we create a `Bird` class and give protected access to its members:

```
package pond.shore;
public class Bird {
 protected String text = "floating"; // protected access
 protected void floatInWater() { // protected access
 System.out.println(text);
 } }

```

Next we create a subclass:

```
package pond.goose;
import pond.shore.Bird; // in a different package
public class Gosling extends Bird { // extends means create subclass
 public void swim() {
 floatInWater(); // calling protected member
 System.out.println(text); // calling protected member
 } }

```

This is a very simple subclass. It extends the `Bird` class. Extending means creating a subclass that has access to any protected or public members of the parent class. Running this code prints `floating` twice: once from calling `floatInWater()` and once from the `print` statement in `swim()`. Since `Gosling` is a subclass of `Bird`, it can access these members even though it is in a different package.

Remember that protected also gives us access to everything that default access does. This means that a class in the same package as `Bird` can access its protected members.

```
package pond.shore; // same package as Bird
public class BirdWatcher {
 public void watchBird() {
 Bird bird = new Bird();
 bird.floatInWater(); // calling protected member
 System.out.println(bird.text); // calling protected member
 } }

```

Since `Bird` and `BirdWatcher` are in the same package, `BirdWatcher` can access members of the `bird` variable. The definition of protected allows access to subclasses and classes in the same package. This example uses the same package part of that definition.

Now let's try the same thing from a different package:

```
package pond.inland;
import pond.shore.Bird; // different package than Bird
public class BirdWatcherFromAfar {
 public void watchBird() {
 Bird bird = new Bird();
 bird.floatInWater(); // DOES NOT COMPILE
 } }

```

```

 System.out.println(bird.text); // DOES NOT COMPILE
} }

```

`BirdWatcherFromAfar` is not in the same package as `Bird` and it doesn't inherit from `Bird`. This means that it is not allowed to access protected members of `Bird`.

Got that? Subclasses and classes in the same package are the only ones allowed to access protected members.

There is one gotcha for protected access. Consider this class:

```

1: package pond.swan;
2: import pond.shore.Bird; // in different package than Bird
3: public class Swan extends Bird { // but subclass of bird
4: public void swim() {
5: floatInWater(); // package access to superclass
6: System.out.println(text); // package access to superclass
7: }
8: public void helpOtherSwanSwim() {
9: Swan other = new Swan();
10: other.floatInWater(); // package access to superclass
11: System.out.println(other.text); // package access to superclass
12: }
13: public void helpOtherBirdSwim() {
14: Bird other = new Bird();
15: other.floatInWater(); // DOES NOT COMPILE
16: System.out.println(other.text); // DOES NOT COMPILE
17: }
18: }

```

Take a deep breath. This is interesting. `Swan` is not in the same package as `Bird`, but does extend it—which implies it has access to the protected members of `Bird` since it is a subclass. And it does. Lines 5 and 6 refer to protected members via inheriting them.

Lines 10 and 11 also successfully use protected members of `Bird`. This is allowed because these lines refer to a `Swan` object. `Swan` inherits from `Bird` so this is okay. It is sort of a two-phase check. The `Swan` class is allowed to use protected members of `Bird` and we are referring to a `Swan` object. Granted, it is a `Swan` object created on line 9 rather than an inherited one, but it is still a `Swan` object.

Lines 15 and 16 do *not* compile. Wait a minute. They are almost exactly the same as lines 10 and 11! There's one key difference. This time a `Bird` reference is used. It is created on line 14. `Bird` is in a different package, and this code isn't inheriting from `Bird`, so it doesn't get to use protected members. Say what now? We just got through saying repeatedly that `Swan` inherits from `Bird`. And it does. However, the variable reference isn't a `Swan`. The code just happens to be in the `Swan` class.



It's okay to be confused. This is arguably one of the most confusing points on the exam. Looking at it a different way, **the protected rules apply under two scenarios:**

- A member is used without referring to a variable. This is the case on lines 5 and 6. In this case, we are taking advantage of inheritance and protected access is allowed.
- A member is used through a variable. This is the case on lines 10, 11, 15, and 16. In this case, the rules for the reference type of the variable are what matter. If it is a subclass, protected access is allowed. This works for references to the same class or a subclass.

We're going to try this again to make sure you understand what is going on. Can you figure out why these examples don't compile?

```
package pond.goose;
import pond.shore.Bird;
public class Goose extends Bird {
 public void helpGooseSwim() {
 Goose other = new Goose();
 other.floatInWater();
 System.out.println(other.text);
 }
 public void helpOtherGooseSwim() {
 Bird other = new Goose();
 other.floatInWater(); // DOES NOT COMPILE
 System.out.println(other.text); // DOES NOT COMPILE
 } }
```

The first method is fine. In fact, it is equivalent to the Swan example. Goose extends Bird. Since we are in the Goose subclass and referring to a Goose reference, it can access protected members. The second method is a problem. Although the object happens to be a Goose, it is stored in a Bird reference. We are not allowed to refer to members of the Bird class since we are not in the same package and Bird is not a subclass of Bird.

What about this one?

```
package pond.duck;
import pond.goose.Goose;
public class GooseWatcher {
 public void watch() {
 Goose goose = new Goose();
 goose.floatInWater(); // DOES NOT COMPILE
 } }
```

This code doesn't compile because we are not in the Goose class. The floatInWater() method is declared in Bird. GooseWatcher is not in the same package as Bird, nor does it

extend Bird. Goose extends Bird. That only lets Goose refer to floatInWater() and not callers of Goose.

If this is still puzzling, try it out. Type in the code and try to make it compile. Then reread this section. Don't worry—it wasn't obvious to the authors the first time either!

## Public Access

Protected access was a tough concept. Luckily, the last type of access modifier is easy: public means anyone can access the member from anywhere.

```
package pond.duck;
public class DuckTeacher {
 public String name = "helpful"; // public access
 public void swim() { // public access
 System.out.println("swim");
 }
}
```

DuckTeacher allows access to any class that wants it. Now we can try it out:

```
package pond.goose;
import pond.duck.DuckTeacher;
public class LostDuckling {
 public void swim() {
 DuckTeacher teacher = new DuckTeacher();
 teacher.swim(); // allowed
 System.out.println("Thanks" + teacher.name); // allowed
 }
}
```

LostDuckling is able to refer to swim() and name on DuckTeacher because they are public. The story has a happy ending. LostDuckling has learned to swim and can find its parents—all because DuckTeacher made them public.

To review access modifiers, make sure you know why everything in Table 4.2 is true. Remember that a member is a method or field.

**TABLE 4.2** Access modifiers

| Can access                              | If that member is private? | If that member has default (package private) access? | If that member is protected? | If that member is public? |
|-----------------------------------------|----------------------------|------------------------------------------------------|------------------------------|---------------------------|
| Member in the same class                | Yes                        | Yes                                                  | Yes                          | Yes                       |
| Member in another class in same package | No                         | Yes                                                  | Yes                          | Yes                       |

| Can access                                                    | If that member is private? | If that member has default (package private) access? | If that member is protected? | If that member is public? |
|---------------------------------------------------------------|----------------------------|------------------------------------------------------|------------------------------|---------------------------|
| Member in a superclass in a different package                 | No                         | No                                                   | Yes                          | Yes                       |
| Method/field in a non-superclass class in a different package | No                         | No                                                   | No                           | Yes                       |

## Designing Static Methods and Fields

Except for the `main()` method, we've been looking at instance methods. Static methods don't require an instance of the class. They are shared among all users of the class. You can think of statics as being a member of the single class object that exist independently of any instances of that class.

### Does Each Class Have Its Own Copy of the Code?

Each class has a copy of the instance variables. There is only one copy of the code for the instance methods. Each instance of the class can call it as many times as it would like. However, each call of an instance method (or any method) gets space on the stack for method parameters and local variables.

The same thing happens for static methods. There is one copy of the code. Parameters and local variables go on the stack.

Just remember that only data gets its "own copy." There is no need to duplicate copies of the code itself.

We have seen one static method since Chapter 1. The `main()` method is a static method. That means you can call it by the classname.

```
public class Koala {
 public static int count = 0; // static variable
 public static void main(String[] args) { // static method
 System.out.println(count);
 }
}
```

We said that the JVM basically calls `Koala.main()` to get the program started. You can do this too. We can have a `KoalaTester` that does nothing but call the `main()` method.

```
public class KoalaTester {
 public static void main(String[] args) {
 Koala.main(new String[0]); // call static method
 }
}
```

Quite a complicated way to print 0, isn't it? When we run `KoalaTester`, it makes a call to the `main()` method of `Koala`, which prints the value of `count`. The purpose of all these examples is to show that `main()` can be called just like any other static method.

In addition to `main()` methods, static methods have two main purposes:

- For utility or helper methods that don't require any object state. Since there is no need to access instance variables, having static methods eliminates the need for the caller to instantiate the object just to call the method.
- For state that is shared by all instances of a class, like a counter. All instances must share the same state. Methods that merely use that state should be static as well.

Let's look at some examples covering other static concepts.

## Calling a Static Variable or Method

Usually, accessing a static member is easy. You just put the classname before the method or variable and you are done. For example:

```
System.out.println(Koala.count);
Koala.main(new String[0]);
```

Both of these are nice and easy. There is one rule that is trickier. You can use an instance of the object to call a static method. The compiler checks for the type of the reference and uses that instead of the object—which is sneaky of Java. This code is perfectly legal:

```
5: Koala k = new Koala();
6: System.out.println(k.count); // k is a Koala
7: k = null;
8: System.out.println(k.count); // k is still a Koala
```

Believe it or not, this code outputs 0 twice. Line 6 sees that `k` is a `Koala` and `count` is a static variable, so it reads that static variable. Line 8 does the same thing. Java doesn't care that `k` happens to be `null`. Since we are looking for a static, it doesn't matter.



Remember to look at the reference type for a variable when you see a static method or variable. The exam creators will try to trick you into thinking a `NullPointerException` is thrown because the variable happens to be `null`. Don't be fooled!

One more time because this is really important: what does the following output?

```
Koala.count = 4;
Koala koala1 = new Koala();
Koala koala2 = new Koala();
koala1.count = 6;
koala2.count = 5;
System.out.println(Koala.count);
```

Hopefully, you answered 5. There is only one *count* variable since it is static. It is set to 4, then 6, and finally winds up as 5. All the *Koala* variables are just distractions.

## Static vs. Instance

There's another way the exam creators will try to trick you regarding static and instance members. (Remember that “member” means field or method.) A static member cannot call an instance member. This shouldn't be a surprise since static doesn't require any instances of the class to be around.

The following is a common mistake for rookie programmers to make:

```
public class Static {
 private String name = "Static class";
 public static void first() { }
 public static void second() { }
 public void third() { System.out.println(name); }
 public static void main(String args[]) {
 first();
 second();
 third(); // DOES NOT COMPILE
 } }
}
```

The compiler will give you an error about making a static reference to a nonstatic method. If we fix this by adding *static* to *third()*, we create a new problem. Can you figure out what it is?

All this does is move the problem. Now, *third()* is referring to nonstatic *name*. Adding *static* to *name* as well would solve the problem. Another solution would have been to call *third* as an instance method—for example, *new Static().third()*;

The exam creators like this topic. A static method or instance method can call a static method because static methods don't require an object to use. Only an instance method can call another instance method on the same class without using a reference variable, because instance methods do require an object. Similar logic applies for the instance and static variables. Make sure you understand Table 4.3 before continuing.

**TABLE 4.3** Static vs. instance calls

| Type            | Calling                             | Legal? | How?                                        |
|-----------------|-------------------------------------|--------|---------------------------------------------|
| Static method   | Another static method or variable   | Yes    | Using the classname                         |
| Static method   | An instance method or variable      | No     |                                             |
| Instance method | A static method or variable         | Yes    | Using the classname or a reference variable |
| Instance method | Another instance method or variable | Yes    | Using a reference variable                  |

Let’s try one more example so you have more practice at recognizing this scenario. Do you understand why the following lines fail to compile?

```
1: public class Gorilla {
2: public static int count;
3: public static void addGorilla() { count++; }
4: public void babyGorilla() { count++; }
5: public void announceBabies() {
6: addGorilla();
7: babyGorilla();
8: }
9: public static void announceBabiesToEveryone() {
10: addGorilla();
11: babyGorilla(); // DOES NOT COMPILE
12: }
13: public int total;
14: public static average = total / count; // DOES NOT COMPILE
15: }
```

Lines 3 and 4 are fine because both static and instance methods can refer to a static variable. Lines 5–8 are fine because an instance method can call a static method. Line 11 doesn’t compile because a static method cannot call an instance method. Similarly, line 14 doesn’t compile because a static variable is trying to use an instance variable.

A common use for static variables is counting the number of instances:

```
public class Counter {
 private static int count;
 public Counter() { count++; }
```

```

public static void main(String[] args) {
 Counter c1 = new Counter();
 Counter c2 = new Counter();
 Counter c3 = new Counter();
 System.out.println(count); // 3
}
}

```

Each time the constructor gets called, it increments *count* by 1. This example relies on the fact that static (and instance) variables are automatically initialized to the default value for that type, which is 0 for `int`. See Chapter 1 to review the default values.

Also notice that we didn't write `Counter.count`. We could have. It isn't necessary because we are already in that class so the compiler can infer it.

## Static Variables

Some static variables are meant to change as the program runs. Counters are a common example of this. We want the count to increase over time. Just as with instance variables, you can initialize a static variable on the line it is declared:

```

public class Initializers {
 private static int counter = 0; // initialization
}

```

Other static variables are meant to never change during the program. This type of variable is known as a *constant*. It uses the `final` modifier to ensure the variable never changes. `static final` constants use a different naming convention than other variables. They use all uppercase letters with underscores between “words.” For example:

```

public class Initializers {
 private static final int NUM_BUCKETS = 45;
 public static void main(String[] args) {
 NUM_BUCKETS = 5; // DOES NOT COMPILE
 } }

```

The compiler will make sure that you do not accidentally try to update a `final` variable. This can get interesting. Do you think the following compiles?

```

private static final ArrayList<String> values = new ArrayList<>();
public static void main(String[] args) {
 values.add("changed");
}

```

It actually does compile. *values* is a reference variable. We are allowed to call methods on reference variables. All the compiler can do is check that we don't try to reassign the `final values` to point to a different object.

## Static Initialization

In Chapter 1, we covered instance initializers that looked like unnamed methods. Just code inside braces. Static initializers look similar. They add the `static` keyword to specify they should be run when the class is first used. For example:

```
private static final int NUM_SECONDS_PER_HOUR;
static {
 int numSecondsPerMinute = 60;
 int numMinutesPerHour = 60;
 NUM_SECONDS_PER_HOUR = numSecondsPerMinute * numMinutesPerHour;
}
```

The static initializer runs when the class is first used. The statements in it run and assign any static variables as needed. There is something interesting about this example. We just got through saying that final variables aren't allowed to be reassigned. The key here is that the static initializer is the first assignment. And since it occurs up front, it is okay.

Let's try another example to make sure you understand the distinction:

```
14: private static int one;
15: private static final int two;
16: private static final int three = 3;
17: private static final int four; // DOES NOT COMPILE
18: static {
19: one = 1;
20: two = 2;
21: three = 3; // DOES NOT COMPILE
22: two = 4; // DOES NOT COMPILE
23: }
```

Line 14 declares a static variable that is not final. It can be assigned as many times as we like. Line 15 declares a final variable without initializing it. This means we can initialize it exactly once in a static block. Line 22 doesn't compile because this is the second attempt. Line 16 declares a final variable and initializes it at the same time. We are not allowed to assign it again, so line 21 doesn't compile. Line 17 declares a final variable that never gets initialized. The compiler gives a compiler error because it knows that the static blocks are the only place the variable could possibly get initialized. Since the programmer forgot, this is clearly an error.

### Try to Avoid Static and Instance Initializers

Using static and instance initializers can make your code much harder to read. Everything that could be done in an instance initializer could be done in a constructor instead. The constructor approach is easier to read.



There is a common case to use a static initializer: when you need to initialize a static field and the code to do so requires more than one line. This often occurs when you want to initialize a collection like an `ArrayList`. When you do need to use a static initializer, put all the static initialization in the same block. That way, the order is obvious.

## Static Imports

Back in Chapter 1, you saw that we could import a specific class or all the classes in a package:

```
import java.util.ArrayList;
import java.util.*;
```

We could use this technique to import:

```
import java.util.List;
import java.util.Arrays;
public class Imports {
 public static void main(String[] args) {
 List<String> list = Arrays.asList("one", "two");
 }
}
```

Imports are convenient because you don't need to specify where each class comes from each time you use it. There is another type of import called a static import. Regular imports are for importing classes. Static imports are for importing static members of classes. Just like regular imports, you can use a wildcard or import a specific member. The idea is that you shouldn't have to specify where each static method or variable comes from each time you use it. An example of when static interfaces shine are when you are referring to a lot of constants in another class.



In a large program, static imports can be overused. When importing from too many places, it can be hard to remember where each static member comes from.

The previous method has one static method call: `Arrays.asList`. Rewriting the code to use a static import yields the following:

```
import java.util.List;
import static java.util.Arrays.asList; // static import
public class StaticImports {
 public static void main(String[] args) {
 List<String> list = asList("one", "two"); // no Arrays.
 } }
```

In this example, we are specifically importing the `asList` method. This means that any time we refer to `asList` in the class, it will call `Arrays.asList()`.

An interesting case is what would happen if we created an `asList` method in our `StaticImports` class. Java would give it preference over the imported one and the method we coded would be used.

The exam will try to trick you with misusing static imports. This example shows almost everything you can do wrong. Can you figure out what is wrong with each one?

```
1: import static java.util.Arrays; // DOES NOT COMPILE
2: import static java.util.Arrays.asList;
3: static import java.util.Arrays.*; // DOES NOT COMPILE
4: public class BadStaticImports {
5: public static void main(String[] args) {
6: Arrays.asList("one"); // DOES NOT COMPILE
7: } }
```

Line 1 tries to use a static import to import a class. Remember that static imports are only for importing static members. Regular imports are for importing a class. Line 3 tries to see if you are paying attention to the order of keywords. The syntax is `import static` and not vice versa. Line 6 is sneaky. We imported the `asList` method on line 2. However, we did not import the `Arrays` class anywhere. This makes it okay to write `asList("one");` but not `Arrays.asList("one");`.

There's only one more scenario with static imports. In Chapter 1, you learned that importing two classes with the same name gives a compiler error. This is true of static imports as well. The compiler will complain if you try to explicitly do a static import of two methods with the same name or two static variables with the same name. For example:

```
import static statics.A.TYPE;
import static statics.B.TYPE; // DOES NOT COMPILE
```

Luckily when this happens, we can just refer to the static members via their classname in the code instead of trying to use a static import.

## Passing Data Among Methods

Java is a “pass-by-value” language. This means that a copy of the variable is made and the method receives that copy. Assignments made in the method do not affect the caller. Let's look at an example:

```
2: public static void main(String[] args) {
3: int num = 4;
4: newNumber(5);
```

```
5: System.out.println(num); // 4
6: }
7: public static void newNumber(int num) {
8: num = 8;
9: }
```

On line 3, `num` is assigned the value of 4. On line 4, we call a method. On line 8, the `num` parameter in the method gets set to 8. Although this parameter has the same name as the variable on line 3, this is a coincidence. The name could be anything. The exam will often use the same name to try to confuse you. The variable on line 3 never changes because no assignments are made to it.

Now that you've seen primitives, let's try an example with a reference type. What do you think is output by the following code?


```
public static void main(String[] args) {
 String name = "Webby";
 speak(name);
 System.out.println(name);
}
public static void speak(String name) {
 name = "Sparky";
}
```

The correct answer is `Webby`. Just as in the primitive example, the variable assignment is only to the method parameter and doesn't affect the caller.

Notice how we keep talking about variable assignments. This is because we can call methods on the parameters. As an example, we have code that calls a method on the `StringBuilder` passed into the method:

```
public static void main(String[] args) {
 StringBuilder name = new StringBuilder();
 speak(name);
 System.out.println(name); // Webby
}
public static void speak(StringBuilder s) {
 s.append("Webby");
}
```

In this case, the output is `Webby` because the method merely calls a method on the parameter. It doesn't reassign `name` to a different object. In Figure 4.4, you can see how pass-by-value is still used. `s` is a copy of the variable `name`. Both point to the same `StringBuilder`, which means that changes made to the `StringBuilder` are available to both references.

**FIGURE 4.4** Copying a reference with pass-by-value

## Real World Scenario

### Pass-by-Value vs. Pass-by-Reference

Different languages handle parameters in different ways. Pass-by-value is used by many languages, including Java. In this example, the swap method does not change the original values. It only changes *a* and *b* within the method.

```
public static void main(String[] args) {
 int original1 = 1;
 int original2 = 2;
 swap(original1, original2);
 System.out.println(original1); // 1
 System.out.println(original2); // 2
}

public static void swap(int a, int b) {
 int temp = a;
 a = b;
 b = temp;
}
```

The other approach is pass-by-reference. It is used by default in a few languages, such as Perl. We aren't going to show you Perl code here because you are studying for the Java exam and we don't want to confuse you. The following example is in a made-up language that shows pass-by-reference:

```
original1 = 1;
original2 = 2;
swapByReference(original1, original2);
print(original1); // 2 (not in Java)
print(original2); // 1 (not in Java)

swapByReference(a, b) {
 temp = a;
 a = b;
 b = temp;
}
```

See the difference? In our made-up language, the caller is affected by variable assignments made in the method.

To review, Java uses pass-by-value to get data into a method. Assigning a new primitive or reference to a parameter doesn't change the caller. Calling methods on a reference to an object does affect the caller.

Getting data back from a method is easier. A copy is made of the primitive or reference and returned from the method. Most of the time, this returned value is used. For example, it might be stored in a variable. If the returned value is not used, the result is ignored. Watch for this on the exam. Ignored returned values are tricky.

Let's try an example. Pay attention to the return types.

```

1: public class ReturningValues {
2: public static void main(String[] args) {
3: int number = 1; // 1
4: String letters = "abc"; // abc
5: number(number); // 1
6: letters = letters(letters); // abcd
7: System.out.println(number + letters); // 1abcd
8: }
9: public static int number(int number) {
10: number++;
11: return number;
12: }
13: public static String letters(String letters) {
14: letters += "d";
15: return letters;
16: }
17: }

```

This is a tricky one because there is a lot to keep track of. When you see such questions on the exam, write down the values of each variable. Lines 3 and 4 are straightforward assignments. Line 5 calls a method. Line 10 increments the method parameter to 2 but leaves the *numbers* variable in the *main()* method as 1. While line 11 returns the value, the caller ignores it. The method call on line 6 doesn't ignore the result so *letters* becomes "abcd". Remember that this is happening because of the returned value and not the method parameter.

## Overloading Methods

Now that you are familiar with the rules for declaring methods, it is time to look at creating methods with the same signature in the same class. *Method overloading* occurs when there are different method signatures with the same name but different type parameters.

We've been calling overloaded methods for a while. `System.out.println` and `StringBuilder`'s `append` methods provide many overloaded versions so you can pass just about anything to them without having to think about it. In both of these examples, the

only change was the type of the parameter. Overloading also allows different numbers of parameters.

Everything other than the method signature can vary for overloaded methods. This means there can be different access modifiers, specifiers (like static), return types, and exception lists.

These are all valid overloaded methods:

```
public void fly(int numMiles) { }
public void fly(short numFeet) { }
public boolean fly() { return false; }
void fly(int numMiles, short numFeet) { }
public void fly(short numFeet, int numMiles) throws Exception { }
```

As you can see, we can overload by changing anything in the parameter list. We can have a different type, more types, or the same types in a different order. Also notice that the access modifier and exception list are irrelevant to overloading.

Now let's look at an example that is not valid overloading:

```
public void fly(int numMiles) { }
public int fly(int numMiles) { } // DOES NOT COMPILE
```

This method doesn't compile because it only differs from the original by return type. The parameter lists are the same so they are duplicate methods as far as Java is concerned.

What about these two? Why does the second not compile?

```
public void fly(int numMiles) { }
public static void fly(int numMiles) { } // DOES NOT COMPILE
```

Again, the parameter list is the same. The only difference is that one is an instance method and one is a static method.

Calling overloaded methods is easy. You just write code and Java calls the right one. For example, look at these two methods:

```
public void fly(int numMiles) {
 System.out.println("short");
}
public void fly(short numFeet) {
 System.out.println("short");
}
```

The call `fly((short) 1);` prints short. It looks for matching types and calls the appropriate method. Of course, it can be more complicated than this.

Now that you know the basics of overloading, let's look at some more complex scenarios that you may encounter on the exam.

## Overloading and Varargs

Which method do you think is called if we pass an `int[]`?

```
public void fly(int[] lengths) { }
public void fly(int... lengths) { } // DOES NOT COMPILE
```

Trick question! Remember that Java treats varargs as if they were an array. This means that the method signature is the same for both methods. Since we are not allowed to overload methods with the same parameter list, this code doesn't compile. Even though the code doesn't look the same, it compiles to the same parameter list.

Now that we've just gotten through explaining that they are the same, it is time to mention how they are not the same. It shouldn't be a surprise that you can call either method by passing an array:

```
fly(new int[] { 1, 2, 3 });
```

However, you can only call the varargs version with stand-alone parameters:

```
fly(1, 2, 3);
```

Obviously, this means they don't compile *exactly* the same. The parameter list is the same, though, and that is what you need to know with respect to overloading for the exam.

## Autoboxing

In the previous chapter, you saw how Java will convert a primitive `int` to an object `Integer` to add it to an `ArrayList` through the wonders of autoboxing. This works for code you write too.

```
public void fly(Integer numMiles) { }
```

This means calling `fly(3);` will call the previous method as expected. However, what happens if we have both a primitive and an integer version?

```
public void fly(int numMiles) { }
public void fly(Integer numMiles) { }
```

Java will match the `int numMiles` version. Java tries to use the most specific parameter list it can find. When the primitive `int` version isn't present, it will autobox. However, when the primitive `int` version is provided, there is no reason for Java to do the extra work of autoboxing.

## Reference Types

Given the rule about Java picking the most specific version of a method that it can, what do you think this code outputs?

```
public class ReferenceTypes {
 public void fly(String s) {
 System.out.print("string ");
 }

 public void fly(Object o) {
 System.out.print("object ");
 }

 public static void main(String[] args) {
 ReferenceTypes r = new ReferenceTypes();
```

```
 r.fly("test");
 r.fly(56);
} }
```

The answer is "string object". The first call is a `String` and finds a direct match. There's no reason to use the `Object` version when there is a nice `String` parameter list just waiting to be called. The second call looks for an `int` parameter list. When it doesn't find one, it autoboxes to `Integer`. Since it still doesn't find a match, it goes to the `Object` one.

## Primitives

Primitives work in a way similar to reference variables. Java tries to find the most specific matching overloaded method. What do you think happens here?

```
public class Plane {
 public void fly(int i) {
 System.out.print("int ");
 }
 public void fly(long l) {
 System.out.print("long ");
 }
 public static void main(String[] args) {
 Plane p = new Plane();
 p.fly(123);
 p.fly(123L);
 } }
```

The answer is `int long`. The first call passes an `int` and sees an exact match. The second call passes a `long` and also sees an exact match. If we comment out the overloaded method with the `int` parameter list, the output becomes `long long`. Java has no problem calling a larger primitive. However, it will not do so unless a better match is not found.

Note that Java can only accept wider types. An `int` can be passed to a method taking a `long` parameter. Java will not automatically convert to a narrower type. If you want to pass a `long` to a method taking an `int` parameter, you have to add a cast to explicitly say narrowing is okay.

## Putting It All Together

So far, all the rules for when an overloaded method is called should be logical. Java calls the most specific method it can. When some of the types interact, the Java rules focus on backward compatibility. In Java 1.4 and earlier, autoboxing and varargs didn't exist. Although that was a long time ago, old code still needs to work—which means autoboxing and varargs come last when Java looks at overloaded methods. Ready for the official order? Table 4.4 lays it out for you.



**TABLE 4.4** Order Java uses to choose the right overloaded method

| Rule                  | Example of what will be chosen for <code>glide(1,2)</code> |
|-----------------------|------------------------------------------------------------|
| Exact match by type   | <code>public String glide(int i, int j) {}</code>          |
| Larger primitive type | <code>public String glide(long i, long j) {}</code>        |
| Autoboxed type        | <code>public String glide(Integer i, Integer j) {}</code>  |
| Varargs               | <code>public String glide(int... nums) {}</code>           |

Let's give this a practice run using the rules in Table 4.4. What do you think this outputs?

```
public class Glider2 {
 public static String glide(String s) {
 return "1";
 }
 public static String glide(String... s) {
 return "2";
 }
 public static String glide(Object o) {
 return "3";
 }
 public static String glide(String s, String t) {
 return "4";
 }
 public static void main(String[] args) {
 System.out.print(glide("a"));
 System.out.print(glide("a", "b"));
 System.out.print(glide("a", "b", "c"));
 } }
```

It prints out 142. The first call matches the signature taking a single `String` because that is the most specific match. The second call matches the signature, taking two `String` parameters since that is an exact match. It isn't until the third call that the varargs version is used since there are no better matches.

As accommodating as Java is with trying to find a match, it will only do one conversion:

```
public class TooManyConversions {
 public static void play(Long l) { }
```

```
public static void play(Long... l) { }
public static void main(String[] args) {
 play(4); // DOES NOT COMPILE
 play(4L); // calls the Long version
} }
```

Here we have a problem. Java is happy to convert the `int 4` to a `long 4` or an `Integer 4`. It cannot handle converting in two steps to a `long` and then to a `Long`. If we had `public static void play(Object o) { }`, it would match because only one conversion would be necessary: from `int` to `Integer`. An `Integer` is an `Object`, as you'll see in Chapter 5.

## Creating Constructors

As you learned in Chapter 1, a constructor is a special method that matches the name of the class and has no return type. Here's an example:

```
public class Bunny {
 public Bunny() {
 System.out.println("constructor");
 }
}
```

The name of the constructor, `Bunny`, matches the name of the class, `Bunny`, and there is no return type, not even `void`. That makes this a constructor. Can you tell why these two are not valid constructors for the `Bunny` class?

```
public bunny() { } // DOES NOT COMPILE
public void Bunny() { }
```

The first one doesn't match the classname because Java is case sensitive. Since it doesn't match, Java knows it can't be a constructor and is supposed to be a regular method. However, it is missing the return type and doesn't compile. The second method is a perfectly good method, but is not a constructor because it has a return type.

Constructors are used when creating a new object. This process is called *instantiation* because it creates a new instance of the class. A constructor is called when we write `new` followed by the name of the class we want to instantiate. For example:

```
new Bunny()
```

When Java sees the `new` keyword, it allocates memory for the new object. Java also looks for a constructor and calls it.

A constructor is typically used to initialize instance variables. The `this` keyword tells Java you want to reference an instance variable. Most of the time, `this` is optional. The problem is that sometimes there are two variables with the same name. In a constructor, one is a parameter and one is an instance variable. If you don't say otherwise, Java gives

you the one with the most granular scope, which is the parameter. Using `this.name` tells Java you want the instance variable.

Here's a common way of writing a constructor:

```
1: public class Bunny {
2: private String color;
3: public Bunny(String color) {
4: this.color = color;
5: } }
```

On line 4, we assign the parameter `color` to the instance variable `color`. The right side of the assignment refers to the parameter because we don't specify anything special. The left side of the assignment uses `this` to tell Java we want it to use the instance variable.

Now let's look at some examples that aren't common but that you might see on the exam:

```
1: public class Bunny {
2: private String color;
3: private int height;
4: private int length;
5: public Bunny(int length, int theHeight) {
6: length = this.length; // backwards - no good!
7: height = theHeight; // fine because a different name
8: this.color = "white"; // fine, but redundant
9: }
10: public static void main(String[] args) {
11: Bunny b = new Bunny(1, 2);
12: System.out.println(b.length + " " + b.height + " " + b.color);
13: } }
```

Line 6 is incorrect and you should watch for it on the exam. The instance variable `length` starts out with a 0 value. That 0 is assigned to the method parameter `length`. The instance variable stays at 0. Line 7 is more straightforward. The parameter `theHeight` and instance variable `height` have different names. Since there is no naming collision, `this` is not required. Finally, line 8 shows that it is allowed to use `this` even when there is no duplication of variable names.

In this section, we'll look at default constructors, overloading constructors, final fields, and the order of initialization in a class.

## Default Constructor

Every class in Java has a constructor whether you code one or not. If you don't include any constructors in the class, Java will create one for you without any parameters.

This Java-created constructor is called the *default constructor*. Sometimes we call it the default no-arguments constructor for clarity. Here's an example:

```
public class Rabbit {
 public static void main(String[] args) {
 Rabbit rabbit = new Rabbit(); // Calls default constructor
 }
}
```

In the `Rabbit` class, Java sees no constructor was coded and creates one. This default constructor is equivalent to typing this:

```
public Rabbit() {}
```

The default constructor has an empty parameter list and an empty body. It is fine for you to type this in yourself. However, since it doesn't do anything, Java is happy to supply it for you and save you some typing.

We keep saying *generated*. This happens during the compile step. If you look at the file with the `.java` extension, the constructor will still be missing. It is only in the compiled file with the `.class` extension that it makes an appearance.

Remember that a default constructor is only supplied if there are no constructors present. Which of these classes do you think has a default constructor?

```
class Rabbit1 {
}
class Rabbit2 {
 public Rabbit2() { }
}
class Rabbit3 {
 public Rabbit3(boolean b) { }
}
class Rabbit4 {
 private Rabbit4() { }
}
```

Only `Rabbit1` gets a default no-argument constructor. It doesn't have a constructor coded so Java generates a default no-argument constructor. `Rabbit2` and `Rabbit3` both have public constructors already. `Rabbit4` has a private constructor. Since these three classes have a constructor defined, the default no-argument constructor is not inserted for you.

Let's take a quick look at how to call these constructors:

```
1: public class RabbitsMultiply {
2: public static void main(String[] args) {
```

```

3: Rabbit1 r1 = new Rabbit1();
4: Rabbit2 r2 = new Rabbit2();
5: Rabbit3 r3 = new Rabbit3(true);
6: Rabbit4 r4 = new Rabbit4(); // DOES NOT COMPILE
7: } }

```

Line 3 calls the generated default no-argument constructor. Lines 4 and 5 call the user-provided constructors. Line 6 does not compile. `Rabbit4` made the constructor private so that other classes could not call it.

Having a private constructor in a class tells the compiler not to provide a default no-argument constructor. It also prevents other classes from instantiating the class. This is useful when a class only has static methods or the class wants to control all calls to create new instances of itself.

## Overloading Constructors

Up to now, you’ve only seen one constructor per class. You can have multiple constructors in the same class as long as they have different method signatures. When overloading methods, the method name and parameter list needed to match. With constructors, the name is always the same since it has to be the same as the name of the class. This means constructors must have different parameters in order to be overloaded.

This example shows two constructors:

```

public class Hamster {
 private String color;
 private int weight;
 public Hamster(int weight) { // first constructor
 this.weight = weight;
 color = "brown";
 }
 public Hamster(int weight, String color) { // second constructor
 this.weight = weight;
 this.color = color;
 }
}

```

One of the constructors takes a single `int` parameter. The other takes an `int` and a `String`. These parameter lists are different, so the constructors are successfully overloaded.

There is a problem here, though. There is a bit of duplication. In programming, even a bit of duplication tends to turn into a lot of duplication as we keep adding “just one more

thing.” What we really want is for the first constructor to call the second constructor with two parameters. You might be tempted to write this:

```
public Hamster(int weight) {
 Hamster(weight, "brown"); // DOES NOT COMPILE
}
```

This will not work. Constructors can be called only by writing `new` before the name of the constructor. They are not like normal methods that you can just call. What happens if we stick `new` before the constructor name?

```
public Hamster(int weight) {
 new Hamster(weight, "brown"); // Compiles but does not do what we want
}
```

This attempt does compile. It doesn't do what we want, though. When the constructor with one parameter is called, it creates an object with the default weight and color. It then constructs a different object with the desired weight and color and ignores the new object. That's not what we want. We want weight and color set on the object we are trying to instantiate in the first place.

Java provides a solution: `this`—yes, the same keyword we used to refer to instance variables. When `this` is used as if it were a method, Java calls another constructor on the same instance of the class.

```
public Hamster(int weight) {
 this(weight, "brown");
}
```

Success! Now Java calls the constructor that takes two parameters. `weight` and `color` get set on this instance.

`this()` has one special rule you need to know. If you choose to call it, the `this()` call must be the first noncommented statement in the constructor.

```
3: public Hamster(int weight) {
4: System.out.println("in constructor");
5: // ready to call this
6: this(weight, "brown"); // DOES NOT COMPILE
7: }
```

Even though a print statement on line 4 doesn't change any variables, it is still a Java statement and is not allowed to be inserted before the call to `this()`. The comment on line 5 is just fine. Comments don't run Java statements and are allowed anywhere.



## Real World Scenario

### Constructor Chaining

Overloaded constructors often call each other. One common technique is to have each constructor add one parameter until getting to the constructor that does all the work. This approach is called *constructor chaining*. In this example, all three constructors are chained.

```
public class Mouse {
 private int numTeeth;
 private int numWhiskers;
 private int weight;

 public Mouse(int weight) {
 this(weight, 16); // calls constructor with 2 parameters
 }
 public Mouse(int weight, int numTeeth) {
 this(weight, numTeeth, 6); // calls constructor with 3 parameters
 }
 public Mouse(int weight, int numTeeth, int numWhiskers) {
 this.weight = weight;
 this.numTeeth = numTeeth;
 this.numWhiskers = numWhiskers;
 }
 public void print() {
 System.out.println(weight + " " + numTeeth + " " + numWhiskers);
 }
 public static void main(String[] args) {
 Mouse mouse = new Mouse(15);
 mouse.print();
 }
}
```

This code prints 15 16 6. The `main()` method calls the constructor with one parameter. That constructor adds a second hard-coded value and calls the constructor with two parameters. That constructor adds one more hard-coded value and calls the constructor with three parameters. The three-parameter constructor assigns the instance variables.

## Final Fields

As you saw earlier in the chapter, `final` instance variables must be assigned a value exactly once. We saw this happen in the line of the declaration and in an instance initializer. There is one more location this assignment can be done: in the constructor.

```
public class MouseHouse {
 private final int volume;
 private final String name = "The Mouse House";
 public MouseHouse(int length, int width, int height) {
 volume = length * width * height;
 }
}
```

The constructor is part of the initialization process, so it is allowed to assign `final` instance variables in it. By the time the constructor completes, all `final` instance variables must have been set.

## Order of Initialization

Chapter 1 covered the order of initialization. Now that you’ve learned about static initializers, it is time to revisit that. Unfortunately, you do have to memorize this list. We’ll give you lots of practice, but you do need to know this order by heart.

1. If there is a superclass, initialize it first (we’ll cover this rule in the next chapter. For now, just say “no superclass” and go on to the next rule.)
2. Static variable declarations and static initializers in the order they appear in the file.
3. Instance variable declarations and instance initializers in the order they appear in the file.
4. The constructor.

Let’s try the first example:

```
1: public class InitializationOrderSimple {
2: private String name = "Torchie";
3: { System.out.println(name); }
4: private static int COUNT = 0;
5: static { System.out.println(COUNT); }
6: static { COUNT += 10; System.out.println(COUNT); }
7: public InitializationOrderSimple() {
8: System.out.println("constructor");
9: } }
```

```
1: public class CallInitializationOrderSimple {
2: public static void main(String[] args) {
```



```

3: InitializationOrderSimple init = new InitializationOrderSimple();
4: } }

```

The output is:

```

0
10
Torchie
constructor

```

Let's look at why. Rule 1 doesn't apply because there is no superclass. Rule 2 says to run the static variable declarations and static initializers—in this case, lines 5 and 6, which output 0 and 10. Rule 3 says to run the instance variable declarations and instance initializers—here, lines 2 and 3, which output Torchie. Finally, rule 4 says to run the constructor—here, lines 7–9, which output constructor.

The next example is a little harder. Keep in mind that the four rules apply only if an object is instantiated. If the class is referred to without a new call, only rules 1 and 2 apply. The other two rules relate to instances and constructors. They have to wait until there is code to instantiate the object.

What do you think happens here?

```

1: public class InitializationOrder {
2: private String name = "Torchie";
3: { System.out.println(name); }
4: private static int COUNT = 0;
5: static { System.out.println(COUNT); }
6: { COUNT++; System.out.println(COUNT); }
7: public InitializationOrder() {
8: System.out.println("constructor");
9: }
10: public static void main(String[] args) {
11: System.out.println("read to construct");
12: new InitializationOrder();
13: }
14: }

```

The output looks like this:

```

0
read to construct
Torchie
1
constructor

```

Again, rule 1 doesn't apply because there is no superclass. Rule 2 tells us to look at the static variables and static initializers—lines 4 and 5, in that order. Line 5 outputs 0. Now that the statics are out of the way, the `main()` method can run. Next, we can use rule 3 to run the instance variables and instance initializers. Here that is lines 2 and 3, which output Torchie. Finally, rule 4 says to run the constructor—in this case, lines 7–9, which output constructor.

We are going to try one more example. This one is as hard as it gets. If you understand the output of this next one, congratulations on a job well done; if not, don't worry. Write some programs to play with this. Try typing in the examples in this section and making minor changes to see what happens. For example, you might try commenting out part of the code. This will give you a better feel for what is going on. Then come back and reread this section to try the examples.

Ready for the tough example? Here it is:

```
1: public class YetMoreInitializationOrder {
2: static { add(2); }
3: static void add(int num) { System.out.print(num + " "); }
4: YetMoreInitializationOrder() { add(5); }
5: static { add(4); }
6: { add(6); }
7: static { new YetMoreInitializationOrder(); }
8: { add(8); }
9: public static void main(String[] args) { } }
```

The correct answer is 2 4 6 8 5. Let's walk through why that is. There is no superclass, so we jump right to rule 2—the statics. There are three static blocks: on lines 2, 5, and 7. They run in that order. The static block on line 2 calls the `add()` method, which prints 2. The static block on line 5 calls the `add()` method, which prints 4. The last static block, on line 7, calls `new` to instantiate the object. This means we can go on to rule 3 to look at the instance variables and instance initializers. There are two of those: on lines 6 and 8. They both call the `add()` method and print 6 and 8, respectively. Finally, we go on to rule 4 and call the constructor, which calls the `add()` method one more time and prints 5.

This example is tricky for a few reasons. There's a lot to keep track of. Also, the question has a lot of one-line code blocks and methods, making it harder to visualize which is a block. Luckily, questions like this are rare on the exam. If you see one, just write down what is going on as you read the code.

# Encapsulating Data

In Chapter 1, we had an example of a class with a field that wasn't private:

```
public class Swan {
 int numberEggs; // instance variable
}
```

Why do we care? Since there is default (package private) access, that means any class in the package can set *numberEggs*. We no longer have control of what gets set in our own class. A caller could even write this:

```
mother.numberEggs = -1;
```

This is clearly no good. We do not want the mother Swan to have a negative number of eggs!

*Encapsulation* to the rescue. Encapsulation means we set up the class so only methods in the class with the variables can refer to the instance variables. Callers are required to use these methods. Let's take a look at our newly encapsulated Swan class:

```
1: public class Swan {
2: private int numberEggs; // private
3: public int getNumberEggs() { // getter
4: return numberEggs;
5: }
6: public void setNumberEggs(int numberEggs) { // setter
7: if (numberEggs >= 0) // guard condition
8: this.numberEggs = numberEggs;
9: } }
```

Note *numberEggs* is now private on line 2. This means only code within the class can read or write the value of *numberEggs*. Since we wrote the class, we know better than to set a negative number of eggs. We added a method on lines 3–5 to read the value, which is called an *accessor method* or a *getter*. We also added a method on lines 6–9 to update the value, which is called a *mutator method* or a *setter*. The setter has an *if* statement in this example to prevent setting the instance variable to an invalid value. This guard condition protects the instance variable.

On line 8, we used the *this* keyword that we saw in constructors to differentiate between the method parameter *numberEggs* and the instance variable *numberEggs*.

For encapsulation, remember that data (an instance variable) is private and getters/setters are public. Java defines a naming convention that is used in *JavaBeans*. JavaBeans are reusable software components. JavaBeans call an instance variable a *property*. The only thing you need to know about JavaBeans for the exam is the naming conventions listed in Table 4.5.

**TABLE 4.5** Rules for JavaBeans naming conventions

| Rule                                                                                                                                                                  | Example                                                                            |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------|
| Properties are private.                                                                                                                                               | <pre>private int numEggs;</pre>                                                    |
| Getter methods begin with <code>is</code> if the property is a boolean.                                                                                               | <pre>public boolean isHappy() {<br/>    return happy;<br/>}</pre>                  |
| Getter methods begin with <code>get</code> if the property is not a boolean.                                                                                          | <pre>public int getNumEggs() {<br/>    return numEggs;<br/>}</pre>                 |
| Setter methods begin with <code>set</code> .                                                                                                                          | <pre>public void setHappy(boolean happy) {<br/>    this.happy = happy;<br/>}</pre> |
| The method name must have a prefix of <code>set/get/is</code> , followed by the first letter of the property in uppercase, followed by the rest of the property name. | <pre>public void setNumEggs(int num) {<br/>    numEggs = num;<br/>}</pre>          |

From the last example in Table 4.5, you noticed that you can name the method parameter to set anything you want. Only the method name and property name have naming conventions here.

It's time for some practice. See if you can figure out which lines follow JavaBeans naming conventions:

```
12: private boolean playing;
13: private String name;
14: public boolean getPlaying() { return playing; }
15: public boolean isPlaying() { return playing; }
16: public String name() { return name; }
17: public void updateName(String n) { name = n; }
18: public void setname(String n) { name = n; }
```

Lines 12 and 13 are good. They are private instance variables. Line 14 doesn't follow the JavaBeans naming conventions. Since *playing* is a boolean, the getter must begin with *is*. Line 15 is a correct getter for *playing*. Line 16 doesn't follow the JavaBeans naming conventions because it should be called *getName*. Lines 17 and 18 do not follow the JavaBeans naming conventions because they should be named *setName*. Remember that Java is case sensitive, so *setname* is not adequate to meet the naming convention.

## Creating Immutable Classes

Encapsulating data is helpful because it prevents callers from making uncontrolled changes to your class. Another common technique is making classes immutable so they cannot be changed at all.

Immutable classes are helpful because you know they will always be the same. You can pass them around your application with a guarantee that the caller didn't change anything. This helps make programs easier to maintain. It also helps with performance by limiting the number of copies, as you saw with `String` in Chapter 3, "Core Java APIs."

One step in making a class immutable is to omit the setters. But wait: we still want the caller to be able to specify the initial value—we just don't want it to change after the object is created. Constructors to the rescue!

```
public class ImmutableSwan {
 private int numberEggs;
 public ImmutableSwan(int numberEggs) {
 this.numberEggs = numberEggs;
 }
 public int getNumberEggs() {
 return numberEggs;
 } }

```

In this example, we don't have a setter. We do have a constructor that allows a value to be set. Remember, immutable is only measured after the object is constructed. Immutable classes are allowed to have values. They just can't change after instantiation.



### Real World Scenario

#### Return Types in Immutable Classes

When you are writing an immutable class, be careful about the return types. On the surface, this class appears to be immutable since there is no setter:

```
public class NotImmutable {
 private StringBuilder builder;
 public NotImmutable(StringBuilder b) {
 builder = b;
 }
 public StringBuilder getBuilder() {
 return builder;
 } }

```

The problem is that it isn't really. Consider this code snippet:

*continues*

*continued*

```
StringBuilder sb = new StringBuilder("initial");
NotImmutable problem = new NotImmutable(sb);
sb.append(" added");
StringBuilder gotBuilder = problem.getBuilder();
gotBuilder.append(" more");
System.out.println(problem.getBuilder());
```

This outputs "initial added more"—clearly not what we were intending. The problem is that we are just passing the same `StringBuilder` all over. The caller has a reference since it was passed to the constructor. Anyone who calls the getter gets a reference too. A solution is to make a copy of the mutable object. This is called a defensive copy.

```
public Mutable(StringBuilder b) {
 builder = new StringBuilder(b);
}
public StringBuilder getBuilder() {
 return new StringBuilder(builder);
}
```

Now the caller can make changes to the initial `sb` object and it is fine. `Mutable` no longer cares about that object after the constructor gets run. The same goes for the getter: callers can change their `StringBuilder` without affecting `Mutable`.

Another approach for the getter is to return an immutable object:

```
public String getValue() {
 return builder.toString();
}
```

There's no rule that says we have to return the same type as we are storing. `String` is safe to return because it is immutable in the first place.

To review, encapsulation refers to preventing callers from changing the instance variables directly. Immutability refers to preventing callers from changing the instance variables at all.

## Writing Simple Lambdas

Java is an object-oriented language at heart. You've seen plenty of objects by now. In Java 8, the language added the ability to write code using another style. *Functional programming* is a way of writing code more declaratively. You specify what you want to do rather than dealing with the state of objects. You focus more on expressions than loops.

Functional programming uses lambda expressions to write code. A *lambda expression* is a block of code that gets passed around. You can think of a lambda expression as an

anonymous method. It has parameters and a body just like full-fledged methods do, but it doesn't have a name like a real method. Lambda expressions are often referred to as lambdas for short. You might also know them as closures if Java isn't your first language. If you had a bad experience with closures in the past, don't worry. They are far simpler in Java.

In other words, a lambda expression is like a method that you can pass as if it were a variable. For example, there are different ways to calculate age. One human year is equivalent to seven dog years. You want to write a method that takes an `age()` method as input. To do this in an object-oriented program, you'd need to define a `Human` subclass and a `Dog` subclass. With lambdas, you can just pass in the relevant expression to calculate age.

Lambdas allow you to write powerful code in Java. **Only the simplest lambda expressions are on the OCA exam.** The goal is to get you comfortable with the syntax and the concepts. This means you aren't truly doing functional programming yet. You'll see lambdas again on the OCP exam.

In this section, we'll cover an example of why lambdas are helpful, the syntax of lambdas, and how to use predicates.

## Lambda Example

Our goal is to print out all the animals in a list according to some criteria. We'll show you how to do this without lambdas to illustrate how lambdas are useful. We start out with the `Animal` class:

```
public class Animal {
 private String species;
 private boolean canHop;
 private boolean canSwim;
 public Animal(String speciesName, boolean hopper, boolean swimmer) {
 species = speciesName;
 canHop = hopper;
 canSwim = swimmer;
 }
 public boolean canHop() { return canHop; }
 public boolean canSwim() { return canSwim; }
 public String toString() { return species; }
}
```

The `Animal` class has three instance variables, which are set in the constructor. It has two methods that get the state of whether the animal can hop or swim. It also has a `toString()` method so we can easily identify the `Animal` in programs.

We plan to write a lot of different checks, so we want an interface. You'll learn more about interfaces in the next chapter. For now, it is enough to remember that an interface specifies the methods that our class needs to implement:

```
public interface CheckTrait {
 boolean test(Animal a);
}
```

The first thing we want to check is whether the `Animal` can hop. We provide a class that can check this:

```
public class CheckIfHopper implements CheckTrait {
 public boolean test(Animal a) {
 return a.canHop();
 }
}
```

This class may seem simple—and it is. This is actually part of the problem that lambdas solve. Just bear with us for a bit. Now we have everything that we need to write our code to find the `Animals` that hop:

```
1: public class TraditionalSearch {
2: public static void main(String[] args) {
3: List<Animal> animals = new ArrayList<Animal>(); // list of animals
4: animals.add(new Animal("fish", false, true));
5: animals.add(new Animal("kangaroo", true, false));
6: animals.add(new Animal("rabbit", true, false));
7: animals.add(new Animal("turtle", false, true));
8:
9: print(animals, new CheckIfHopper()); // pass class that does check
10: }
11: private static void print(List<Animal> animals, CheckTrait checker) {
12: for (Animal animal : animals) {
13: if (checker.test(animal)) // the general check
14: System.out.print(animal + " ");
15: }
16: System.out.println();
17: }
18: }
```

The `print()` method on line 11 method is very general—it can check for any trait. This is good design. It shouldn't need to know what specifically we are searching for in order to print a list of animals.

Now what happens if we want to print the `Animals` that swim? Sigh. We need to write another class `CheckIfSwims`. Granted, it is only a few lines. Then we need to add a new line under line 9 that instantiates that class. That's two things just to do another check.

Why can't we just specify the logic we care about right here? Turns out that we can with lambda expressions. We could repeat that whole class here and make you find the one line that changed. Instead, we'll just show you. We could replace line 9 with the following, which uses a lambda:

```
9: print(animals, a -> a.canHop());
```



Don't worry that the syntax looks a little funky. You'll get used to it and we'll describe it in the next section. We'll also explain the bits that look like magic. For now, just focus on how easy it is to read. We are telling Java that we only care about `Animals` that can hop.

It doesn't take much imagination to figure how we would add logic to get the `Animals` that can swim. We only have to add one line of code—no need for an extra class to do something simple. Here's that other line:

```
print(animals, a -> a.canSwim());
```

How about `Animals` that cannot swim?

```
print(animals, a -> ! a.canSwim());
```

The point here is that it is really easy to write code that uses lambdas once you get the basics in place. This code uses a concept called *deferred execution*. *Deferred execution* means that code is specified now but will run later. In this case, later is when the `print()` method calls it.

## Lambda Syntax

One of the simplest lambda expressions you can write is the one you just saw:

```
a -> a.canHop();
```

This means that Java should call a method with an `Animal` parameter that returns a boolean value that's the result of `a.canHop()`. We know all this because we wrote the code. But how does Java know?

Java replies on context when figuring out what lambda expressions mean. We are passing this lambda as the second parameter of the `print()` method. That method expects a `CheckTrait` as the second parameter. Since we are passing a lambda instead, Java tries to map our lambda to that interface:

```
boolean test(Animal a);
```


Since that interface's method takes an `Animal`, that means the lambda parameter has to be an `Animal`. And since that interface's method returns a boolean, we know the lambda returns a boolean.

The syntax of lambdas is tricky because many parts are optional. These two lines do the exact same thing:

```
a -> a.canHop()
(Animal a) -> { return a.canHop(); }
```


Let's look at what is going on here. The first example, shown in Figure 4.5, has three parts:

- Specify a single parameter with the name `a`
- The arrow operator to separate the parameter and body
- A body that calls a single method and returns the result of that method

**FIGURE 4.5** Lambda syntax omitting optional parts

The second example also has three parts; it's just more verbose (see Figure 4.6):

- Specify a single parameter with the name `a` and stating the type is `Animal`
- The arrow operator to separate the parameter and body
- A body that has one or more lines of code, including a semicolon and a return statement

**FIGURE 4.6** Lambda syntax, including optional parts

The parentheses can only be omitted if there is a single parameter and its type is not explicitly stated. Java does this because developers commonly use lambda expressions this way and they can do as little typing as possible.

It shouldn't be news to you that we can omit braces when we only have a single statement. We did this with `if` statements and loops already. What is different here is that the rules change when you omit the braces. Java doesn't require you to type `return` or use a semicolon when no braces are used. This special shortcut doesn't work when we have two or more statements. At least this is consistent with using `{}` to create blocks of code elsewhere.

Let's look at some examples of valid lambdas. Pretend that there are valid interfaces that can consume a lambda with zero, one, or two `String` parameters.

```
3: print(() -> true); // 0 parameters
4: print(a -> a.startsWith("test")); // 1 parameter
5: print((String a) -> a.startsWith("test")); // 1 parameter
6: print((a, b) -> a.startsWith("test")); // 2 parameters
7: print((String a, String b) -> a.startsWith("test")); // 2 parameters
```

Notice that all of these examples have parentheses around the parameter list except the one that takes only one parameter and doesn't specify the type. Line 3 takes 0 parameters and always returns the Boolean true. Line 4 takes one parameter and calls a method on it, returning the result. Line 5 does the same except that it explicitly defines the type of the variable. Lines 6 and 7 take two parameters and ignore one of them—there isn't a rule that says you must use all defined parameters.

Now let's make sure you can identify invalid syntax. Do you see what's wrong with each of these?

```
print(a, b -> a.startsWith("test")); // DOES NOT COMPILE
print(a -> { a.startsWith("test"); }); // DOES NOT COMPILE
print(a -> { return a.startsWith("test") }); // DOES NOT COMPILE
```

The first line needs parentheses around the parameter list. **Remember that the parentheses are *only* optional when there is one parameter and it doesn't have a type declared.** The second line is missing the return keyword. The last line is missing the semicolon.

You might have noticed all of our lambdas return a boolean. **That is because the scope for the OCA exam limits what you need to learn.**



## Real World Scenario

### What Variables Can My Lambda Access?

**Lambdas are allowed to access variables. This topic isn't on the OCA exam, but you may come across it when practicing.** Lambdas are allowed to access variables. Here's an example:

```
boolean wantWhetherCanHop = true;
print(animals, a -> a.canHop() == wantWhetherCanHop);
```

The trick is that they cannot access all variables. Instance and static variables are okay. Method parameters and local variables are fine if they are not assigned new values.

There is one more issue you might see with lambdas. We've been defining an argument list in our lambda expressions. Since Java doesn't allow us to redeclare a local variable, the following is an issue:

```
(a, b) -> { int a = 0; return 5; } // DOES NOT COMPILE
```

We tried to redeclare `a`, which is not allowed. By contrast, the following line is okay because it uses a different variable name:

```
(a, b) -> { int c = 0; return 5; }
```

## Predicates

In our earlier example, we created an interface with one method:

```
boolean test(Animal a);
```

**Lambdas work with interfaces that have only one method. These are called functional interfaces—interfaces that can be used with functional programming. (It's actually more complicated than this, but for the OCA exam this definition is fine.)**

You can imagine that we'd have to create lots of interfaces like this to use lambdas. We want to test `Animals` and `Strings` and `Plants` and anything else that we come across.

Luckily, Java recognizes that this is a common problem and provides such an interface for us. It's in the package `java.util.function` and the gist of it is as follows:

```
public interface Predicate<T> {
 boolean test(T t);
}
```

That looks a lot like our method. The only difference is that it uses this type `T` instead of `Animal`. That's the syntax for generics. It's like when we created an `ArrayList` and got to specify any type that goes in it.

This means we don't need our own interface anymore and can put everything related to our search in one class:

```
1: import java.util.*;
2: import java.util.function.*;
3: public class PredicateSearch {
4: public static void main(String[] args) {
5: List<Animal> animals = new ArrayList<Animal>();
6: animals.add(new Animal("fish", false, true));
7:
8: print(animals, a -> a.canHop());
9: }
10: private static void print(List<Animal> animals, Predicate<Animal>
 checker) {
11: for (Animal animal : animals) {
12: if (checker.test(animal))
13: System.out.print(animal + " ");
14: }
15: System.out.println();
16: }
17: }
```

This time, line 10 is the only one that changed. We expect to have a `Predicate` passed in that uses type `Animal`. Pretty cool. We can just use it without having to write extra code.

Java 8 even integrated the `Predicate` interface into some existing classes. There is only one you need to know for the exam. `ArrayList` declares a `removeIf()` method that takes a `Predicate`. Imagine we have a list of names for pet bunnies. We decide we want to remove all of the bunny names that don't begin with the letter h because our little cousin really wants us to choose an H name. We could solve this problem by writing a loop. Or we could solve it in one line:

```
3: List<String> bunnies = new ArrayList<>();
4: bunnies.add("long ear");
5: bunnies.add("floppy");
6: bunnies.add("hoppy");
7: System.out.println(bunnies); // [long ear, floppy, hoppy]
8: bunnies.removeIf(s -> s.charAt(0) != 'h');
9: System.out.println(bunnies); // [hoppy]
```

Line 8 takes care of everything for us. It defines a predicate that takes a `String` and returns a `boolean`. The `removeIf()` method does the rest.

For the OCA exam, you only need to know how to implement lambda expressions that use the `Predicate` interface. Remember the one method in the interface called `test()`? It takes any one reference type parameter and returns a `boolean`. Functional programming is a large topic and just the basics are covered. On the OCP exam, you'll learn how to get rid of the loop entirely for more than just `removeIf()`. You'll also learn the rules for implementing your own functional interfaces as we did with `CheckTrait`.

## Summary

As you learned in this chapter, Java methods start with an access modifier of `public`, `private`, `protected` or blank (default access). This is followed by an optional specifier such as `static`, `final`, or `abstract`. Next comes the return type, which is `void` or a Java type. The method name follows, using standard Java identifier rules. Zero or more parameters go in parentheses as the parameter list. Next come any optional exception types. Finally, zero or more statements go in braces to make up the method body.

Using the `private` keyword means the code is only available from within the same class. Default (package private) access means the code is only available from within the same package. Using the `protected` keyword means the code is available from the same package or subclasses. Using the `public` keyword means the code is available from anywhere. Static methods and static variables are shared by the class. When referenced from outside the class, they are called using the classname—for example, `StaticClass.method()`. Instance members are allowed to call static members, but static members are not allowed to call instance members. Static imports are used to import static members.

Java uses pass-by-value, which means that calls to methods create a copy of the parameters. Assigning new values to those parameters in the method doesn't affect the caller's variables.

Calling methods on objects that are method parameters changes the state of those objects and is reflected in the caller.

Overloaded methods are methods with the same name but a different parameter list. Java calls the most specific method it can find. Exact matches are preferred, followed by wider primitives. After that comes autoboxing and finally varargs.

Constructors are used to instantiate new objects. The default no-argument constructor is called when no constructor is coded. Multiple constructors are allowed and can call each other by writing `this()`. If `this()` is present, it must be the first statement in the constructor. Constructors can refer to instance variables by writing `this` before a variable name to indicate they want the instance variable and not the method parameter with that name. The order of initialization is the superclass (which we will cover in Chapter 5); static variables and static initializers in the order they appear; instance variables and instance initializers in the order they appear; and finally the constructor.

Encapsulation refers to preventing callers from changing the instance variables directly. This is done by making instance variables private and getters/setters public. Immutability refers to preventing callers from changing the instance variables at all. This uses several techniques, including removing setters. JavaBeans use methods beginning with `is` and `get` for boolean and non-boolean property types, respectively. Methods beginning with `set` are used for setters.

Lambda expressions, or lambdas, allow passing around blocks of code. The full syntax looks like `(String a, String b) -> { return a.equals(b); }`. The parameter types can be omitted. When only one parameter is specified without a type, the parentheses can also be omitted. The braces and return statement can be omitted for a single statement, making the short form `(a -> a.equals(b))`. Lambdas are passed to a method expecting an interface with one method. Predicate is a common interface. It has one method named `test` that returns a boolean and takes any type. The `removeIf()` method on `ArrayList` takes a Predicate.

## Exam Essentials

**Be able to identify correct and incorrect method declarations.** A sample method signature is `public static void method(String... args) throws Exception {}`.

**Identify when a method or field is accessible.** Recognize when a method or field is accessed when the access modifier (private, protected, public, or default access) does not allow it.

**Recognize valid and invalid uses of static imports.** Static imports import static members. They are written as `import static`, not *static import*. Make sure they are importing static methods or variables rather than classnames.

**State the output of code involving methods.** Identify when to call static rather than instance methods based on whether the classname or object comes before the method.

Recognize the correct overloaded method. Exact matches are used first, followed by wider primitives, followed by autoboxing, followed by varargs. Assigning new values to method parameters does not change the caller, but calling methods on them does.

**Evaluate code involving constructors.** Constructors can call other constructors by calling `this()` as the first line of the constructor. Recognize when the default constructor is provided. Remember the order of initialization is the superclass, static variables/initializers, instance variables/initializers, and the constructor.

**Be able to recognize when a class is properly encapsulated.** Look for `private` instance variables and `public` getters and setters when identifying encapsulation.

**Write simple lambda expressions.** Look for the presence or absence of optional elements in lambda code. Parameter types are optional. Braces and the `return` keyword are optional when the body is a single statement. Parentheses are optional when only one parameter is specified and the type is implicit. The `Predicate` interface is commonly used with lambdas because it declares a single method called `test()`, which takes one parameter.

## Review Questions

1. Which of the following can fill in the blank in this code to make it compile? (Choose all that apply)

```
public class Ant {
 _____ void method() { }
}
```

- A. default
  - B. final
  - C. private
  - D. Public
  - E. String
  - F. zzz:
2. Which of the following compile? (Choose all that apply)
- A. `final static void method4() { }`
  - B. `public final int void method() { }`
  - C. `private void int method() { }`
  - D. `static final void method3() { }`
  - E. `void final method() {}`
  - F. `void public method() { }`
3. Which of the following methods compile? (Choose all that apply)
- A. `public void methodA() { return;}`
  - B. `public void methodB() { return null;}`
  - C. `public void methodD() {}`
  - D. `public int methodD() { return 9;}`
  - E. `public int methodE() { return 9.0;}`
  - F. `public int methodF() { return;}`
  - G. `public int methodG() { return null;}`
4. Which of the following compile? (Choose all that apply)
- A. `public void moreA(int... nums) {}`
  - B. `public void moreB(String values, int... nums) {}`
  - C. `public void moreC(int... nums, String values) {}`
  - D. `public void moreD(String... values, int... nums) {}`
  - E. `public void moreE(String[] values, ...int nums) {}`
  - F. `public void moreF(String... values, int[] nums) {}`
  - G. `public void moreG(String[] values, int[] nums) {}`



5. Given the following method, which of the method calls return 2? (Choose all that apply)

```
public int howMany(boolean b, boolean... b2) {
 return b2.length;
}
```

- A. `howMany();`
- B. `howMany(true);`
- C. `howMany(true, true);`
- D. `howMany(true, true, true);`
- E. `howMany(true, {true});`
- F. `howMany(true, {true, true});`
- G. `howMany(true, new boolean[2]);`

6. Which of the following are true? (Choose all that apply)

- A. Package private access is more lenient than protected access.
- B. A public class that has private fields and package private methods is not visible to classes outside the package.
- C. You can use access modifiers so only some of the classes in a package see a particular package private class.
- D. You can use access modifiers to allow read access to all methods, but not any instance variables.
- E. You can use access modifiers to restrict read access to all classes that begin with the word `Test`.

7. Given the following `my.school.ClassRoom` and `my.city.School` class definitions, which line numbers in `main()` generate a compiler error? (Choose all that apply)

```
1: package my.school;
2: public class Classroom {
3: private int roomNumber;
4: protected String teacherName;
5: static int globalKey = 54321;
6: public int floor = 3;
7: Classroom(int r, String t) {
8: roomNumber = r;
9: teacherName = t; } }
```

```
1: package my.city;
2: import my.school.*;
3: public class School {
4: public static void main(String[] args) {
5: System.out.println(Classroom.globalKey);
6: Classroom room = new Classroom(101, ""Mrs. Anderson");
```

```
7: System.out.println(room.roomNumber);
8: System.out.println(room.floor);
9: System.out.println(room.teacherName); } }
```

- A. None, the code compiles fine.
  - B. Line 5
  - C. Line 6
  - D. Line 7
  - E. Line 8
  - F. Line 9
8. Which of the following are true? (Choose all that apply)
- A. Encapsulation uses package private instance variables.
  - B. Encapsulation uses private instance variables.
  - C. Encapsulation allows setters.
  - D. Immutability uses package private instance variables.
  - E. Immutability uses private instance variables.
  - F. Immutability allows setters.
9. Which are methods using JavaBeans naming conventions for accessors and mutators? (Choose all that apply)
- A. `public boolean getCanSwim() { return canSwim;}`
  - B. `public boolean canSwim() { return numberWings;}`
  - C. `public int getNumWings() { return numberWings;}`
  - D. `public int numWings() { return numberWings;}`
  - E. `public void setCanSwim(boolean b) { canSwim = b;}`
10. What is the output of the following code?
- ```
1: package rope;
2: public class Rope {
3:     public static int LENGTH = 5;
4:     static {
5:         LENGTH = 10;
6:     }
```

```
7: public static void swing() {  
8:     System.out.print("swing ");  
9: }  
10: }
```

```
1: import rope.*;  
2: import static rope.Rope.*;  
3: public class Chimp {  
4:     public static void main(String[] args) {  
5:         Rope.swing();  
6:         new Rope().swing();  
7:         System.out.println(LENGTH);  
8:     }  
9: }
```

- A. swing swing 5
- B. swing swing 10
- C. Compiler error on line 2 of Chimp.
- D. Compiler error on line 5 of Chimp.
- E. Compiler error on line 6 of Chimp.
- F. Compiler error on line 7 of Chimp.

11. Which are true of the following code? (Choose all that apply)

```
1: public class Rope {  
2:     public static void swing() {  
3:         System.out.print("swing ");  
4:     }  
5:     public void climb() {  
6:         System.out.println("climb ");  
7:     }  
8:     public static void play() {  
9:         swing();  
10:        climb();  
11:    }  
12:    public static void main(String[] args) {  
13:        Rope rope = new Rope();  
14:        rope.play();  
15:        Rope rope2 = null;  
16:        rope2.play();  
17:    }  
18: }
```

- A. The code compiles as is.
- B. There is exactly one compiler error in the code.
- C. There are exactly two compiler errors in the code.
- D. If the lines with compiler errors are removed, the output is `climb climb`.
- E. If the lines with compiler errors are removed, the output is `swing swing`.
- F. If the lines with compile errors are removed, the code throws a `NullPointerException`.

12. What is the output of the following code?

```
import rope.*;
import static rope.Rope.*;
public class RopeSwing {
    private static Rope rope1 = new Rope();
    private static Rope rope2 = new Rope();
    {
        System.out.println(rope1.length);
    }
    public static void main(String[] args) {
        rope1.length = 2;
        rope2.length = 8;
        System.out.println(rope1.length);
    }
}
```

```
package rope;
public class Rope {
    public static int length = 0;
}
```

- A. 02
- B. 08
- C. 2
- D. 8
- E. The code does not compile.
- F. An exception is thrown.

13. How many compiler errors are in the following code?

```
1: public class RopeSwing {
2:     private static final String leftRope;
3:     private static final String rightRope;
4:     private static final String bench;
5:     private static final String name = "name";
```

```
6:  static {
7:      leftRope = "left";
8:      rightRope = "right";
9:  }
10: static {
11:     name = "name";
12:     rightRope = "right";
13: }
14: public static void main(String[] args) {
15:     bench = "bench";
16: }
17: }
```

- A. 0
- B. 1
- C. 2
- D. 3
- E. 4
- F. 5

14. Which of the following can replace line 2 to make this code compile? (Choose all that apply)

```
1: import java.util.*;
2: // INSERT CODE HERE
3: public class Imports {
4:     public void method(ArrayList<String> list) {
5:         sort(list);
6:     }
7: }
```

- A. `import static java.util.Collections;`
- B. `import static java.util.Collections.*;`
- C. `import static java.util.Collections.sort(ArrayList<String>);`
- D. `static import java.util.Collections;`
- E. `static import java.util.Collections.*;`
- F. `static import java.util.Collections.sort(ArrayList<String>);`

15. What is the result of the following statements?

```
1: public class Test {
2:     public void print(byte x) {
3:         System.out.print("byte");
4:     }
5:     public void print(int x) {
6:         System.out.print("int");
```

```
7:     }
8:     public void print(float x) {
9:         System.out.print("float");
10:    }
11:    public void print(Object x) {
12:        System.out.print("Object");
13:    }
14:    public static void main(String[] args) {
15:        Test t = new Test();
16:        short s = 123;
17:        t.print(s);
18:        t.print(true);
19:        t.print(6.789);
20:    }
21: }
```

- A. byteFloatObject
- B. intFloatObject
- C. byteObjectFloat
- D. intObjectFloat
- E. intObjectObject
- F. byteObjectObject

16. What is the result of the following program?

```
1: public class Squares {
2:     public static long square(int x) {
3:         long y = x * (long) x;
4:         x = -1;
5:         return y;
6:     }
7:     public static void main(String[] args) {
8:         int value = 9;
9:         long result = square(value);
10:        System.out.println(value);
11:    } }
```

- A. -1
- B. 9
- C. 81
- D. Compiler error on line 9.
- E. Compiler error on a different line.

17. Which of the following are output by the following code? (Choose all that apply)

```
public class StringBuilders {
    public static StringBuilder work(StringBuilder a,
    StringBuilder b) {
        a = new StringBuilder("a");
        b.append("b");
        return a;
    }
    public static void main(String[] args) {
        StringBuilder s1 = new StringBuilder("s1");
        StringBuilder s2 = new StringBuilder("s2");
        StringBuilder s3 = work(s1, s2);
        System.out.println("s1 = " + s1);
        System.out.println("s2 = " + s2);
        System.out.println("s3 = " + s3);
    }
}
```

- A. s1 = a
 - B. s1 = s1
 - C. s2 = s2
 - D. s2 = s2b
 - E. s3 = a
 - F. s3 = null
 - G. The code does not compile.
18. Which of the following are true? (Choose 2)
- A. this() can be called from anywhere in a constructor.
 - B. this() can be called from any instance method in the class.
 - C. this.variableName can be called from any instance method in the class.
 - D. this.variableName can be called from any static method in the class.
 - E. You must include a default constructor in the code if the compiler does not include one.
 - F. You can call the default constructor written by the compiler using this().
 - G. You can access a private constructor with the main() method.
19. Which of these classes compile and use a default constructor? (Choose all that apply)
- A. public class Bird { }
 - B. public class Bird { public bird() {} }
 - C. public class Bird { public bird(String name) {} }
 - D. public class Bird { public Bird() {} }

- E. `public class Bird { Bird(String name) {} }`
- F. `public class Bird { private Bird(int age) {} }`
- G. `public class Bird { void Bird() { } }`

20. Which code can be inserted to have the code print 2?

```
public class BirdSeed {
    private int numberBags;
    boolean call;

    public BirdSeed() {
        // LINE 1
        call = false;
        // LINE 2
    }
    public BirdSeed(int numberBags) {
        this.numberBags = numberBags;
    }
    public static void main(String[] args) {
        BirdSeed seed = new BirdSeed();
        System.out.println(seed.numberBags);
    } }
```

- A. Replace line 1 with `BirdSeed(2);`
- B. Replace line 2 with `BirdSeed(2);`
- C. Replace line 1 with `new BirdSeed(2);`
- D. Replace line 2 with `new BirdSeed(2);`
- E. Replace line 1 with `this(2);`
- F. Replace line 2 with `this(2);`

21. Which of the following complete the constructor so that this code prints out 50? (Choose all that apply)

```
public class Cheetah {
    int numSpots;
    public Cheetah(int numSpots) {
        // INSERT CODE HERE
    }
    public static void main(String[] args) {
        System.out.println(new Cheetah(50).numSpots);
    }
}
```


- A. numSpots = numSpots;
- B. numSpots = this.numSpots;
- C. this.numSpots = numSpots;
- D. numSpots = super.numSpots;
- E. super.numSpots = numSpots;
- F. None of the above.

22. What is the result of the following?

```
1: public class Order {  
2:     static String result = "";  
3:     { result += "c"; }  
4:     static  
5:     { result += "u"; }  
6:     { result += "r"; }  
7: }
```

```
1: public class OrderDriver {  
2:     public static void main(String[] args) {  
3:         System.out.print(Order.result + " ");  
4:         System.out.print(Order.result + " ");  
5:         new Order();  
6:         new Order();  
7:         System.out.print(Order.result + " ");  
8:     }  
9: }
```

- A. curur
- B. ucr cr
- C. u ucr cr
- D. u u cur cur
- E. u u ucr cr
- F. ur ur urc
- G. The code does not compile.

23. What is the result of the following?

```
1: public class Order {  
2:     String value = "t";  
3:     { value += "a"; }  
4:     { value += "c"; }  
5:     public Order() {
```

```
6:     value += "b";
7: }
8: public Order(String s) {
9:     value += s;
10: }
11: public static void main(String[] args) {
12:     Order order = new Order("f");
13:     order = new Order();
14:     System.out.println(order.value);
15: } }
```

- A. tacb
- B. tacf
- C. tacbf
- D. tacfb
- E. tacftacb
- F. The code does not compile.
- G. An exception is thrown.

24. Which of the following will compile when inserted in the following code? (Choose all that apply)

```
public class Order3 {
    final String value1 = "1";
    static String value2 = "2";
    String value3 = "3";
    {
        // CODE SNIPPET 1
    }
    static {
        // CODE SNIPPET 2
    }
}
```

- A. value1 = "d"; instead of // CODE SNIPPET 1
- B. value2 = "e"; instead of // CODE SNIPPET 1
- C. value3 = "f"; instead of // CODE SNIPPET 1
- D. value1 = "g"; instead of // CODE SNIPPET 2
- E. value2 = "h"; instead of // CODE SNIPPET 2
- F. value3 = "i"; instead of // CODE SNIPPET 2

25. Which of the following are true about the following code? (Choose all that apply)

```
public class Create {
    Create() {
        System.out.print("1 ");
    }
    Create(int num) {
        System.out.print("2 ");
    }
    Create(Integer num) {
        System.out.print("3 ");
    }
    Create(Object num) {
        System.out.print("4 ");
    }
    Create(int... nums) {
        System.out.print("5 ");
    }
    public static void main(String[] args) {
        new Create(100);
        new Create(1000L);
    }
}
```

- A. The code prints out 2 4.
 - B. The code prints out 3 4.
 - C. The code prints out 4 2.
 - D. The code prints out 4 4.
 - E. The code prints 3 4 if you remove the constructor `Create(int num)`.
 - F. The code prints 4 4 if you remove the constructor `Create(int num)`.
 - G. The code prints 5 4 if you remove the constructor `Create(int num)`.
26. What is the result of the following class?
- ```
1: import java.util.function.*;
2:
3: public class Panda {
4: int age;
5: public static void main(String[] args) {
6: Panda p1 = new Panda();
7: p1.age = 1;
8: check(p1, p -> p.age < 5);
```

```
9: }
10: private static void check(Panda panda, Predicate<Panda> pred) {
11: String result = pred.test(panda) ? "match" : "not match";
12: System.out.print(result);
13: } }
```

- A. match
- B. not match
- C. Compiler error on line 8.
- D. Compiler error on line 10.
- E. Compiler error on line 11.
- F. A runtime exception is thrown.

27. What is the result of the following code?

```
1: interface Climb {
2: boolean isTooHigh(int height, int limit);
3: }
4:
5: public class Climber {
6: public static void main(String[] args) {
7: check((h, l) -> h.append(l).isEmpty(), 5);
8: }
9: private static void check(Climb climb, int height) {
10: if (climb.isTooHigh(height, 10))
11: System.out.println("too high");
12: else
13: System.out.println("ok");
14: }
15: }
```

- A. ok
- B. too high
- C. Compiler error on line 7.
- D. Compiler error on line 10.
- E. Compiler error on a different line.
- F. A runtime exception is thrown.

28. Which of the following lambda expressions can fill in the blank? (Choose all that apply)

```
List<String> list = new ArrayList<>();
list.removeIf(_____);
```

- A. `s -> s.isEmpty()`
- B. `s -> {s.isEmpty()}`
- C. `s -> {s.isEmpty();}`
- D. `s -> {return s.isEmpty();}`
- E. `String s -> s.isEmpty()`
- F. `(String s) -> s.isEmpty()`

29. Which lambda can replace the `MySecret` class to return the same value? (Choose all that apply)

```
interface Secret {
 String magic(double d);
}
```

```
class MySecret implements Secret {
 public String magic(double d) {
 return "Poof";
 }
}
```

- A. `caller((e) -> "Poof");`
- B. `caller((e) -> {"Poof"});`
- C. `caller((e) -> { String e = ""; "Poof" });`
- D. `caller((e) -> { String e = ""; return "Poof"; });`
- E. `caller((e) -> { String e = ""; return "Poof" });`
- F. `caller((e) -> { String f = ""; return "Poof"; });`



# Chapter 5

## Class Design

---

### OCA EXAM OBJECTIVES COVERED IN THIS CHAPTER:

#### ✓ Working with Inheritance

- Describe inheritance and its benefits
- Develop code that demonstrates the use of polymorphism; including overriding and object type versus reference type
- Determine when casting is necessary
- Use super and this to access objects and constructors
- Use abstract classes and interfaces





In Chapter 1, “Java Building Blocks,” we introduced the basic definition for a class in Java. In Chapter 4, “Methods and Encapsulation,” we delved into constructors, methods, and modifiers, and showed how you can use them to build more structured classes. In this chapter, we’ll take things one step further and show how class structure is one of the most powerful features in the Java language.

At its core, proper Java class design is about code reusability, increased functionality, and standardization. For example, by creating a new class that extends an existing class, you may gain access to a slew of inherited primitives, objects, and methods. Alternatively, by designing a standard interface for your application, you ensure that any class that implements the interface has certain required methods defined. Finally, by creating abstract class definitions, you’re defining a platform that other developers can extend and build on top of.

## Introducing Class Inheritance

When creating a new class in Java, you can define the class to inherit from an existing class. *Inheritance* is the process by which the new child subclass automatically includes any public or protected primitives, objects, or methods defined in the parent class.

For illustrative purposes, we refer to any class that inherits from another class as a *child class*, or a descendent of that class. Alternatively, we refer to the class that the child inherits from as the *parent class*, or an ancestor of the class. If child X inherits from class Y, which in turn inherits from class Z, then X would be considered an indirect child, or descendent, of class Z.

Java supports *single inheritance*, by which a class may inherit from only one direct parent class. Java also supports multiple levels of inheritance, by which one class may extend another class, which in turn extends another class. You can extend a class any number of times, allowing each descendent to gain access to its ancestor’s members.


To truly understand single inheritance, it may helpful to contrast it with *multiple inheritance*, by which a class may have multiple direct parents. By design, Java doesn’t support multiple inheritance in the language because studies have shown that multiple inheritance can lead to complex, often difficult-to-maintain code. Java does allow one exception to the single inheritance rule: classes may implement multiple interfaces, as you’ll see later in this chapter.

Figure 5.1 illustrates the various types of inheritance models. The items on the left are considered single inheritance because each child has exactly one parent. You may notice that single inheritance doesn’t preclude parents from having multiple children. The right



side shows items that have multiple inheritance. For example, a dog object has multiple parent designations. Part of what makes multiple inheritance complicated is determining which parent to inherit values from in case of a conflict. For example, if you have an object or method defined in all of the parents, which one does the child inherit? There is no natural ordering for parents in this example, which is why Java avoids these issues by disallowing multiple inheritance altogether.

**FIGURE 5.1** Types of inheritance



It is possible in Java to prevent a class from being extended by marking the class with the `final` modifier. If you try to define a class that inherits from a `final` class, the compiler will throw an error and not compile. Unless otherwise specified, throughout this chapter you can assume the classes we work with are not marked as `final`.

## Extending a Class

In Java, you can extend a class by adding the parent class name in the definition using the `extends` keyword. The syntax of defining and extending a class is shown in Figure 5.2.

**FIGURE 5.2** Defining and extending a class

```

public or default access modifier class name
 | |
 v v
abstract or final keyword (optional)
 |
 v
class keyword (required)
 |
 v
public abstract class ElephantSeal extends Seal {
 // Methods and Variables defined here
}

```

We'll discuss what it means for a class to be `abstract` and `final`, as well as the class access modifiers, later in this chapter.

Because Java allows only one public class per file, we can create two files, `Animal.java` and `Lion.java`, in which the `Lion` class extends the `Animal` class. Assuming they are in the same package, an import statement is not required in `Lion.java` to access the `Animal` class.

Here are the contents of `Animal.java`:

```
public class Animal {
 private int age;
 public int getAge() {
 return age;
 }
 public void setAge(int age) {
 this.age = age;
 }
}
```

And here are the contents of `Lion.java`:

```
public class Lion extends Animal {
 private void roar() {
 System.out.println("The "+getAge()+" year old lion says: Roar!");
 }
}
```

Notice the use of the `extends` keyword in `Lion.java` to indicate that the `Lion` class extends from the `Animal` class. In this example, we see that `getAge()` and `setAge()` are accessible by subclass `Lion`, because they are marked as `public` in the parent class. The primitive `age` is marked as `private` and therefore not accessible from the subclass `Lion`, as the following would not compile:

```
public class Lion extends Animal {
 private void roar() {
 System.out.println("The "+age+" year old lion says: Roar!");
 // DOES NOT COMPILE
 }
}
```

Despite the fact that `age` is inaccessible by the child class, if we have an instance of a `Lion` object, there is still an `age` value that exists within the instance. The `age` value just cannot be directly referenced by the child class nor any instance of the class. In this manner, the `Lion` object is actually “bigger” than the `Animal` object in the sense that it includes all the properties of the `Animal` object (although not all of those properties may be directly accessible) along with its own set of `Lion` attributes.

## Applying Class Access Modifiers

As discussed in Chapter 4, you can apply access modifiers (`public`, `private`, `protected`, `default`) to both class methods and variables. It probably comes as no surprise that you can also apply access modifiers to class definitions, since we have been adding the `public` access modifier to nearly every class up to now.



For the OCA exam, you should only be familiar with `public` and `default` package-level class access modifiers, because these are the only ones that can be applied to top-level classes within a Java file. The `protected` and `private` modifiers can only be applied to inner classes, which are classes that are defined within other classes, but this is well out of scope for the OCA exam.

The `public` access modifier applied to a class indicates that it can be referenced and used in any class. The `default` package private modifier, which is the lack of any access modifier, indicates the class can be accessed only by a subclass or class within the same package.

As you know, a Java file can have many classes but at most one `public` class. In fact, it may have no `public` class at all. One feature of using the `default` package private modifier is that you can define many classes within the same Java file. For example, the following definition could appear in a single Java file named `Groundhog.java`, since it contains only one `public` class:

```
class Rodent {}
```

```
public class Groundhog extends Rodent {}
```

If we were to update the `Rodent` class with the `public` access modifier, the `Groundhog.java` file would not compile unless the `Rodent` class was moved to its own `Rodent.java` file.

The rules for applying class access modifiers are identical for interfaces. There can be at most one `public` class or interface in a Java file. Like classes, top-level interfaces can also be declared with the `public` or `default` modifiers. We'll discuss interfaces in detail later in this chapter.



For simplicity, any time you see multiple `public` classes or interfaces defined in the same code block in this chapter, assume each class is defined in its own Java file.

## Creating Java Objects

Throughout our discussion of Java in this book, we have thrown around the word *object* numerous times—and with good reason. In Java, all classes inherit from a single class,

`java.lang.Object`. Furthermore, `java.lang.Object` is the only class that doesn't have any parent classes.

You might be wondering, “None of the classes I've written so far extend `java.lang.Object`, so how do all classes inherit from it?” The answer is that the compiler has been automatically inserting code into any class you write that doesn't extend a specific class. For example, consider the following two equivalent class definitions:


```
public class Zoo {
}
```

```
public class Zoo extends java.lang.Object {
}
```

The key is that when Java sees you define a class that doesn't extend another class, it immediately adds the syntax `extends java.lang.Object` to the class definition.

If you define a new class that extends an existing class, Java doesn't add this syntax, although the new class still inherits from `java.lang.Object`. Since all classes inherit from `java.lang.Object`, extending an existing class means the child automatically inherits from `java.lang.Object` by construction. This means that if you look at the inheritance structure of any class, it will always end with `java.lang.Object` on the top of the tree, as shown in Figure 5.3.

**FIGURE 5.3** Java object inheritance



## Defining Constructors

As you may recall from Chapter 4, **every class has at least one constructor. In the case that no constructor is declared, the compiler will automatically insert a default no-argument constructor.** In the case of extending a class, though, things are a bit more interesting.

In Java, the first statement of every constructor is either a call to another constructor within the class, using `this()`, or a call to a constructor in the direct parent class, using

`super()`. If a parent constructor takes arguments, the `super` constructor would also take arguments. For simplicity in this section, we refer to the `super()` command as any parent constructor, even those that take an argument. Notice the user of both `super()` and `super(age)` in the following example:

```
public class Animal {
 private int age;
 public Animal(int age) {
 super();
 this.age = age;
 }
}

public class Zebra extends Animal {
 public Zebra(int age) {
 super(age);
 }
 public Zebra() {
 this(4);
 }
}
```

In the first class, `Animal`, the first statement of the constructor is a call to the parent constructor defined in `java.lang.Object`, which takes no arguments. In the second class, `Zebra`, the first statement of the first constructor is a call to `Animal`'s constructor, which takes a single argument. The class `Zebra` also includes a second no-argument constructor that doesn't call `super()` but instead calls the other constructor within the `Zebra` class using `this(4)`.

Like the `this()` command that you saw in Chapter 4, **the `super()` command may only be used as the first statement of the constructor**. For example, the following two class definitions will not compile:

```
public class Zoo {
 public Zoo() {
 System.out.println("Zoo created");
 super(); // DOES NOT COMPILE
 }
}

public class Zoo {
 public Zoo() {
 super();
 System.out.println("Zoo created");
 }
}
```

```
 super(); // DOES NOT COMPILE
}
}
```

The first class will not compile because the call to the parent constructor must be the first statement of the constructor, not the second statement. In the second code snippet, `super()` is the first statement of the constructor, but it also used as the third statement. Since `super()` can only be used as the first statement of the constructor, the code will likewise not compile.

If the parent class has more than one constructor, the child class may use any valid parent constructor in its definition, as shown in the following example:

```
public class Animal {
 private int age;
 private String name;
 public Animal(int age, String name) {
 super();
 this.age = age;
 this.name = name;
 }
 public Animal(int age) {
 super();
 this.age = age;
 this.name = null;
 }
}

public class Gorilla extends Animal {
 public Gorilla(int age) {
 super(age, "Gorilla");
 }
 public Gorilla() {
 super(5);
 }
}
```

In this example, the first child constructor takes one argument, `age`, and calls the parent constructor, which takes two arguments, `age` and `name`. The second child constructor takes no arguments, and it calls the parent constructor, which takes one argument, `age`. In this example, notice that the child constructors are not required to call matching parent constructors. Any valid parent constructor is acceptable as long as the appropriate input parameters to the parent constructor are provided.

## Understanding Compiler Enhancements

Up to now, we defined numerous classes that did not explicitly call the parent constructor via the `super()` keyword, so why did the code compile? The answer is that the Java compiler automatically inserts a call to the no-argument constructor `super()` if the first statement is not a call to the parent constructor. For example, the following three class and constructor definitions are equivalent, because the compiler will automatically convert them all to the last example:

```
public class Donkey {
}
```

```
public class Donkey {
 public Donkey() {
 }
}
```

```
public class Donkey {
 public Donkey() {
 super();
 }
}
```

Make sure you understand the differences between these three `Donkey` class definitions and why Java will automatically convert them all to the last definition. Keep the process the Java compile performs in mind as we discuss the next few examples.

What happens if the parent class doesn't have a no-argument constructor? Recall that the no-argument constructor is not required and only inserted if there is no constructor defined in the class. In this case, the Java compiler will not help and you must create at least one constructor in your child class that explicitly calls a parent constructor via the `super()` command. For example, the following code will not compile:

```
public class Mammal {
 public Mammal(int age) {
 }
}
```

```
public class Elephant extends Mammal { // DOES NOT COMPILE
}
```

In this example no constructor is defined within the `Elephant` class, so the compiler tries to insert a default no-argument constructor with a `super()` call, as it did in the `Donkey` example. The compiler stops, though, when it realizes there is no parent constructor that takes no arguments. In this example, we must explicitly define at least one constructor, as in the following code:

```
public class Mammal {
 public Mammal(int age) {
 }
}

public class Elephant extends Mammal {
 public Elephant() { // DOES NOT COMPILE
 }
}
```

This code still doesn't compile, though, because the compiler tries to insert the no-argument `super()` as the first statement of the constructor in the `Elephant` class, and there is no such constructor in the parent class. We can fix this, though, by adding a call to a parent constructor that takes a fixed argument:

```
public class Mammal {
 public Mammal(int age) {
 }
}

public class Elephant extends Mammal {
 public Elephant() {
 super(10);
 }
}
```

This code will compile because we have added a constructor with an explicit call to a parent constructor. Note that the class `Elephant` now has a no-argument constructor even though its parent class `Mammal` doesn't. Subclasses may define no-argument constructors even if their parent classes do not, provided the constructor of the child maps to a parent constructor via an explicit call of the `super()` command.

You should be wary of any exam question in which the parent class defines a constructor that takes arguments and doesn't define a no-argument constructor. Be sure to check that the code compiles before answering a question about it.

## Reviewing Constructor Rules

Let's review the rules we covered in this section.

### Constructor Definition Rules:

1. The first statement of every constructor is a call to another constructor within the class using `this()`, or a call to a constructor in the direct parent class using `super()`.
2. The `super()` call may not be used after the first statement of the constructor.



3. If no `super()` call is declared in a constructor, Java will insert a no-argument `super()` as the first statement of the constructor.
4. If the parent doesn't have a no-argument constructor and the child doesn't define any constructors, the compiler will throw an error and try to insert a default no-argument constructor into the child class.
5. If the parent doesn't have a no-argument constructor, the compiler requires an explicit call to a parent constructor in each child constructor.

Make sure you understand these rules; the exam will often provide code that breaks one or many of these rules and therefore doesn't compile.

## Calling Constructors

Now that we have covered how to define a valid constructor, we'll show you how Java calls the constructors. In Java, the parent constructor is always executed before the child constructor. For example, try to determine what the following code outputs:

```
class Primate {
 public Primate() {
 System.out.println("Primate");
 }
}

class Ape extends Primate {
 public Ape() {
 System.out.println("Ape");
 }
}

public class Chimpanzee extends Ape {
 public static void main(String[] args) {
 new Chimpanzee();
 }
}
```

The compiler first inserts the `super()` command as the first statement of both the `Primate` and `Ape` constructors. Next, the compiler inserts a default no-argument constructor in the `Chimpanzee` class with `super()` as the first statement of the constructor. The code will execute with the parent constructors called first and yields the following output:

```
Primate
Ape
```

The exam creators are fond of questions similar to the previous one that try to get you to determine the output of statements involving constructors. Just remember to “think like the compiler” as much as possible and insert the missing constructors or statements as needed.

## Calling Inherited Class Members

Java classes may use any public or protected member of the parent class, including methods, primitives, or object references. If the parent class and child class are part of the same package, the child class may also use any default members defined in the parent class. Finally, a child class may never access a private member of the parent class, at least not through any direct reference. As you saw in the first example in this chapter, a private member age may be accessed indirectly via a public or protected method.

To reference a member in a parent class, you can just call it directly, as in the following example with the output function `displaySharkDetails()`:

```
class Fish {
 protected int size;
 private int age;

 public Fish(int age) {
 this.age = age;
 }

 public int getAge() {
 return age;
 }
}

public class Shark extends Fish {

 private int numberOfFins = 8;

 public Shark(int age) {
 super(age);
 this.size = 4;
 }

 public void displaySharkDetails() {
 System.out.print("Shark with age: "+getAge());
 }
}
```

```
 System.out.print(" and "+size+" meters long");
 System.out.print(" with "+numberOfFins+" fins");
 }
}
```

In the child class, we use the public method `getAge()` and protected member `size` to access values in the parent class.

As you may remember from Chapter 4, you can use the keyword `this` to access a member of the class. You may also use `this` to access members of the parent class that are accessible from the child class, since a child class inherits all of its parent members. Consider the following alternative definition to the `displaySharkDetails()` method in the previous example:

```
public void displaySharkDetails() {
 System.out.print("Shark with age: "+this.getAge());
 System.out.print(" and "+this.size+" meters long");
 System.out.print(" with "+this.numberOfFins+" fins");
}
```

In Java, you can explicitly reference a member of the parent class by using the `super` keyword, as in the following alternative definition of `displaySharkDetails()`:

```
public void displaySharkDetails() {
 System.out.print("Shark with age: "+super.getAge());
 System.out.print(" and "+super.size+" meters long");
 System.out.print(" with "+this.numberOfFins+" fins");
}
```

In the previous example, we could use `this` or `super` to access a member of the parent class, but is the same true for a member of the child class? Consider this example:

```
public void displaySharkDetails() {
 System.out.print("Shark with age: "+super.getAge());
 System.out.print(" and "+super.size+" meters long");
 System.out.print(" with "+super.numberOfFins+" fins"); // DOES NOT COMPILE
}
```

This code will not compile because `numberOfFins` is only a member of the current class, not the parent class. In other words, we see that `this` and `super` may both be used for methods or variables defined in the parent class, but only `this` may be used for members defined in the current class.

As you'll see in the next section, if the child class overrides a member of the parent class, `this` and `super` could have very different effects when applied to a class member.

***super()* vs. *super***

As discussed in Chapter 4, `this()` and `this` are unrelated in Java. Likewise, `super()` and `super` are quite different but may be used in the same methods on the exam. The first, `super()`, is a statement that explicitly calls a parent constructor and may only be used in the first line of a constructor of a child class. The second, `super`, is a keyword used to reference a member defined in a parent class and may be used throughout the child class.

The exam may try to trick you by using both `super()` and `super` in a constructor. For example, consider the following code:

```
public Rabbit(int age) {
 super();
 super.setAge(10);
}
```

The first statement of the constructor calls the parent's constructor, whereas the second statement calls a function defined in the parent class. Contrast this with the following code, which doesn't compile:

```
public Rabbit(int age) {
 super; // DOES NOT COMPILE
 super().setAge(10); // DOES NOT COMPILE
}
```

This code looks similar to the previous example, but neither line of the constructor will compile since they are using the keywords incorrectly. When you see `super()` or `super` on the exam, be sure to check that they are being used correctly.

## Inheriting Methods

Inheriting a class grants us access to the public and protected members of the parent class, but also sets the stage for collisions between methods defined in both the parent class and the subclass. In this section, we'll review the rules for method inheritance and how Java handles such scenarios.

### Overriding a Method

What if there is a method defined in both the parent and child class? For example, you may want to define a new version of an existing method in a child class that makes use of the definition in the parent class. In this case, you can *override* a method a method by declaring a new method with the signature and return type as the method in the parent class. As you may recall from Chapter 4, the method signature includes the name and list of input parameters.

When you override a method, you may reference the parent version of the method using the `super` keyword. In this manner, the keywords `this` and `super` allow you to select between the current and parent version of a method, respectively. We illustrate this with the following example:

```
public class Canine {
 public double getAverageWeight() {
 return 50;
 }
}

public class Wolf extends Canine {
 public double getAverageWeight() {
 return super.getAverageWeight()+20;
 }
 public static void main(String[] args) {
 System.out.println(new Canine().getAverageWeight());
 System.out.println(new Wolf().getAverageWeight());
 }
}
```

In this example, in which the child class `Wolf` overrides the parent class `Canine`, the method `getAverageWeight()` runs without issue and outputs the following:

```
50.00
70.00
```

You might be wondering, was the use of `super` in the child's method required? For example, what would the following code output if we removed the `super` keyword in the `getAverageWeight()` method of the `Wolf` class?

```
public double getAverageWeight() {
 return getAverageWeight()+20; // INFINITE LOOP
}
```

In this example, the compiler would not call the parent `Canine` method; it would call the current `Wolf` method since it would think you were executing a recursive call. A recursive function is one that calls itself as part of execution, and it is common in programming. A recursive function must have a termination condition. In this example, there is no termination condition; therefore, the application will attempt to call itself infinitely and produce a stack overflow error at runtime.

Overriding a method is not without limitations, though. The compiler performs the following checks when you override a nonprivate method:

1. The method in the child class must have the same signature as the method in the parent class.

2. The method in the child class must be at least as accessible or more accessible than the method in the parent class.
3. The method in the child class may not throw a checked exception that is new or broader than the class of any exception thrown in the parent class method.
4. If the method returns a value, it must be the same or a subclass of the method in the parent class, known as *covariant return types*.

The first rule of overriding a method is somewhat self-explanatory. If two methods have the same name but different signatures, the methods are overloaded, not overridden. As you may recall from our discussion of overloaded methods in Chapter 4, the methods are unrelated to each other and do not share any properties.

### Overloading vs. Overriding

Overloading a method and overriding a method are similar in that they both involve redefining a method using the same name. **They differ in that an overloaded method will use a different signature than an overridden method.** This distinction allows overloaded methods a great deal more freedom in syntax than an overridden method would have. For example, take a look at the following code sample:

```
public class Bird {
 public void fly() {
 System.out.println("Bird is flying");
 }
 public void eat(int food) {
 System.out.println("Bird is eating "+food+" units of food");
 }
}

public class Eagle extends Bird {
 public int fly(int height) {
 System.out.println("Bird is flying at "+height+" meters");
 return height;
 }
 public int eat(int food) { // DOES NOT COMPILE
 System.out.println("Bird is eating "+food+" units of food");
 return food;
 }
}
```

The first method, `fly()`, is overloaded in the subclass `Eagle`, since the signature changes from a no-argument constructor to a constructor with one `int` argument. Because the method is being overloaded and not overridden, the return type can be changed from `void` to `int` without issue.

The second method, `eat()`, is overridden in the subclass `Eagle`, since the signature is the same as it is in the parent class `Bird`—they both take a single argument `int`. Because the method is being overridden, the return type of the method in `Eagle` must be a subclass of the return type of the method in `Bird`. In this example, the return type `void` is not a subclass of `int`; therefore, the compiler will throw an exception on this method definition.

Any time you see a method on the exam with the same name as a method in the parent class, determine whether the method is being overloaded or overridden first; doing so will help you with questions about whether the code will compile.

Let's review some examples of the last three rules of overriding methods so you can learn to spot the issues when they arise:

```
public class Camel {
 protected String getNumberOfHumps() {
 return "Undefined";
 }
}

public class BactrianCamel extends Camel {
 private int getNumberOfHumps() { // DOES NOT COMPILE
 return 2;
 }
}
```

In this example, the method in the child class doesn't compile for two reasons. First, it violates the second rule of overriding methods: the child method must be at least as accessible as the parent. In the example, the parent method uses the `protected` modifier, but the child method uses the `private` modifier, making it less accessible in the child method than in the parent method. It also violates the fourth rule of overriding methods: the return type of the parent method and child method must be covariant. In this example, the return type of the parent method is `String`, whereas the return type of the child method is `int`, neither of which is covariant with each other.

Any time you see a method that appears to be overridden on the example, first check to make sure it is truly being overridden and not overloaded. Once you have confirmed it is being overridden, check that the access modifiers, return types, and any exceptions defined in the method are compatible with one another. Let's take a look at some example methods that use exceptions:

```

public class InsufficientDataException extends Exception {}

public class Reptile {
 protected boolean hasLegs() throws InsufficientDataException {
 throw new InsufficientDataException();
 }
 protected double getWeight() throws Exception {
 return 2;
 }
}

public class Snake extends Reptile {
 protected boolean hasLegs() {
 return false;
 }
 protected double getWeight() throws InsufficientDataException{
 return 2;
 }
}

```

In this example, both parent and child classes define two methods, `hasLegs()` and `getWeight()`. The first method, `hasLegs()`, throws an exception `InsufficientDataException` in the parent class but doesn't throw an exception in the child class. This does not violate the third rule of overriding methods, though, as no new exception is defined. In other words, **a child method may hide or eliminate a parent method's exception without issue.**

The second method, `getWeight()`, throws `Exception` in the parent class and `InsufficientDataException` in the child class. This is also permitted, as `InsufficientDataException` is a subclass of `Exception` by construction.

Neither of the methods in the previous example violates the third rule of overriding methods, so the code compiles and runs without issue. Let's review some examples that do violate the third rule of overriding methods:

```

public class InsufficientDataException extends Exception {}

public class Reptile {
 protected double getHeight() throws InsufficientDataException {
 return 2;
 }
 protected int getLength() {
 return 10;
 }
}

```



```
public class Snake extends Reptile {
 protected double getHeight() throws Exception { // DOES NOT COMPILE
 return 2;
 }
 protected int getLength() throws InsufficientDataException { // DOES NOT COMPILE
 return 10;
 }
}
```

Unlike the earlier example, neither of the methods in the child class of this code will compile. The `getHeight()` method in the parent class throws an `InsufficientDataException`, whereas the method in the child class throws an `Exception`. Since `Exception` is not a subclass of `InsufficientDataException`, the third rule of overriding methods is violated and the code will not compile. Coincidentally, `Exception` is a superclass of `InsufficientDataException`.

Next, the `getLength()` method doesn't throw an exception in the parent class, but it does throw an exception, `InsufficientDataException`, in the child class. In this manner, the child class defines a new exception that the parent class did not, which is a violation of the third rule of overriding methods.

The last three rules of overriding a method may seem arbitrary or confusing at first, but as you'll see later in this chapter when we discuss polymorphism, they are needed for consistency of the language. Without these rules in place, it is possible to create contradictions within the Java language.

## Redeclaring *private* Methods

The previous section defined the behavior if you override a public or protected method in the class. Now let's expand our discussion to private methods. **In Java, it is not possible to override a private method in a parent class since the parent method is not accessible from the child class.** Just because a child class doesn't have access to the parent method, doesn't mean the child class can't define its own version of the method. It just means, strictly speaking, that the new method is not an overridden version of the parent class's method.

Java permits you to redeclare a new method in the child class with the same or modified signature as the method in the parent class. This method in the child class is a separate and independent method, unrelated to the parent version's method, so none of the rules for overriding methods are invoked. For example, let's return to the `Camel` example we used in the previous section and show two related classes that define the same method:

```
public class Camel {
 private String getNumberOfHumps() {
 return "Undefined";
 }
}
```

```
public class BactrianCamel extends Camel {
 private int getNumberOfHumps() {
 return 2;
 }
}
```

This code compiles without issue. Notice that the return type differs in the child method from `String` to `int`. In this example, the method `getNumberOfHumps()` in the parent class is hidden, so the method in the child class is a new method and not an override of the method in the parent class. As you saw in the previous section, if the method in the parent class were `public` or `protected`, the method in the child class would not compile because it would violate two rules of overriding methods. The parent method in this example is `private`, so there are no such issues.

## Hiding Static Methods

A *hidden method* occurs when a child class defines a static method with the same name and signature as a static method defined in a parent class. Method hiding is similar but not exactly the same as method overriding. First, the four previous rules for overriding a method must be followed when a method is hidden. In addition, a new rule is added for hiding a method, namely that the usage of the `static` keyword must be the same between parent and child classes. **The following list summarizes the five rules for hiding a method:**

1. The method in the child class must have the same signature as the method in the parent class.
2. The method in the child class must be at least as accessible or more accessible than the method in the parent class.
3. The method in the child class may not throw a checked exception that is new or broader than the class of any exception thrown in the parent class method.
4. If the method returns a value, it must be the same or a subclass of the method in the parent class, known as *covariant return types*.
5. The method defined in the child class must be marked as `static` if it is marked as `static` in the parent class (method hiding). Likewise, the method must not be marked as `static` in the child class if it is not marked as `static` in the parent class (method overriding).

Note that the first four are the same as the rules for overriding a method.

Let's review some examples of the new rule:

```
public class Bear {
 public static void eat() {
 System.out.println("Bear is eating");
 }
}
```

```
public class Panda extends Bear {
 public static void eat() {
 System.out.println("Panda bear is chewing");
 }
 public static void main(String[] args) {
 Panda.eat();
 }
}
```

In this example, the code compiles and runs without issue. The `eat()` method in the child class hides the `eat()` method in the parent class. Because they are both marked as `static`, this is not considered an overridden method. Let's contrast this with examples that violate the fifth rule:

```
public class Bear {
 public static void sneeze() {
 System.out.println("Bear is sneezing");
 }
 public void hibernate() {
 System.out.println("Bear is hibernating");
 }
}

public class Panda extends Bear {
 public void sneeze() { // DOES NOT COMPILE
 System.out.println("Panda bear sneezes quietly");
 }
 public static void hibernate() { // DOES NOT COMPILE
 System.out.println("Panda bear is going to sleep");
 }
}
```

In this example, `sneeze()` is marked as `static` in the parent class but not in the child class. The compiler detects that you're trying to override a method that should be hidden and generates a compiler error. In the second method, `hibernate()` is an instance member in the parent class but a static method in the child class. In this scenario, the compiler thinks that you're trying to hide a method that should be overridden and also generates a compiler error.



As you saw in the previous example, hiding static methods is fraught with pitfalls and potential problems and as a practice should be avoided. Though you might see questions on the exam that contain hidden static methods that are syntactically correct, avoid hiding static methods in your own work, since it tends to lead to confusing and difficult-to-read code. You should not reuse the name of a static method in your class if it is already used in the parent class.

## Overriding vs. Hiding Methods

In our description of hiding of static methods, we indicated there was a distinction between overriding and hiding methods. **Unlike overriding a method, in which a child method replaces the parent method in calls defined in both the parent and child, hidden methods only replace parent methods in the calls defined in the child class.**

**At runtime the child version of an overridden method is always executed for an instance regardless of whether the method call is defined in a parent or child class method. In this manner, the parent method is never used unless an explicit call to the parent method is referenced, using the syntax *ParentClassName.method()*. Alternatively, at runtime the parent version of a hidden method is always executed if the call to the method is defined in the parent class.** Let's take a look at an example:

```
public class Marsupial {
 public static boolean isBiped() {
 return false;
 }
 public void getMarsupialDescription() {
 System.out.println("Marsupial walks on two legs: "+isBiped());
 }
}

public class Kangaroo extends Marsupial {
 public static boolean isBiped() {
 return true;
 }
 public void getKangarooDescription() {
 System.out.println("Kangaroo hops on two legs: "+isBiped());
 }
 public static void main(String[] args) {
 Kangaroo joey = new Kangaroo();
 joey.getMarsupialDescription();
 joey.getKangarooDescription();
 }
}
```

In this example, the code compiles and runs without issue, outputting the following:

```
Marsupial walks on two legs: false
Kangaroo hops on two legs: true
```

Notice that `isBiped()` returns `false` in the parent class and `true` in the child class. In the first method call, the parent method `getMarsupialDescription()` is used. The `Marsupial` class only knows about `isBiped()` from its own class definition, so it outputs `false`. In the second method call, the child executes a method of `isBiped()`, which hides the parent method's version and returns `true`.

Contrast this first example with the following example, which uses an overridden version of `isBiped()` instead of a hidden version:

```
class Marsupial {
 public boolean isBiped() {
 return false;
 }
 public void getMarsupialDescription() {
 System.out.println("Marsupial walks on two legs: "+isBiped());
 }
}

public class Kangaroo extends Marsupial {
 public boolean isBiped() {
 return true;
 }
 public void getKangarooDescription() {
 System.out.println("Kangaroo hops on two legs: "+isBiped());
 }
 public static void main(String[] args) {
 Kangaroo joey = new Kangaroo();
 joey.getMarsupialDescription();
 joey.getKangarooDescription();
 }
}
```

This code also compiles and runs without issue, but it outputs slightly different text:

```
Marsupial walks on two legs: true
Kangaroo hops on two legs: true
```

In this example, the `isBiped()` method is overridden, not hidden, in the child class. Therefore, it is replaced at runtime in the parent class with the call to the child class's method.

Make sure you understand these examples as they show how hidden and overridden methods are fundamentally different. This example makes use of polymorphism, which we'll discuss later in this chapter.

## Creating *final* methods

We conclude our discussion of method inheritance with a somewhat self-explanatory rule: **final methods cannot be overridden**. If you recall our discussion of modifiers from Chapter 4, you can create a method with the `final` keyword. By doing so, though, you forbid a child class from overriding this method. **This rule is in place both when you override a method and when you hide a method. In other words, you cannot hide a static method in a parent class if it is marked as final.**

Let's take a look at an example:

```
public class Bird {
 public final boolean hasFeathers() {
 return true;
 }
}

public class Penguin extends Bird {
 public final boolean hasFeathers() { // DOES NOT COMPILE
 return false;
 }
}
```

In this example, the method `hasFeathers()` is marked as `final` in the parent class `Bird`, so the child class `Penguin` cannot override the parent method, resulting in a compiler error. Note that whether or not the child method used the `final` keyword is irrelevant—the code will not compile either way.

### Why Mark a Method as *final*?

Although marking methods as `final` prevents them from being overridden, it does have advantages in practice. For example, you'd mark a method as `final` when you're defining a parent class and want to guarantee certain behavior of a method in the parent class, regardless of which child is invoking the method.

For example, in the previous example with `Birds`, the author of the parent class may want to ensure the method `hasFeathers()` always returns `true`, regardless of the child class instance on which it is invoked. The author is confident that there is no example of a `Bird` in which feathers are not present.

The reason methods are not commonly marked as `final` in practice, though, is that it may be difficult for the author of a parent class method to consider all of the possible ways her child class may be used. For example, although all adult birds have feathers, a baby chick doesn't; therefore, if you have an instance of a `Bird` that is a chick, it would not have feathers. In short, the `final` modifier is only used on methods when the author of the parent method wants to guarantee very precise behavior.

## Inheriting Variables

As you saw with method overriding, there were a lot of rules when two methods have the same signature and are defined in both the parent and child classes. Luckily, the rules for variables with the same name in the parent and child classes are a lot simpler, because Java doesn't allow variables to be overridden but instead hidden.

## Hiding Variables

When you hide a variable, you define a variable with the same name as a variable in a parent class. This creates two copies of the variable within an instance of the child class: one instance defined for the parent reference and another defined for the child reference.

As when hiding a static method, you can't override a variable; you can only hide it. Also similar to hiding a static method, the rules for accessing the parent and child variables are quite similar. If you're referencing the variable from within the parent class, the variable defined in the parent class is used. Alternatively, if you're referencing the variable from within a child class, the variable defined in the child class is used. Likewise, you can reference the parent value of the variable with an explicit use of the `super` keyword. Consider the following example:

```
public class Rodent {
 protected int tailLength = 4;
 public void getRodentDetails() {
 System.out.println("[parentTail="+tailLength+"]");
 }
}

public class Mouse extends Rodent {
 protected int tailLength = 8;
 public void getMouseDetails() {
 System.out.println("[tail="+tailLength +",parentTail="+super.tailLength+"]");
 }
 public static void main(String[] args) {
 Mouse mouse = new Mouse();
 mouse.getRodentDetails();
 mouse.getMouseDetails();
 }
}
```

This code compiles without issue and outputs the following when executed:

```
[parentTail=4]
[tail=8,parentTail=4]
```

Notice that the instance of `Mouse` contains two copies of the `tailLength` variables: one defined in the parent and one defined in the child. These instances are kept separate from each other, allowing our instance of `Mouse` to reference both `tailLength` values independently. In the first method call, `getRodentDetails()`, the parent method outputs the parent value of the `tailLength` variable. In the second method call, `getMouseDetails()`, the child method outputs both the child and parent version of the `tailLength` variables, using the `super` keyword to access the parent variable's value.

The important thing to remember is that there is no notion of overriding a member variable. For example, there is no code change that could be made to cause Java to override the value of `tailLength`, making it the same in both parent and child. These rules are the same regardless of whether the variable is an instance variable or a static variable.



## Real World Scenario

### Don't Hide Variables in Practice

Although Java allows you to hide a variable defined in a parent class with one defined in a child class, it is considered an extremely poor coding practice. For example, take a look at the following code, which uses a hidden variable `length`, marked as `public` in both parent and child classes.

```
public class Animal {
 public int length = 2;
}

public class Jellyfish extends Animal {
 public int length = 5;
 public static void main(String[] args) {
 Jellyfish jellyfish = new Jellyfish();
 Animal animal = new Jellyfish();
 System.out.println(jellyfish.length);
 System.out.println(animal.length);
 }
}
```

This code compiles without issue. Here's the output:

```
5
2
```



Notice the same type of object was created twice, but the reference to the object determines which value is seen as output. If the object `Jellyfish` was passed to a method by an `Animal` reference, as you'll see in the section "Understanding Polymorphism," later in this chapter, the wrong value might be used.

Hiding variables makes the code very confusing and difficult to read, especially if you start modifying the value of the variable in both the parent and child methods, since it may not be clear which variable you're updating.

When defining a new variable in a child class, it is considered good coding practice to select a name for the variable that is not already a public, protected, or default variable in use in a parent class. Hiding private variables is considered less problematic because the child class did not have access to the variable in the parent class to begin with.

## Creating Abstract Classes

Let's say you want to define a parent class that other developers are going to subclass. Your goal is to provide some reusable variables and methods to developers in the parent class, whereas the developers provide specific implementations or overrides of other methods in the child classes. Furthermore, let's say you also don't want an instance of the parent class to be instantiated unless it is an instance of the child class.

For example, you might define an `Animal` parent class that a number of classes extend from and use but for which an instance of `Animal` itself cannot be instantiated. All subclasses of the `Animal` class, such as `Swan`, are required to implement a `getName()` method, but there is no implementation for the method in the parent `Animal` class. How do you ensure all classes that extend `Animal` provide an implementation for this method?

In Java, you can accomplish this task by using an abstract class and abstract method. **An *abstract class* is a class that is marked with the `abstract` keyword and cannot be instantiated. An *abstract method* is a method marked with the `abstract` keyword defined in an abstract class, for which no implementation is provided in the class in which it is declared.**

The following code is based on our `Animal` and `Swan` description:

```
public abstract class Animal {
 protected int age;
 public void eat() {
 System.out.println("Animal is eating");
 }
 public abstract String getName();
}
```

```
public class Swan extends Animal {
 public String getName() {
 return "Swan";
 }
}
```

The first thing to notice about this sample code is that the `Animal` class is declared abstract and `Swan` is not. Next, the member `age` and the method `eat()` are marked as protected and public, respectively; therefore, they are inherited in subclasses such as `Swan`. Finally, the abstract method `getName()` is terminated with a semicolon and doesn't provide a body in the parent class `Animal`. This method is implemented with the same name and signature as the parent method in the `Swan` class.

## Defining an Abstract Class

The previous sample code illustrates a number of important rules about abstract classes. For example, **an abstract class may include nonabstract methods and variables**, as you saw with the variable `age` and the method `eat()`. In fact, an abstract class is not required to include any abstract methods. For example, the following code compiles without issue even though it doesn't define any abstract methods:

```
public abstract class Cow {
}
```

Although an abstract class doesn't have to implement any abstract methods, **an abstract method may only be defined in an abstract class**. For example, the following code won't compile because an abstract method is not defined within an abstract class:

```
public class Chicken {
 public abstract void peck(); // DOES NOT COMPILE
}
```

The exam creators are fond of questions like this one, which mixes nonabstract classes with abstract methods. They are also fond of questions with methods marked as abstract for which an implementation is also defined. For example, neither method in the following code will compile because the methods are marked as abstract:

```
public abstract class Turtle {
 public abstract void swim() {} // DOES NOT COMPILE
 public abstract int getAge() { // DOES NOT COMPILE
 return 10;
 }
}
```

The first method, `swim()`, doesn't compile because two brackets are provided instead of a semicolon, and Java interprets this as providing a body to an abstract method. The second method, `getAge()`, doesn't compile because it also provides a body to an abstract method. Pay close attention to `swim()`, because you'll likely see a question like this on the exam.

### Default Method Implementations in Abstract Classes

Although you can't provide a default implementation to an abstract method in an abstract class, **you can still define a method with a body—you just can't mark it as abstract. As long as you do not mark it as `final`, the subclass still has the option to override it**, as explained in the previous section.

Next, we note that an abstract class cannot be marked as `final` for a somewhat obvious reason. **By definition, an abstract class is one that must be extended by another class to be instantiated, whereas a `final` class can't be extended by another class. By marking an abstract class as `final`, you're saying the class can never be instantiated, so the compiler refuses to process the code.** For example, the following code snippet will not compile:

```
public final abstract class Tortoise { // DOES NOT COMPILE
}
```

Likewise, **an abstract method may not be marked as `final` for the same reason that an abstract class may not be marked as `final`.** Once marked as `final`, the method can never be overridden in a subclass, making it impossible to create a concrete instance of the abstract class.

```
public abstract class Goat {
 public abstract final void chew(); // DOES NOT COMPILE
}
```

Finally, **a method may not be marked as both `abstract` and `private`.** This rule makes sense if you think about it. How would you define a subclass that implements a required method if the method is not accessible by the subclass itself? The answer is you can't, which is why the compiler will complain if you try to do the following:

```
public abstract class Whale {
 private abstract void sing(); // DOES NOT COMPILE
}

public class HumpbackWhale extends Whale {
 private void sing() {
 System.out.println("Humpback whale is singing");
 }
}
```

In this example, the abstract method `sing()` defined in the parent class `Whale` is not visible to the subclass `HumpbackWhale`. Even though `HumpbackWhale` does provide an implementation, it is not considered an override of the abstract method since the abstract method is unreachable. The compiler recognizes this in the parent class and throws an exception as soon as `private` and `abstract` are applied to the same method.

If we changed the access modified from `private` to `protected` in the parent class `Whale`, would the code compile? Let's take a look:

```
public abstract class Whale {
 protected abstract void sing();
}

public class HumpbackWhale extends Whale {
 private void sing() { // DOES NOT COMPILE
 System.out.println("Humpback whale is singing");
 }
}
```

In this modified example, the code will still not compile but for a completely different reason. **If you remember the rules earlier in this chapter for overriding a method, the subclass cannot reduce the visibility of the parent method, `sing()`. Because the method is declared `protected` in the parent class, it must be marked as `protected` or `public` in the child class. Even with abstract methods, the rules for overriding methods must be followed.**

## Creating a Concrete Class

When working with abstract classes, it is important to remember that by themselves, they cannot be instantiated and therefore do not do much other than define static variables and methods. For example, the following code will not compile as it is an attempt to instantiate an abstract class.

```
public abstract class Eel {
 public static void main(String[] args) {
 final Eel eel = new Eel(); // DOES NOT COMPILE
 }
}
```

An abstract class becomes useful when it is extended by a concrete subclass. **A concrete class is the first nonabstract subclass that extends an abstract class and is required to implement all inherited abstract methods.** When you see a concrete class extending an abstract class on the exam, check that it implements all of the required abstract methods. Let's review this with the following example.

```
public abstract class Animal {
 public abstract String getName();
}
```

```
public class Walrus extends Animal { // DOES NOT COMPILE
}
```

First, note that `Animal` is marked as abstract and `Walrus` is not. In this example, `Walrus` is considered the first concrete subclass of `Animal`. Second, since `Walrus` is the first concrete subclass, it must implement all inherited abstract methods, `getName()` in this example. Because it doesn't, the compiler rejects the code.

Notice that when we define a concrete class as the “first” nonabstract subclass, we include the possibility that another nonabstract class may extend an existing nonabstract class. **The key point is that the first class to extend the nonabstract class must implement all inherited abstract methods.** For example, the following variation will also not compile:

```
public abstract class Animal {
 public abstract String getName();
}

public class Bird extends Animal { // DOES NOT COMPILE
}

public class Flamingo extends Bird {
 public String getName() {
 return "Flamingo";
 }
}
```

Even though a second subclass `Flamingo` implements the abstract method `getName()`, the first concrete subclass `Bird` doesn't; therefore, the `Bird` class will not compile.

## Extending an Abstract Class

Let's expand our discussion of abstract classes by introducing the concept of extending an abstract class with another abstract. We'll repeat our previous `Walrus` example with one minor variation:

```
public abstract class Animal {
 public abstract String getName();
}

public class Walrus extends Animal { // DOES NOT COMPILE
}

public abstract class Eagle extends Animal {
}
```

In this example, we again have an abstract class `Animal` with a concrete subclass `Walrus` that doesn't compile since it doesn't implement a `getName()` method. We also have an abstract class `Eagle`, which like `Walrus` extends `Animal` and doesn't provide an implementation for `getName()`. In this situation, `Eagle` does compile because it is marked as abstract. Be sure you understand why `Walrus` doesn't compile and `Eagle` does in this example.

As you saw in this example, **abstract classes can extend other abstract classes and are not required to provide implementations for any of the abstract methods. It follows, then, that a concrete class that extends an abstract class must implement all inherited abstract methods.** For example, the following concrete class `Lion` must implement two methods, `getName()` and `roar()`:

```
public abstract class Animal {
 public abstract String getName();
}

public abstract class BigCat extends Animal {
 public abstract void roar();
}

public class Lion extends BigCat {
 public String getName() {
 return "Lion";
 }
 public void roar() {
 System.out.println("The Lion lets out a loud ROAR!");
 }
}
```

In this sample code, `BigCat` extends `Animal` but is marked as abstract; therefore, it is not required to provide an implementation for the `getName()` method. The class `Lion` is not marked as abstract, and as the first concrete subclass, it must implement all inherited abstract methods not defined in a parent class.

There is one exception to the rule for abstract methods and concrete classes: a concrete subclass is not required to provide an implementation for an abstract method if an intermediate abstract class provides the implementation. For example, take a look at the following variation on our previous example:

```
public abstract class Animal {
 public abstract String getName();
}

public abstract class BigCat extends Animal {
 public String getName() {
```

```
 return "BigCat";
 }
 public abstract void roar();
}

public class Lion extends BigCat {
 public void roar() {
 System.out.println("The Lion lets out a loud ROAR!");
 }
}
```

In this example, BigCat provides an implementation for the abstract method getName() defined in the abstract Animal class. Therefore, Lion inherits only one abstract method, roar(), and is not required to provide an implementation for the method getName().

Here's one way to think about this: **if an intermediate class provides an implementation for an abstract method, that method is inherited by subclasses as a concrete method, not as an abstract one.** In other words, the subclasses do not consider it an inherited abstract method because it is no longer abstract by the time it reaches the subclasses.

The following are lists of rules for abstract classes and abstract methods that we have covered in this section. Review and understand these rules before taking the exam.

#### **Abstract Class Definition Rules:**

- 1. Abstract classes cannot be instantiated directly.**
- 2. Abstract classes may be defined with any number, including zero, of abstract and non-abstract methods.**
- 3. Abstract classes may not be marked as private or final.**
- 4. An abstract class that extends another abstract class inherits all of its abstract methods as its own abstract methods.**
- 5. The first concrete class that extends an abstract class must provide an implementation for all of the inherited abstract methods.**


#### **Abstract Method Definition Rules:**

- 1. Abstract methods may only be defined in abstract classes.**
- 2. Abstract methods may not be declared private or final.**
- 3. Abstract methods must not provide a method body/implementation in the abstract class for which is it declared.**
- 4. Implementing an abstract method in a subclass follows the same rules for overriding a method. For example, the name and signature must be the same, and the visibility of the method in the subclass must be at least as accessible as the method in the parent class.**


# Implementing Interfaces

Although Java doesn't allow multiple inheritance, it does allow classes to implement any number of interfaces. **An *interface* is an abstract data type that defines a list of abstract public methods that any class implementing the interface must provide.** An interface can also include a list of constant variables and default methods, which we'll cover in this section. In Java, an interface is defined with the `interface` keyword, analogous to the `class` keyword used when defining a class. A class invokes the interface by using the `implements` keyword in its class definition. Refer to Figures 5.4 and 5.5 for proper syntax usage.

**FIGURE 5.4** Defining an interface



**FIGURE 5.5** Implementing an interface



As you see in this example, an interface is not declared an abstract class, although it has many of the same properties of abstract class. Notice that the method modifiers in this



example, `abstract` and `public`, are assumed. In other words, whether or not you provide them, the compiler will automatically insert them as part of the method definition.

A class may implement multiple interfaces, each separated by a comma, such as in the following example:

```
public class Elephant implements WalksOnFourLegs, HasTrunk, Herbivore {
}
```

In the example, if any of the interfaces defined abstract methods, the concrete class `Elephant` would be required to implement those methods.

New to Java 8 is the notion of default and static interface methods, which we'll cover at the end of this section.

## Defining an Interface

It may be helpful to think of an interface as a specialized kind of abstract class, since it shares many of the same properties and rules as an abstract class. The following is a list of rules for creating an interface, many of which you should recognize as adaptations of the rules for defining abstract classes.

1. Interfaces cannot be instantiated directly.
2. An interface is not required to have any methods.
3. An interface may not be marked as `final`.
4. All top-level interfaces are assumed to have `public` or default access, and they must include the `abstract` modifier in their definition. Therefore, marking an interface as `private`, `protected`, or `final` will trigger a compiler error, since this is incompatible with these assumptions.
5. All nondefault methods in an interface are assumed to have the modifiers `abstract` and `public` in their definition. Therefore, marking a method as `private`, `protected`, or `final` will trigger compiler errors as these are incompatible with the `abstract` and `public` keywords.

The fourth rule doesn't apply to inner interfaces, although inner classes and interfaces are not in scope for the OCA exam. The first three rules are identical to the first three rules for creating an abstract class. Imagine we have an interface `WalksOnTwoLegs`, defined as follows:

```
public interface WalksOnTwoLegs {}
```

It compiles without issue, since interfaces are not required to define any methods. Now consider the following two examples, which do not compile:

```
public class TestClass {
 public static void main(String[] args) {
 WalksOnTwoLegs example = new WalksOnTwoLegs(); // DOES NOT COMPILE
 }
}
```

```

 }
}

public final interface WalksOnEightLegs { // DOES NOT COMPILE
}

```

The first example doesn't compile, as `WalksOnTwoLegs` is an interface and cannot be instantiated directly. The second example, `WalksOnEightLegs`, doesn't compile since **interfaces may not be marked as final for the same reason that abstract classes cannot be marked as final.**

The fourth and fifth rule about “assumed keywords” might be new to you, but you should think of these in the same light as the compiler inserting a default no-argument constructor or `super()` statement into your constructor. You may provide these modifiers yourself, although the compiler will insert them automatically if you do not. For example, the following two interface definitions are equivalent, as the compiler will convert them both to the second example:

```

public interface CanFly {
 void fly(int speed);
 abstract void takeoff();
 public abstract double dive();
}

public abstract interface CanFly {
 public abstract void fly(int speed);
 public abstract void takeoff();
 public abstract double dive();
}

```

In this example, the `abstract` keyword is first automatically added to the interface definition. Then, each method is prepended with `abstract` and `public` keywords. If the method already has either of these keywords, then no change is required. Let's take a look at an example that violates the assumed keywords:

```

private final interface CanCrawl { // DOES NOT COMPILE
 private void dig(int depth); // DOES NOT COMPILE
 protected abstract double depth(); // DOES NOT COMPILE
 public final void surface(); // DOES NOT COMPILE
}

```

Every single line of this example doesn't compile. The first line doesn't compile for two reasons. First, it is marked as `final`, which cannot be applied to an interface since it conflicts with the assumed `abstract` keyword. Next, it is marked as `private`, which conflicts with the `public` or default required access for interfaces. The second and third line do not compile because all interface methods are assumed to be `public` and marking them

as `private` or `protected` throws a compiler error. Finally, the last line doesn't compile because the method is marked as `final` and since interface methods are assumed to be abstract, the compiler throws an exception for using both `abstract` and `final` keywords on a method.



Adding the assumed keywords to an interface is a matter of personal preference, although it is considered good coding practice to do so. Code with the assumed keywords written out tends to be easier and clearer to read, and leads to fewer potential conflicts, as you saw in the previous examples.

Be sure to review the previous example and understand why each of the lines doesn't compile. There will likely be at least one question on the exam in which an interface or interface method uses an invalid modifier.

## Inheriting an Interface

There are two inheritance rules you should keep in mind when extending an interface:

1. An interface that extends another interface, as well as an abstract class that implements an interface, inherits all of the abstract methods as its own abstract methods.
2. The first concrete class that implements an interface, or extends an abstract class that implements an interface, must provide an implementation for all of the inherited abstract methods.

Like an abstract class, an interface may be extended using the `extends` keyword. In this manner, the new child interface inherits all the abstract methods of the parent interface. Unlike an abstract class, though, an interface may extend multiple interfaces. Consider the following example:

```
public interface HasTail {
 public int getTailLength();
}

public interface HasWhiskers {
 public int getNumberOfWhiskers();
}

public interface Seal extends HasTail, HasWhiskers {
}
```

Any class that implements the `Seal` interface must provide an implementation for all methods in the parent interfaces—in this case, `getTailLength()` and `getNumberOfWhiskers()`.

What about an abstract class that implements an interface? In this scenario, the abstract class is treated in the same way as an interface extending another interface. In other words, the abstract class inherits the abstract methods of the interface but is not required to implement them. That said, like an abstract class, the first concrete class to extend the abstract class must implement all the inherited abstract methods of the interface. We illustrate this in the following example:

```
public interface HasTail {
 public int getTailLength();
}

public interface HasWhiskers {
 public int getNumberOfWhiskers();
}

public abstract class HarborSeal implements HasTail, HasWhiskers {
}

public class LeopardSeal implements HasTail, HasWhiskers { // DOES NOT COMPILE
}
```

In this example, we see that HarborSeal is an abstract class and compiles without issue. Any class that extends HarborSeal will be required to implement all of the methods in the HasTail and HasWhiskers interface. Alternatively, LeopardSeal is not an abstract class, so it must implement all the interface methods within its definition. In this example, LeopardSeal doesn't provide an implementation for the interface methods, so the code doesn't compile.

## Classes, Interfaces, and Keywords

The exam creators are fond of questions that mix class and interface terminology. **Although a class can implement an interface, a class cannot extend an interface. Likewise, whereas an interface can extend another interface, an interface cannot implement another interface.** The following examples illustrate these principles:

```
public interface CanRun {}

public class Cheetah extends CanRun {} // DOES NOT COMPILE

public class Hyena {}

public interface HasFur extends Hyena {} // DOES NOT COMPILE
```

The first example shows a class trying to extend an interface that doesn't compile. The second example shows an interface trying to extend a class, which also doesn't compile.

Be wary of examples on the exam that mix class and interface definitions. Make sure the only connection between a class and an interface is with the *class implements interface* syntax.

## Abstract Methods and Multiple Inheritance

Since Java allows for multiple inheritance via interfaces, you might be wondering what will happen if you define a class that inherits from two interfaces that contain the same abstract method:

```
public interface Herbivore {
 public void eatPlants();
}
```

```
public interface Omnivore {
 public void eatPlants();
 public void eatMeat();
}
```

In this scenario, the signatures for the two interface methods `eatPlants()` are compatible, so you can define a class that fulfills both interfaces simultaneously:

```
public class Bear implements Herbivore, Omnivore {
 public void eatMeat() {
 System.out.println("Eating meat");
 }
 public void eatPlants() {
 System.out.println("Eating plants");
 }
}
```

Why does this work? Remember that interface methods in this example are abstract and define the “behavior” that the class implementing the interface must have. If two abstract interface methods have identical behaviors—or in this case the same method signature—creating a class that implements one of the two methods automatically implements the second method. In this manner, the interface methods are considered duplicates since they have the same signature.

What happens if the two methods have different signatures? If the method name is the same but the input parameters are different, there is no conflict because this is considered a method overload. We demonstrate this principle in the following example:

```
public interface Herbivore {
 public int eatPlants(int quantity);
}
```

```
public interface Omnivore {
 public void eatPlants();
}

public class Bear implements Herbivore, Omnivore {
 public int eatPlants(int quantity) {
 System.out.println("Eating plants: "+quantity);
 return quantity;
 }
 public void eatPlants() {
 System.out.println("Eating plants");
 }
}
```

In this example, we see that the class that implements both interfaces must provide implements of both versions of `eatPlants()`, since they are considered separate methods. Notice that it doesn't matter if the return type of the two methods is the same or different, because the compiler treats these methods as independent.

Unfortunately, **if the method name and input parameters are the same but the return types are different between the two methods, the class or interface attempting to inherit both interfaces will not compile.** The reason the code doesn't compile has less to do with interfaces and more to do with class design, as discussed in Chapter 4. It is not possible in Java to define two methods in a class with the same name and input parameters but different return types. Given the following two interface definitions for `Herbivore` and `Omnivore`, the following code will not compile:

```
public interface Herbivore {
 public int eatPlants();
}

public interface Omnivore {
 public void eatPlants();
}

public class Bear implements Herbivore, Omnivore {
 public int eatPlants() { // DOES NOT COMPILE
 System.out.println("Eating plants: 10");
 return 10;
 }
 public void eatPlants() { // DOES NOT COMPILE
 System.out.println("Eating plants");
 }
}
```

The code doesn't compile, as the class defines two methods with the same name and input parameters but different return types. If we were to remove either definition of `eatPlants()`, the compiler would stop because the definition of `Bear` would be missing one of the required methods. In other words, there is no implementation of the `Bear` class that inherits from `Herbivore` and `Omnivore` that the compiler would accept.

The compiler would also throw an exception if you define an interface or abstract class that inherits from two conflicting interfaces, as shown here:

```
public interface Herbivore {
 public int eatPlants();
}

public interface Omnivore {
 public void eatPlants();
}

public interface Supervore extends Herbivore, Omnivore {} // DOES NOT COMPILE

public abstract class AbstractBear implements Herbivore, Omnivore {}

// DOES NOT COMPILE
```

Even without implementation details, the compiler detects the problem with the abstract definition and prevents compilation.

This concludes our discussion of abstract interface methods and multiple inheritance. We'll return to this discussion shortly after we introduce default interface methods. You'll see that things work a bit differently with default interface methods.

## Interface Variables

Let's expand our discussion of interfaces to include interface variables, which can be defined within an interface. Like interface methods, interface variables are assumed to be public. Unlike interface methods, though, interface variables are also assumed to be static and final.

Here are two interface variables rules:

1. Interface variables are assumed to be public, static, and final. Therefore, marking a variable as private or protected will trigger a compiler error, as will marking any variable as abstract.
2. The value of an interface variable must be set when it is declared since it is marked as final.

In this manner, interface variables are essentially constant variables defined on the interface level. Because they are assumed to be static, they are accessible even without

an instance of the interface. Like our earlier CanFly example, the following two interface definitions are equivalent, because the compiler will automatically convert them both to the second example:

```
public interface CanSwim {
 int MAXIMUM_DEPTH = 100;
 final static boolean UNDERWATER = true;
 public static final String TYPE = "Submersible";
}
```

```
public interface CanSwim {
 public static final int MAXIMUM_DEPTH = 100;
 public static final boolean UNDERWATER = true;
 public static final String TYPE = "Submersible";
}
```

As we see in this example, the compiler will automatically insert `public static final` to any constant interface variables it finds missing those modifiers. Also note that it is a common coding practice to use uppercase letters to denote constant values within a class.

Based on these rules, it should come as no surprise that the following entries will not compile:

```
public interface CanDig {
 private int MAXIMUM_DEPTH = 100; // DOES NOT COMPILE
 protected abstract boolean UNDERWATER = false; // DOES NOT COMPILE
 public static String TYPE; // DOES NOT COMPILE
}
```

The first example, `MAXIMUM_DEPTH`, doesn't compile because the `private` modifier is used, and all interface variables are assumed to be `public`. The second line, `UNDERWATER`, doesn't compile for two reasons. It is marked as `protected`, which conflicts with the assumed modifier `public`, and it is marked as `abstract`, which conflicts with the assumed modifier `final`. Finally, the last example, `TYPE`, doesn't compile because it is missing a value. Unlike the other examples, the modifiers are correct, but as you may remember from Chapter 4, you must provide a value to a `static final` member of the class when it is defined.

## Default Interface Methods

With the release of Java 8, the authors of Java have introduced a new type of method to an interface, referred to as a default method. **A default method is a method defined within an interface with the default keyword in which a method body is provided.** Contrast default methods with “regular” methods in an interface, which are assumed to be abstract and may not have a method body.



A default method within an interface defines an abstract method with a default implementation. In this manner, classes have the option to override the default method if they need to, but they are not required to do so. If the class doesn't override the method, the default implementation will be used. In this manner, the method definition is concrete, not abstract.

The purpose of adding default methods to the Java language was in part to help with code development and backward compatibility. Imagine you have an interface that is shared among dozens or even hundreds of users that you would like to add a new method to. If you just update the interface with the new method, the implementation would break among all of your subscribers, who would then be forced to update their code. In practice, this might even discourage you from making the change altogether. By providing a default implementation of the method, though, the interface becomes backward compatible with the existing codebase, while still providing those individuals who do want to use the new method with the option to override it.

The following is an example of a default method defined in an interface:

```
public interface ISwarmBlooded {
 boolean hasScales();
 public default double getTemperature() {
 return 10.0;
 }
}
```

This example defines two interface methods, one is a normal abstract method and the other a default method. Note that both methods are assumed to be public, as all methods of an interface are all public. The first method is terminated with a semicolon and doesn't provide a body, whereas the second default method provides a body. Any class that implements ISwarmBlooded may rely on the default implementation of getTemperature() or override the method and create its own version.

Note that the default access modifier as defined in Chapter 4 is completely different from the default method defined in this chapter. We defined a default access modifier in Chapter 4 as lack of an access modifier, which indicated a class may access a class, method, or value within another class if both classes are within the same package. In this chapter, we are specifically talking about the keyword default as applied to a method within an interface. Because all methods within an interface are assumed to be public, the access modifier for a default method is therefore public.

The following are the default interface method rules you need to be familiar with:

1. A default method may only be declared within an interface and not within a class or abstract class.
2. A default method must be marked with the default keyword. If a method is marked as default, it must provide a method body.
3. A default method is not assumed to be static, final, or abstract, as it may be used or overridden by a class that implements the interface.

**4. Like all methods in an interface, a default method is assumed to be public and will not compile if marked as private or protected.**

The first rule should give you some comfort in that you'll only see default methods in interfaces. If you see them in a class on the exam, assume the code will not compile. The second rule just denotes syntax, as default methods must use the default keyword. For example, the following code snippets will not compile:

```
public interface Carnivore {
 public default void eatMeat(); // DOES NOT COMPILE
 public int getRequiredFoodAmount() { // DOES NOT COMPILE
 return 13;
 }
}
```

In this example, the first method, `eatMeat()`, doesn't compile because it is marked as default but doesn't provide a method body. The second method, `getRequiredFoodAmount()`, also doesn't compile because it provides a method body but is not marked with the default keyword.

Unlike interface variables, which are assumed static class members, default methods cannot be marked as `static` and require an instance of the class implementing the interface to be invoked. They can also not be marked as `final` or `abstract`, because they are allowed to be overridden in subclasses but are not required to be overridden.

When an interface extends another interface that contains a default method, it may choose to ignore the default method, in which case the default implementation for the method will be used. Alternatively, **the interface may override the definition of the default method using the standard rules for method overriding, such as not limiting the accessibility of the method and using covariant returns**. Finally, the interface may redeclare the method as `abstract`, requiring classes that implement the new interface to explicitly provide a method body. Analogous options apply for an abstract class that implements an interface.

For example, the following class overrides one default interface method and redeclares a second interface method as `abstract`:

```
public interface HasFins {
 public default int getNumberOfFins() {
 return 4;
 }
 public default double getLongestFinLength() {
 return 20.0;
 }
 public default boolean doFinsHaveScales() {
 return true;
 }
}
```

```
public interface SharkFamily extends HasFins {
 public default int getNumberOfFins() {
 return 8;
 }
 public double getLongestFinLength();
 public boolean doFinsHaveScales() { // DOES NOT COMPILE
 return false;
 }
}
```

In this example, the first interface, `HasFins`, defines three default methods: `getNumberOfFins()`, `getLongestFinLength()`, and `doFinsHaveScales()`. The second interface, `SharkFamily`, extends `HasFins` and overrides the default method `getNumberOfFins()` with a new method that returns a different value. Next, the `SharkFamily` interface replaces the default method `getLongestFinLength()` with a new abstract method, forcing any class that implements the `SharkFamily` interface to provide an implementation of the method. Finally, the `SharkFamily` interface overrides the `doFinsHaveScales()` method but doesn't mark the method as default. Since interfaces may only contain methods with a body that are marked as default, the code will not compile.

Because default methods are new to Java 8, there will probably be a few questions on the exam about them, although they likely will not be any more difficult than the previous example.

## Default Methods and Multiple Inheritance

You may have realized that by allowing default methods in interfaces, coupled with the fact a class may implement multiple interfaces, Java has essentially opened the door to multiple inheritance problems. For example, what value would the following code output?

```
public interface Walk {
 public default int getSpeed() {
 return 5;
 }
}

public interface Run {
 public default int getSpeed() {
 return 10;
 }
}

public class Cat implements Walk, Run { // DOES NOT COMPILE
 public static void main(String[] args) {
 System.out.println(new Cat().getSpeed());
 }
}
```

In this example, `Cat` inherits the two default methods for `getSpeed()`, so which does it use? Since `Walk` and `Run` are considered siblings in terms of how they are used in the `Cat` class, it is not clear whether the code should output 5 or 10. The answer is that the code outputs neither value—it fails to compile.

If a class implements two interfaces that have default methods with the same name and signature, the compiler will throw an error. There is an exception to this rule, though: if the subclass overrides the duplicate default methods, the code will compile without issue—the ambiguity about which version of the method to call has been removed. For example, the following modified implementation of `Cat` will compile and output 1:

```
public class Cat implements Walk, Run {
 public int getSpeed() {
 return 1;
 }

 public static void main(String[] args) {
 System.out.println(new Cat().getSpeed());
 }
}
```

You can see that having a class that implements or inherits two duplicate default methods forces the class to implement a new version of the method, or the code will not compile. This rule holds true even for abstract classes that implement multiple interfaces, because the default method could be called in a concrete method within the abstract class.

## Static Interface Methods

Java 8 also now includes support for static methods within interfaces. These methods are defined explicitly with the `static` keyword and function nearly identically to static methods defined in classes, as discussed in Chapter 4. In fact, there is really only one distinction between a static method in a class and an interface. A static method defined in an interface is not inherited in any classes that implement the interface.

Here are the static interface method rules you need to be familiar with:

1. Like all methods in an interface, a static method is assumed to be `public` and will not compile if marked as `private` or `protected`.
2. To reference the static method, a reference to the name of the interface must be used.

The following is an example of a static method defined in an interface:

```
public interface Hop {
 static int getJumpHeight() {
 return 8;
 }
}
```

The method `getJumpHeight()` works just like a static method as defined in a class. In other words, it can be accessed without an instance of the class using the `Hop.getJumpHeight()` syntax. Also, note that the compiler will automatically insert the access modifier `public` since all methods in interfaces are assumed to be `public`.

The following is an example of a class `Bunny` that implements `Hop`:

```
public class Bunny implements Hop {
 public void printDetails() {
 System.out.println(getJumpHeight()); // DOES NOT COMPILE
 }
}
```

As you can see, without an explicit reference to the name of the interface the code will not compile, even though `Bunny` implements `Hop`. In this manner, the static interface methods are not inherited by a class implementing the interface. The following modified version of the code resolves the issue with a reference to the interface name `Hop`:

```
public class Bunny implements Hop {
 public void printDetails() {
 System.out.println(Hop.getJumpHeight());
 }
}
```

It follows, then, that a class that implements two interfaces containing static methods with the same signature will still compile at runtime, because the static methods are not inherited by the subclass and must be accessed with a reference to the interface name. Contrast this with the behavior you saw for default interface methods in the previous section: the code would compile if the subclass overrode the default methods and would fail to compile otherwise. You can see that static interface methods have none of the same multiple inheritance issues and rules as default interface methods do.

## Understanding Polymorphism

Java supports *polymorphism*, the property of an object to take on many different forms. To put this more precisely, a Java object may be accessed using a reference with the same type as the object, a reference that is a superclass of the object, or a reference that defines an interface the object implements, either directly or through a superclass. Furthermore, a cast is not required if the object is being reassigned to a super type or interface of the object.

Let's illustrate this polymorphism property with the following example:

```
public class Primate {
 public boolean hasHair() {
 return true;
 }
}
```

```
}
}

public interface HasTail {
 public boolean isTailStriped();
}

public class Lemur extends Primate implements HasTail {
 public boolean isTailStriped() {
 return false;
 }
 public int age = 10;
 public static void main(String[] args) {
 Lemur lemur = new Lemur();
 System.out.println(lemur.age);

 HasTail hasTail = lemur;
 System.out.println(hasTail.isTailStriped());

 Primate primate = lemur;
 System.out.println(primate.hasHair());
 }
}
```

This code compiles and executes without issue and yields the following output:

```
10
false
true
```

The most important thing to note about this example is that only one object, `Lemur`, is created and referenced. The ability of an instance of `Lemur` to be passed as an instance of an interface it implements, `HasTail`, as well as an instance of one of its superclasses, `Primate`, is the nature of polymorphism.

Once the object has been assigned a new reference type, only the methods and variables available to that reference type are callable on the object without an explicit cast. For example, the following snippets of code will not compile:

```
HasTail hasTail = lemur;
System.out.println(hasTail.age); // DOES NOT COMPILE
```

```
Primate primate = lemur;
System.out.println(primate.isTailStriped()); // DOES NOT COMPILE
```

In this example, the reference `hasTail` has direct access only to methods defined with the `HasTail` interface; therefore, it doesn't know the variable `age` is part of the object. Likewise, the reference `primate` has access only to methods defined in the `Primate` class, and it doesn't have direct access to the `isTailStriped()` method.

## Object vs. Reference

In Java, all objects are accessed by reference, so as a developer you never have direct access to the object itself. Conceptually, though, you should consider the object as the entity that exists in memory, allocated by the Java runtime environment. Regardless of the type of the reference you have for the object in memory, the object itself doesn't change. For example, since all objects inherit `java.lang.Object`, they can all be reassigned to `java.lang.Object`, as shown in the following example:

```
Lemur lemur = new Lemur();
```

```
Object lemurAsObject = lemur;
```

Even though the `Lemur` object has been assigned a reference with a different type, the object itself has not changed and still exists as a `Lemur` object in memory. What has changed, then, is our ability to access methods within the `Lemur` class with the `lemurAsObject` reference. Without an explicit cast back to `Lemur`, as you'll see in the next section, we no longer have access to the `Lemur` properties of the object.

We can summarize this principle with the following two rules:


1. The type of the object determines which properties exist within the object in memory.
2. The type of the reference to the object determines which methods and variables are accessible to the Java program.

It therefore follows that successfully changing a reference of an object to a new reference type may give you access to new properties of the object, but those properties existed before the reference change occurred.

Let's illustrate this property using the previous example in Figure 5.6. As you can see in the figure, the same object exists in memory regardless of which reference is pointing to it. Depending on the type of the reference, we may only have access to certain methods. For example, the `hasTail` reference has access to the method `isTailStriped()` but doesn't have access to the variable `age` defined in the `Lemur` class. As you'll learn in the next section, it is

possible to reclaim access to the variable `age` by explicitly casting the `hasTail` reference to a reference of type `Lemur`.

**FIGURE 5.6** Object vs. reference



## Casting Objects

In the previous example, we created a single instance of a `Lemur` object and accessed it via superclass and interface references. Once we changed the reference type, though, we lost access to more specific methods defined in the subclass that still exist within the object. We can reclaim those references by casting the object back to the specific subclass it came from:

```
Primate primate = lemur;
Lemur lemur2 = primate; // DOES NOT COMPILE
```

```
Lemur lemur3 = (Lemur)primate;
System.out.println(lemur3.age);
```

In this example, we first try to convert the `primate` reference back to a `lemur` reference, `lemur2`, without an explicit cast. The result is that the code will not compile. In the second example, though, we explicitly cast the object to a subclass of the object `Primate` and we gain access to all the methods available to the `Lemur` class.

**Here are some basic rules to keep in mind when casting variables:**

1. Casting an object from a subclass to a superclass doesn't require an explicit cast.
2. Casting an object from a superclass to a subclass requires an explicit cast.
3. The compiler will not allow casts to unrelated types.
4. Even when the code compiles without issue, an exception may be thrown at runtime if the object being cast is not actually an instance of that class.



The third rule is important; the exam may try to trick you with a cast that the compiler doesn't allow. For example, we were able to cast a `Primate` reference to a `Lemur` reference, because `Lemur` is a subclass of `Primate` and therefore related.

Consider this example:

```
public class Bird {}

public class Fish {
 public static void main(String[] args) {
 Fish fish = new Fish();
 Bird bird = (Bird)fish; // DOES NOT COMPILE
 }
}
```

In this example, the classes `Fish` and `Bird` are not related through any class hierarchy that the compiler is aware of; therefore, the code will not compile.

Casting is not without its limitations. Even though two classes share a related hierarchy, that doesn't mean an instance of one can automatically be cast to another. Here's an example:

```
public class Rodent {
}

public class Capybara extends Rodent {
 public static void main(String[] args) {
 Rodent rodent = new Rodent();
 Capybara capybara = (Capybara)rodent; // Throws ClassCastException at runtime
 }
}
```

This code creates an instance of `Rodent` and then tries to cast it to a subclass of `Rodent`, `Capybara`. Although this code will compile without issue, it will throw a `ClassCastException` at runtime since the object being referenced is not an instance of the `Capybara` class. The thing to keep in mind in this example is the object that was created is not related to the `Capybara` class in any way.



**Although this topic is out of scope for the OCA exam, keep in mind that the `instanceof` operator can be used to check whether an object belongs to a particular class and to prevent `ClassCastException`s at runtime.** Unlike the previous example, the following code snippet doesn't throw an exception at runtime and performs the cast only if the `instanceof` operator returns true.

```
if(rodent instanceof Capybara) {
 Capybara capybara = (Capybara)rodent;
}
```

When reviewing a question on the exam that involves casting and polymorphism, be sure to remember what the instance of the object actually is. Then, focus on whether the compiler will allow the object to be referenced with or without explicit casts.

## Virtual Methods

The most important feature of polymorphism—and one of the primary reasons we have class structure at all—is to support virtual methods. **A virtual method is a method in which the specific implementation is not determined until runtime. In fact, all non-final, non-static, and non-private Java methods are considered virtual methods, since any of them can be overridden at runtime.** What makes a virtual method special in Java is that if you call a method on an object that overrides a method, you get the overridden method, even if the call to the method is on a parent reference or within the parent class.

We'll illustrate this principle with the following example:

```
public class Bird {
 public String getName() {
 return "Unknown";
 }
 public void displayInformation() {
 System.out.println("The bird name is: "+getName());
 }
}

public class Peacock extends Bird {
 public String getName() {
 return "Peacock";
 }
 public static void main(String[] args) {
 Bird bird = new Peacock();
 bird.displayInformation();
 }
}
```

This code compiles and executes without issue and outputs the following:

**The bird name is: Peacock**

As you saw in similar examples in the section “Overriding a Method,” the method `getName()` is overridden in the child class `Peacock`. More importantly, though, the value of the `getName()` method at runtime in the `displayInformation()` method is replaced with the value of the implementation in the subclass `Peacock`.

In other words, even though the parent class `Bird` defines its own version of `getName()` and doesn't know anything about the `Peacock` class during compile-time, at runtime the instance uses the overridden version of the method, as defined on the instance of the object. We emphasize this point by using a reference to the `Bird` class in the `main()` method, although the result would have been the same if a reference to `Peacock` was used.

You now know the true purpose of overriding a method and how it relates to polymorphism. The nature of the polymorphism is that an object can take on many different forms. By combining your understanding of polymorphism with method overriding, you see that objects may be interpreted in vastly different ways at runtime, especially in methods defined in the superclasses of the objects.

## Polymorphic Parameters

One of the most useful applications of polymorphism is the ability to pass instances of a subclass or interface to a method. For example, you can define a method that takes an instance of an interface as a parameter. In this manner, any class that implements the interface can be passed to the method. Since you're casting from a subtype to a supertype, an explicit cast is not required. This property is referred to as *polymorphic parameters* of a method, and we demonstrate it in the following example:

```
public class Reptile {
 public String getName() {
 return "Reptile";
 }
}

public class Alligator extends Reptile {
 public String getName() {
 return "Alligator";
 }
}

public class Crocodile extends Reptile {
 public String getName() {
 return "Crocodile";
 }
}

public class ZooWorker {
 public static void feed(Reptile reptile) {
 System.out.println("Feeding reptile "+reptile.getName());
 }
}
```

```
public static void main(String[] args) {
 feed(new Alligator());
 feed(new Crocodile());
 feed(new Reptile());
}
}
```

This code compiles and executes without issue, yielding the following output:

```
Feeding: Alligator
Feeding: Crocodile
Feeding: Reptile
```

Let's focus on the `feed(Reptile reptile)` method in this example. As you can see, that method was able to handle instances of `Alligator` and `Crocodile` without issue, because both are subclasses of the `Reptile` class. It was also able to accept a matching type `Reptile` class. If we had tried to pass an unrelated class, such as the previously defined `Rodent` or `Capybara` classes, or a superclass such as `java.lang.Object`, to the `feed()` method, the code would not have compiled.

### Polymorphic Parameters and Code Reusability

If you're defining a method that will be accessible outside the current class, either to subclasses of the current class or publicly to objects outside the current class, it is considered good coding practice to use the superclass or interface type of input parameters whenever possible.

As you may remember from Chapter 3, "Core Java APIs," the type `java.util.List` is an interface, not a class. Although there are many classes that implement `java.util.List`, such as `java.util.ArrayList` and `java.util.Vector`, when you're passing an existing `List` you're not usually interested in the particular subclass of the `List`. In this manner, a method that passes a `List` should use the interface type `java.util.List` as the polymorphic parameter type, rather than a specific class that implements `List`, as the code will be more reusable for other types of lists.

For example, it is common to see code such as the following that uses the interface reference type over the class type for greater reusability:

```
java.util.List list = new java.util.ArrayList();
```

## Polymorphism and Method Overriding

Let's conclude this chapter by returning to the last three rules for method overriding to demonstrate how polymorphism requires them to be included as part of the Java specification. You'll see that without such rules in place, it is easy to construct an example with polymorphism in Java.

The first rule is that an overridden method must be at least as accessible as the method it is overriding. Let's assume this rule is not necessary and consider the following example:

```
public class Animal {
 public String getName() {
 return "Animal";
 }
}

public class Gorilla extends Animal {
 protected String getName() { // DOES NOT COMPILE
 return "Gorilla";
 }
}

public class ZooKeeper {
 public static void main(String[] args) {
 Animal animal = new Gorilla();
 System.out.println(animal.getName());
 }
}
```

For the purpose of this discussion, we'll ignore the fact that the implementation of `getName()` in the `Gorilla` class doesn't compile because it is less accessible than the version it is overriding in the `Animal` class.

As you can see, this example creates an ambiguity problem in the `ZooKeeper` class. The reference `animal.getName()` is allowed because the method is public in the `Animal` class, but due to polymorphism, the `Gorilla` object itself has been overridden with a less accessible version, not available to the `ZooKeeper` class. This creates a contradiction in that the compiler should not allow access to this method, but because it is being referenced as an instance of `Animal`, it is allowed. Therefore, **Java eliminates this contradiction, thus disallowing a method from being overridden by a less accessible version of the method.**

**Likewise, a subclass cannot declare an overridden method with a new or broader exception than in the superclass, since the method may be accessed using a reference to the superclass.** For example, if an instance of the subclass is passed to a method using a superclass reference, then the enclosing method would not know about any new checked exceptions that exist on methods for this object, potentially leading to compiled code with

“unchecked” checked exceptions. Therefore, the Java compiler disallows overriding methods with new or broader exceptions.

Finally, overridden methods must use covariant return types for the same kinds of reasons as just discussed. If an object is cast to a superclass reference and the overridden method is called, the return type must be compatible with the return type of the parent method. If the return type in the child is too broad, it will result in an inherent cast exception when accessed through the superclass reference.

For example, if the return type of a method is `Double` in the parent class and is overridden in a subclass with a method that returns `Number`, a superclass of `Double`, then the subclass method would be allowed to return any valid `Number`, including `Integer`, another subclass of `Number`. If we are using the object with a reference to the superclass, that means an `Integer` could be returned when a `Double` was expected. Since `Integer` is not a subclass of `Double`, this would lead to an implicit cast exception as soon as the value was referenced. Java solves this problem by only allowing covariant return types for overridden methods.

## Summary

This chapter took the basic class structure we presented in Chapter 4 and expanded it by introducing the notion of inheritance. Java classes follow a multilevel single-inheritance pattern in which every class has exactly one direct parent class, with all classes eventually inheriting from `java.lang.Object`. Java interfaces simulate a limited form of multiple inheritance, since Java classes may implement multiple interfaces.

Inheriting a class gives you access to all of the public and protected methods of the class, but special rules for constructors and overriding methods must be followed or the code will not compile. For example, if the parent class doesn't include a no-argument constructor, an explicit call to a parent constructor must be provided in the child's constructors. Pay close attention on the exam to any class that defines a constructor with arguments and doesn't define a no-argument constructor.

We reviewed overloaded, overridden, and hidden methods and showed how they differ, especially in terms of polymorphism. We also introduced the notion of hiding variables, although we strongly discourage this in practice as it often leads to confusing, difficult-to-maintain code.

We introduced abstract classes and interfaces and showed how you can use them to define a platform for other developers to interact with. By definition, an abstract type cannot be instantiated directly and requires a concrete subclass for the code to be used. Since default and static interface methods are new to Java 8, expect to see at least one question on them on the exam.

Finally, this chapter introduced the concept of polymorphism, central to the Java language, and showed how objects can be accessed in a variety of forms. Make sure you understand when casts are needed for accessing objects, and be able to spot the difference between compile-time and runtime cast problems.

## Exam Essentials

**Be able to write code that extends other classes.** A Java class that extends another class inherits all of its public and protected methods and variables. The first line of every constructor is a call to another constructor within the class using `this()` or a call to a constructor of the parent class using the `super()` call. If the parent class doesn't contain a no-argument constructor, an explicit call to the parent constructor must be provided. Parent methods and objects can be accessed explicitly using the `super` keyword. Finally, all classes in Java extend `java.lang.Object` either directly or from a superclass.

**Understand the rules for method overriding.** The Java compiler allows methods to be overridden in subclasses if certain rules are followed: a method must have the same signature, be at least as accessible as the parent method, must not declare any new or broader exceptions, and must use covariant return types.

**Understand the rules for hiding methods and variables.** When a static method is re-created in a subclass, it is referred to as method hiding. Likewise, variable hiding is when a variable name is reused in a subclass. In both situations, the original method or variable still exists and is used in methods that reference the object in the parent class. For method hiding, the use of `static` in the method declaration must be the same between the parent and child class. Finally, variable and method hiding should generally be avoided since it leads to confusing and difficult-to-follow code.

**Recognize the difference between method overriding and method overloading.** Both method overloading and overriding involve creating a new method with the same name as an existing method. When the method signature is the same, it is referred to as method overriding and must follow a specific set of override rules to compile. When the method signature is different, with the method taking different inputs, it is referred to as method overloading and none of the override rules are required.

**Be able to write code that creates and extends abstract classes.** In Java, classes and methods can be declared as abstract. Abstract classes cannot be instantiated and require a concrete subclass to be accessed. Abstract classes can include any number, including zero, of abstract and nonabstract methods. Abstract methods follow all the method override rules and may only be defined within abstract classes. The first concrete subclass of an abstract class must implement all the inherited methods. Abstract classes and methods may not be marked as `final` or `private`.

**Be able to write code that creates, extends, and implements interfaces.** Interfaces are similar to a specialized abstract class in which only abstract methods and constant static final variables are allowed. New to Java 8, an interface can also define default and static methods with method bodies. All members of an interface are assumed to be public. Methods are assumed to be abstract if not explicitly marked as default or static. An interface that extends another interface inherits all its abstract methods. An interface cannot extend a class, nor can a class extend an interface. Finally, classes may implement any number of interfaces.

**Be able to write code that uses default and static interface methods.** A default method allows a developer to add a new method to an interface used in existing implementations, without forcing other developers using the interface to recompile their code. A developer using the interface may override the default method or use the provided one. A static method in an interface follows the same rules for a static method in a class.

**Understand polymorphism.** An object in Java may take on a variety of forms, in part depending on the reference used to access the object. Methods that are overridden will be replaced everywhere they are used, whereas methods and variables that are hidden will only be replaced in the classes and subclasses that they are defined. It is common to rely on polymorphic parameters—the ability of methods to be automatically passed as a superclass or interface reference—when creating method definitions.

**Recognize valid reference casting.** An instance can be automatically cast to a superclass or interface reference without an explicit cast. Alternatively, an explicit cast is required if the reference is being narrowed to a subclass of the object. The Java compiler doesn't permit casting to unrelated types. You should be able to discern between compiler-time casting errors and those that will not occur until runtime and that throw a `CastClassException`.



# Review Questions

1. What modifiers are implicitly applied to all interface methods? (Choose all that apply)

- A. protected
- B. public
- C. static
- D. void
- E. abstract
- F. default

2. What is the output of the following code?

```
1: class Mammal {
2: public Mammal(int age) {
3: System.out.print("Mammal");
4: }
5: }
6: public class Platypus extends Mammal {
7: public Platypus() {
8: System.out.print("Platypus");
9: }
10: public static void main(String[] args) {
11: new Mammal(5);
12: }
13: }
```

- A. Platypus
- B. Mammal
- C. PlatypusMammal
- D. MammalPlatypus
- E. The code will not compile because of line 8.
- F. The code will not compile because of line 11.

3. Which of the following statements can be inserted in the blank line so that the code will compile successfully? (Choose all that apply)

```
public interface CanHop {}
public class Frog implements CanHop {
 public static void main(String[] args) {
 _____ frog = new TurtleFrog();
 }
}
```

```
public class BrazilianHornedFrog extends Frog {}
public class TurtleFrog extends Frog {}
```

- A. Frog
- B. TurtleFrog
- C. BrazilianHornedFrog
- D. CanHop
- E. Object
- F. Long

4. Which statement(s) are correct about the following code? (Choose all that apply)

```
public class Rodent {
 protected static Integer chew() throws Exception {
 System.out.println("Rodent is chewing");
 return 1;
 }
}

public class Beaver extends Rodent {
 public Number chew() throws RuntimeException {
 System.out.println("Beaver is chewing on wood");
 return 2;
 }
}
```

- A. It will compile without issue.
  - B. It fails to compile because the type of the exception the method throws is a subclass of the type of exception the parent method throws.
  - C. It fails to compile because the return types are not covariant.
  - D. It fails to compile because the method is protected in the parent class and public in the subclass.
  - E. It fails to compile because of a static modifier mismatch between the two methods.
5. Which of the following may only be hidden and not overridden? (Choose all that apply)
- A. private instance methods
  - B. protected instance methods
  - C. public instance methods
  - D. static methods
  - E. public variables
  - F. private variables

6. Choose the correct statement about the following code:

```
1: interface HasExoskeleton {
2: abstract int getNumberOfSections();
3: }
4: abstract class Insect implements HasExoskeleton {
5: abstract int getNumberOfLegs();
6: }
7: public class Beetle extends Insect {
8: int getNumberOfLegs() { return 6; }
9: }
```

- A. It compiles and runs without issue.
  - B. The code will not compile because of line 2.
  - C. The code will not compile because of line 4.
  - D. The code will not compile because of line 7.
  - E. It compiles but throws an exception at runtime.
7. Which of the following statements about polymorphism are true? (Choose all that apply)
- A. A reference to an object may be cast to a subclass of the object without an explicit cast.
  - B. If a method takes a superclass of three objects, then any of those classes may be passed as a parameter to the method.
  - C. A method that takes a parameter with type `java.lang.Object` will take any reference.
  - D. All cast exceptions can be detected at compile-time.
  - E. By defining a public instance method in the superclass, you guarantee that the specific method will be called in the parent class at runtime.

8. Choose the correct statement about the following code:

```
1: public interface Herbivore {
2: int amount = 10;
3: public static void eatGrass();
4: public int chew() {
5: return 13;
6: }
7: }
```

- A. It compiles and runs without issue.
- B. The code will not compile because of line 2.
- C. The code will not compile because of line 3.
- D. The code will not compile because of line 4.
- E. The code will not compile because of lines 2 and 3.
- F. The code will not compile because of lines 3 and 4.

9. Choose the correct statement about the following code:
- ```
1: public interface CanFly {  
2:     void fly();  
3: }  
4: interface HasWings {  
5:     public abstract Object getWindSpan();  
6: }  
7: abstract class Falcon implements CanFly, HasWings {  
8: }
```
- A. It compiles without issue.
 - B. The code will not compile because of line 2.
 - C. The code will not compile because of line 4.
 - D. The code will not compile because of line 5.
 - E. The code will not compile because of lines 2 and 5.
 - F. The code will not compile because the class `Falcon` doesn't implement the interface methods.
10. Which statements are true for both abstract classes and interfaces? (Choose all that apply)
- A. All methods within them are assumed to be abstract.
 - B. Both can contain `public static final` variables.
 - C. Both can be extended using the `extend` keyword.
 - D. Both can contain default methods.
 - E. Both can contain static methods.
 - F. Neither can be instantiated directly.
 - G. Both inherit `java.lang.Object`.
11. What modifiers are assumed for all interface variables? (Choose all that apply)
- A. `public`
 - B. `protected`
 - C. `private`
 - D. `static`
 - E. `final`
 - F. `abstract`
12. What is the output of the following code?
- ```
1: interface Nocturnal {
2: default boolean isBlind() { return true; }
3: }
4: public class Owl implements Nocturnal {
```

```
5: public boolean isBlind() { return false; }
6: public static void main(String[] args) {
7: Nocturnal nocturnal = (Nocturnal)new Owl();
8: System.out.println(nocturnal.isBlind());
9: }
10: }
```

- A. true
- B. false
- C. The code will not compile because of line 2.
- D. The code will not compile because of line 5.
- E. The code will not compile because of line 7.
- F. The code will not compile because of line 8.

**13.** What is the output of the following code?

```
1: class Arthropod
2: public void printName(double input) { System.out
 .print("Arthropod"); }
3: }
4: public class Spider extends Arthropod {
5: public void printName(int input) { System.out.print("Spider"); }
6: public static void main(String[] args) {
7: Spider spider = new Spider();
8: spider.printName(4);
9: spider.printName(9.0);
10: }
11: }
```

- A. SpiderArthropod
- B. ArthropodSpider
- C. SpiderSpider
- D. ArthropodArthropod
- E. The code will not compile because of line 5.
- F. The code will not compile because of line 9.

**14.** Which statements are true about the following code? (Choose all that apply)

```
1: interface HasVocalCords {
2: public abstract void makeSound();
3: }
4: public interface CanBark extends HasVocalCords {
5: public void bark();
6: }
```

- A. The CanBark interface doesn't compile.
  - B. A class that implements HasVocalCords must override the makeSound() method.
  - C. A class that implements CanBark inherits both the makeSound() and bark() methods.
  - D. A class that implements CanBark only inherits the bark() method.
  - E. An interface cannot extend another interface.
15. Which of the following is true about a concrete subclass? (Choose all that apply)
- A. A concrete subclass can be declared as abstract.
  - B. A concrete subclass must implement all inherited abstract methods.
  - C. A concrete subclass must implement all methods defined in an inherited interface.
  - D. A concrete subclass cannot be marked as final.
  - E. Abstract methods cannot be overridden by a concrete subclass.
16. What is the output of the following code?
- ```
1: abstract class Reptile {  
2:     public final void layEggs() { System.out.println("Reptile laying eggs");  
    }  
3:     public static void main(String[] args) {  
4:         Reptile reptile = new Lizard();  
5:         reptile.layEggs();  
6:     }  
7: }  
8: public class Lizard extends Reptile {  
9:     public void layEggs() { System.out.println("Lizard laying eggs"); }  
10: }
```
- A. Reptile laying eggs
 - B. Lizard laying eggs
 - C. The code will not compile because of line 4.
 - D. The code will not compile because of line 5.
 - E. The code will not compile because of line 9.
17. What is the output of the following code?
- ```
1: public abstract class Whale {
2: public abstract void dive() {};
3: public static void main(String[] args) {
4: Whale whale = new Orca();
5: whale.dive();
6: }
7: }
```

```
8: class Orca extends Whale {
9: public void dive(int depth) { System.out.println("Orca diving"); }
10: }
```

- A. Orca diving
- B. The code will not compile because of line 2.
- C. The code will not compile because of line 8.
- D. The code will not compile because of line 9.
- E. The output cannot be determined from the code provided.

18. What is the output of the following code? (Choose all that apply)

```
1: interface Aquatic {
2: public default int getNumberOfGills(int input) { return 2; }
3: }
4: public class ClownFish implements Aquatic {
5: public String getNumberOfGills() { return "4"; }
6: public String getNumberOfGills(int input) { return "6"; }
7: public static void main(String[] args) {
8: System.out.println(new ClownFish().getNumberOfGills(-1));
9: }
10: }
```

- A. 2
- B. 4
- C. 6
- D. The code will not compile because of line 5.
- E. The code will not compile because of line 6.
- F. The code will not compile because of line 8.

19. Which of the following statements can be inserted in the blank so that the code will compile successfully? (Choose all that apply)

```
public class Snake {}
public class Cobra extends Snake {}
public class GardenSnake {}
public class SnakeHandler {
 private Snake snake;
 public void setSnake(Snake snake) { this.snake = snake; }
 public static void main(String[] args) {
 new SnakeHandler().setSnake(_____);
 }
}
```

- A. `new Cobra()`
- B. `new GardenSnake()`
- C. `new Snake()`
- D. `new Object()`
- E. `new String("Snake")`
- F. `null`

20. What is the result of the following code?

```
1: public abstract class Bird {
2: private void fly() { System.out.println("Bird is flying"); }
3: public static void main(String[] args) {
4: Bird bird = new Pelican();
5: bird.fly();
6: }
7: }
8: class Pelican extends Bird {
9: protected void fly() { System.out.println("Pelican is flying"); }
10: }
```

- A. Bird is flying
- B. Pelican is flying
- C. The code will not compile because of line 4.
- D. The code will not compile because of line 5.
- E. The code will not compile because of line 9.



# Chapter 6

## Exceptions

---

### OCA EXAM OBJECTIVES COVERED IN THIS CHAPTER:

#### ✓ Handling Exceptions

- Differentiate among checked exceptions, unchecked exceptions and Errors
- Create a try-catch block and determine how exceptions alter normal program flow
- Describe the advantages of Exception handling
- Create and invoke a method that throws an exception
- Recognize common exception classes (such as `NullPointerException`, `ArithmeticException`, `ArrayIndexOutOfBoundsException`, `ClassCastException`)





Many things can go wrong in a program. Java uses exceptions to deal with some of these scenarios. The OCA exam covers only the basics of working with exceptions. The rest are on the OCP exam.

## Understanding Exceptions

A program can fail for just about any reason. Here are just a few possibilities:

- The code tries to connect to a website, but the Internet connection is down.
- You made a coding mistake and tried to access an invalid index in an array.
- One method calls another with a value that the method doesn't support.

As you can see, some of these are coding mistakes. Others are completely beyond your control. Your program can't help it if the Internet connection goes down. What it *can* do is deal with the situation.

First, we'll look at the role of exceptions. Then we'll cover the various types of exceptions, followed by an explanation of how to throw an exception in Java.

### The Role of Exceptions

An *exception* is Java's way of saying, "I give up. I don't know what to do right now. You deal with it." When you write a method, you can either deal with the exception or make it the calling code's problem.

As an example, think of Java as a child who visits the zoo. The *happy path* is when nothing goes wrong. The child continues to look at the animals until the program nicely ends. Nothing went wrong and there were no exceptions to deal with.

This child's younger sister doesn't experience the happy path. In all the excitement she trips and falls. Luckily, it isn't a bad fall. The little girl gets up and proceeds to look at more animals. She has handled the issue all by herself. Unfortunately, she falls again later in the day and starts crying. This time, she has declared she needs help by crying. The story ends well. Her daddy rubs her knee and gives her a hug. Then they go back to seeing more animals and enjoy the rest of the day.

These are the two approaches Java uses when dealing with exceptions. A method can handle the exception case itself or make it the caller's responsibility. You saw both in the trip to the zoo.

You saw an exception in Chapter 1, “Java Building Blocks,” with a very simple Zoo example. You wrote a class that printed out the name of the zoo:

```
1: public class Zoo {
2: public static void main(String[] args) {
3: System.out.println(args[0]);
4: System.out.println(args[1]);
5: } }
```

Then you tried to call it without enough arguments:

```
$ javac Zoo.java
$ java Zoo Zoo
```

On line 4, Java realized there’s only one element in the array and index 1 is not allowed. Java threw up its hands in defeat and threw an exception. It didn’t try to handle the exception. It just said, “I can’t deal with it” and the exception was displayed:

```
ZooException in thread "main"
java.lang.ArrayIndexOutOfBoundsException: 1
at mainmethod.Zoo.main(Zoo.java:7)
```

Exceptions can and do occur all the time, even in solid program code. In our example, toddlers falling is a fact of life. When you write more advanced programs, you’ll need to deal with failures in accessing files, networks, and outside services. On the OCA exam, exceptions deal largely with mistakes in programs. For example, a program might try to access an invalid position in an array. The key point to remember is that exceptions alter the program flow.



## Real World Scenario

### Return Codes vs. Exceptions

Exceptions are used when “something goes wrong.” However, the word “wrong” is subjective. The following code returns –1 instead of throwing an exception if no match is found:

```
public int indexOf(String[] names, String name) {
 for (int i = 0; i < names.length; i++) {
 if (names[i].equals(name)) { return i; }
 }
 return -1;
}
```

*continues*

*continued*


This approach is common when writing a method that does a search. For example, imagine being asked to find the name Joe in the array. It is perfectly reasonable that Joe might not appear in the array. When this happens, a special value is returned. An exception should be reserved for exceptional conditions like *names* being null.

In general, try to avoid return codes. Return codes are commonly used in searches, so programmers are expecting them. In other methods, you will take your callers by surprise by returning a special value. An exception forces the program to deal with them or end with the exception if left unhandled, whereas a return code could be accidentally ignored and cause problems later in the program. An exception is like shouting, “Deal with me!”

## Understanding Exception Types

As we’ve explained, an exception is an event that alters program flow. Java has a `Throwable` superclass for all objects that represent these events. Not all of them have the word exception in their classname, which can be confusing. Figure 6.1 shows the key subclasses of `Throwable`.

**FIGURE 6.1** Categories of exception



Error means something went so horribly wrong that your program should not attempt to recover from it. For example, the disk drive “disappeared.” These are abnormal conditions that you aren’t likely to encounter.

**A runtime exception is defined as the `RuntimeException` class and its subclasses. Runtime exceptions tend to be unexpected but not necessarily fatal.** For example, accessing an invalid array index is unexpected. Runtime exceptions are also known as *unchecked exceptions*.

### Runtime vs. at the Time the Program is Run

A runtime (unchecked) exception is a specific type of exception. All exceptions occur at the time that the program is run. (The alternative is compile time, which would be a compiler error.) People don't refer to them as run time exceptions because that would be too easy to confuse with runtime! **When you see runtime, it means unchecked.**

**A checked exception includes Exception and all subclasses that do not extend RuntimeException. Checked exceptions tend to be more anticipated**—for example, trying to read a file that doesn't exist.

Checked exceptions? What are we checking? Java has a rule called the *handle or declare rule*. For checked exceptions, Java requires the code to either handle them or declare them in the method signature.

For example, this method declares that it might throw an exception:

```
void fall() throws Exception {
 throw new Exception();
}
```

Notice that you're using two different keywords here. `throw` tells Java that you want to throw an `Exception`. `throws` simply declares that the method might throw an `Exception`. It also might not. You will see the `throws` keyword more later in the chapter.

Because checked exceptions tend to be anticipated, Java enforces that the programmer do something to show the exception was thought about. Maybe it was handled in the method. Or maybe the method declares that it can't handle the exception and someone else should.

An example of a runtime exception is a `NullPointerException`, which happens when you try to call a member on a null reference. This can occur in any method. If you had to declare runtime exceptions everywhere, every single method would have that clutter!

### Checked vs. Unchecked (Runtime) Exceptions

In the past, developers used checked exceptions more often than they do now. According to Oracle, they are intended for issues a programmer "might reasonably be expected to recover from." Then developers started writing code where a chain of methods kept declaring the same exception and nobody actually handled it. Some libraries started using runtime exceptions for issues a programmer might reasonably be expected to recover from. Many programmers can hold a debate with you on which approach is better. For the OCA exam, you need to know the rules for how checked versus unchecked exceptions function. You don't have to decide philosophically whether an exception should be checked or unchecked.

## Throwing an Exception

Any Java code can throw an exception; this includes code you write. For the OCP exam, you'll learn how to create your own exception classes. The OCA exam is limited to exceptions that someone else has created. Most likely, they will be exceptions that are provided with Java. You might encounter an exception that was made up for the exam. This is fine. The question will make it obvious that these are exceptions by having the classname end with exception. For example, "MyMadeUpException" is clearly an exception.

On the exam, you will see two types of code that result in an exception. The first is code that's wrong. For example:

```
String[] animals = new String[0];
System.out.println(animals[0]);
```

This code throws an `ArrayIndexOutOfBoundsException`. That means questions about exceptions can be hidden in questions that appear to be about something else.



On the OCA exam, the vast majority of questions have a choice about not compiling and about throwing an exception. Pay special attention to code that calls a method on a null or that references an invalid array or `ArrayList` index. If you spot this, you know the correct answer is that the code throws an exception.

The second way for code to result in an exception is to explicitly request Java to throw one. Java lets you write statements like these:

```
throw new Exception();
throw new Exception("Ow! I fell.");
throw new RuntimeException();
throw new RuntimeException("Ow! I fell.");
```

The `throw` keyword tells Java you want some other part of the code to deal with the exception. This is the same as the young girl crying for her daddy. Someone else needs to figure out what to do about the exception.

When creating an exception, you can usually pass a `String` parameter with a message or you can pass no parameters and use the defaults. We say *usually* because this is a convention. Someone could create an exception class that does not have a constructor that takes a message. The first two examples create a new object of type `Exception` and throw it. The last two show that the code looks the same regardless of which type of exception you throw.

These rules are very important. Be sure to closely study everything in Table 6.1.


**TABLE 6.1** Types of exceptions

| Type              | How to recognize                                                                     | Okay for program to catch? | Is program required to handle or declare? |
|-------------------|--------------------------------------------------------------------------------------|----------------------------|-------------------------------------------|
| Runtime exception | Subclass of <code>RuntimeException</code>                                            | Yes                        | No                                        |
| Checked exception | Subclass of <code>Exception</code> but not subclass of <code>RuntimeException</code> | Yes                        | Yes                                       |
| Error             | Subclass of <code>Error</code>                                                       | No                         | No                                        |

## Using a *try* Statement

Now that you know what exceptions are, let's explore how to handle them. Java uses a *try statement* to separate the logic that might throw an exception from the logic to handle that exception. Figure 6.2 shows the syntax of a try statement.

**FIGURE 6.2** The syntax of a *try* statement



The code in the `try` block is run normally. If any of the statements throw an exception that can be caught by the exception type listed in the `catch` block, the `try` block stops running and execution goes to the `catch` statement. If none of the statements in the `try` block throw an exception that can be caught, the *catch clause* is not run.

You probably noticed the words “block” and “clause” used interchangeably. The exam does this as well, so we are getting you used to it. Both are correct. “Block” is correct because there are braces present. “Clause” is correct because they are part of a `try` statement.

There aren’t a ton of syntax rules here. The curly braces are required for the `try` and `catch` blocks.

In our example, the little girl gets up by herself the first time she falls. Here’s what this looks like:

```
3: void explore() {
4: try {
5: fall();
6: System.out.println("never get here");
7: } catch (RuntimeException e) {
8: getUp();
9: }
10: seeAnimals();
11: }
12: void fall() { throw new RuntimeException(); }
```

First, line 5 calls the `fall()` method. Line 12 throws an exception. This means Java jumps straight to the `catch` block, skipping line 6. The girl gets up on line 8. Now the `try` statement is over and execution proceeds normally with line 10.

Now let’s look at some invalid `try` statements that the exam might try to trick you with. Do you see what’s wrong with this one?

```
try // DOES NOT COMPILE
 fall();
catch (Exception e)
 System.out.println("get up");
```

The problem is that the braces are missing. It needs to look like this:

```
try {
 fall();
} catch (Exception e) {
 System.out.println("get up");
}
```



*try* statements are like methods in that the curly braces are required even if there is only one statement inside the code blocks. *if* statements and loops are special in this respect as they allow you to omit the curly braces.

What about this one?

```
try { // DOES NOT COMPILE
 fall();
}
```


This code doesn't compile because the *try* block doesn't have anything after it. Remember, the point of a *try* statement is for something to happen if an exception is thrown. Without another clause, the *try* statement is lonely.

Now that you know the basics, let's start adding more features to exceptions. The following sections show you how to add a *finally* clause to a *try* statement and catch different types of exceptions and describe what happens if an exception is thrown in *catch* or *finally*.

## Adding a *finally* Block

The *try* statement also lets you run code at the end with a *finally clause* regardless of whether an exception is thrown. Figure 6.3 shows the syntax of a *try* statement with this extra functionality.

**FIGURE 6.3** The syntax of a *try* statement with *finally*



There are two paths through code with both a *catch* and a *finally*. If an exception is thrown, the *finally* block is run after the *catch* block. If no exception is thrown, the *finally* block is run after the *try* block completes.

Let's go back to our young girl example, this time with finally:

```
12: void explore() {
13: try {
14: seeAnimals();
15: fall();
16: } catch (Exception e) {
17: getHugFromDaddy();
18: } finally {
19: seeMoreAnimals();
20: }
21: goHome();
22: }
```

The girl falls on line 15. If she gets up by herself, the code goes on to the finally block and runs line 19. Then the try statement is over and the code proceeds on line 21. If the girl doesn't get up by herself, she throws an exception. The catch block runs and she gets a hug on line 17. Then the try statement is over and the code proceeds on line 21. Either way, the ending is the same. The finally block is executed and the try statement ends.



On the OCA exam, a try statement must have catch and/or finally. Having both is fine. Having neither is a problem. On the OCP exam, you'll learn about a special syntax for a try statement called try-with-resources that allows neither a catch nor a finally block. On the OCA exam, you get to assume a try statement is just a regular try statement and not a try-with-resources statement.

The exam will try to trick you with missing clauses or clauses in the wrong order. Do you see why the following do or do not compile?

```
25: try { // DOES NOT COMPILE
26: fall();
27: } finally {
28: System.out.println("all better");
29: } catch (Exception e) {
30: System.out.println("get up");
31: }
32:
33: try { // DOES NOT COMPILE
34: fall();
35: }
36:
37: try {
```

```
38: fall();
39: } finally {
40: System.out.println("all better");
41: }
```

The first example (lines 25–31) does not compile because the `catch` and `finally` blocks are in the wrong order. The second example (lines 33–35) does not compile because there must be a `catch` or `finally` block. The third example (lines 37–41) is just fine. `catch` is not required if `finally` is present.

One problem with `finally` is that any realistic uses for it are out of the scope of the OCA exam. `finally` is typically used to close resources such as files or databases—both of which are topics on the OCP exam. This means most of the examples you encounter on the OCA exam with `finally` are going to look contrived. For example, you'll get asked questions such as what this code outputs:

```
String s = "";
try {
 s += "t";
} catch(Exception e) {
 s += "c";
} finally {
 s += "f";
}
s += "a";
System.out.print(s);
```

The answer is `tfa`. The `try` block is executed. Since no exception is **thrown**, Java goes straight to the `finally` block. Then the code after the `try` statement is run. We know; this is a silly example. Expect to see examples like this on the OCA exam.

### ***System.exit***

There is one exception to “the `finally` block always runs after the `catch` block” rule: Java defines a method that you call as `System.exit(0)`; . The integer parameter is the error code that gets returned. `System.exit` tells Java, “Stop. End the program right now. Do not pass go. Do not collect \$200.” When `System.exit` is called in the `try` or `catch` block, `finally` does not run.

## **Catching Various Types of Exceptions**

So far, you have been catching only one type of exception. Now let's see what happens when different types of exceptions can be thrown from the same method.

Creating your own exceptions is not on the OCA exam, but it is on the OCP exam. However, the OCA exam can define basic exceptions to show you the hierarchy. You only need to do two things with this information. First, you must be able to recognize if the exception is a checked or an unchecked exception. Second, you need to determine if any of the exceptions are subclasses of the others.

```
class AnimalsOutForAWalk extends RuntimeException { }
class ExhibitClosed extends RuntimeException { }
class ExhibitClosedForLunch extends ExhibitClosed { }
```

In this example, there are three custom exceptions. All are unchecked exceptions because they directly or indirectly extend `RuntimeException`. Now we catch both types of exceptions and handle them by printing out the appropriate message:

```
public void visitPorcupine() {
 try {
 seeAnimal();
 } catch (AnimalsOutForAWalk e) { // first catch block
 System.out.print("try back later");
 } catch (ExhibitClosed e) { // second catch block
 System.out.print("not today");
 }
}
```

There are three possibilities for when this code is run. If `seeAnimal()` doesn't throw an exception, nothing is printed out. If the animal is out for a walk, only the first catch block runs. If the exhibit is closed, only the second catch block runs.

A rule exists for the order of the catch blocks. Java looks at them in the order they appear. If it is impossible for one of the catch blocks to be executed, a compiler error about unreachable code occurs. This happens when a superclass is caught before a subclass. Remember, we warned you to pay attention to any subclass exceptions.

In the porcupine example, the order of the catch blocks could be reversed because the exceptions don't inherit from each other. And yes, we have seen a porcupine be taken for a walk on a leash.

The following example shows exception types that do inherit from each other:

```
public void visitMonkeys() {
 try {
 seeAnimal();
 } catch (ExhibitClosedForLunch e) { // subclass exception
 System.out.print("try back later");
 } catch (ExhibitClosed e) { // superclass exception
 System.out.print("not today");
 }
}
```

If the more specific `ExhibitClosedForLunch` exception is thrown, the first catch block runs. If not, Java checks if the superclass `ExhibitClosed` exception is thrown and catches it. This time, the order of the catch blocks does matter. The reverse does not work.

```
public void visitMonkeys() {
 try {
 seeAnimal();
 } catch (ExhibitClosed e) {
 System.out.print("not today");
 } catch (ExhibitClosedForLunch e) { // DOES NOT COMPILE
 System.out.print("try back later");
 }
}
```

This time, if the more specific `ExhibitClosedForLunch` exception is thrown, the catch block for `ExhibitClosed` runs—which means there is no way for the second catch block to ever run. Java correctly tells us there is an unreachable catch block.

Let's try this one more time. Do you see why this code doesn't compile?

```
public void visitSnakes() {
 try {
 seeAnimal();
 } catch (RuntimeException e) {
 System.out.print("runtime exception");
 } catch (ExhibitClosed e) { // DOES NOT COMPILE
 System.out.print("not today");
 } catch (Exception e) {
 System.out.print("exception");
 }
}
```

It's the same problem. `ExhibitClosed` is a `RuntimeException`. If it is thrown, the first catch block takes care of it, making sure there no way to get to the second catch block.

To review catching multiple exceptions, **remember that at most one catch block will run and it will be the first catch block that can handle it.**

## Throwing a Second Exception

So far, we've limited ourselves to one try statement in each example. **However, a catch or finally block can have any valid Java code in it—including another try statement.**



Even though the topic of reading files is on the OCP exam, the OCA exam may ask you about exception handling with those classes. This is actually a gift. When you see such a question, you know the problem has to be about basic Java syntax or exception handling!

The following code tries to read a file:

```
16: public static void main(String[] args) {
17: FileReader reader = null;
18: try {
19: reader = read();
20: } catch (IOException e) {
21: try {
22: if (reader != null) reader.close();
23: } catch (IOException inner) {
24: }
25: }
26: }
27: private static FileReader read() throws IOException {
28: // CODE GOES HERE
29: }
```

The easiest case is if line 28 doesn't throw an exception. Then the entire catch block on lines 20–25 is skipped. Next, consider if line 28 throws a `NullPointerException`. That isn't an `IOException`, so the catch block on lines 20–25 will still be skipped.

If line 28 does throw an `IOException`, the catch block on lines 20–25 does get run. Line 22 tries to close the reader. If that goes well, the code completes and the `main()` method ends normally. If the `close()` method does throw an exception, Java looks for more catch blocks. There aren't any, so the `main` method throws that new exception. Regardless, the exception on line 28 is handled. A different exception might be thrown, but the one from line 28 is done.

Most of the examples you see with exception handling on the exam are abstract. They use letters or numbers to make sure you understand the flow. This one shows that only the last exception to be thrown matters. (This is true for the OCA exam. It will change a bit on the OCP exam.)

```
26: try {
27: throw new RuntimeException();
28: } catch (RuntimeException e) {
29: throw new RuntimeException();
30: } finally {
31: throw new Exception();
32: }
```

Line 27 throws an exception, which is caught on line 28. The catch block then throws an exception on line 29. If there were no `finally` block, the exception from line 29 would

be thrown. However, the `finally` block runs after the `try` block. Since the `finally` block throws an exception of its own on line 31, this one gets thrown. The exception from the `catch` block gets forgotten about. This is why you often see another `try/catch` inside a `finally` block—to make sure it doesn't mask the exception from the `catch` block.

Next we are going to show you the hardest example you can be asked related to exceptions. What do you think this method returns? Go slowly. It's tricky.

```
30: public String exceptions() {
31: String result = "";
32: String v = null;
33: try {
34: try {
35: result += "before";
36: v.length();
37: result += "after";
38: } catch (NullPointerException e) {
39: result += "catch";
40: throw new RuntimeException();
41: } finally {
42: result += "finally";
43: throw new Exception();
44: }
45: } catch (Exception e) {
46: result += "done";
47: }
48: return result;
49: }
```

The correct answer is `before catch finally done`. Everything is normal up until line 35, when "before" is added. Line 36 throws a `NullPointerException`. Line 37 is skipped as Java goes straight to the `catch` block. Line 38 does catch the exception, and "catch" is added on line 39. Then line 40 throws a `RuntimeException`. The `finally` block runs after the `catch` regardless of whether an exception is thrown; it adds "finally" to result. At this point, we have completed the inner `try` statement that ran on lines 34–44. The outer `catch` block then sees an exception was thrown and catches it on line 45; it adds "done" to result.

## Recognizing Common Exception Types

You need to recognize three types of exceptions for the OCA exam: runtime exceptions, checked exceptions, and errors. We'll look at common examples of each type. For the exam, you'll need to recognize which type of an exception it is and whether it's thrown by the JVM or a programmer. So you can recognize them, we'll show you some code examples for those exceptions.

## Runtime Exceptions

Runtime exceptions extend `RuntimeException`. They don't have to be handled or declared. They can be thrown by the programmer or by the JVM. Common runtime exceptions include the following:

**ArithmeticException** Thrown by the JVM when code attempts to divide by zero

**ArrayIndexOutOfBoundsException** Thrown by the JVM when code uses an illegal index to access an array

**ClassCastException** Thrown by the JVM when an attempt is made to cast an exception to a subclass of which it is not an instance

**IllegalArgumentException** Thrown by the programmer to indicate that a method has been passed an illegal or inappropriate argument

**NullPointerException** Thrown by the JVM when there is a null reference where an object is required

**NumberFormatException** Thrown by the programmer when an attempt is made to convert a string to a numeric type but the string doesn't have an appropriate format

### *ArithmeticException*

Trying to divide an `int` by zero gives an undefined result. When this occurs, the JVM will throw an `ArithmeticException`:

```
int answer = 11 / 0;
```

Running this code results in the following output:

```
Exception in thread "main" java.lang.ArithmeticException: / by zero
```

Java doesn't spell out the word "divide." That's okay, though, because we know that `/` is the division operator and that Java is trying to tell us division by zero occurred.

The thread "main" is telling us the code was called directly or indirectly from a program with a `main` method. On the OCA exam, this is all the output we will see. Next comes the name of the exception, followed by extra information (if any) that goes with the exception.

### *ArrayIndexOutOfBoundsException*

You know by now that array indexes start with 0 and go up to 1 less than the length of the array—which means this code will throw an `ArrayIndexOutOfBoundsException`:

```
int[] countsOfMoose = new int[3];
System.out.println(countsOfMoose[-1]);
```

This is a problem because there's no such thing as a negative array index. Running this code yields the following output:

```
Exception in thread "main" java.lang.ArrayIndexOutOfBoundsException: -1
```

At least Java tells us what index was invalid. Can you see what's wrong with this one?



```
int total = 0;
int[] countsOfMoose = new int[3];
for (int i = 0; i <= countsOfMoose.length; i++)
 total += countsOfMoose[i];
```

The problem is that the for loop should have < instead of <=. On the final iteration of the loop, Java tries to call `countsOfMoose[3]`, which is invalid. The array includes only three elements, making 2 the largest possible index. The output looks like this:

```
Exception in thread "main" java.lang.ArrayIndexOutOfBoundsException: 3
```

## ***ClassCastException***

Java tries to protect you from impossible casts. This code doesn't compile because `Integer` is not a subclass of `String`:

```
String type = "moose";
Integer number = (Integer) type; // DOES NOT COMPILE
```

More complicated code thwarts Java's attempts to protect you. When the cast fails at runtime, Java will throw a `ClassCastException`:

```
String type = "moose";
Object obj = type;
Integer number = (Integer) obj;
```

The compiler sees a cast from `Object` to `Integer`. This could be okay. The compiler doesn't realize there's a `String` in that `Object`. When the code runs, it yields the following output:

```
Exception in thread "main" java.lang.ClassCastException: java.lang.String
cannot be cast to java.lang.Integer
```

Java tells us both types that were involved in the problem, making it apparent what's wrong.

## ***IllegalArgumentException***

**`IllegalArgumentException` is a way for your program to protect itself.** We first saw the following setter method in the `Swan` class in Chapter 4, "Methods and Encapsulation."

```
6: public void setNumberEggs(int numberEggs) { // setter
7: if (numberEggs >= 0) // guard condition
8: this.numberEggs = numberEggs;
9: }
```

This code works, but we don't really want to ignore the caller's request when they tell us a `Swan` has -2 eggs. We want to tell the caller that something is wrong—preferably in a very obvious way that the caller can't ignore so that the programmer will fix the problem. Exceptions are an efficient way to do this. Seeing the code end with an exception is a great reminder that something is wrong:

```
public static void setNumberEggs(int numberEggs) {
 if (numberEggs < 0)
 throw new IllegalArgumentException(
 "# eggs must not be negative");
 this.numberEggs = numberEggs;
}
```

The program throws an exception when it's not happy with the parameter values. The output looks like this:

```
Exception in thread "main" java.lang.IllegalArgumentException: # eggs must not
be negative
```

Clearly this is a problem that must be fixed if the programmer wants the program to do anything useful.

### ***NullPointerException***

Instance variables and methods must be called on a non-null reference. If the reference is null, the JVM will throw a `NullPointerException`. It's usually subtle, such as this example, which checks whether you remember instance variable references default to null.

```
String name;

public void printLength() throws NullPointerException {
 System.out.println(name.length());
}
```

Running this code results in this output:

```
Exception in thread "main" java.lang.NullPointerException
```

### ***NumberFormatException***

Java provides methods to convert strings to numbers. When these are passed an invalid value, they throw a `NumberFormatException`. The idea is similar to `IllegalArgumentException`. Since this is a common problem, Java gives it a separate class.

In fact, `NumberFormatException` is a subclass of `IllegalArgumentException`. Here's an example of trying to convert something non-numeric into an `int`:

```
Integer.parseInt("abc");
```

The output looks like this:

```
Exception in thread "main" java.lang.NumberFormatException: For input string:
"abc"
```

## Checked Exceptions

Checked exceptions have `Exception` in their hierarchy but not `RuntimeException`. They must be handled or declared. They can be thrown by the programmer or by the JVM.

Common runtime exceptions include the following:

**`FileNotFoundException`** Thrown programmatically when code tries to reference a file that does not exist

**`IOException`** Thrown programmatically when there's a problem reading or writing a file

For the OCA exam, you only need to know that these are checked exceptions. Also keep in mind that `FileNotFoundException` is a subclass of `IOException`, although the exam will remind you of that fact if it comes up. You'll see these two exceptions in more detail on the OCP exam.

## Errors

Errors extend the `Error` class. They are thrown by the JVM and should not be handled or declared. Errors are rare, but you might see these:

**`ExceptionInInitializerError`** Thrown by the JVM when a static initializer throws an exception and doesn't handle it

**`StackOverflowError`** Thrown by the JVM when a method calls itself too many times (this is called *infinite recursion* because the method typically calls itself without end)

**`NoClassDefFoundError`** Thrown by the JVM when a class that the code uses is available at compile time but not runtime

### ***`ExceptionInInitializerError`***

Java runs static initializers the first time a class is used. If one of the static initializers throws an exception, Java can't start using the class. It declares defeat by throwing an `ExceptionInInitializerError`. This code shows an `ArrayIndexOutOfBoundsException` in a static initializer:

```
static {
 int[] countsOfMoose = new int[3];
 int num = countsOfMoose[-1];
}
public static void main(String[] args) { }
```

This code yields information about two exceptions:

```
Exception in thread "main" java.lang.ExceptionInInitializerError
Caused by: java.lang.ArrayIndexOutOfBoundsException: -1
```

We get the `ExceptionInInitializerError` because the error happened in a static initializer. That information alone wouldn't be particularly useful in fixing the problem. Therefore, Java also tells us the original cause of the problem: the `ArrayIndexOutOfBoundsException` that we need to fix.

The `ExceptionInInitializerError` is an error because Java failed to load the whole class. This failure prevents Java from continuing.

## ***StackOverflowError***

When Java calls methods, it puts parameters and local variables on the stack. After doing this a very large number of times, the stack runs out of room and overflows. This is called a `StackOverflowError`. Most of the time, this error occurs when a method calls itself.

```
public static void doNotCodeThis(int num) {
 doNotCodeThis(1);
}
```

The output contains this line:

```
Exception in thread "main" java.lang.StackOverflowError
```

Since the method calls itself, it will never end. Eventually, Java runs out of room on the stack and throws the error. This is called infinite recursion. It is better than an infinite loop because at least Java will catch it and throw the error. With an infinite loop, Java just uses all your CPU until you can kill it.

## ***NoClassDefFoundError***

This error won't show up in code on the exam—you just need to know that it is an error. `NoClassDefFoundError` occurs when Java can't find the class at runtime.

# Calling Methods That Throw Exceptions

When you're calling a method that throws an exception, the rules are the same as within a method. Do you see why the following doesn't compile?

```
class NoMoreCarrotsException extends Exception {}
public class Bunny {
 public static void main(String[] args) {
 eatCarrot(); // DOES NOT COMPILE
 }
 private static void eatCarrot() throws NoMoreCarrotsException {
 }
}
```

The problem is that `NoMoreCarrotsException` is a checked exception. Checked exceptions must be handled or declared. The code would compile if we changed the `main()` method to either of these:

```
public static void main(String[] args)
 throws NoMoreCarrotsException { // declare exception
 eatCarrot();
}
public static void main(String[] args) {
 try {
 eatCarrot();
 } catch (NoMoreCarrotsException e) { // handle exception
 System.out.print("sad rabbit");
 }
}
```

You might have noticed that `eatCarrot()` didn't actually throw an exception; it just declared that it could. This is enough for the compiler to require the caller to handle or declare the exception.

The compiler is still on the lookout for unreachable code. Declaring an unused exception isn't considered unreachable code. It gives the method the option to change the implementation to throw that exception in the future. Do you see the issue here?

```
public void bad() {
 try {
 eatCarrot();
 } catch (NoMoreCarrotsException e) { // DOES NOT COMPILE
 System.out.print("sad rabbit");
 }
}
public void good() throws NoMoreCarrotsException {
 eatCarrot();
}
private static void eatCarrot() { }
```

**Java knows that `eatCarrot()` can't throw a checked exception—which means there's no way for the catch block in `bad()` to be reached. In comparison, `good()` is free to declare other exceptions.**

## Subclasses

Now that you have a deeper understanding of exceptions, let's look at overriding methods with exceptions in the method declaration. **When a class overrides a method from a**

superclass or implements a method from an interface, it's not allowed to add new checked exceptions to the method signature. For example, this code isn't allowed:

```
class CanNotHopException extends Exception { }
class Hopper {
 public void hop() { }
}
class Bunny extends Hopper {
 public void hop() throws CanNotHopException { } // DOES NOT COMPILE
}
```

Java knows `hop()` isn't allowed to throw any checked exceptions because the superclass `Hopper` doesn't declare any. Imagine what would happen if subclasses could add checked exceptions—you could write code that calls `Hopper`'s `hop()` method and not handle any exceptions. Then if `Bunny` was used in its place, the code wouldn't know to handle or declare `CanNotHopException`.

A subclass is allowed to declare fewer exceptions than the superclass or interface. This is legal because callers are already handling them.

```
class Hopper {
 public void hop() throws CanNotHopException { }
}
class Bunny extends Hopper {
 public void hop() { }
}
```

A subclass not declaring an exception is similar to a method declaring it throws an exception that it never actually throws. This is perfectly legal.

Similarly, a class is allowed to declare a subclass of an exception type. The idea is the same. The superclass or interface has already taken care of a broader type. Here's an example:

```
class Hopper {
 public void hop() throws Exception { }
}
class Bunny extends Hopper {
 public void hop() throws CanNotHopException { }
}
```

`Bunny` could declare that it throws `Exception` directly, or it could declare that it throws a more specific type of `Exception`. It could even declare that it throws nothing at all.

This rule applies only to checked exceptions. The following code is legal because it has a runtime exception in the subclass's version:

```
class Hopper {
 public void hop() { }
}
```

```
class Bunny extends Hopper {
 public void hop() throws IllegalStateException { }
}
```

The reason that it's okay to declare new runtime exceptions in a subclass method is that the declaration is redundant. Methods are free to throw any runtime exceptions they want without mentioning them in the method declaration.

## Printing an Exception

There are three ways to print an exception. You can let Java print it out, print just the message, or print where the stack trace comes from. This example shows all three approaches:

```
5: public static void main(String[] args) {
6: try {
7: hop();
8: } catch (Exception e) {
9: System.out.println(e);
10: System.out.println(e.getMessage());
11: e.printStackTrace();
12: }
13: }
14: private static void hop() {
15: throw new RuntimeException("cannot hop");
16: }
```


This code results in the following output:

```
java.lang.RuntimeException: cannot hop
cannot hop
java.lang.RuntimeException: cannot hop
at trycatch.Handling.hop(Handling.java:15)
at trycatch.Handling.main(Handling.java:7)
```

The first line shows what Java prints out by default: the exception type and message. The second line shows just the message. The rest shows a stack trace.

The stack trace is usually the most helpful one because it shows where the exception occurred in each method that it passed through. On the OCA exam, you will mostly see the first approach. This is because the exam often shows code snippets.

The stack trace shows all the methods on the stack. Figure 6.4 shows what the stack looks like for this code. Every time you call a method, Java adds it to the stack until it completes. When an exception is thrown, it goes through the stack until it finds a method that can handle it or it runs out of stack.

**FIGURE 6.4** A method stack

## Real World Scenario

### Why Swallowing Exception Is Bad

Because checked exceptions require you to handle or declare them, there is a temptation to catch them so they “go away.” But doing so can cause problems. In the following code, there’s a problem reading in the file:

```
public static void main(String[] args) {
 String textInFile = null;
 try {
 readInFile();
 } catch (IOException e) {
 // ignore exception
 }
 // imagine many lines of code here
 System.out.println(textInFile.replace(" ", ""));
}

private static void readInFile() throws IOException {
 throw new IOException();
}
```

The code results in a `NullPointerException`. Java doesn’t tell you anything about the original `IOException` because it was handled. Granted, it was handled poorly, but it *was* handled.

When writing your own code, print out a stack trace or at least a message when catching an exception. Also, consider whether continuing is the best course of action. In our example, the program can’t do anything after it fails to read in the file. It might as well have just thrown the `IOException`.



# Summary

An exception indicates something unexpected happened. A method can handle an exception by catching it or declaring it for the caller to deal with. Many exceptions are thrown by Java libraries. You can throw your own exception with code such as `throw new Exception()`.

Subclasses of `java.lang.Error` are exceptions that a programmer should not attempt to handle. Subclasses of `java.lang.RuntimeException` are runtime (unchecked) exceptions. Subclasses of `java.lang.Exception`, but not `java.lang.RuntimeException` are checked exceptions. Java requires checked exceptions to be handled or declared.

If a `try` statement has multiple catch blocks, at most one catch block can run. Java looks for an exception that can be caught by each catch block in the order they appear, and the first match is run. Then execution continues after the `try` statement. If both `catch` and `finally` throw an exception, the one from `finally` gets thrown.

Common runtime exceptions include:

- `ArithmeticException`
- `ArrayIndexOutOfBoundsException`
- `ClassCastException`
- `IllegalArgumentException`
- `NullPointerException`
- `NumberFormatException`

`IllegalArgumentException` and `NumberFormatException` are typically thrown by the programmer, whereas the others are typically thrown by the JVM.

Common checked exceptions include:

- `IOException`
- `FileNotFoundException`

Common errors include:

- `ExceptionInInitializerError`
- `StackOverflowError`
- `NoClassDefFoundError`

When a method overrides a method in a superclass or interface, it is not allowed to add checked exceptions. It is allowed to declare fewer exceptions or declare a subclass of a declared exception. Methods declare exceptions with the keyword `throws`.

## Exam Essentials

**Differentiate between checked and unchecked exceptions.** Unchecked exceptions are also known as runtime exceptions and are subclasses of `java.lang.RuntimeException`. All other subclasses of `java.lang.Exception` are checked exceptions.

**Understand the flow of a try statement.** A try statement must have a catch or a finally block. Multiple catch blocks are also allowed, provided no superclass exception type appears in an earlier catch block than its subclass. The finally block runs last regardless of whether an exception is thrown.

**Identify whether an exception is thrown by the programmer or the JVM.** `IllegalArgumentException` and `NumberFormatException` are commonly thrown by the programmer. Most of the other runtime exceptions are typically thrown by the JVM.

**Declare methods that declare exceptions.** The `throws` keyword is used in a method declaration to indicate an exception might be thrown. When overriding a method, the method is allowed to throw fewer exceptions than the original version.

**Recognize when to use `throw` versus `throws`.** The `throw` keyword is used when you actually want to throw an exception—for example, `throw new RuntimeException()`. The `throws` keyword is used in a method declaration.

# Review Questions

1. Which of the following statements are true? (Choose all that apply)
  - A. Runtime exceptions are the same thing as checked exceptions.
  - B. Runtime exceptions are the same thing as unchecked exceptions.
  - C. You can declare only checked exceptions.
  - D. You can declare only unchecked exceptions.
  - E. You can handle only Exception subclasses.
2. Which of the following pairs fill in the blanks to make this code compile? (Choose all that apply)

```
7: public void ohNo() _____ Exception {
8: _____ Exception();
9: }
```

  - A. On line 7, fill in throw
  - B. On line 7, fill in throws
  - C. On line 8, fill in throw
  - D. On line 8, fill in throw new
  - E. On line 8, fill in throws
  - F. On line 8, fill in throws new
3. When are you required to use a finally block in a regular try statement (not a try-with-resources)?
  - A. Never.
  - B. When the program code doesn't terminate on its own.
  - C. When there are no catch blocks in a try statement.
  - D. When there is exactly one catch block in a try statement.
  - E. When there are two or more catch blocks in a try statement.
4. Which exception will the following throw?

```
Object obj = new Integer(3);
String str = (String) obj;
System.out.println(str);
```

  - A. ArrayIndexOutOfBoundsException
  - B. ClassCastException
  - C. IllegalArgumentException
  - D. NumberFormatException
  - E. None of the above.

5. Which of the following exceptions are thrown by the JVM? (Choose all that apply)
- A. `ArrayIndexOutOfBoundsException`
  - B. `ExceptionInInitializerError`
  - C. `java.io.IOException`
  - D. `NullPointerException`
  - E. `NumberFormatException`
6. What will happen if you add the statement `System.out.println(5 / 0);` to a working `main()` method?
- A. It will not compile.
  - B. It will not run.
  - C. It will run and throw an `ArithmeticException`.
  - D. It will run and throw an `IllegalArgumentException`.
  - E. None of the above.
7. What is printed besides the stack trace caused by the `NullPointerException` from line 16?

```
1: public class DoSomething {
2: public void go() {
3: System.out.print("A");
4: try {
5: stop();
6: } catch (ArithmeticException e) {
7: System.out.print("B");
8: } finally {
9: System.out.print("C");
10: }
11: System.out.print("D");
12: }
13: public void stop() {
14: System.out.print("E");
15: Object x = null;
16: x.toString();
17: System.out.print("F");
18: }
19: public static void main(String[] args) {
20: new DoSomething().go();
21: }
22: }
```

- A. AE
- B. AEBCD

- C. AEC
  - D. AECD
  - E. No output appears other than the stack trace.
8. What is the output of the following snippet, assuming a and b are both 0?
- ```
3:     try {
4:         return a / b;
5:     } catch (RuntimeException e) {
6:         return -1;
7:     } catch (ArithmeticException e) {
8:         return 0;
9:     } finally {
10:        System.out.print("done");
11:    }
```
- A. -1
 - B. 0
 - C. done-1
 - D. done0
 - E. The code does not compile.
 - F. An uncaught exception is thrown.
9. What is the output of the following program?
- ```
1: public class Laptop {
2: public void start() {
3: try {
4: System.out.print("Starting up ");
5: throw new Exception();
6: } catch (Exception e) {
7: System.out.print("Problem ");
8: System.exit(0);
9: } finally {
10: System.out.print("Shutting down ");
11: }
12: }
13: public static void main(String[] args) {
14: new Laptop().start();
15: } }
```
- A. Starting up
  - B. Starting up Problem
  - C. Starting up Problem Shutting down

- D. Starting up Shutting down
- E. The code does not compile.
- F. An uncaught exception is thrown.

10. What is the output of the following program?

```
1: public class Dog {
2: public String name;
3: public void parseName() {
4: System.out.print("1");
5: try {
6: System.out.print("2");
7: int x = Integer.parseInt(name);
8: System.out.print("3");
9: } catch (NumberFormatException e) {
10: System.out.print("4");
11: }
12: }
13: public static void main(String[] args) {
14: Dog leroy = new Dog();
15: leroy.name = "Leroy";
16: leroy.parseName();
17: System.out.print("5");
18: } }
```

- A. 12
- B. 1234
- C. 1235
- D. 124
- E. 1245
- F. The code does not compile.
- G. An uncaught exception is thrown.

11. What is the output of the following program?

```
1: public class Cat {
2: public String name;
3: public void parseName() {
4: System.out.print("1");
5: try {
6: System.out.print("2");
7: int x = Integer.parseInt(name);
8: System.out.print("3");
```

```
9: } catch (NullPointerException e) {
10: System.out.print("4");
11: }
12: System.out.print("5");
13: }
14: public static void main(String[] args) {
15: Cat leo = new Cat();
16: leo.name = "Leo";
17: leo.parseName();
18: System.out.print("6");
19: }
20: }
```

- A. 12, followed by a stack trace for a `NumberFormatException`
- B. 124, followed by a stack trace for a `NumberFormatException`
- C. 12456
- D. 12456
- E. 1256, followed by a stack trace for a `NumberFormatException`
- F. The code does not compile.
- G. An uncaught exception is thrown.

**12.** What is printed by the following? (Choose all that apply)

```
1: public class Mouse {
2: public String name;
3: public void run() {
4: System.out.print("1");
5: try {
6: System.out.print("2");
7: name.toString();
8: System.out.print("3");
9: } catch (NullPointerException e) {
10: System.out.print("4");
11: throw e;
12: }
13: System.out.print("5");
14: }
15: public static void main(String[] args) {
16: Mouse jerry = new Mouse();
17: jerry.run();
18: System.out.print("6");
19: } }
```

- A. 1
  - B. 2
  - C. 3
  - D. 4
  - E. 5
  - F. 6
  - G. The stack trace for a `NullPointerException`
13. Which of the following statements are true? (Choose all that apply)
- A. You can declare a method with `Exception` as the return type.
  - B. You can declare any subclass of `Error` in the throws part of a method declaration.
  - C. You can declare any subclass of `Exception` in the throws part of a method declaration.
  - D. You can declare any subclass of `Object` in the throws part of a method declaration.
  - E. You can declare any subclass of `RuntimeException` in the throws part of a method declaration.
14. Which of the following can be inserted on line 8 to make this code compile? (Choose all that apply)
- ```
7: public void ohNo() throws IOException {  
8:     // INSERT CODE HERE  
9: }
```
- A. `System.out.println("it's ok");`
 - B. `throw new Exception();`
 - C. `throw new IllegalArgumentException();`
 - D. `throw new java.io.IOException();`
 - E. `throw new RuntimeException();`
15. Which of the following are unchecked exceptions? (Choose all that apply)
- A. `ArrayIndexOutOfBoundsException`
 - B. `IllegalArgumentException`
 - C. `IOException`
 - D. `NumberFormatException`
 - E. Any exception that extends `RuntimeException`
 - F. Any exception that extends `Exception`
16. Which scenario is the best use of an exception?
- A. An element is not found when searching a list.
 - B. An unexpected parameter is passed into a method.

- C. The computer caught fire.
- D. You want to loop through a list.
- E. You don't know how to code a method.

17. Which of the following can be inserted into `Lion` to make this code compile? (Choose all that apply)

```
class HasSoreThroatException extends Exception {}
class TiredException extends RuntimeException {}
interface Roar {
    void roar() throws HasSoreThroatException;
}
class Lion implements Roar { // INSERT CODE HERE
}
```

- A. `public void roar() {}`
- B. `public void roar() throws Exception {}`
- C. `public void roar() throws HasSoreThroatException {}`
- D. `public void roar() throws IllegalArgumentException {}`
- E. `public void roar() throws TiredException {}`

18. Which of the following are true? (Choose all that apply)

- A. Checked exceptions are allowed to be handled or declared.
- B. Checked exceptions are required to be handled or declared.
- C. Errors are allowed to be handled or declared.
- D. Errors are required to be handled or declared.
- E. Runtime exceptions are allowed to be handled or declared.
- F. Runtime exceptions are required to be handled or declared.

19. Which of the following can be inserted in the blank to make the code compile? (Choose all that apply)

```
public static void main(String[] args) {
    try {
        System.out.println("work real hard");
    } catch (_____ e) {
    } catch (RuntimeException e) {
    }
}
```

- A. `Exception`
- B. `IOException`
- C. `IllegalArgumentException`
- D. `RuntimeException`

- E. StackOverflowError
- F. None of the above.

20. What does the output of the following contain? (Choose all that apply)

```
12: public static void main(String[] args) {  
13:     System.out.print("a");  
14:     try {  
15:         System.out.print("b");  
16:         throw new IllegalArgumentException();  
17:     } catch (IllegalArgumentException e) {  
18:         System.out.print("c");  
19:         throw new RuntimeException("1");  
20:     } catch (RuntimeException e) {  
21:         System.out.print("d");  
22:         throw new RuntimeException("2");  
23:     } finally {  
24:         System.out.print("e");  
25:         throw new RuntimeException("3");  
26:     }  
27: }
```

- A. abce
- B. abde
- C. An exception with the message set to "1"
- D. An exception with the message set to "2"
- E. An exception with the message set to "3"
- F. Nothing; the code does not compile.

Appendix

A

Answers to Review Questions



Chapter 1: Java Building Blocks

1. A, B, E. Option A is valid because you can use the dollar sign in identifiers. Option B is valid because you can use an underscore in identifiers. Option C is not a valid identifier because `true` is a Java reserved word. Option D is not valid because the dot (`.`) is not allowed in identifiers. Option E is valid because Java is case sensitive, so `Public` is not a reserved word and therefore a valid identifier. Option F is not valid because the first character is not a letter, `$`, or `_`.
2. D. Boolean fields initialize to `false` and references initialize to `null`, so `empty` is `false` and `brand` is `null`. `Brand = null` is output.
3. B, D, E. Option A (line 4) compiles because `short` is an integral type. Option B (line 5) generates a compiler error because `int` is an integral type, but `5.6` is a floating-point type. Option C (line 6) compiles because it is assigned a `String`. Options D and E (lines 7 and 8) do not compile because `short` and `int` are primitives. Primitives do not allow methods to be called on them. Option F (line 9) compiles because `length()` is defined on `String`.
4. A, B. Adding the variable at line 2 makes `result` an instance variable. Since instance variables are in scope for the entire life of the object, option A is correct. Option B is correct because adding the variable at line 4 makes `result` a local variable with a scope of the whole method. Adding the variable at line 6 makes `result` a local variable with a scope of lines 6–7. Since it is out of scope on line 8, the `println` does not compile and option C is incorrect. Adding the variable at line 9 makes `result` a local variable with a scope of lines 9 and 10. Since line 8 is before the declaration, it does not compile and option D is incorrect. Finally, option E is incorrect because the code can be made to compile.
5. C, D. Option C is correct because it imports `Jelly` by classname. Option D is correct because it imports all the classes in the `jellies` package, which includes `Jelly`. Option A is incorrect because it only imports classes in the `aquarium` package—`Tank` in this case—and not those in lower-level packages. Option B is incorrect because you cannot use wildcards anywhere other than the end of an import statement. Option E is incorrect because you cannot import parts of a class with a regular import statement. Option F is incorrect because options C and D do make the code compile.
6. E. The first two imports can be removed because `java.lang` is automatically imported. The second two imports can be removed because `Tank` and `Water` are in the same package, making the correct answer E. If `Tank` and `Water` were in different packages, one of these two imports could be removed. In that case, the answer would be option D.
7. A, B, C. Option A is correct because it imports all the classes in the `aquarium` package including `aquarium.Water`. Options B and C are correct because they import `Water` by classname. Since importing by classname takes precedence over wildcards, these compile. Option D is incorrect because Java doesn't know which of the two wildcard `Water`

classes to use. Option E is incorrect because you cannot specify the same classname in two imports.

8. B. Option B is correct because arrays start counting from zero and strings with spaces must be in quotes. Option A is incorrect because it outputs Blue. C is incorrect because it outputs Jay. Option D is incorrect because it outputs Sparrow. Options E and F are incorrect because they output Error: Could not find or load main class Bird-Display.class.
9. A, C, D, E. Option A is correct because it is the traditional `main()` method signature and variables may begin with underscores. Options C and D are correct because the array operator may appear after the variable name. Option E is correct because varargs are allowed in place of an array. Option B is incorrect because variables are not allowed to begin with a digit. Option F is incorrect because the argument must be an array or varargs. Option F is a perfectly good method. However, it is not one that can be run from the command line because it has the wrong parameter type.
10. E. Option E is the canonical `main()` method signature. You need to memorize it. Option A is incorrect because the `main()` method must be public. Options B and F are incorrect because the `main()` method must have a void return type. Option C is incorrect because the `main()` method must be static. Option D is incorrect because the `main()` method must be named `main`.
11. C, D. Option C is correct because all non-primitive values default to `null`. Option D is correct because float and double primitives default to `0.0`. Options B and E are incorrect because `int` primitives default to `0`.
12. G. Option G is correct because local variables do not get assigned default values. The code fails to compile if a local variable is not explicitly initialized. If this question were about instance variables, options D and F would be correct. A `boolean` primitive defaults to `false` and a `float` primitive defaults to `0.0`.
13. A, D. Options A and D are correct because `boolean` primitives default to `false` and `int` primitives default to `0`.
14. D. The package name represents any folders underneath the current path, which is named `.A` in this case. Option B is incorrect because package names are case sensitive, just like variable names and other identifiers.
15. A, E. Underscores are allowed as long as they are directly between two other digits. This means options A and E are correct. Options B and C are incorrect because the underscore is adjacent to the decimal point. Option D is incorrect because the underscore is the last character.
16. B, C, D. `0b` is the prefix for a binary value and is correct. `0x` is the prefix for a hexadecimal value. This value can be assigned to many primitive types, including `int` and `double`, making options C and D correct. Option A is incorrect because `9L` is a long value. `long amount = 9L` would be allowed. Option E is incorrect because the underscore is immediately before the decimal. Option F is incorrect because the underscore is the very last character.

17. A, E. Bunny is a class, which can be seen from the declaration: `public class Bunny`. `bun` is a reference to an object. `main()` is a method.
18. C, D, E. `package` and `import` are both optional. If both are present, the order must be `package`, then `import`, then `class`. Option A is incorrect because `class` is before `package` and `import`. Option B is incorrect because `import` is before `package`. Option F is incorrect because `class` is before `package`. Option G is incorrect because `class` is before `import`.
19. B, D. The `Rabbit` object from line 3 has two references to it: one and three. The references are nulled out on lines 6 and 8, respectively. Option B is correct because this makes the object eligible for garbage collection after line 8. Line 7 sets the reference four to the now null one, which means it has no effect on garbage collection. The `Rabbit` object from line 4 only has a single reference to it: two. Option D is correct because this single reference becomes null on line 9. The `Rabbit` object declared on line 10 becomes eligible for garbage collection at the end of the method on line 12. Calling `System.gc()` has no effect on eligibility for garbage collection.
20. B, E. Calling `System.gc()` suggests that Java might wish to run the garbage collector. Java is free to ignore the request, making option E correct. `finalize()` runs if an object attempts to be garbage collected, making option B correct.
21. A. While the code on line 3 does compile, it is not a constructor because it has a return type. It is a method that happens to have the same name as the class. When the code runs, the default constructor is called and `count` has the default value (0) for an `int`.
22. B, E. C++ has operator overloading and pointers. Java made a point of not having either. Java does have references to objects, but these are pointing to an object that can move around in memory. Option B is correct because Java is platform independent. Option E is correct because Java is object oriented. While it does support some parts of functional programming, these occur within a class.
23. C, D. Java puts source code in `.java` files and bytecode in `.class` files. It does not use a `.bytecode` file. When running a Java program, you pass just the name of the class without the `.class` extension.

Chapter 2: Operators and Statements

1. A, D. Option A is the equality operator and can be used on numeric primitives, boolean values, and object references. Options B and C are both **arithmetic operators and cannot be applied to a boolean value**. Option D is the logical complement operator and is used exclusively with boolean values. Option E is **the modulus operator, which can only be used with numeric primitives**. Finally, option F is a relational operator that compares the values of two numbers.

2. A, B, D. The value `x + y` is automatically promoted to `int`, so `int` and data types that can be promoted automatically from `int` will work. Options A, B, D are such data types. Option C will not work because `boolean` is not a numeric data type. Options E and F will not work without an explicit cast to a smaller data type.
3. F. In this example, the ternary operator has two expressions, one of them a `String` and the other a `boolean` value. The ternary operator is permitted to have expressions that don't have matching types, but the key here is the assignment to the `String` reference. The compiler knows how to assign the first expression value as a `String`, but the second `boolean` expression cannot be set as a `String`; therefore, this line will not compile.
4. B, C, D, F. The code will not compile as is, so option A is not correct. The value `2 * x` is automatically promoted to `long` and cannot be automatically stored in `y`, which is in an `int` value. Options B, C, and D solve this problem by reducing the `long` value to `int`. Option E does not solve the problem and actually makes it worse by attempting to place the value in a smaller data type. Option F solves the problem by increasing the data type of the assignment so that `long` is allowed.
5. C. This code does not contain any compilation errors or an infinite loop, so options D, E, and F are incorrect. The `break` statement on line 8 causes the loop to execute once and finish, so option C is the correct answer.
6. F. The code does not compile because two `else` statements cannot be chained together without additional `if-then` statements, so the correct answer is option F. Option E is incorrect as Line 6 by itself does not cause a problem, only when it is paired with Line 7. One way to fix this code so it compiles would be to add an `if-then` statement on line 6. The other solution would be to remove line 7.
7. D. As you learned in the section “Ternary Operator,” although parentheses are not required, they do greatly increase code readability, such as the following equivalent statement:

```
System.out.println((x > 2) ? ((x < 4) ? 10 : 8) : 7)
```

We apply the outside ternary operator first, as it is possible the inner ternary expression may never be evaluated. Since `(x > 2)` is true, this reduces the problem to:

```
System.out.println((x < 4) ? 10 : 8)
```

Since `x` is greater than 2, the answer is 8, or option D in this case.

8. B. This example is tricky because of the second assignment operator embedded in line 5. The expression `(z=false)` assigns the value `false` to `z` and returns `false` for the entire expression. Since `y` does not equal 10, the left-hand side returns `true`; therefore, the exclusive or (`^`) of the entire expression assigned to `x` is `true`. The output reflects these assignments, with no change to `y`, so option B is the only correct answer. The code compiles and runs without issue, so option F is not correct.
9. F. In this example, the update statement of the `for` loop is missing, which is fine as the statement is optional, so option D is incorrect. The expression inside the loop increments `i` but then assigns `i` the old value. Therefore, `i` ends the loop with the same value

that it starts with: 0. The loop will repeat infinitely, outputting the same statement over and over again because `i` remains 0 after every iteration of the loop.

10. D. Line 4 generates a possible loss of precision compiler error. The cast operator has the highest precedence, so it is evaluated first, casting `a` to a byte. Then, the addition is evaluated, causing both `a` and `b` to be promoted to `int` values. The value 90 is an `int` and cannot be assigned to the byte `sum` without an explicit cast, so the code does not compile. The code could be corrected with parentheses around `(a + b)`, in which case option C would be the correct answer.
11. A. The `*` and `%` have the same operator precedence, so the expression is evaluated from left-to-right. The result of `5 * 4` is 20, and `20 % 3` is 2 (20 divided by 3 is 18, the remainder is 2). The output is 2 and option A is the correct answer.
12. D. The variable `x` is an `int` and `s` is a reference to a `String` object. The two data types are incomparable because neither variable can be converted to the other variable's type. The compiler error occurs on line 5 when the comparison is attempted, so the answer is option D.
13. A. The code compiles successfully, so options C and D are incorrect. The value of `b` after line 4 is `false`. However, the `if-then` statement on line 5 contains an assignment, not a comparison. The variable `b` is assigned `true` on line 3, and the assignment operator returns `true`, so line 5 executes and displays `Success`, so the answer is option A.
14. C. The code compiles successfully, so option F is incorrect. On line 5, the pre-increment operator is used, so `c` is incremented to 4 and the new value is returned to the expression. The value of `result` is computed by adding 4 to the original value of 8, resulting in a new value of 12, which is output on line 6. Therefore, option C is the correct answer.
15. E. This is actually a much simpler problem than it appears to be. The `while` statement on line 4 is missing parentheses, so the code will not compile, and option E is the correct answer. If the parentheses were added, though, option F would be the correct answer since the loop does not use curly braces to include `x++` and the boolean expression never changes. Finally, if curly braces were added around both expressions, the output would be 10, 6 and option B would be correct.
16. D. The variable `y` is declared within the body of the `do-while` statement, so it is out of scope on line 6. Line 6 generates a compiler error, so option D is the correct answer.
17. D. The code compiles without issue, so option F is incorrect. After the first execution of the loop, `i` is decremented to 9 and `result` to 13. Since `i` is not 8, `keepGoing` is `false`, and the loop continues. On the next iteration, `i` is decremented to 8 and `result` to 11. On the second execution, `i` does equal 8, so `keepGoing` is set to `false`. At the conclusion of the loop, the loop terminates since `keepGoing` is no longer `true`. The value of `result` is 11, and the correct answer is option D.
18. A. The expression on line 5 is `true` when `row * col` is an even number. On the first iteration, `row = 1` and `col = 1`, so the expression on line 6 is `false`, the `continue` is skipped, and `count` is incremented to 1. On the second iteration, `row = 1` and

`col = 2`, so the expression on line 6 is true and the `continue` ends the outer loop with count still at 1. On the third iteration, `row = 2` and `col = 1`, so the expression on line 6 is true and the `continue` ends the outer loop with count still at 1. On the fourth iteration, `row = 3` and `col = 1`, so the expression on line 6 is false, the `continue` is skipped, and count is incremented to 2. Finally, on the fifth and final iteration, `row = 3` and `col = 2`, so the expression on line 6 is true and the `continue` ends the outer loop with count still at 2. **The result of 2 is displayed, so the answer is option B.**

19. D. Prior to the first iteration, `m = 9`, `n = 1`, and `x = 0`. After the iteration of the first loop, `m` is updated to 8, `n` to 3, and `x` to the sum of the new values for `m + n`, `0 + 11 = 11`. After the iteration of the second loop, `m` is updated to 7, `n` to 5, and `x` to the sum of the new values for `m + n`, `11 + 12 = 23`. After the iteration of the third loop, `m` is updated to 6, `n` to 7, and `x` to the sum of the new values for `m + n`, `23 + 13 = 36`. On the fourth iteration of the loop, `m > n` evaluates to false, as `6 < 7` is not true. The loop ends and the most recent value of `x`, 36, is output, so the correct answer is option D.
20. B. The code compiles and runs without issue, so options C, D, and E are not correct. The value of `grade` is 'B' and there is a matching case statement that will cause "great" to be printed. There is no `break` statement after the case, though, so the next case statement will be reached, and "good" will be printed. There is a `break` after this case statement, though, so the `switch` statement will end. The correct answer is thus option B.

Chapter 3: Core Java APIs

1. G. Line 5 does not compile. This question is checking to see if you are paying attention to the types. `numFish` is an `int` and 1 is an `int`. Therefore, we use numeric addition and get 5. The problem is that we can't store an `int` in a `String` variable. Supposing line 5 said `String anotherFish = numFish + 1 + ""`; In that case, the answer would be options A and D. The variable defined on line 5 would be the string "5", and both output statements would use concatenation.
2. A, C, D. The code compiles fine. Line 3 points to the `String` in the string pool. Line 4 calls the `String` constructor explicitly and is therefore a different object than `s`. Lines 5 and 7 check for object equality, which is true, and so print one and three. Line 6 uses object reference equality, which is not true since we have different objects. Line 7 also compares references but is true since both references point to the object from the string pool. Finally, line 8 compares one object from the string pool with one that was explicitly constructed and returns false.
3. B, C, E. Immutable means the state of an object cannot change once it is created. Immutable objects can be garbage collected just like mutable objects. `String` is immutable. `StringBuilder` can be mutated with methods like `append()`. Although

StringBuffer isn't on the exam, you should know about it anyway in case older questions haven't been removed.

4. B. This example uses method chaining. After the call to `append()`, `sb` contains "aaa". That result is passed to the first `insert()` call, which inserts at index 1. At this point `sb` contains `abbbaa`. That result is passed to the final `insert()`, which inserts at index 4, resulting in `abbacca`.
5. F. The question is trying to distract you into paying attention to logical equality versus object reference equality. It is hoping you will miss the fact that line 4 does not compile. Java does not allow you to compare `String` and `StringBuilder` using `==`.
6. B. A `String` is immutable. Calling `concat()` returns a new `String` but does not change the original. A `StringBuilder` is mutable. Calling `append()` adds characters to the existing character sequence along with returning a reference to the same object.
7. B, D, E. `length()` is simply a count of the number of characters in a `String`. In this case, there are six characters. `charAt()` returns the character at that index. Remember that indexes are zero based, which means that index 3 corresponds to `d` and index 6 corresponds to 1 past the end of the array. A `StringIndexOutOfBoundsException` is thrown for the last line.
8. A, D, E. `substring()` has two forms. The first takes the index to start with and the index to stop immediately before. The second takes just the index to start with and goes to the end of the `String`. Remember that indexes are zero based. The first call starts at index 1 and ends with index 2 since it needs to stop before index 3. The second call starts at index 7 and ends in the same place, resulting in an empty `String`. This prints out a blank line. The final call starts at index 7 and goes to the end of the `String`.
9. C. This question is trying to see if you know that `String` objects are immutable. Line 4 returns "PURR" but the result is ignored and not stored in `s`. Line 5 returns "purrr" since there is no whitespace present but the result is again ignored. Line 6 returns "ur" because it starts with index 1 and ends before index 3 using zero-based indexes. The result is ignored again. Finally, on line 6 something happens. We concatenate four new characters to `s` and now have a `String` of length 8.
10. F. `a += 2` expands to `a = a + 2`. A `String` concatenated with any other type gives a `String`. Lines 14, 15, and 16 all append to `a`, giving a result of "2cfalse". The `if` statement on line 18 returns `false` because the values of the two `String` objects are the same using object equality. The `if` statement on line 17 returns `false` because the two `String` objects are not the same in memory. One comes directly from the string pool and the other comes from building using `String` operations.
11. E. Line 6 adds 1 to `total` because `substring()` includes the starting index but not the ending index. Line 7 adds 0 to `total`. Line 8 is a problem: Java does not allow the indexes to be specified in reverse order and the code throws a `StringIndexOutOfBoundsException`.

12. A. First, we delete the characters at index 2 until the character one before index 8. At this point, 0189 is in numbers. The following line uses method chaining. It appends a dash to the end of the characters sequence, resulting in 0189-, and then inserts a plus sign at index 2, resulting in 01+89-.
13. F. This is a trick question. The first line does not compile because you cannot assign a String to a StringBuilder. If that line were `StringBuilder b = new StringBuilder("rumble")`, the code would compile and print rum4. Watch out for this sort of trick on the exam. You could easily spend a minute working out the character positions for no reason at all.
14. A, C. The `reverse()` method is the easiest way of reversing the characters in a `StringBuilder`; therefore, option A is correct. Option B is a nice distraction—it does in fact return "avaJ". However, `substring()` returns a `String`, which is not stored anywhere. Option C uses method chaining. First it creates the value "JavavaJ\$". Then it removes the first three characters, resulting in "avaJ\$". Finally, it removes the last character, resulting in "avaJ". Option D throws an exception because you cannot delete the character after the last index. Remember that `deleteCharAt()` uses indexes that are zero based and `length()` counts starting with 1.
15. C, E, F. Option C uses the variable name as if it were a type, which is clearly illegal. Options E and F don't specify any size. **Although it is legal to leave out the size for later dimensions of a multidimensional array, the first one is required.** Option A declares a legal 2D array. Option B declares a legal 3D array. Option D declares a legal 2D array. Remember that it is normal to see on the exam types you might not have learned. You aren't expected to know anything about them.
16. C. **Arrays define a property called length. It is not a method, so parentheses are not allowed.**
17. F. **The ArrayList class defines a method called size().**
18. A, C, D, E. **An array is not able to change size and can have multiple dimensions. Both an array and ArrayList are ordered and have indexes. Neither is immutable. The elements can change in value.**
19. B, C. **An array does not override equals() and so uses object equality. ArrayList does override equals() and defines it as the same elements in the same order.** The compiler does not know when an index is out of bounds and thus can't give you a compiler error. The code will throw an exception at runtime, though.
20. D. The code does not compile because `list` is instantiated using generics. Only `String` objects can be added to `list` and 7 is an `int`.
21. C. After line 4, `values` has one element (4). After line 5, `values` has two elements (4, 5). After line 6, `values` has two elements (4, 6) because `set()` does a replace. After line 7, `values` has only one element (6).
22. D. The code compiles and runs fine. However, an array must be sorted for `binarySearch()` to return a meaningful result.

23. A. Line 4 creates a fixed size array of size 4. Line 5 sorts it. Line 6 converts it back to an array. The brackets aren't in the traditional place, but they are still legal. Line 7 prints the first element, which is now -1.
24. C. Converting from an array to an ArrayList uses `Arrays.asList(names)`. There is no `asList()` method on an array instance. If this code were corrected to compile, the answer would be option A.
25. D. After sorting, `hex` contains `[30, 3A, 8, FF]`. Remember that numbers sort before letters and strings sort alphabetically. This makes 30 come before 8. A binary search correctly finds 8 at index 2 and 3A at index 1. It cannot find 4F but notices it should be at index 2. The rule when an item isn't found is to negate that index and subtract 1. Therefore, we get $-2-1$, which is -3 .
26. A, B, D. Lines 5 and 7 use autoboxing to convert an `int` to an `Integer`. Line 6 does not because `valueOf()` returns an `Integer`. Line 8 does not because `null` is not an `int`. The code does not compile. However, when the for loop tries to unbox `null` into an `int`, it fails and throws a `NullPointerException`.
27. B. The first `if` statement is false because the variables do not point to the same object. The second `if` statement is true because `ArrayList` implements equality to mean the same elements in the same order.
28. D, F. Options A and B are incorrect because `LocalDate` does not have a public constructor. Option C is incorrect because months start counting with 1 rather than 0. Option E is incorrect because it uses the old pre-Java 8 way of counting months, again beginning with 0. Options D and F are both correct ways of specifying the desired date.
29. D. A `LocalDate` does not have a time element. Therefore, it has no method to add hours and the code does not compile.
30. F. Java throws an exception if invalid date values are passed. There is no 40th day in April—or any other month for that matter.
31. B. The date starts out as April 30, 2018. Since dates are immutable and the plus methods have their return values ignored, the result is unchanged. Therefore, option B is correct.
32. E. Even though `d` has both date and time, the formatter only outputs time.
33. B. `Period` does not allow chaining. Only the last `Period` method called counts, so only the two years are subtracted.

Chapter 4: Methods and Encapsulation

1. B, C. `void` is a return type. Only the access modifier or optional specifiers are allowed before the return type. Option C is correct, creating a method with private access. Option B is correct, creating a method with default access and the optional specifier `final`. Since default access does not require a modifier, we get to jump right to `final`.

Option A is incorrect because default access omits the access modifier rather than specifying default. Option D is incorrect because Java is case sensitive. It would have been correct if `public` were the choice. Option E is incorrect because the method already has a `void` return type. Option F is incorrect because labels are not allowed for methods.

2. A, D. Options A and D are correct because the optional specifiers are allowed in any order. Options B and C are incorrect because they each have two return types. Options E and F are incorrect because the return type is before the optional specifier and access modifier, respectively.
3. A, C, D. Options A and C are correct because a `void` method is allowed to have a return statement as long as it doesn't try to return a value. Options B and G do not compile because `null` requires a reference object as the return type. `void` is not a reference object since it is a marker for no return type. `int` is not a reference object since it is a primitive. Option D is correct because it returns an `int` value. Option E does not compile because it tries to return a `double` when the return type is `int`. Since a `double` cannot be assigned to an `int`, it cannot be returned as one either. Option F does not compile because no value is actually returned.
4. A, B, G. Options A and B are correct because the single `vararg` parameter is the last parameter declared. Option G is correct because it doesn't use any `vararg` parameters at all. Options C and F are incorrect because the `vararg` parameter is not last. Option D is incorrect because two `vararg` parameters are not allowed in the same method. Option E is incorrect because the `...` for a `vararg` must be after the type, not before it.
5. D, G. Option D passes the initial parameter plus two more to turn into a `vararg` array of size 2. Option G passes the initial parameter plus an array of size 2. Option A does not compile because it does not pass the initial parameter. Options E and F do not compile because they do not declare an array properly. It should be `new boolean[] {true}`. Option B creates a `vararg` array of size 0 and option C creates a `vararg` array of size 1.
6. D. Option D is correct. This is the common implementation for encapsulation by setting all fields to be `private` and all methods to be `public`. Option A is incorrect because `protected` access allows everything that `package private` access allows and additionally allows subclasses access. Option B is incorrect because the class is `public`. This means that other classes can see the class. However, they cannot call any of the methods or read any of the fields. It is essentially a useless class. Option C is incorrect because `package private` access applies to the whole package. Option E is incorrect because Java has no such capability.
7. B, C, D, F. The two classes are in different packages, which means `private` access and default (`package private`) access will not compile. Additionally, `protected` access will not compile since `School` does not inherit from `Classroom`. Therefore, only line 8 will compile because it uses `public` access.
8. B, C, E. Encapsulation requires using methods to get and set instance variables so other classes are not directly using them. Instance variables must be `private` for this to work. Immutability takes this a step further, allowing only getters, so the instance variables do not change state.

9.  C, E. Option A is incorrect because the property is of type boolean and getters must begin with `is` for booleans. Options B and D are incorrect because they don't follow the naming convention of beginning with `get/is/set`. Options C and E follow normal getter and setter conventions.
10. B. Rope runs line 3, setting `LENGTH` to 5, then immediately after runs the static initializer, which sets it to 10. Line 5 calls the static method normally and prints `swing`. Line 6 also calls the static method. Java allows calling a static method through an instance variable. Line 7 uses the static import on line 2 to reference `LENGTH`.
11. B, E. Line 10 does not compile because **static methods are not allowed to call instance methods**. Even though we are calling `play()` as if it were an instance method and an instance exists, Java knows `play()` is really a static method and treats it as such. If line 10 is removed, the code works. It does not throw a `NullPointerException` on line 16 because `play()` is a static method. Java looks at the type of the reference for `rope2` and translates the call to `Rope.play()`.
12. D. There are two details to notice in this code. First, note that `RopeSwing` has an instance initializer and not a static initializer. Since `RopeSwing` is never constructed, the instance initializer does not run. The other detail is that `length` is static. Changes from one object update this common static variable.
13. E. static final variables must be set exactly once, and it must be in the declaration line or in a static initialization block. Line 4 doesn't compile because `bench` is not set in either of these locations. Line 15 doesn't compile because final variables are not allowed to be set after that point. Line 11 doesn't compile because `name` is set twice: once in the declaration and again in the static block. Line 12 doesn't compile because `rightRope` is set twice as well. Both are in static initialization blocks.
14. B. The two valid ways to do this are `import static java.util.Collections.*;` and `import static java.util.Collections.sort;`. Option A is incorrect because you can only do a static import on static members. Classes such as `Collections` require a regular import. Option C is nonsense as method parameters have no business in an import. Options D, E, and F try to trick you into reversing the syntax of `import static`.
15. E. The argument on line 17 is a short. It can be promoted to an int, so `print()` on line 5 is invoked. The argument on line 18 is a boolean. It can be autoboxed to a boolean, so `print()` on line 11 is invoked. The argument on line 19 is a double. It can be autoboxed to a double, so `print()` on line 11 is invoked. Therefore, the output is `intObjectObject` and the correct answer is option E.
16. B. Since Java is pass-by-value and the variable on line 8 never gets reassigned, it stays as 9. In the method `square`, `x` starts as 9. `y` becomes 81 and then `x` gets set to -1. Line 9 does set `result` to 81. However, we are printing out `value` and that is still 9.
17. B, D, E. Since Java is pass-by-reference, assigning a new object to `a` does not change the caller. Calling `append()` does affect the caller because both the method parameter and

caller have a reference to the same object. Finally, returning a value does pass the reference to the caller for assignment to `s3`.

18. C, G. **Since the `main()` method is in the same class, it can call private methods in the class.** `this()` may only be called as the first line of a constructor. `this.variableName` can be called from any instance method to refer to an instance variable. It cannot be called from a static method because there is no instance of the class to refer to. Option F is tricky. **The default constructor is only written by the compiler if no user-defined constructors were provided.** `this()` can only be called from a constructor in the same class. Since there can be no user-defined constructors in the class if a default constructor was created, it is impossible for option F to be true.
19. A, G. Options B and C don't compile because the constructor name must match the classname. Since Java is case sensitive, these don't match. Options D, E, and F all compile and provide one user-defined constructor. Since a constructor is coded, a default constructor isn't supplied. Option G defines a method, but not a constructor. Option A does not define a constructor, either. Since no constructor is coded, a default constructor is provided for options A and G.
20. E. Options A and B will not compile because constructors cannot be called without `new`. Options C and D will compile but will create a new object rather than setting the fields in this one. Option F will not compile because `this()` must be the first line of a constructor. Option E is correct.
21. C. Within the constructor `numSpots` refers to the constructor parameter. The instance variable is hidden because they have the same name. `this.numSpots` tells Java to use the instance variable. In the `main()` method, `numSpots` refers to the instance variable. Option A sets the constructor parameter to itself, leaving the instance variable as 0. Option B sets the constructor parameter to the value of the instance variable, making them both 0. Option C is correct, setting the instance variable to the value of the constructor parameter. Options D and E do not compile.
22. E. On line 3 of `OrderDriver`, we refer to `Order` for the first time. At this point the statics in `Order` get initialized. In this case, the statics are the static declaration of `result` and the static initializer. `result` is `u` at this point. On line 4, `result` is the same because the static initialization is only run once. On line 5, we create a new `Order`, which triggers the instance initializers in the order they appear in the file. Now `result` is `ucr`. Line 6 creates another `Order`, triggering another set of initializers. Now `result` is `ucrcr`. Notice how the static is on a different line than the initialization code in lines 4–5 of `Order`. The exam may try to trick you by formatting the code like this to confuse you.
23. A. Line 4 instantiates an `Order`. Java runs the declarations and instance initializers first in the order they appear. This sets `value` to `tacf`. Line 5 creates another `Order` and initializes `value` to `tacb`. The object on line 5 is stored in the same variable line 4 used. This makes the object created on line 4 unreachable. When `value` is printed, it is the instance variable in the object created on line 5.

24. B, C, E. *value1* is a final instance variable. It can only be set once: in the variable declaration, an instance initializer, or a constructor. Option A does not compile because the final variable was already set in the declaration. *value2* is a static variable. Both instance and static initializers are able to access static variables, making options B and E correct. *value3* is an instance variable. Options D and F do not compile because a static initializer does not have access to instance variables.
25. A, E. The 100 parameter is an int and so calls the matching int constructor. When this constructor is removed, Java looks for the next most specific constructor. **Java prefers autoboxing to varargs, and so chooses the Integer constructor.** The 100L parameter is a long. Since it can't be converted into a smaller type, it is autoboxed into a Long and then the constructor for Object is called.
26. A. This code is correct. Line 8 creates a lambda expression that checks if the age is less than 5. **Since there is only one parameter and it does not specify a type, the parentheses around the type parameter are optional.** Line 10 uses the Predicate interface, which declares a test() method.
27. C. The interface takes two int parameters. The code on line 7 attempts to use them as if one is a StringBuilder. It is tricky to use types in a lambda when they are implicitly specified. Remember to check the interface for the real type.
28. A, D, F. removeIf() expects a Predicate, which takes a parameter list of one parameter using the specified type. Options B and C are incorrect because they do not use the return keyword. It is required inside braces for lambda bodies. Option E is incorrect because it is missing the parentheses around the parameter list. This is only optional for a single parameter with an inferred type.
29. A, F. Option B is incorrect because it does not use the return keyword. Options C, D, and E are incorrect because the variable *e* is already in use from the lambda and cannot be redefined. Additionally, option C is missing the return keyword and option E is missing the semicolon.

Chapter 5: Class Design

1. B. **All interface methods are implicitly public,** so option B is correct and option A is not. Interface methods may be declared as static or default but are never implicitly added, so options C and F are incorrect. Option D is incorrect—void is not a modifier; it is a return type. Option E is a tricky one, because prior to Java 8 all interface methods would be assumed to be abstract. **Since Java 8 now includes default and static methods and they are never abstract, you cannot assume the abstract modifier will be implicitly applied to all methods by the compiler.**
2. E. The code will not compile because the parent class Mammal doesn't define a no-argument constructor, so the first line of a Platypus constructor should be an explicit call to super(int age). If there was such a call, then the output would be MammalPlatypus, since the super constructor is executed before the child constructor.

3. A, B, D, E. The blank can be filled with any class or interface that is a supertype of `TurtleFrog`. Option A is a superclass of `TurtleFrog`, and option B is the same class, so both are correct. `BrazilianHornedFrog` is not a superclass of `TurtleFrog`, so option C is incorrect. `TurtleFrog` inherits the `CanHope` interface, so option D is correct. All classes inherit `Object`, so option E is correct. Finally, `Long` is an unrelated class that is not a superclass of `TurtleFrog`, and is therefore incorrect.
4. C, E. The code doesn't compile, so option A is incorrect. Option B is also not correct because **the rules for overriding a method allow a subclass to define a method with an exception that is a subclass of the exception in the parent method**. Option C is correct because the return types are not covariant; in particular, `Number` is not a subclass of `Integer`. Option D is incorrect because the subclass defines a method that is more accessible than the method in the parent class, which is allowed. Finally, option E is correct because the method is declared as `static` in the parent class and not so in the child class. **For nonprivate methods in the parent class, both methods must use `static` (hide) or neither should use `static` (override).**
5. A, D, E, F. First off, options B and C are incorrect because protected and public methods may be overridden, not hidden. Option A is correct because **private methods are always hidden in a subclass**. Option D is also correct because **static methods cannot be overridden, only hidden**. Options E and F are correct because **variables may only be hidden, regardless of the access modifier**.
6. D. The code fails to compile because `Beetle`, the first concrete subclass, doesn't implement `getNumberOfSections()`, which is inherited as an abstract method; therefore, option D is correct. Option B is incorrect because there is nothing wrong with this interface method definition. Option C is incorrect because an abstract class is not required to implement any abstract methods, including those inherited from an interface. Option E is incorrect because the code fails at compilation-time.
7. B, C. **A reference to an object requires an explicit cast if referenced with a subclass, so option A is incorrect. If the cast is to a superclass reference, then an explicit cast is not required. Because of polymorphic parameters, if a method takes the superclass of an object as a parameter, then any subclass references may be used without a cast,** so option B is correct. All objects extend `java.lang.Object`, so if a method takes that type, any valid object, including `null`, may be passed; therefore, option C is correct. Some cast exceptions can be detected as errors at compile-time, but others can only be detected at runtime, so D is incorrect. Due to the nature of polymorphism, a public instance method can be overridden in a subclass and calls to it will be replaced even in the superclass it was defined, so E is incorrect.
8. F. The interface variable amount is correctly declared, with `public` and `static` being assumed and automatically inserted by the compiler, so option B is incorrect. The method declaration for `eatGrass()` on line 3 is incorrect because the method has been marked as `static` but no method body has been provided. The method declaration for `chew()` on line 4 is also incorrect, since an interface method that provides a body must be marked as `default` or `static` explicitly. Therefore, option F is the correct answer since this code contains two compile-time errors.

9. A. Although the definition of methods on lines 2 and 5 vary, both will be converted to public abstract by the compiler. Line 4 is fine, because an interface can have public or default access. Finally, the class Falcon doesn't need to implement the interface methods because it is marked as abstract. Therefore, the code will compile without issue.
10. B, C, E, F. Option A is wrong, because an abstract class may contain concrete methods. Since Java 8, interfaces may also contain concrete methods in form of static or default methods. **Although all variables in interfaces are assumed to be public static final, abstract classes may contain them as well**, so option B is correct. Both abstract classes and interfaces can be extended with the extends keyword, so option C is correct. **Only interfaces can contain default methods**, so option D is incorrect. **Both abstract classes and interfaces can contain static methods**, so option E is correct. **Both structures require a concrete subclass to be instantiated**, so option F is correct. Finally, though an instance of an object that implements an interface inherits java.lang.Object, the interface itself doesn't; otherwise, Java would support multiple inheritance for objects, which it doesn't. Therefore, option G is incorrect.
11. A, D, E. **Interface variables are assumed to be public static final**; therefore, options A, D, and E are correct. Options B and C are incorrect because interface variables must be public—interfaces are implemented by classes, not inherited by interfaces. Option F is incorrect because variables can never be abstract.
12. B. This code compiles and runs without issue, outputting false, so option B is the correct answer. The first declaration of isBlind() is as a default interface method, assumed public. The second declaration of isBlind() correctly overrides the default interface method. Finally, the newly created Owl instance may be automatically cast to a Nocturnal reference without an explicit cast, although adding it doesn't break the code.
13. A. The code compiles and runs without issue, so options E and F are incorrect. The printName() method is an overload in Spider, not an override, so both methods may be called. The call on line 8 references the version that takes an int as input defined in the Spider class, and the call on line 9 references the version in the Arthropod class that takes a double. Therefore, SpiderArthropod is output and option A is the correct answer.
14. C. The code compiles without issue, so option A is wrong. Option B is incorrect, since an abstract class could implement HasVocalCords without the need to override the makeSound() method. Option C is correct; any class that implements CanBark automatically inherits its methods, as well as any inherited methods defined in the parent interface. Because option C is correct, it follows that option D is incorrect. Finally, an interface can extend multiple interfaces, so option E is incorrect.
15. B. Concrete classes are, by definition, not abstract, so option A is incorrect. A concrete class must implement all inherited abstract methods, so option B is correct. Option C is incorrect; a superclass may have already implemented an inherited interface, so the concrete subclass would not need to implement the method. Concrete classes can be both final and not final, so option D is incorrect. Finally, abstract methods must be overridden by a concrete subclass, so option E is incorrect.

16. E. The code doesn't compile, so options A and B are incorrect. The issue with line 9 is that `layEggs()` is marked as `final` in the superclass `Reptile`, which means it cannot be overridden. There are no errors on any other lines, so options C and D are incorrect.
17. B. This may look like a complex question, but it is actually quite easy. Line 2 contains an invalid definition of an abstract method. Abstract methods cannot contain a body, so the code will not compile and option B is the correct answer. If the body `{}` was removed from line 2, the code would still not compile, although it would be line 8 that would throw the compilation error. Since `dive()` in `Whale` is abstract and `Orca` extends `Whale`, then it must implement an overridden version of `dive()`. The method on line 9 is an overloaded version of `dive()`, not an overridden version, so `Orca` is an invalid subclass and will not compile.
18. E. The code doesn't compile because line 6 contains an incompatible override of the `getNumberOfGills(int input)` method defined in the `Aquatic` interface. In particular, `int` and `String` are not covariant return types, since `int` is not a subclass of `String`. Note that line 5 compiles without issue; `getNumberOfGills()` is an overloaded method that is not related to the parent interface method that takes an `int` value.
19. A, C, F. First off, `Cobra` is a subclass of `Snake`, so option A can be used. `GardenSnake` is not defined as a subclass of `Snake`, so it cannot be used and option B is incorrect. The class `Snake` is not marked as abstract, so it can be instantiated and passed, so option C is correct. Next, `Object` is a superclass of `Snake`, not a subclass, so it also cannot be used and option D is incorrect. The class `String` is unrelated in this example, so option E is incorrect. Finally, a `null` value can always be passed as an object value, regardless of type, so option F is correct.
20. A. The code compiles and runs without issue, so options C, D, and E are incorrect. The trick here is that the method `fly()` is marked as `private` in the parent class `Bird`, which means it may only be hidden, not overridden. With hidden methods, the specific method used depends on where it is referenced. Since it is referenced within the `Bird` class, the method declared on line 2 was used, and option A is correct. Alternatively, if the method was referenced within the `Pelican` class, or if the method in the parent class was marked as `protected` and overridden in the subclass, then the method on line 9 would have been used.

Chapter 6: Exceptions

1. B. Runtime exceptions are also known as unchecked exceptions. They are allowed to be declared, but they don't have to be. Checked exceptions must be handled or declared. Legally, you can handle `java.lang.Error` subclasses, but it's not a good idea.
2. B, D. In a method declaration, the keyword `throws` is used. To actually throw an exception, the keyword `throw` is used and a new exception is created.
3. C. A `try` statement is required to have a `catch` clause and/or `finally` clause. If it goes the `catch` route, it is allowed to have multiple `catch` clauses.

4. B. The second line tries to cast an Integer to a String. Since String does not extend Integer, this is not allowed and a ClassCastException is thrown.
5. A, B, D. `java.io.IOException` is thrown by many methods in the `java.io` package, but it is always thrown programmatically. The same is true for `NumberFormatException`; it is thrown programmatically by the wrapper classes of `java.lang`. The other three exceptions are all thrown by the JVM when the corresponding problem arises.
6. C. The compiler tests the operation for a valid type but not a valid result, so the code will still compile and run. At runtime, evaluation of the parameter takes place before passing it to the `print()` method, so an `ArithmeticException` object is raised.
7. C. The `main()` method invokes `go` and A is printed on line 3. The `stop` method is invoked and E is printed on line 14. Line 16 throws a `NullPointerException`, so `stop` immediately ends and line 17 doesn't execute. The exception isn't caught in `go`, so the `go` method ends as well, but not before its `finally` block executes and C is printed on line 9. Because `main()` doesn't catch the exception, the stack trace displays and no further output occurs, so AEC was the output printed before the stack trace.
8. E. The order of catch blocks is important because they're checked in the order they appear after the try block. Because `ArithmeticException` is a child class of `RuntimeException`, the catch block on line 7 is unreachable. (If an `ArithmeticException` is thrown in try try block, it will be caught on line 5.) Line 7 generates a compiler error because it is unreachable code.
9. B. The `main()` method invokes `start` on a new Laptop object. Line 4 prints Starting up; then line 5 throws an Exception. Line 6 catches the exception, line 7 prints Problem, and then line 8 calls `System.exit`, which terminates the JVM. The `finally` block does not execute because the JVM is no longer running.
10. E. The `parseName` method is invoked within `main()` on a new Dog object. Line 4 prints 1. The try block executes and 2 is printed. Line 7 throws a `NumberFormatException`, so line 8 doesn't execute. The exception is caught on line 9, and line 10 prints 4. Because the exception is handled, execution resumes normally. `parseName` runs to completion, and line 17 executes, printing 5. That's the end of the program, so the output is 1245.
11. A. The `parseName` method is invoked on a new Cat object. Line 4 prints 1. The try block is entered, and line 6 prints 2. Line 7 throws a `NumberFormatException`. It isn't caught, so `parseName` ends. `main()` doesn't catch the exception either, so the program terminates and the stack trace for the `NumberFormatException` is printed.
12. A, B, D, G. The `main()` method invokes `run` on a new Mouse object. Line 4 prints 1 and line 6 prints 2, so options A and B are correct. Line 7 throws a `NullPointerException`, which causes line 8 to be skipped, so C is incorrect. The exception is caught on line 9 and line 10 prints 4, so option D is correct. Line 11 throws the exception again, which causes `run()` to immediately end, so line 13 doesn't execute and option E is incorrect. The `main()` method doesn't catch the exception either, so line 18 doesn't execute and option F is incorrect. The uncaught `NullPointerException` causes the stack trace to be printed, so option G is correct.

13. A, B, C, E. Classes listed in the throws part of a method declaration must extend `java.lang.Throwable`. This includes `Error`, `Exception`, and `RuntimeException`. Arbitrary classes such as `String` can't go there. Any Java type, including `Exception`, can be declared as the return type. However, this will simply return the object rather than throw an exception.
14. A, C, D, E. A method that declares an exception isn't required to throw one, making option A correct. Runtime exceptions can be thrown in any method, making options C and E correct. Option D matches the exception type declared and so is also correct. Option B is incorrect because a broader exception is not allowed.
15. A, B, D, E. `ArrayIndexOutOfBoundsException`, `IllegalArgumentException`, and `NumberFormatException` are runtime exceptions. Sorry, you have to memorize them. Any class that extends `RuntimeException` is a runtime (unchecked) exception. Classes that extend `Exception` but not `RuntimeException` are checked exceptions.
16. B. `IllegalArgumentException` is used when an unexpected parameter is passed into a method. Option A is incorrect because returning `null` or `-1` is a common return value for this scenario. Option D is incorrect because a `for` loop is typically used for this scenario. Option E is incorrect because you should find out how to code the method and not leave it for the unsuspecting programmer who calls your method. Option C is incorrect because you should run!
17. A, C, D, E. The method is allowed to throw no exceptions at all, making option A correct. It is also allowed to throw runtime exceptions, making options D and E correct. Option C is also correct since it matches the signature in the interface.
18. A, B, C, E. Checked exceptions are required to be handled or declared. Runtime exceptions are allowed to be handled or declared. Errors are allowed to be handled or declared, but this is bad practice.
19. C, E. Option C is allowed because it is a more specific type than `RuntimeException`. Option E is allowed because it isn't in the same inheritance tree as `RuntimeException`. It's not a good idea to catch either of these. Option B is not allowed because the method called inside the `try` block doesn't declare an `IOException` to be thrown. The compiler realizes that `IOException` would be an unreachable catch block. Option D is not allowed because the same exception can't be specified in two different catch blocks. Finally, option A is not allowed because it's more general than `RuntimeException` and would make that block unreachable.
20. A, E. The code begins normally and prints `a` on line 13, followed by `b` on line 15. On line 16, it throws an exception that's caught on line 17. Remember, only the most specific matching catch is run. Line 18 prints `c`, and then line 19 throws another exception. Regardless, the `finally` block runs, printing `e`. Since the `finally` block also throws an exception, that's the one printed.

Appendix

B

Study Tips





This appendix covers suggestions and recommendations for how you should prepare for the certification exam. If you're an experienced test taker, or you've taken a certification test before, most of this should be common knowledge. For those who are taking the exam for the first time, don't panic! We'll present a number of tips and strategies in this appendix to help you prepare for the exam.

Studying for the Test

Before you even sign up and take the test, you need to study the material. Studying includes the following tasks:

- Create a study plan.
- Read the Study Guide material.
- Create and run sample applications.
- Solve the Review Questions at the end of each chapter.
- Create flashcards and/or use the ones we've provided.
- Take the three practice exams.

The book is divided into chapters with corresponding exam objectives, to make it easier to assimilate. The earlier chapters on syntax and operators are especially important since they are used throughout the code samples on the exam. Unless we explicitly stated something was out of scope for the exam, you will be required to have a strong understanding of all the information in this book.

Creating a Study Plan

Rome wasn't built in a day, so you shouldn't attempt to study for only one day. Even if you have been certified with a previous version of Java, the new test includes features and components unique to Java 8 that are covered in this text.

Once you have decided to take the test, which we assume you have already since you're reading this book, you should construct a study plan that fits with your schedule. We recommend you set aside some amount of time each day, even if it's just a few minutes during lunch, to read or practice for the exam. The idea is to keep your momentum going throughout the exam preparation process. The more consistent you are in how you study, the better prepared you will be for the exam. Try to avoid taking a few days or weeks off from studying, or you're likely to spend a lot of time relearning existing material instead of moving on to new material.

Let's say you begin studying on January 1. Assuming you allot two weeks per chapter, we constructed a study plan in Table B.1 that you can use as a schedule throughout the study process. Of course, if you're new to Java, two weeks per chapter may not be enough; if you're an experienced Java developer, you may only need a few days per chapter.

TABLE B.1 Sample study plan

Date	Task
January 1–January 11	Read Introduction, Appendix B, and Chapter 1
January 12–January 14	Answer Chapter 1 Review Questions
January 15–January 25	Read Chapter 2
January 26–January 28	Answer Chapter 2 Review Questions
January 29–February 8	Read Chapter 3
February 9–February 11	Answer Chapter 3 Review Questions
February 12–February 22	Read Chapter 4
February 23–February 25	Answer Chapter 4 Review Questions
February 26–March 8	Read Chapter 5
March 9–March 11	Answer Chapter 5 Review Questions
March 12–March 22	Read Chapter 6
March 23–March 25	Answer Chapter 6 Review Questions
March 26–April 2	Take practice exams and practice with flashcards
April 3	Take exam

Your own study plan will vary based on your familiarity with Java, your personal and work schedule, and your learning abilities. The idea is to create a plan early on that has self-imposed deadlines that you can follow throughout the studying process. When someone asks how you're doing preparing for the exam, you should have a strong sense of what you've learned so far, what you're currently studying, and how many weeks you need to be prepared to take the exam.

Creating and Running Sample Applications

Although some people can learn Java just by reading a textbook, that's not how we recommend you study for a certification exam. We want you to be writing your own Java sample

applications throughout this book so that you don't just learn the material but you also understand the material. For example, it may not be obvious why the following line of code does not compile, but if you try to compile it yourself, the Java compiler will tell you the problem.

```
float value = 102.0;    // DOES NOT COMPILE
```

In this section, we will discuss how to test Java code and the tools available to assist you in this process.



A lot of people post on the CodeRanch.com forum asking, “Why does this code not compile?” and we encourage you to post the compiler error message anytime you need help. We recommend you also read the compiler message when posting, since it may provide meaningful information about why the code failed to compile.

In the previous example, the compiler failed to compile with the message `Type mismatch: cannot convert from double to float`. This message indicates that we are trying to convert a double value, `102.0`, to a float variable reference using an implicit cast. If we add an explicit cast to `(float)` or change the value to `102.0f`, the code will compile without issue.

Sample Test Class

Throughout this book, we present numerous code snippets and ask you whether they'll compile and what their output will be. These snippets are designed to be placed inside a simple Java application that starts, executes the code, and terminates. As described in Chapter 1, “Java Building Blocks,” you can accomplish this by compiling and running a public class containing a public static void `main(String[] args)` method, such as the following:

```
public class TestClass {  
    public static void main(String[] args) {  
        // Add test code here  
  
        // Add any print statements here  
        System.out.println("Hello World!");  
    }  
}
```

This application isn't particularly interesting—it just outputs “Hello World” and exits. That said, we can insert many of the code snippets present in this book in the `main()` method to determine if the code compiles, as well as what the code outputs when it does compile. We strongly recommend you become familiar with this sample application, so much so that you could write it from memory, without the comments.

We recommend that while reading this book you make note of any sections that you do not fully understand and revisit them when in front of a computer screen with a Java compiler and Java runtime. You should start by copying the code snippet into your test class, and then try experimenting with the code as much as possible. For example, we indicated the previous sample line of code would not compile, but would any of the following compile?

```
float value1 = 102;  
float value2 = (int)102.0;  
float value3 = 1f * 0.0;  
float value4 = 1f * (short)0.0;  
float value5 = 1f * (boolean)0;
```

Try out these samples on your computer and see if the result matches your expectation. Here's a hint: Two of these five lines will not compile.



Real World Scenario

IDE Software

While studying for the exam, you should develop code using a text editor and command-line Java compiler. Some of you may have existing experience with Integrated Development Environments (IDEs) such as Eclipse or IntelliJ. An IDE is a software application that facilitates software development for computer programmers.

Although such tools are extremely valuable in developing software, they can interfere with your ability to readily spot problems on the exam. For example, when a line of code does not compile, the IDE will often underline it in red, whereas on the exam, you'll have to find the line that does not compile, if there is one, on your own.

If you do choose to study with an IDE, make sure you understand everything it is doing in the background for you. For the exam, you'll need to know how to manually compile code from the command line, and this experience is rarely learned using an IDE. You'll also need to understand why the code does not compile without relying on the tips and suggestions provided by the IDE.

Identifying Your Weakest Link

The best advice we can give you to do well on the exam is to practice writing sample applications that push the limits of what you already know, as much and as often as possible. For example, if the previous samples with `float` values were too difficult for you, then you should spend even more time studying numeric promotion and casting expressions.

Prior to taking the OCA exam, you may already be an experienced Java developer, but there is a difference between being able to write Java code and being a certified Java

developer. For example, you might go years without writing a ternary expression or using an abstract class, but that does not mean they are not important features of the Java language. You may also be unaware of some of the more complex features that exist within the Java language. On top of that, there are new features to Java 8, such as lambda expressions and default interface methods, which as of this writing very few professional software developers are using.

The Review Questions in each chapter are designed to help you hone in on those features of the Java language that you may be weak in and that are required knowledge for the exam. For each chapter, you should note which questions you got wrong, understand why you got them wrong, and study those areas even more.

Often, the reason you got a question wrong on the exam is that you did not fully understand the concept. Many topics in Java have subtle rules that you often need to see for yourself to truly understand. For example, you cannot write a class that implements two interfaces that define the same default method unless you override the default method in the class. Writing and attempting to compile your own sample interfaces and classes that reference the default method may illuminate this concept far better than we could ever explain it.

Finally, we find developers who practice writing code while studying for the Java certification tend to write better Java code in their professional career. Anyone can write a Java class that can compile, but just because a class compiles does not mean it is well designed. For example, imagine a class where all class methods and variables were declared `public`, simply because the developer did not understand the other access modifiers, such as `protected` and `private`. Studying for the certification helps you to learn those features that may be applicable in your daily coding experience but that you never knew existed within the Java language.

“Overstudying” Practice Exams

Although we recommend reading this book and writing your own sample applications multiple times, redoing practice exams over and over can have a negative impact in the long run. For example, some individuals study the practice exam questions so much that they end up memorizing them. In this scenario, they can easily become overconfident—they can achieve perfect scores on the practice exams but may fail on the actual exam.

If you get a practice exam question correct, you should move on, and if you get it incorrect you should review the part of the chapter that covers it until you can answer it correctly. Remember that for legal reasons the practice exam questions are not real exam questions, so it is important you learn the material the questions are based on.

On the other hand, we recommend you repeat Review Questions as often as you like to master a chapter. Review Questions are designed to teach you important concepts in the chapter, and you should understand them completely before leaving a section. Furthermore, they help improve your ability to recognize certain types of problems present in many code snippets.

Taking the Test

Studying how to take a test can be just as important as the studying the material itself. For example, you could answer every question correctly, but only make it halfway through the exam, resulting in a failing score! If you're not historically a good test taker, or you've never taken a certification exam before, we recommend you read this section because it contains notes that are relevant to many software certification exams.

Understanding the Question

The majority of questions on the exam will contain code snippets and ask you to answer questions about them. For those containing code snippets, the number one question we recommend you answer before attempting to solve the question is:

Does the code compile?

It sounds simple but many people dive into answering the question without checking whether or not the code actually compiles. If you can determine whether or not a particular set of code compiles, and what line or lines cause it to not compile, answering the question often becomes easy.

Checking the Answers

To determine whether the code will compile, you should briefly review the answer choices to see what options are available. If there are no choices of the form "Code does not compile," then you can be reasonably assured all the lines of the code will compile and you do not need to spend time checking syntax. These questions are often, but not always, among the easiest questions because you can skip determining whether the code compiles and instead focus on what it does.

If the answer choices do include some answers of the form "Does not compile due to line 5," you should immediately focus on those lines and determine whether they compile. For example, take a look at the answer choices for the following question:

18. What is the output of the following code?

- Code Omitted -

- A. Monday
- B. Tuesday
- C. Friday
- D. The code does not compile due to line 4.
- E. The code does not compile due to line 6.

The answer choices act as a guide instructing you to focus on line 4 or 6 for compilation errors. If the question indicates only one answer choice is allowed, it also tells you at most only one line of code contains a compilation problem and the other line is correct. Although the reason line 4 or 6 may not compile could be related to other lines of code, the

key is that those other lines do not throw compiler errors themselves. By quickly browsing the list of answers, you can save time by focusing only on those lines of code that are possible candidates for not compiling.

If you are able to identify a line of code that does not compile, you will be able to finish the question a lot quicker. Often, the most difficult questions are the ones where the code does in fact compile, but one of the answer choices is “Does not compile” without indicating any line numbers. In these situations, you will have to spend extra time verifying that each and every line compiles. If they are taking too much time, we recommend marking these for “Review” and coming back to them later.

Determining What the Question Is Asking

A lot of times, a question may appear to be asking one thing but will actually be asking another. For example, the following question may appear to be asking about method overloading and abstract classes:

12. What is the output of the following code?

```
1: abstract class Mammal {  
2:     protected boolean hasFur() { return false; }  
3: }  
4: class Capybara implements Mammal {  
5:     public boolean hasFur() { return true; }  
6:     public static void main(String[] args) {  
7:         System.out.println(new Capybara().hasFur());  
8:     }  
9: }
```

It turns out this question is a lot simpler than it looks. A class cannot implement another class—it can only extend another class—so line 4 will cause the code to fail to compile. If you notice this compiler problem early on, you’ll likely be able to answer this question quickly and easily.

Taking Advantage of Context Clues

Let’s face it—there will be things you’re likely to forget on the day of the exam. Between being nervous about taking the test and being a bit overtired when you read a particular chapter, you’re likely to encounter at least one question where you do not have a high degree of confidence. Luckily, you do not need to score a perfect 100% to pass.

One advanced test-taking skill that can come in handy is to use information from one question to help answer another. For example, we mentioned in an earlier section that you can assume a question’s code block will compile and run if “Does not compile” and “Throw an exception at runtime” are not available in the list of answers. If you have a piece of code that you know compiles and a related piece of code that you’re not so sure about, you can use information from the former question to help solve the latter question.

Use a similar strategy when a question asks which single line will not compile. If you're able to determine the line that does not compile with some degree of confidence, you can use the remaining code that you know does compile as a guide to help answer other questions.

By using context clues of other questions on the exam, you may be able to more easily solve questions that you are unsure about.

Reviewing Common Compiler Issues

The following is a brief list of common things to look for when trying to determine whether code compiles. Bear in mind that this is only a partial list. We recommend you review each chapter for a comprehensive list of reasons that code will not compile. Also, if you have not finished reading the book, you should set aside this list and return to it when you are preparing to take the exam.

Common Tips to Determine if Code Compiles:

- Keep an eye out for all reserved words. [Chapter 1]
- Verify brackets—{ }—and parentheses—()—are being used correctly. [Chapter 1]
- Verify new is used appropriately for creating objects. [Chapter 1]
- Ignore all line indentation especially with if-then statements that do not use brackets { }. [Chapter 2]
- Make sure operators use compatible data types, such as the logical complement operator (!) only applied to boolean values, and arithmetic operators (+, -, ++, --) only applied to numeric values. [Chapter 2]
- For any numeric operators, check for automatic numeric promotion and order of operation when evaluating an expression. [Chapter 2]
- Verify switch statements use acceptable data types. [Chapter 2]
- Remember == is not the same as equals(). [Chapter 3]
- String values are immutable. [Chapter 3]
- Non-void methods must return a value that matches or is a subclass of the return type of the method. [Chapter 4]
- If two classes are involved, make sure access modifiers allow proper access of variables and methods. [Chapter 4]
- Nonstatic methods and variables require an object instance to access. [Chapter 4]
- If a class is missing a default no-argument constructor or the provided constructors do not explicitly call super(), assume the compiler will automatically insert them. [Chapter 5]
- Make sure abstract methods do not define an implementation, and likewise concrete methods always define an implementation. [Chapter 5]

- You implement an interface and extend a class. [Chapter 5]
- A class can be cast to a reference of any superclass it inherits from or interface it implements. [Chapter 5]
- Checked exceptions must be caught; unchecked exceptions may be caught but do not need to be. [Chapter 6]
- try blocks require a catch and/or finally block for the OCA exam. [Chapter 6]

We have listed the chapter each tip is found in so that you can go back and review any that you do not fully understand. Once you've determined that the code does in fact compile, proceed with tracing the application logic and trying to determine what the code actually does.

Applying Process of Elimination


Although you might not immediately know the correct answer to a question, if you can reduce the question from five answers down to three, your odds of guessing the correct answer will be markedly improved. For example, if you can reduce a question from four answers to two answers, you double your chances of guessing the correct answer. In this section, we will discuss how to apply the process of elimination to help improve your score.

Using the Provided Writing Material

Depending on your particular testing center, you may be provided with a stack of blank paper or a whiteboard to use to help you answer questions. If you sit down and are not provided with anything, please make sure to ask for such materials.

After determining whether a question compiles and what it is asking for, you should then jot down a list of all the answers. You should then proceed to cross out the ones you know are not correct. We provided a sample of what this might look like in Figure B.1.

FIGURE B.1 Eliminating answer choices



If you're using paper and you decide to come back to this question, be sure to write down the question number and save it for later. If you're using a whiteboard and decide to come back to a question later, you may have to redo some of the work, given the limited space on a whiteboard. For those questions you want to come back to later, we suggest jotting down the remaining answer choices on the side of the whiteboard. Some test-taking

software allows you to mark and save which answer choices you've eliminated, although in our experience this does not always work reliably in practice.



Although you aren't allowed to bring any written notes with you into the exam, you're allowed to write things down you remember at the start of the exam on the provided writing material. If there's a particular facet of the Java language that you have difficulty remembering, try memorizing it before the exam and write it down as soon as the exam starts. You can then use it as a guide for the rest of the exam. Of course, this strategy only works for a handful of topics, since there's a limit to what you're likely to remember in a short time.

For example, you may have trouble remembering the list of acceptable data types in `switch` statements. If so, we recommend you memorize that information before the exam and write it down as soon as the exam starts for use in various questions.

Understanding Relationships Between Answers

The exam writers, as well as the writers of this book, are fond of answers that are related to each other. We can apply the process of elimination to remove entire sets of answers from selection, not just a single answer. For example, take a look at the following question:

22. What is the output of the following application?

```
3: int x = 0;
4: while(++x < 5) { x+=1; }
5: String message = x > 5 ? "Greater than" : "Less Than";
6: System.out.println(message+", "+x);
```

- A. Greater than,5
- B. Greater than,6
- C. Greater than,7
- D. Less than,5
- E. Less than,6
- F. Less than,7

In this question, notice that half of the answers output `Greater than`, whereas the other half output `Less than`. Based on the code, as well as the answers available, the question cannot output both values. That means if you can determine what the ternary expression on line 5 evaluates to, you can eliminate half the answers!

You might also notice that this particular question does not include any “Does not compile” or “Code throws an exception at runtime” answers, meaning you can be assured this snippet of code does compile and run without issue. If you have a question similar to this, you can compare the syntax and use this as a guide for solving other related questions.

Guessing the Correct Answer

Unlike with some other standardized tests, there's no penalty for answering a question incorrectly versus leaving it blank. If you're nearly out of time, or you just can't decide on an answer, select a random answer and move on. If you've been able to eliminate even one answer, then your guess will be better than blind luck.

Answer All Questions!

You should set a hard stop of 5 minutes of time remaining on the exam to ensure that you've answered each and every question. Remember, if you fail to answer a question you'll definitely get it wrong and lose points, but if you guess, there's at least a chance you'll be correct. There's no harm in guessing!

When in doubt, we generally recommend picking a random answer that includes "Does not compile" if available, although which choice you select is not nearly as important as making sure to not leave any unanswered questions on the exam!

Optimizing Your Time

One of the most difficult test-taking skills to master is balancing your time on the exam. Although Oracle often varies the precise number of questions on the exam and the amount of time you have to answer them, the general rule of thumb is that you have about one and half minutes per question.

Of course, it can be stressful to frequently look at the time remaining while taking the exam, so the key is pacing yourself. Some questions will take you longer than two minutes to solve, but hopefully others will only take less than a minute. The more time you save on the easier questions, the more time you'll have for the harder questions.

Checking the Time Remaining

The exam software includes a clock that tells you the amount of time you have left on the exam. We don't recommend checking the clock after each and every question to determine your pace. After all, doing such a calculation will waste time and probably make you nervous and stressed out. We do recommend you check the time remaining at certain points while taking the exam to determine whether you should try to increase your pace.

For example, if the exam lasts two hours and is 90 questions long, the following would be a good pace to try to keep.

- 120 Minutes Remaining: Start exam.
- 90 Minutes Remaining: One third of the exam finished.

- 60 Minutes Remaining: Two thirds of the exam finished.
- 30 Minutes Remaining: First pass of all questions complete.
- 5 Minutes Remaining: Finished reviewing all questions marked for “Review.” Select answers to all questions left blank.

As you’re taking the exam you may realize you’re falling behind. In this scenario, you need to start allotting less time per question, which may involve more guessing, or you’ll end up with some questions that you never even answered. As discussed in the previous section, guessing an answer to a question is better than not answering the question at all.

Skipping Hard Questions

If you do find you are having difficulty with a particular set of questions, just skip them. The exam provides a feature to mark questions for “Review” that you can later come back to. Remember that all questions on the exam, easy or difficult, are weighted the same. It is a far better use of your time to spend five minutes answering ten easy questions than the same amount of time answering one difficult question.

You might come to a question that looks difficult and immediately realize it is going to take a lot of time. In this case, skip it before even starting on it. You can save the most difficult problems for the end so that you can get all the easy ones solved early on. Of course, you shouldn’t mark every question for “Review,” so use that sparingly. For example, if you only need 30 more seconds to solve a specific question, it is better to finish it so you do not have to come back to it later. The trick is to not get stuck on a difficult question for a long period of time.

Improving Your Test-Taking Speed

Answering certification exam questions quickly does not come naturally to most people. It takes a bit of practice and skill to look at a question, a code sample, and 4–6 answers, and be able to answer it within a minute or two. The best way to practice is to keep solving the review questions at the end of each chapter until you can read, understand, and answer them in under a minute.

Once you’ve completed all of the material and practiced with the review questions enough that you can answer them quickly and correctly, you should try one of the three 60-question practice exams that come with this Study Guide. You should treat it like the real exam, setting aside two hours and finishing it in one sitting.

Although we recommend you try to avoid taking the practice exams so much that you memorize the questions and answers, we do recommend you keep taking them until you can finish each practice exam in under two hours. Remember not to move on to the next one until you can pass the previous exam in the allotted time. If not, study more and go back to drilling on the Review Questions. The idea is that you want to be good at quickly reading through the question, honing in on the key concept the question is asking, and being able to select the answer that best represents it.

Getting a Good Night's Rest

Although a lot of people are inclined to cram as much material as they can in the hours leading up to the exam, most studies have shown that this is a poor test-taking strategy. The best thing we can recommend you do before the exam is to get a good night's rest!

Given the length of the exam and number of questions, the exam can be quite draining, especially if this is your first time taking a certification exam. You might come in expecting to be done 30 minutes early, only to discover you are only a quarter of the way through the exam with half the time remaining. At some point, you may begin to panic, and it is in these moments that these test-taking skills are most important. Just remember to take a deep breath, stay calm, eliminate as many wrong answers as you can, and make sure to answer each and every question. It is for stressful moments like these that being well rested with a good night's sleep will be most beneficial!



Index

Note to the Reader: Throughout this index **boldfaced** page numbers indicate primary discussions of a topic. *Italicized* page numbers indicate illustrations.

A

- abstract classes
 - concrete classes from, 262–263
 - creating, 259
 - defining, 260–262
 - extending, 263–265
- abstract specifiers
 - interfaces, 267–269
 - methods, 168, 271–273
- access modifiers, 173
 - default, 175
 - description, 7
 - inheritance, 237
 - methods, 166–167
 - private, 173, 174
 - protected, 176, 176
 - public, 180–181
 - static, 181–188
- add() method, 130–131
- addition
 - dates, 142–143
 - precedence, 53–54
- ampersands (&) for logical operators, 64, 64
- ancestors in inheritance, 234
- “and” logical operator, 64, 64
- anonymous arrays, 120
- append() method, 115
- args parameters in main(), 7–8
- arithmetic operators, 53–55
- ArithmeticException class, 314
- ArrayIndexOutOfBoundsException class, 123, 314–315
- ArrayLists, 129
 - arrays from, 136–137
 - autoboxing, 136
 - creating, 129
 - methods, 130–134
 - sorting, 138
 - wrapper classes, 134–135
- arrays, 119
 - ArrayLists, 129
 - declaring, 121
 - indexes, 7, 120, 120

- multidimensional, 126–129, 127–128
 - parameters, 7
 - primitives, 119–121, 119–120
 - reference variables, 121, 122
 - searching, 125–126
 - sorting, 124–125
 - varargs, 126
 - working with, 123–124
- arrow operator (->) for lambda expressions, 211–212, 212
- asList() method, 187–188
- assignment operators
 - compound, 62–63
 - overview, 60
 - precedence, 53
- assumed keywords in interfaces, 268
- asterisks (*)
 - comments, 4–5
 - packages, 10
- asymmetric arrays, 128, 128
- autoboxing
 - overloading methods, 193
 - type conversions, 136

B

- B prefix for binary, 22
- backed lists, 137
- base 10 numbering system, 22
- Beginning Java forum, 6
- binary number format, 22–23
- binary operators
 - arithmetic, 53–55
 - assignment, 60
 - casting, 60–61
 - compound assignment operators, 62–63
 - equality, 65–66
 - logical, 64–65, 64
 - numeric promotion, 55–57
 - relational, 63
- binary searches, 125–126
- bitwise operators, 64, 64

- blocks
 - description, 67
 - initialization order, 19–20
 - instance initializer, 18
 - scope in, 31–34
- bodies
 - lambda expressions, 212, 212
 - methods, 171–172
- boolean type
 - default initialization, 31
 - if statements, 70
 - logical operators, 57–58
 - size and range, 21
 - ternary operators, 71–72
 - while statements, 77
 - wrapper classes, 134–135
- braces ({})
 - blocks, 18–19, 32–33
 - if-then statements, 68
 - lambda expressions, 212, 212
 - methods, 171
 - try statements, 306–307
- brackets ([]) for arrays, 7, 119–122
- break statements
 - in loops, 88–90, 88
 - in switch, 74
- byte type
 - default initialization, 31
 - promotion rules, 56
 - size and range, 21
 - switch statements, 72–73
 - wrapper classes, 134–135
- bytecode, 6

C

- calling
 - constructors, 243–244
 - inherited class members, 244–245
 - static variables and methods, 182–183
- CamelCase names, 28–29
- capacity vs. size, 113–114, 114
- carets (^) for logical operators, 64, 64
- case sensitivity of names, 27
- case statements in switch, 73–76, 73
- casting
 - objects, 282–284
 - operators, 60–61
- catch blocks, 305–309, 305
- chaining
 - constructors, 201
 - methods, 110–111
 - StringBuilder, 112–113
- char type
 - default initialization, 31
 - promotion rules, 56
 - size and range, 21
 - switch statements, 72–73
 - wrapper classes, 134–135
- charAt() method, 106, 114, 135
- checked exceptions, 303, 305, 317
- child classes in inheritance, 234
- child packages, 10
- .class extension, 6
- class variables in default initialization, 30–31
- ClassCastException class, 315
- classes
 - abstract, 259–265
 - concrete, 262–263
 - description, 2
 - element ordering, 34–35
 - extending, 235–236, 235
 - fields and methods, 2–4
 - vs. files, 5
 - immutable, 207–208
 - inheritance. *See* inheritance
 - interfaces. *See* interfaces
 - packages, 9
 - paths, 15
 - subclasses, 319–321
 - wrapper, 134–136
- clear() method, 133
- closures. *See* lambda expressions
- code
 - blocks, 18–19
 - compiling, 6, 14–15
 - exam formatting, 16
- colons (:)
 - dates and times, 150
 - labels, 88
 - paths, 15
 - ternary operators, 71
- commas (,)
 - dates and times, 150
 - exception lists, 171
 - interface implementation, 267
 - parameter lists, 171
 - variable declarations, 26

- comments, 4
- comparators in arrays, 125
- compile-time constant values in switch statements, 73–76
- compiling code
 - extensions for, 6
 - and inheritance, 241–242
 - with packages, 14–15
- compound assignment operators, 62–63
- concatenating strings, 102–104
- concrete classes, 262–263
- conflicts in names, 12–13
- consistency of names, 28
- constant values in switch statements, 73–76
- constructors
 - calling, 243–244
 - chaining, 201
 - creating, 196
 - date and time, 141
 - default, 197–199
 - defining, 238
 - definition rules, 242–243
 - final fields, 202
 - initialization order, 202–204
 - objects, 17
 - overloading, 199–200
 - wrapper classes, 134–135
- contains() method
 - ArrayLists, 133
 - strings, 109–110
- continue statements, 90–91, 90
- control flow statements
 - break, 88–90, 88
 - continue, 90–91, 90
 - do-while, 78–80, 78
 - for, 80–83, 80, 85
 - for-each, 83–86, 83
 - if-then, 67–68, 67
 - if-then-else, 68–70, 69
 - labels, 87–89
 - nested loops, 87
 - switch, 72–76, 73
 - while, 76–77, 77, 79
- converting
 - ArrayLists to arrays, 136–137
 - number systems, 23
 - objects to wrapper classes, 135–136
- covariant return types, 248, 252
- curly braces ({})
 - blocks, 18–19, 32–33
 - if-then statements, 68

- lambda expressions, 212, 212
- methods, 171
- try statements, 306–307

D

- data types
 - in for statements, 82
 - promotion rules, 55–57
 - switch statements, 72–73
- Date class, 12–13
- dates
 - creating, 138–142
 - earlier versions, 141–142, 144, 150
 - formatting, 148–150
 - manipulating, 142–145
 - parsing, 151
 - periods, 145–147
- DateTimeFormatter class, 148–150
- dd format for dates and times, 150
- decimal numbering system, 22
- declare rules for exceptions, 303
- declaring
 - arrays, 121
 - methods, 166, 166
 - variables, 25
- decrement operators, 58
- default statements in switch, 72, 74
- defaults
 - access modifiers, 167, 175
 - constructors, 17, 197–199
 - interface methods, 274–278
 - packages, 13
 - variable initialization, 29
- defining
 - abstract classes, 260–262
 - interfaces, 267–269
- delete() method, 116
- deleteCharAt() method, 116
- descendants in inheritance, 234
- destroying objects
 - finalize(), 38–39
 - garbage collection, 36, 37–38
- diamond operator (<>), 130
- division
 - modulus operation, 54–55
 - precedence, 53–54
- do-while statements, 78–80, 78

- dollar signs (\$) for names, 27, 29
- double quotes (")
 - for spaces, 8
 - strings, 102–103
- double type
 - default initialization, 31
 - size and range, 21
 - wrapper classes, 134–135

E

- else operators, 68–70, 69
- encapsulation
 - description, 39
 - immutable classes, 207–208
 - overview, 205–206
- endsWith() method, 109
- enum type
 - dates, 140
 - support for, 72–73
 - switch statements, 72
- equal signs (=)
 - assignment operators, 60
 - compound assignment operators, 62–63
 - equality, 117–118
 - equality operators, 65–66
 - precedence, 53
 - relational operators, 63
 - unary operator, 57
 - variable declarations, 25
- equality
 - operators, 65–66
 - overview, 117–119
- equals() method
 - ArrayLists, 133–134
 - arrays, 121
 - strings, 109, 118–119
- equalsIgnoreCase() method, 109
- Error class, 302, 302, 305, 317–318
- exams, code formatting in, 16
- exception lists for methods, 171
- ExceptionInInitializerError class, 317
- exceptions
 - checked, 317
 - Error, 317–318
 - finally blocks, 307–309, 307
 - ignoring, 322
 - methods, 318–322
 - multiple, 311–313
 - printing, 321, 322
 - reasons, 300
 - vs. return codes, 301–302
 - roles, 300–301
 - runtime, 314–316
 - subclasses, 319–321
 - throwing, 304–305
 - try statements, 305–307, 305
 - types, 302–303, 302, 309
- exclamation points (!)
 - equality operators, 65–66
 - logical operator, 57–58
 - unary operator, 57
- “exclusive or” operator, 64, 64
- exit() method, 309
- extending
 - abstract classes, 263–265
 - classes, 235–236, 235
 - interfaces, 269

F

- false value in logical operators, 57
- fields
 - in classes, 35, 41
 - final, 202
 - initialization order, 19–20
 - overview, 2–4
 - reading and writing, 18
 - static, 181–182
- FileNotFoundException class, 317
- files vs. classes, 5
- final specifiers
 - abstract classes, 261
 - constants, 74
 - constructor fields, 202
 - interfaces, 267, 273–274
 - methods, 168, 256
- finalize() method, 38
- finally blocks, 307–309, 307
- float type
 - default initialization, 31
 - promotion rules, 55
 - size and range, 21
 - wrapper classes, 134–135
- for statements, 80–83, 80, 85
- for-each statements, 83–86, 83
- format() method, 148
- formatting dates and times, 148–150
- free store, 36

functional interfaces with lambda expressions, 214
 functional programming, 208–209
 functions. *See* methods

G

garbage collection, 36, 37–38
 gc() method, 36
 generics, 130
 getter methods in JavaBeans, 206
 getting variables, 18
 GMT (Greenwich Mean Time), 146
 greater than signs (>)
 diamond operator, 130
 lambda expressions, 211–212, 212
 relational operators, 63
 Greenwich Mean Time (GMT), 146

H

handle rules for exceptions, 303
 heap, 36
 hexadecimal number format, 22–23
 hh format for dates and times, 150
 hiding
 vs. overriding, 254–255
 static methods, 252–254
 variables, 257

I

identifiers, 27
 if-then statements, 67–68, 67
 if-then-else statements, 68–70, 69
 ignored returned values, 191
 IllegalArgumentException class, 315–316
 immutability of strings, 104–105
 immutable classes, creating, 207–208
 imports
 element ordering, 35
 packages, 9, 11–12
 static, 187–188
 “inclusive or” operator, 64, 64
 increment operators, 58
 indentation in if-then statements, 68

indexes
 arrays, 7, 120, 120
 strings, 105–107, 106
 indexOf() method, 106–107, 114
 infinite loops, 78
 infinite recursion, 317
 inheritance
 access modifiers, 237
 calling class members, 244–245
 compiler enhancements, 241–242
 constructors, 238, 242–244
 extending classes, 235–236, 235
 interfaces, 269–273, 277
 methods, 246
 object creation, 237–238, 238
 overview, 234, 235
 variables, 257
 initialization
 constructor order, 202–204
 default variables, 29
 object order, 19–20
 static, 186–187
 variables, 25–26
 initialization blocks in for statements, 80–82, 80
 initializer blocks, instances in, 18–19
 insert() method, 115
 instance members vs. static, 183–185
 instance variables
 default initialization, 30–31
 scope, 33
 instanceof operator, 63, 283
 instances
 creating, 17
 initialization order, 19
 initializer blocks, 18–19
 objects, 2
 instantiation process, 196
 int type
 default initialization, 31
 promotion rules, 55
 size and range, 21–22
 switch statements, 72–73
 wrapper classes, 134–135
 integrated development environment (IDE), 14
 interfaces
 abstract methods, 271–273
 with classes, 270
 default methods, 274–278
 defining, 267–269
 extending, 269
 implementing, 266, 266
 inheriting, 269–273

- lambda expressions, 214
- multiple inheritance, 277
- static methods, 278
- variables, 273–274
- intValue() method, 135
- IOException class, 317
- isEmpty() method, 132–133

J

- JAR files, 15
- Java Development Kit (JDK), 6
- .java extension, 6
- java.lang package, 11
- Java Runtime Environment (JRE), 6
- Java Virtual Machine (JVM), 6
- JavaBeans names, 205–206
- javac command, 15
- Javadoc comments
 - description, 4
 - key benefits, 39–40
- JDK (Java Development Kit), 6
- JRE (Java Runtime Environment), 6
- JVM (Java Virtual Machine), 6

K

- keywords, 3

L

- L suffix for long, 22
- labels, 87–89
- lambda expressions
 - example, 209–211
 - predicates, 214–215
 - syntax, 211–213, 212
 - variable access, 213
 - writing, 208–209
- length() method, 106, 114
- less than signs (<)
 - diamond operator, 130
 - relational operators, 63
- lists. *See* ArrayLists
- literal values
 - primitive types, 22–23
 - strings, 105, 117–118

- switch statements, 73–75
- local variables, 29
- LocalDate, 138–141, 146
- LocalDateTime, 139–141, 146
- LocalTime, 139–140
- logical operators, 64
 - complement, 57
 - overview, 64–65
 - precedence, 53
- long type
 - dates and times, 146
 - default initialization, 31
 - size and range, 21–22
 - wrapper classes, 134–135
- loops
 - arrays, 123
 - break statements, 88–90, 88
 - continue statements, 90–91, 90
 - do-while, 78–80, 78
 - for, 80, 85
 - for-each, 83–86, 83
 - infinite, 78
 - labels, 87–88
 - nested, 87
 - while, 76–77, 77, 79

M

- main() method, 6
- MAX_VALUE constant, 22
- MEDIUM format for dates and times, 149
- members in classes, 2
- memory
 - object destruction, 36, 37–38
 - reference types, 24, 24
- methods
 - abstract, 259, 271–273
 - access modifiers. *See* access modifiers
 - ArrayLists, 130–134
 - bodies, 171–172
 - chaining, 110–111
 - constructors. *See* constructors
 - dates and times, 144
 - declaring, 166, 166
 - defining, 3
 - designing, 166, 166
 - exception lists, 171
 - exceptions thrown by, 318–322
 - final, 256
 - hiding, 252–254
 - interfaces, 271–279

- main(), 6
- names, 170
- optional specifiers, 168–169
- overloading, 191
- overriding, 246, 319–321
- overriding vs. hiding, 254–255
- overview, 2–4
- parameter lists, 171
- passing data, 188–191, 190
- polymorphism, 287–288
- redeclaring, 251–252
- return types, 169–170
- signatures, 3, 7, 166, 166
- static, 181–182
- StringBuilder class, 114–117
- strings, 105–109
- varargs, 172–173
- virtual, 284–285
- minus signs (-)
 - compound assignment operators, 62
 - increment operators, 58–59
 - lambda expressions, 211–212, 212
 - negation operator, 57–58
 - unary operator, 57
- mm format for dates and times, 150
- MMMM format for dates and times, 150
- modulus operator (%)
 - overview, 54–55
 - precedence, 53–54
- multidimensional arrays, 126–129, 127–128
- multiple exceptions, 311–313
- multiple inheritance
 - description, 234–235, 235
 - interfaces, 271–273, 277
- multiple-line comments, 4–5
- multiple variables, declaring, 26
- multiplication, 53–54
- mutability of StringBuilder, 112–113

N

- names
 - conflicts, 12–13
 - constructors, 17
 - identifiers, 27
 - JavaBeans, 205–206
 - methods, 170
 - packages, 10, 12–13
- native specifiers, 168
- negation operator, 57
- nested loops, 87

- new keyword for constructors, 17
- no-argument constructors, 198
- NoClassDefFoundError class, 318
- not equals operators, 65–66
- now() method, 139
- null values
 - autoboxing, 136
 - reference types, 24
- NullPointerException class, 303, 316
- NumberFormatException class, 316
- numbering systems, 22–23
- numeric promotion, 55–57

O

- object-oriented languages, 39
- objects
 - casting, 282–284
 - comparing, 66
 - constructors, 17
 - creating, 16, 237–238, 238
 - description, 2
 - destroying, 36–39, 37–38
 - initialization order, 19–20
 - instance initializer blocks, 18–19
 - polymorphism, 281–282, 282
 - primitive types, 20–23
 - reading and writing fields, 18
 - vs. references, 36, 37
- octal number format, 22
- ofLocalized methods, 149
- operators
 - arithmetic, 53–55
 - assignment, 60, 62
 - casting, 60–61
 - equality, 65–66
 - logical, 64–65, 64
 - numeric promotion, 55–57
 - order, 52
 - overview, 52
 - relational, 63
 - ternary, 71
 - unary, 57–59
- optional specifiers, 168–169
- “or” operator, 64, 64
- order
 - class elements, 34–35
 - constructor initialization, 202–204
 - field initialization, 19–20
 - operators, 52

- overflow, 61
- overloading methods, 191
 - autoboxing, 193
 - constructors, 199–200
 - vs. overriding, 248
 - overview, 194–196
 - primitive types, 194
 - reference types, 193–194
 - varargs, 192–193
- overriding
 - vs. hiding, 254–255
 - methods, 246, 319–321
 - vs. overloading, 248
 - polymorphism, 287–288

P

- package private access modifiers, 167, 175
- packages
 - compiling code with, 14–15
 - creating, 13–14
 - element ordering, 35
 - names, 10, 12–13
 - overview, 9
 - redundant imports, 11–12
 - wildcards, 10–11
- parameters
 - main(), 7
 - methods, 3, 171
 - overloading methods, 192–193, 199
 - polymorphic, 285–286
- parent classes in inheritance, 234
- parentheses ()
 - lambda expressions, 212–213, 212
 - methods, 167
 - precedence, 54
- parse() method, 151
- parseInt() method, 135
- parsing dates and times, 151
- pass-by-reference languages, 190–191
- pass-by-value languages, 188–190, 190
- passing data in methods, 188–191, 190
- paths for classes, 15
- percent signs (%) for modulus, 54
- Period class, 146–147
- periods in dates and times, 145–147
- platform independence, 39–40

- plus methods for dates, 144
- plus signs (+)
 - compound assignment operators, 62
 - increment operators, 58–59
 - string concatenation, 103
 - unary operator, 57
- pointers, 24
- polymorphism
 - casting objects, 282–284
 - method overriding, 287–288
 - objects vs. references, 281–282, 282
 - overview, 279
 - parameters, 285–286
 - virtual methods, 284–285
- pools, strings, 105
- post-decrement operators, 58–59
- post-increment operators, 58–59
- pre-decrement operators, 58–59
- pre-increment operators, 58–59
- precedence of operators, 52
- predicates in lambda expressions, 214–215
- primitive types
 - arrays, 119–121, 119–120
 - casting, 60–61
 - overloading methods, 194
 - overview, 20–23
 - vs. reference, 25
- printing exceptions, 321, 322
- private access modifiers
 - abstract classes, 261
 - inheritance, 237
 - interfaces, 267–269
 - methods, 167
 - overview, 173, 174
- private methods, redeclaring, 251–252
- procedural languages, 39
- procedures. *See* methods
- promotion, numeric, 55–57
- properties in JavaBeans, 205
- protected access modifiers
 - interfaces, 267–269
 - methods, 167
 - overview, 176, 176
- public access modifiers
 - inheritance, 237
 - interfaces, 267–268, 273–274
 - methods, 167
 - overview, 7, 180–181

Q

question marks (?) for ternary operators, 71
 quotes (")
 for spaces, 8
 strings, 102–103

R

ranges of primitive types, 21
 readability
 if-then statements, 68
 literal values, 23
 names, 28
 polymorphic parameters, 286
 reading fields, 18
 recursion, infinite, 317
 recursive functions, 247
 redeclaring private methods, 251–252
 redundant package imports, 11–12
 references
 in arrays, 121, 122
 comparing, 66
 methods, 189–191, 190
 vs. objects, 36, 37
 overloading methods, 194
 overview, 24, 24
 polymorphism, 281–282, 282
 vs. primitive, 25
 relational operators
 overview, 63
 precedence, 53
 remainder operator (%)
 overview, 54–55
 precedence, 53–54
 remove() method, 131–132
 replace() method, 110
 reserved words, 27
 return codes vs. exceptions, 301–302
 return statement in methods, 169
 return types and values
 constructors, 17
 covariant, 248, 252
 description, 7
 ignored, 191
 immutable classes, 207–208

 methods, 169–170
 reverse() method, 116
 robustness, 40
 runtime exceptions, 314–316
 RuntimeException class, 302–303, 302, 305

S

scope
 in for statements, 81
 variables, 31–34
 searching arrays, 125–126
 security, 40
 semicolons (;)
 class paths, 15
 in for statements, 81
 lambda expressions, 212–213, 212
 variable declarations, 27
 set() method, 132
 setter methods in JavaBeans, 206
 setting variables, 18
 shift operators, 53
 short-circuit logical operators
 description, 64–65
 precedence, 53
 SHORT format for dates and times, 149
 short type
 default initialization, 31
 promotion rules, 56
 size and range, 21
 switch statements, 72–73
 wrapper classes, 134–135
 signatures in methods, 3, 7, 166, 166
 simplicity, 40
 single inheritance, 234, 235
 single-line comments, 4
 size
 vs. capacity, 113–114, 114
 primitive types, 21
 size() method, 132–133
 slashes (/)
 comments, 4–5
 division, 53
 precedence, 54
 sorting
 ArrayLists, 138

- arrays, 124–125
- spaces in arguments, 8
- square brackets ([]) for arrays, 7, 119–122
- StackOverflowError class, 318
- startsWith() method, 109
- statements
 - do-while, 78–80, 78
 - for, 80–83, 80, 85
 - for-each, 83–86, 83
 - if-then, 67–68, 67
 - if-then-else, 68–70, 69
 - overview, 66
 - switch, 72–76, 73
 - while, 76–77, 77, 79
- static keyword, 7
- static methods, hiding, 252–254
- static specifiers
 - calling variables and methods, 182–183
 - designing, 181–182
 - imports, 187–188
 - initialization, 186–187
 - vs. instance, 183–185
 - interfaces, 273–274, 278
 - methods, 168
 - variables, 185
- strictfp specifiers, 168
- String class, 102
- StringBuffer class vs, 117
- StringBuilder class, 111–112
 - creating, 113–114
 - methods, 114–117
 - mutability and chaining, 112–113
 - vs. StringBuffer class, 117
- strings
 - as arrays, 119
 - concatenating, 102–104
 - description, 3, 102
 - immutability, 104–105
 - methods, 105–109
 - pools, 105
 - StringBuilder class, 111–117
 - switch statements, 73
- subclasses with exceptions, 319–321
- substring() method, 107–108, 114
- subtraction
 - days, 145
 - precedence, 53–54
- super() keyword, 239–243, 246
- super keyword, 246–247

- switch statements, 72
 - compile-time constant values, 73–76
 - data types, 72–73, 73
- synchronized specifiers, 168
- System.gc() method, 36

T

- ternary operators
 - overview, 71
 - precedence, 53
- this() keyword, 238
- Throwable class, 302, 302
- throwing exceptions, 304–305
- times
 - creating, 138–142
 - formatting, 148–150
 - manipulating, 142–145
 - parsing, 151
 - periods, 145–147
- toEpochDay() method, 146
- toEpochTime() method, 146
- toLowerCase() method, 108
- toString() method, 105, 116–117
- toUpperCase() method, 108
- trim() method, 110
- true value in logical operators, 57
- truth tables, 64, 64
- try statements, 305–307, 305
- two-multidimensional arrays, 127

U

- unary operators
 - increment and decrement, 58–59
 - logical complement and negation, 57–58
 - precedence, 53
 - promotion rules, 56
 - working with, 57
- unchecked exceptions, 302
- underflow, 61
- underscores (_)
 - in literal values, 23
 - names, 27
- Unicode character set, 29
- update statements in for statements, 80–82, 80

V

- valueOf() method, 135
- variable argument (varargs) lists, 7–8
 - ArrayLists, 137
 - arrays, 126
 - methods, 172–173
 - overloading methods, 192–193
- variables
 - assignment operators, 60, 62–63
 - declaring, 25
 - default initialization, 29
 - defining, 3
 - description, 2
 - in for statements, 81–83
 - for-each statements, 83
 - hiding, 257
 - inheriting, 257
 - interfaces, 273–274
 - lambda expression access, 213
 - reading and writing, 18
 - scope, 31–34
 - static, 185
- vertical bars (|) for logical operators, 64, 64
- virtual methods, 284–285
- void type, 7
 - description, 3
 - methods, 169

W

- while statements, 76–77, 77, 79
- wildcards
 - JAR files, 15
 - packages, 10–11
- wrapper classes
 - ArrayLists, 134–135
 - converting to, 135–136
 - switch statements, 73
- “write once, run everywhere”, 40
- writing fields, 18

X

- X for hexadecimal, 22

Z

- ZonedDateTime, 139

Free Online Learning Environment

Register on Sybex.com to gain access to the free online interactive learning environment and test bank to help you study for your OCA Java SE 8 Programmer I certification.

The online test bank includes:

- Assessment Test to help you focus your study to specific objectives
- Chapter Tests to reinforce what you learned
- Practice Exams to test your knowledge of the material
- Electronic Flashcards to reinforce your learning and provide last-minute test prep before the exam
- Searchable Glossary gives you instant access to the key terms you'll need to know for the exam

Go to www.sybex.com/go/ocejavase8 to register and gain access to this comprehensive study tool package.



 **SYBEX**
An Imprint of  **WILEY**

WILEY END USER LICENSE AGREEMENT

Go to www.wiley.com/go/eula to access Wiley's ebook EULA.